Building RAG Agents with LLMs

Jae Y. CHOI

NVIDIA DLI Certified Instructor & University Ambassador

Course Introduction

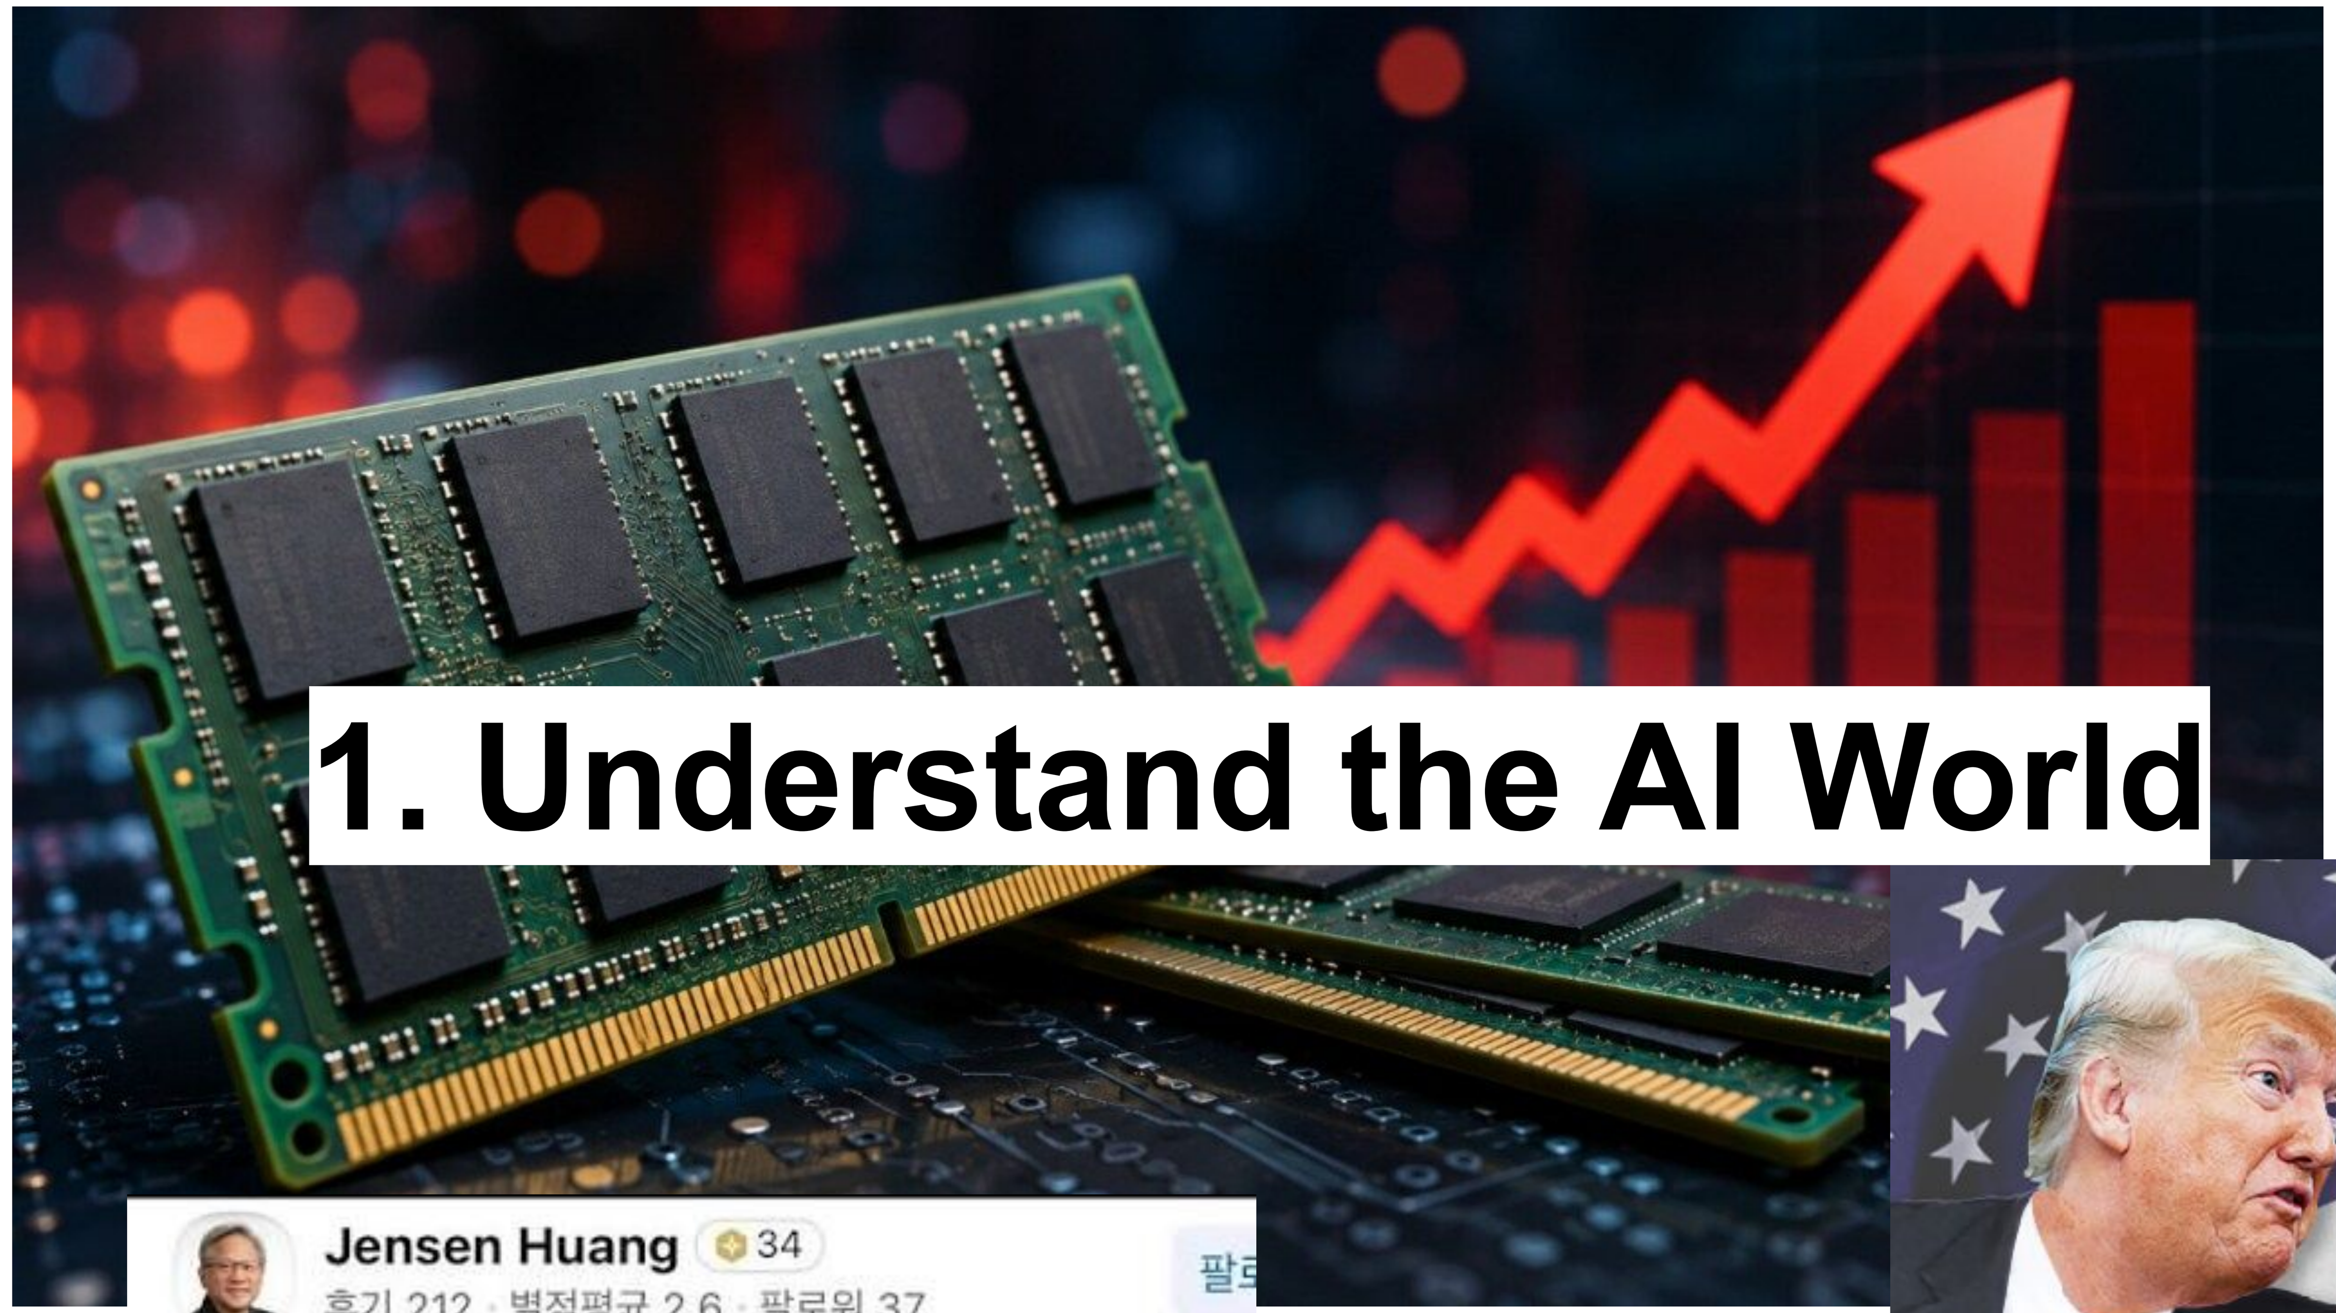

1. Understand the AI World

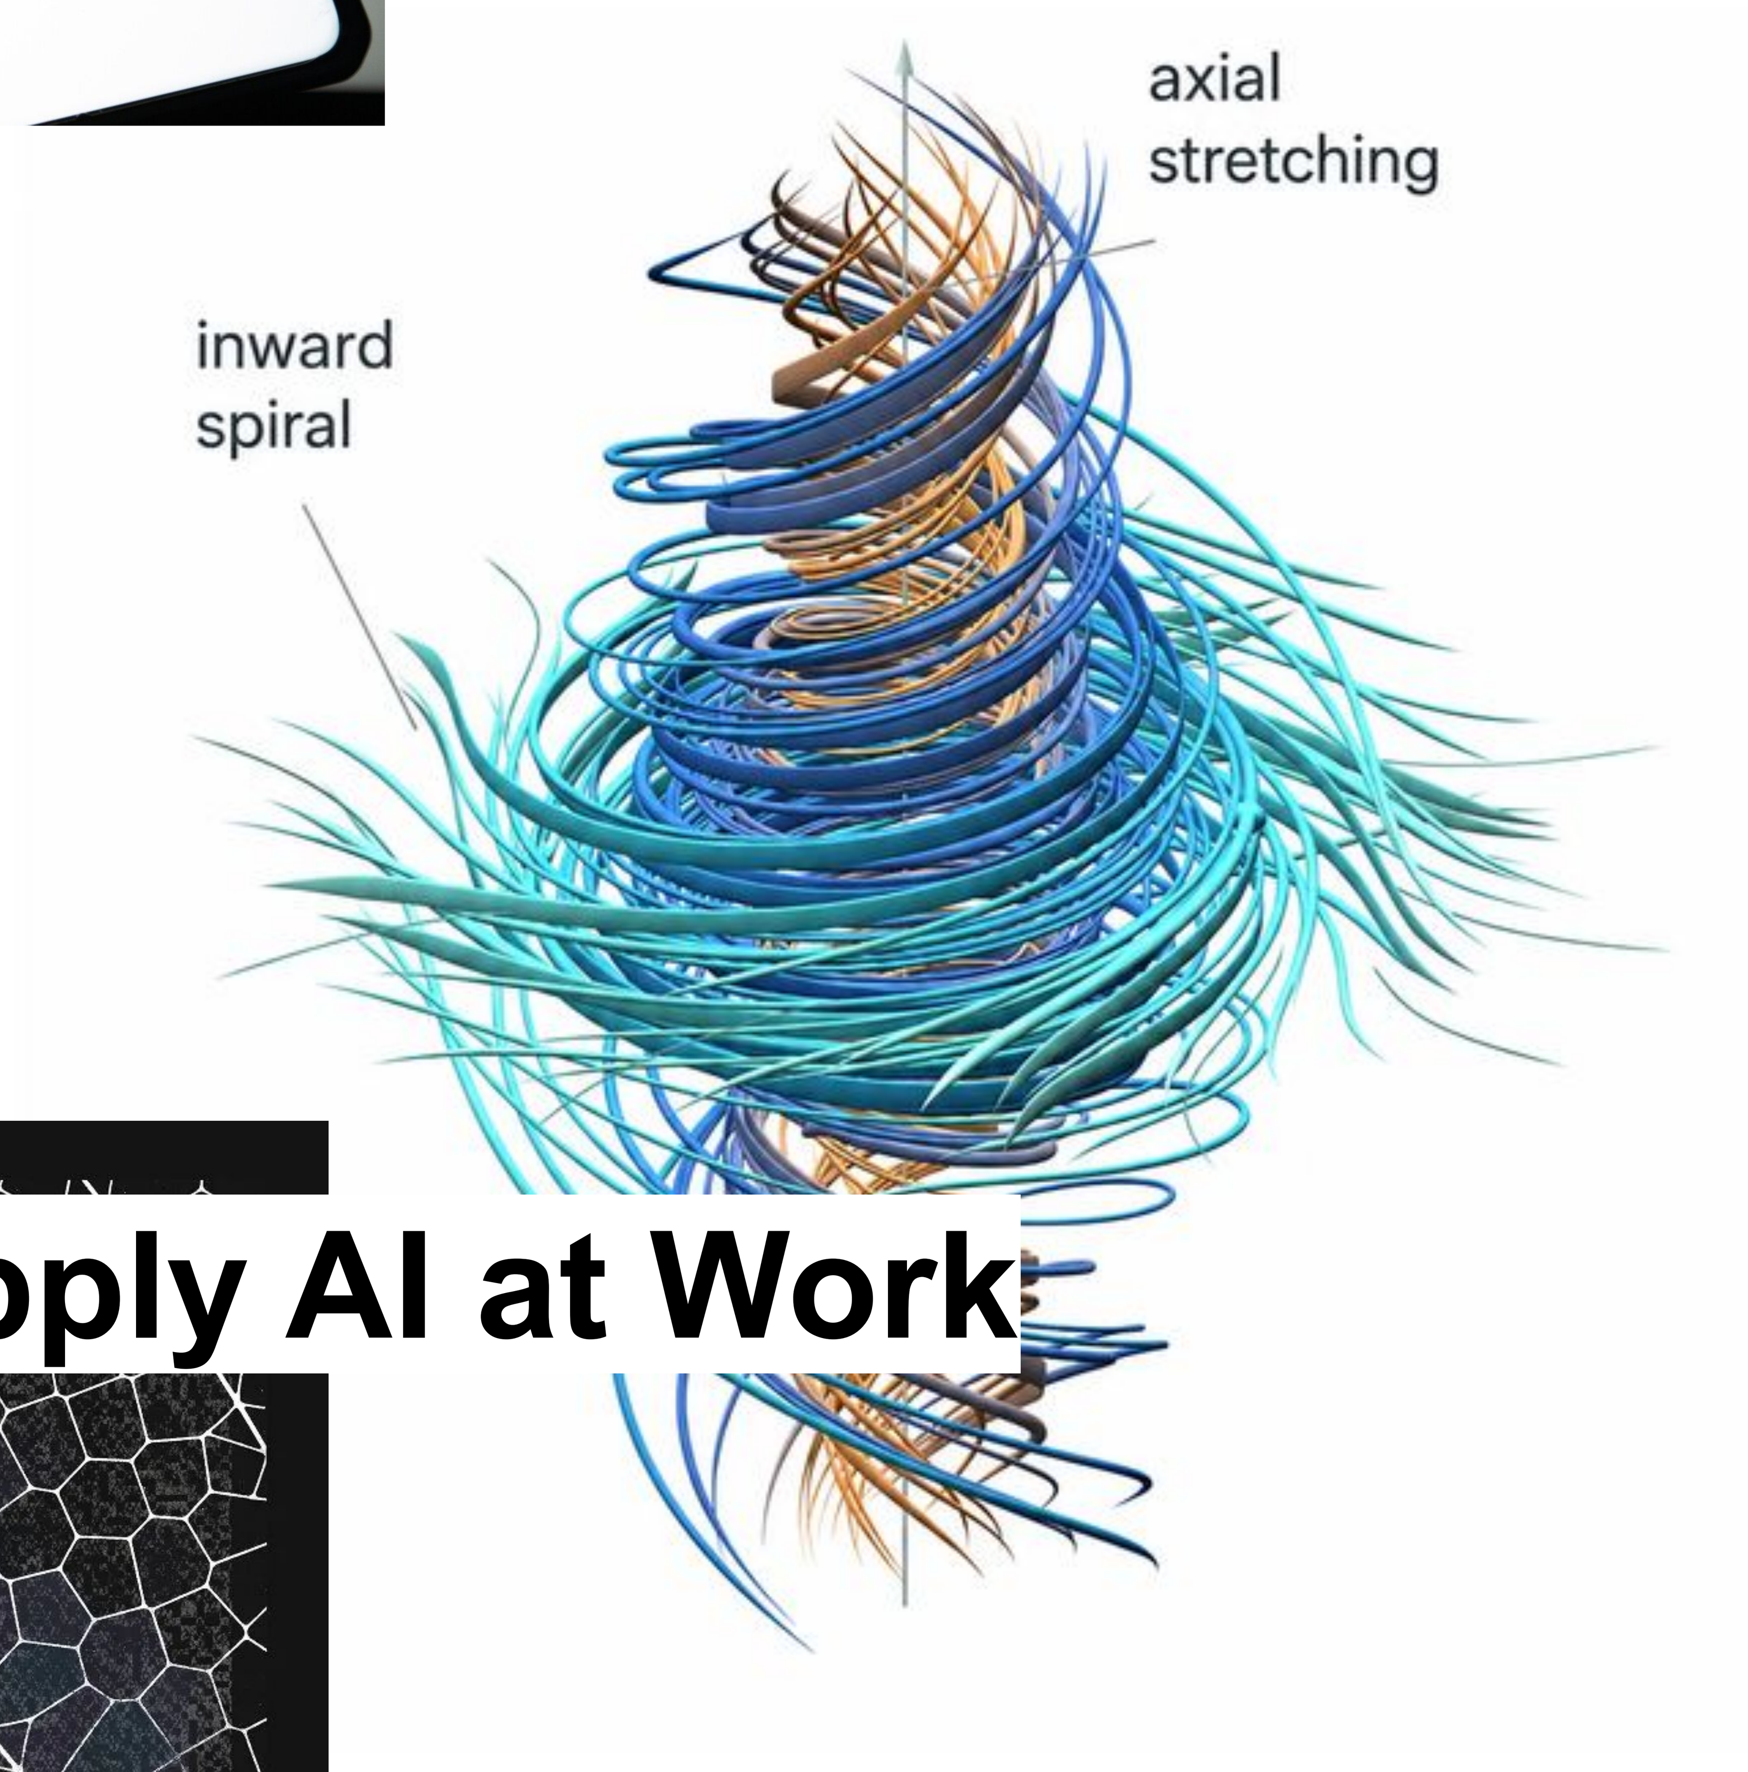

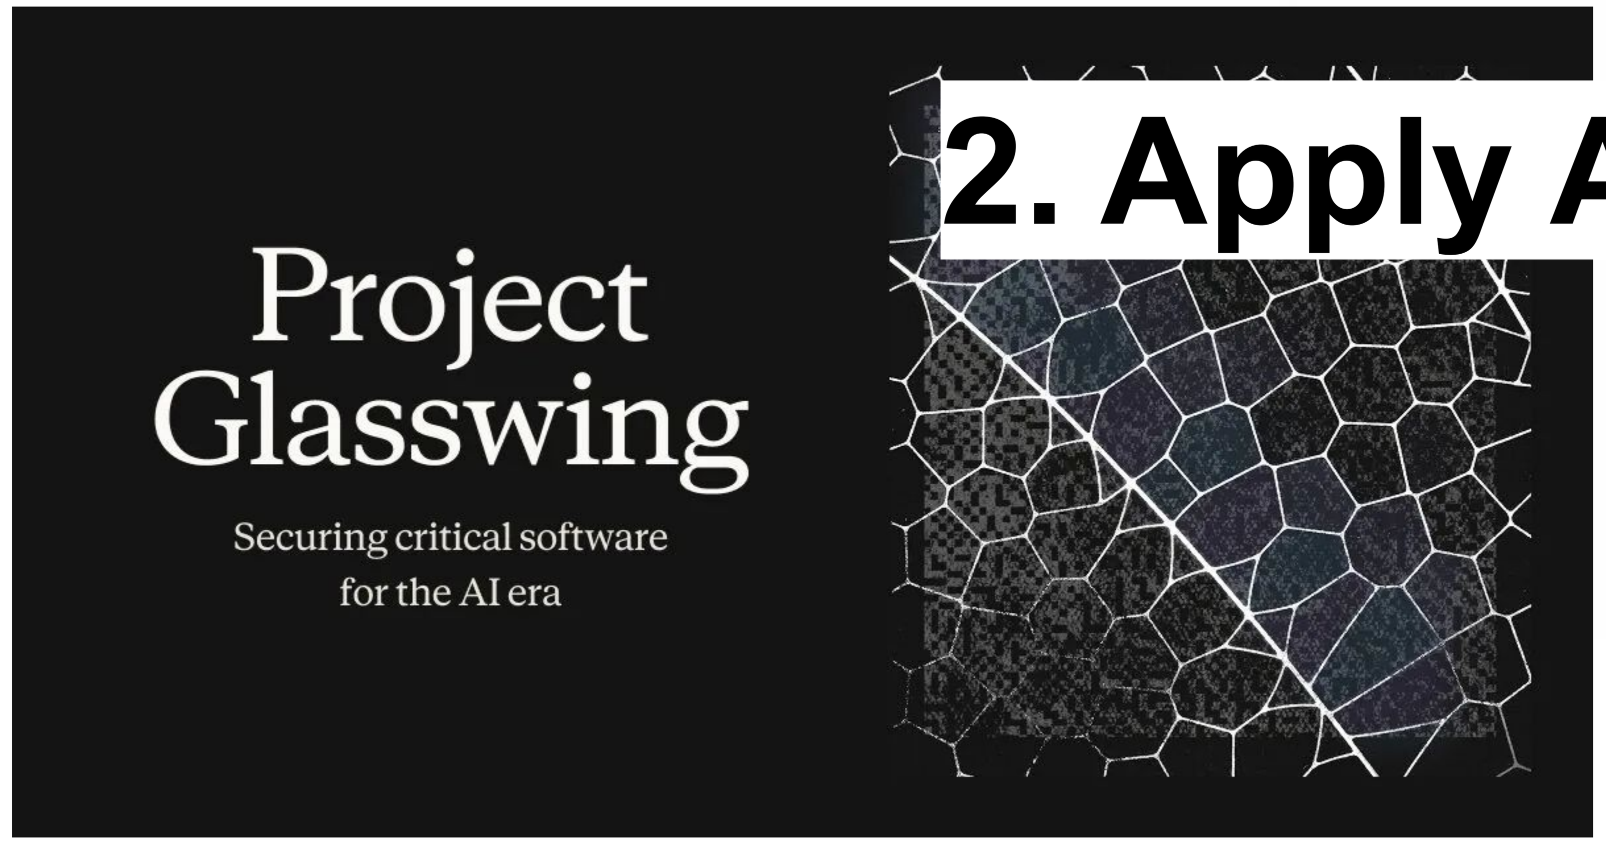

2. Apply AI at Work

Jae Y. Choi

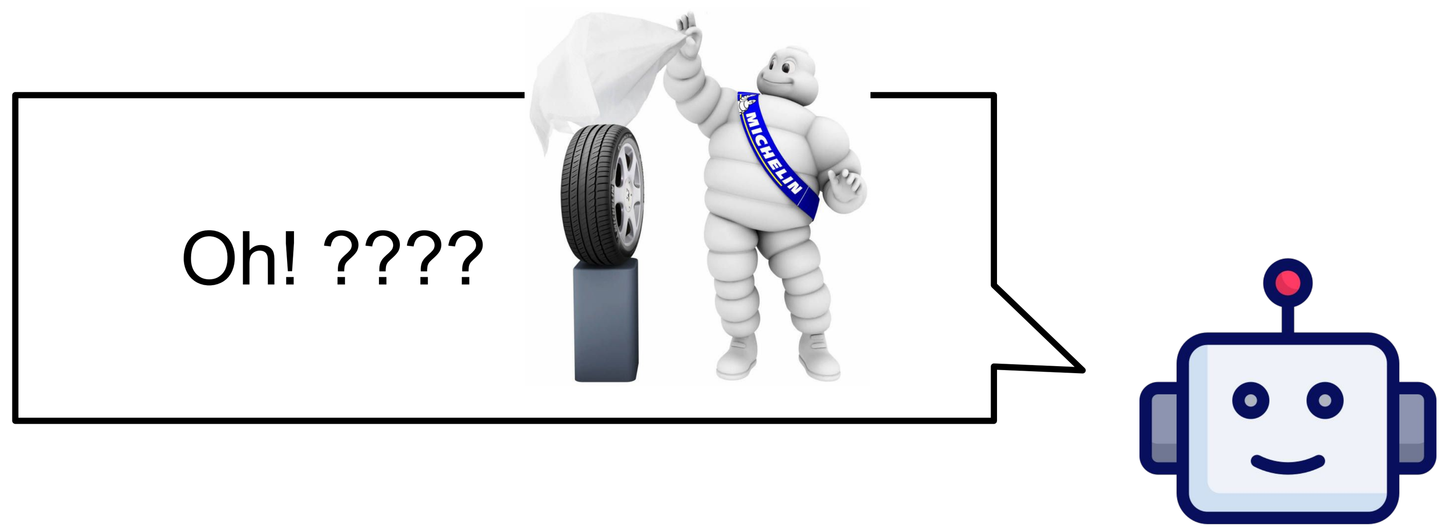

RAG Application in Industrial domain

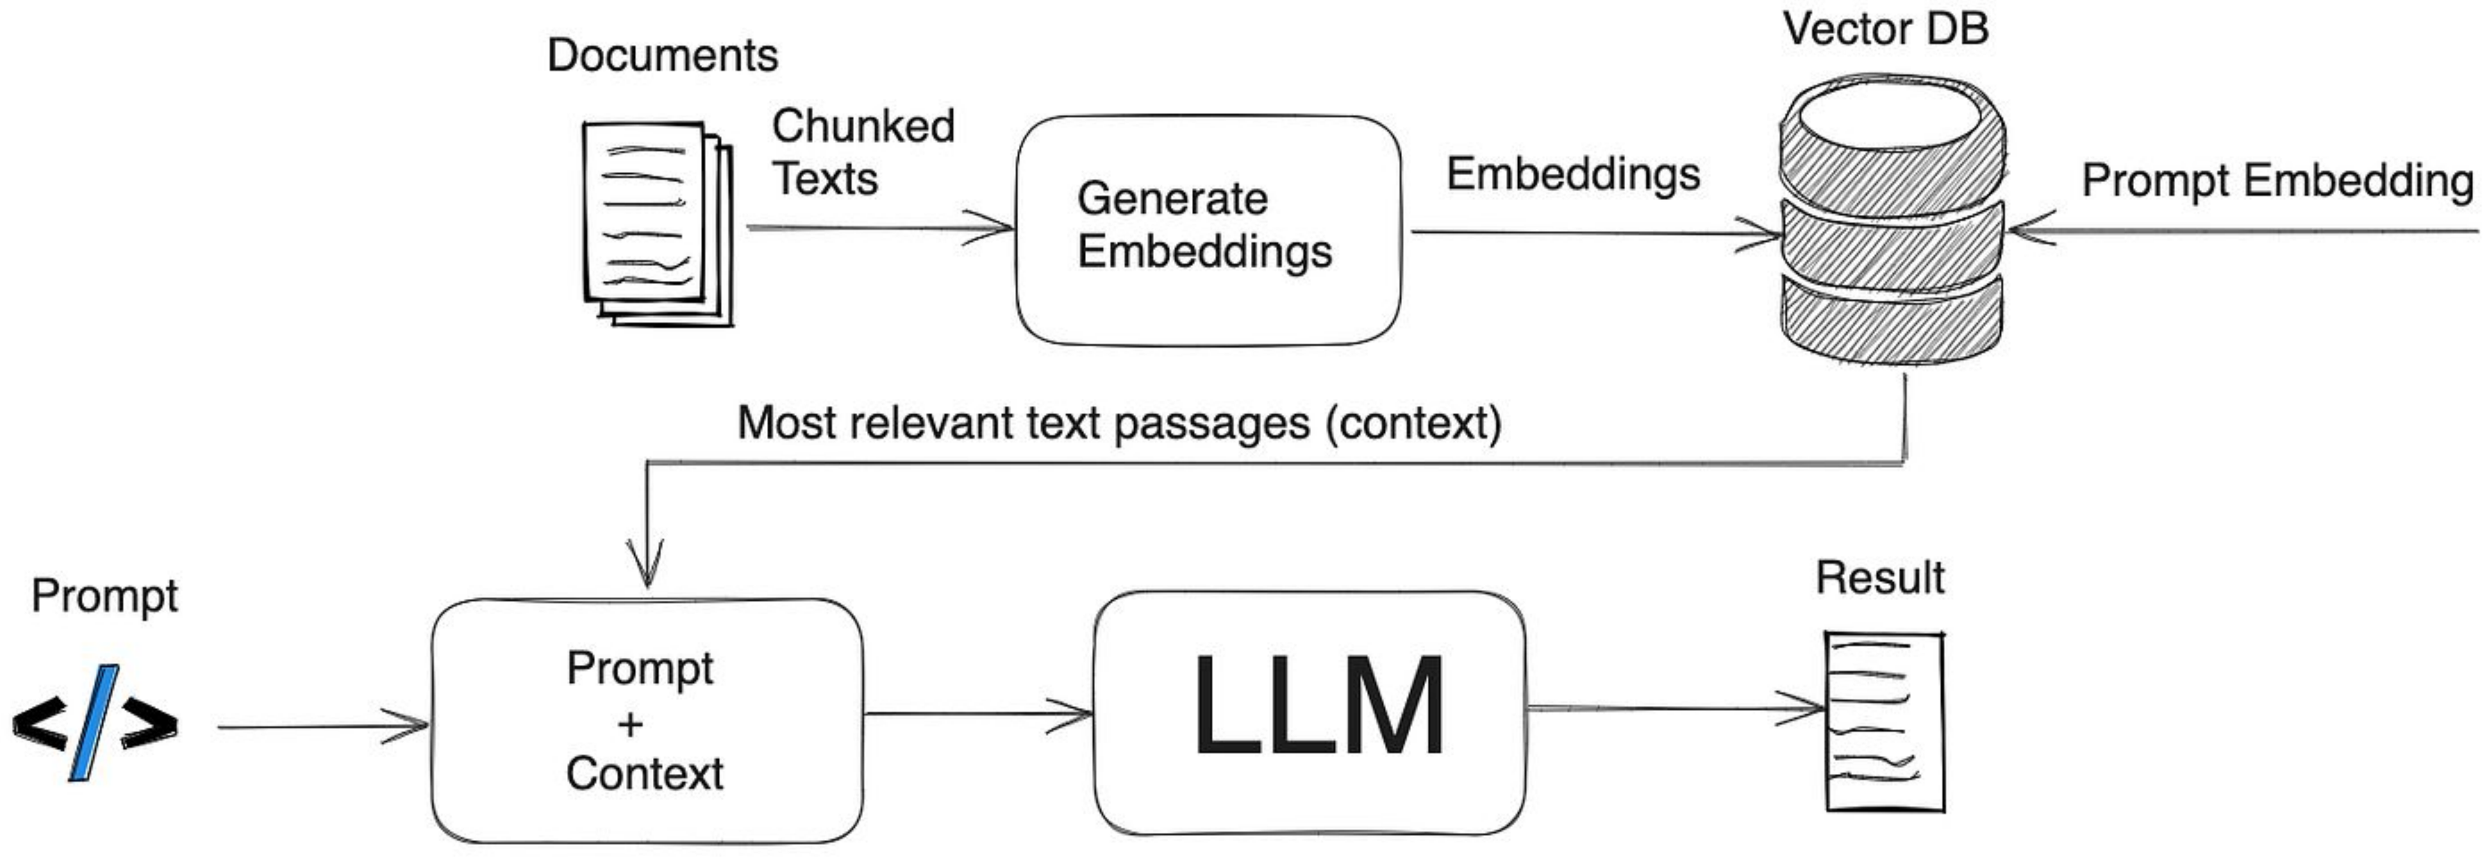

과학기술정보통신부

SEMES

Collaboration among AI agents (MAS) and multimodal AI

1) Video Temporal Grounding(VTG)

VTG is the task of identifying
the start and end timestamps of the video segment
most relevant to a given natural-language query.

Query: "When is the erhu being played?"

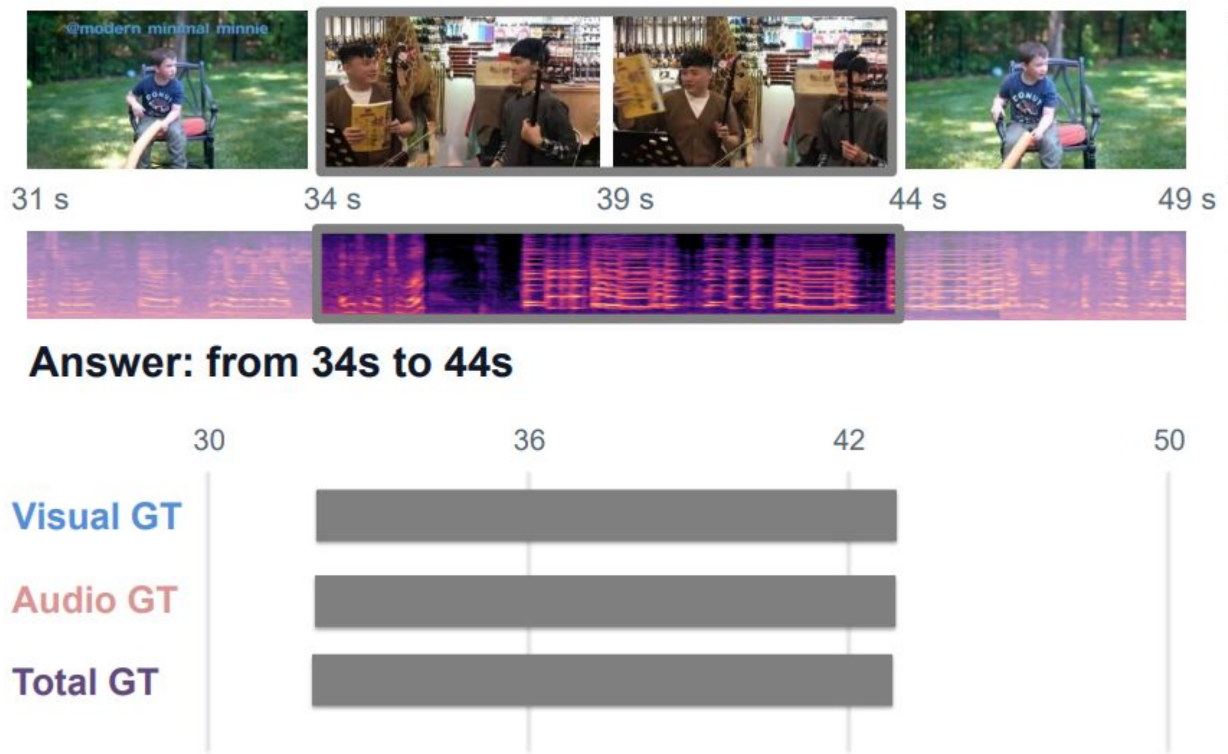

2) Controlled Mismatch Reveals Missing Evidence

Aligned inputs can hide reliance on one modality.
Controlled temporal mismatch tests whether the model covers the complete audible and visible event.

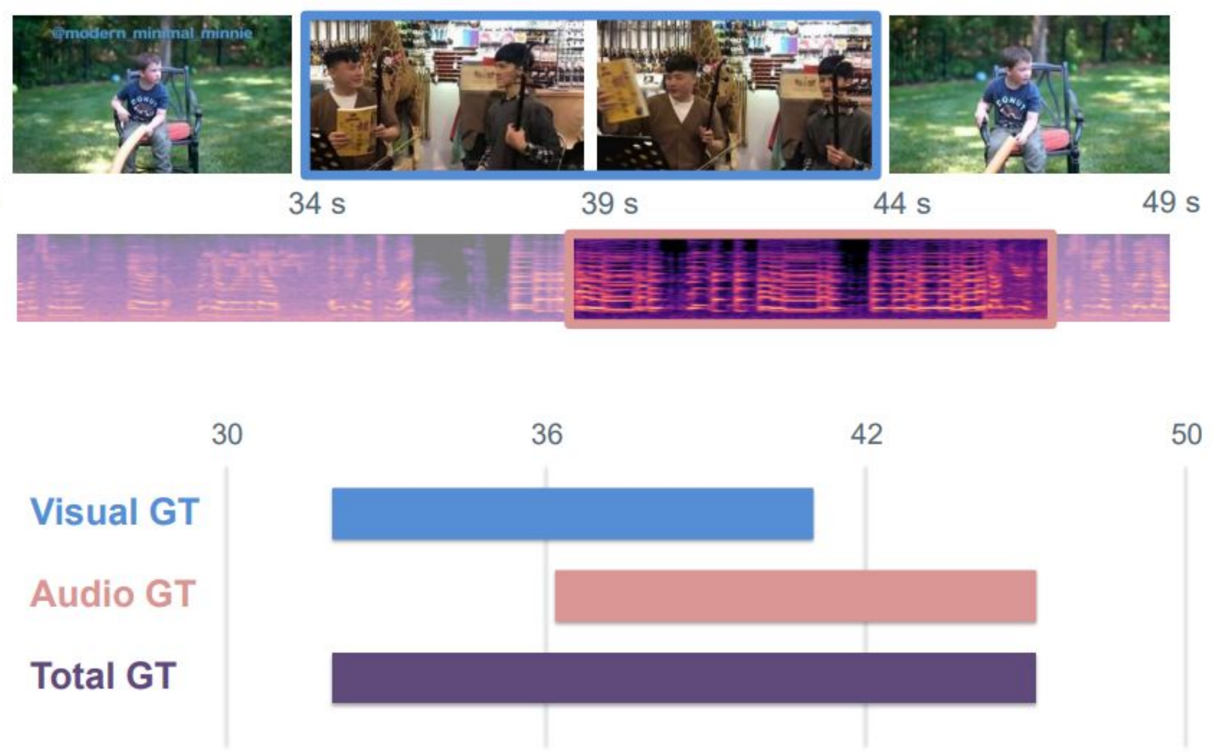

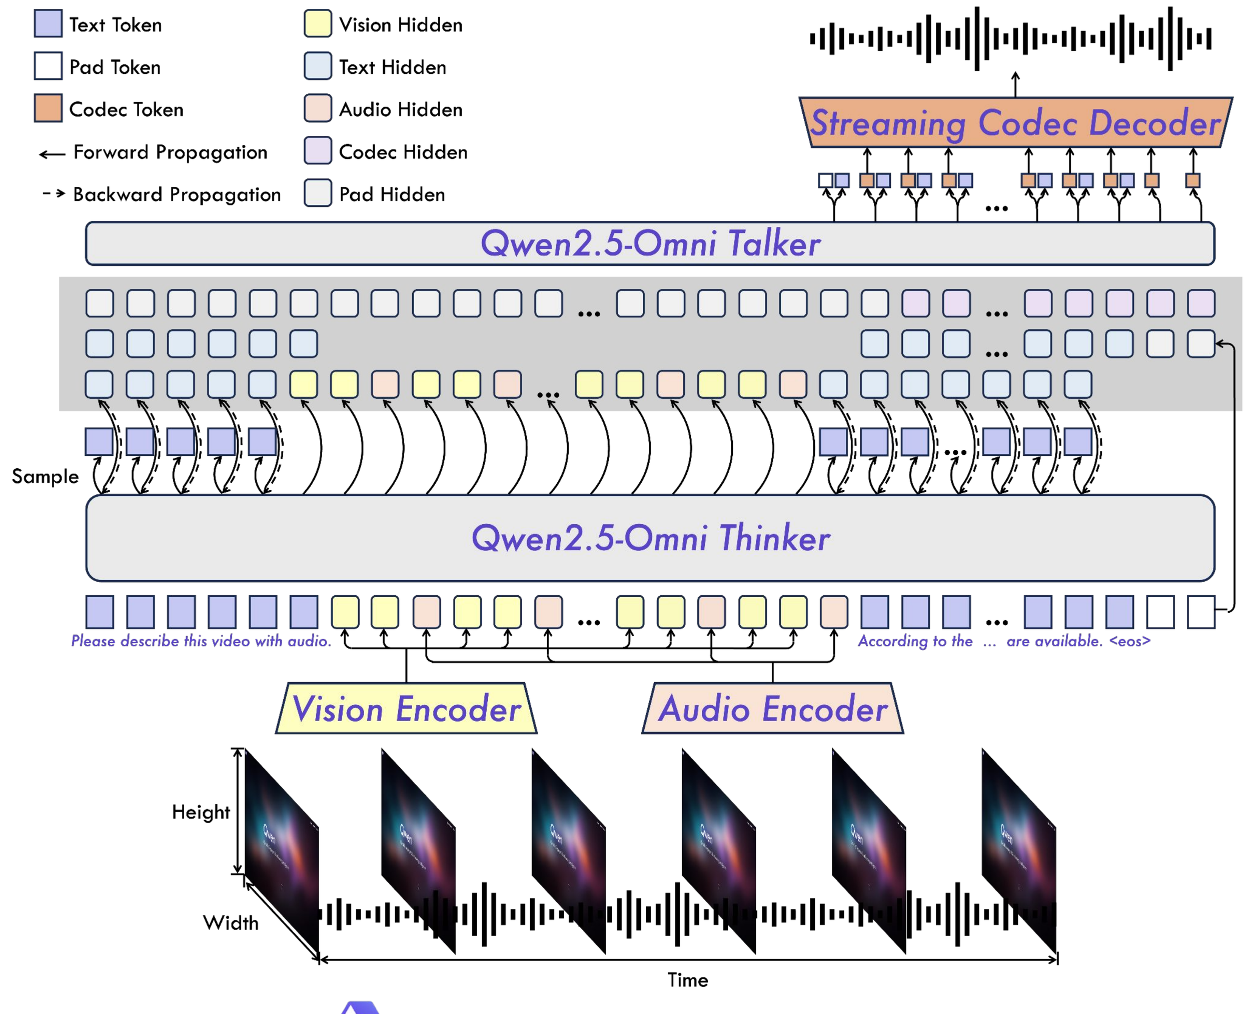

Qwen2.5-Omni

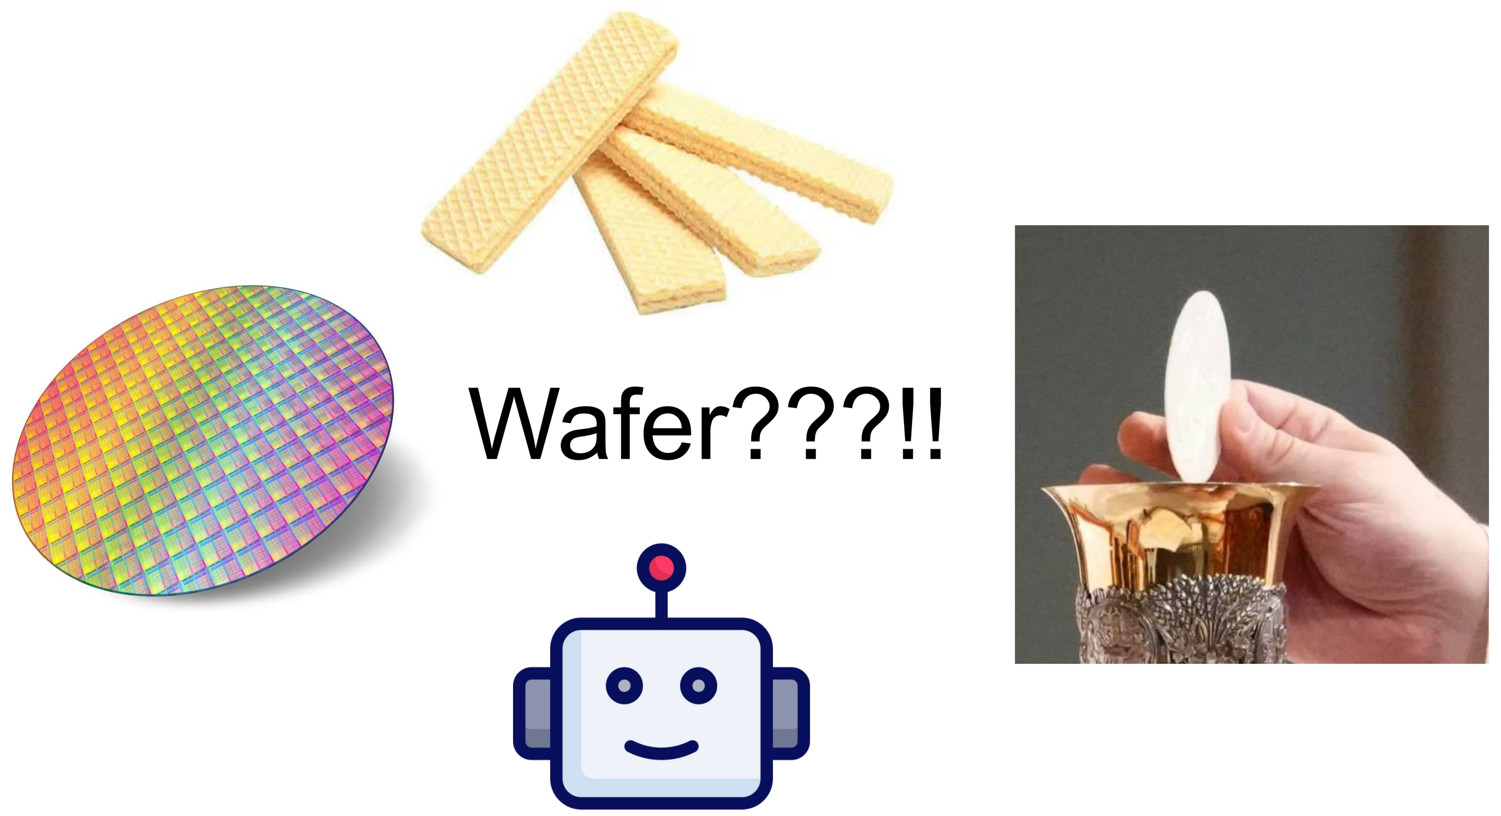

NVIDIA

Course Objectives

Fundamentals of LLM & Agentic AI

- Introduction to LLM
- LLM Reasoning and Agentic AI
- LLM Training
- Evolution of AI Models
- GPU Infrastructure

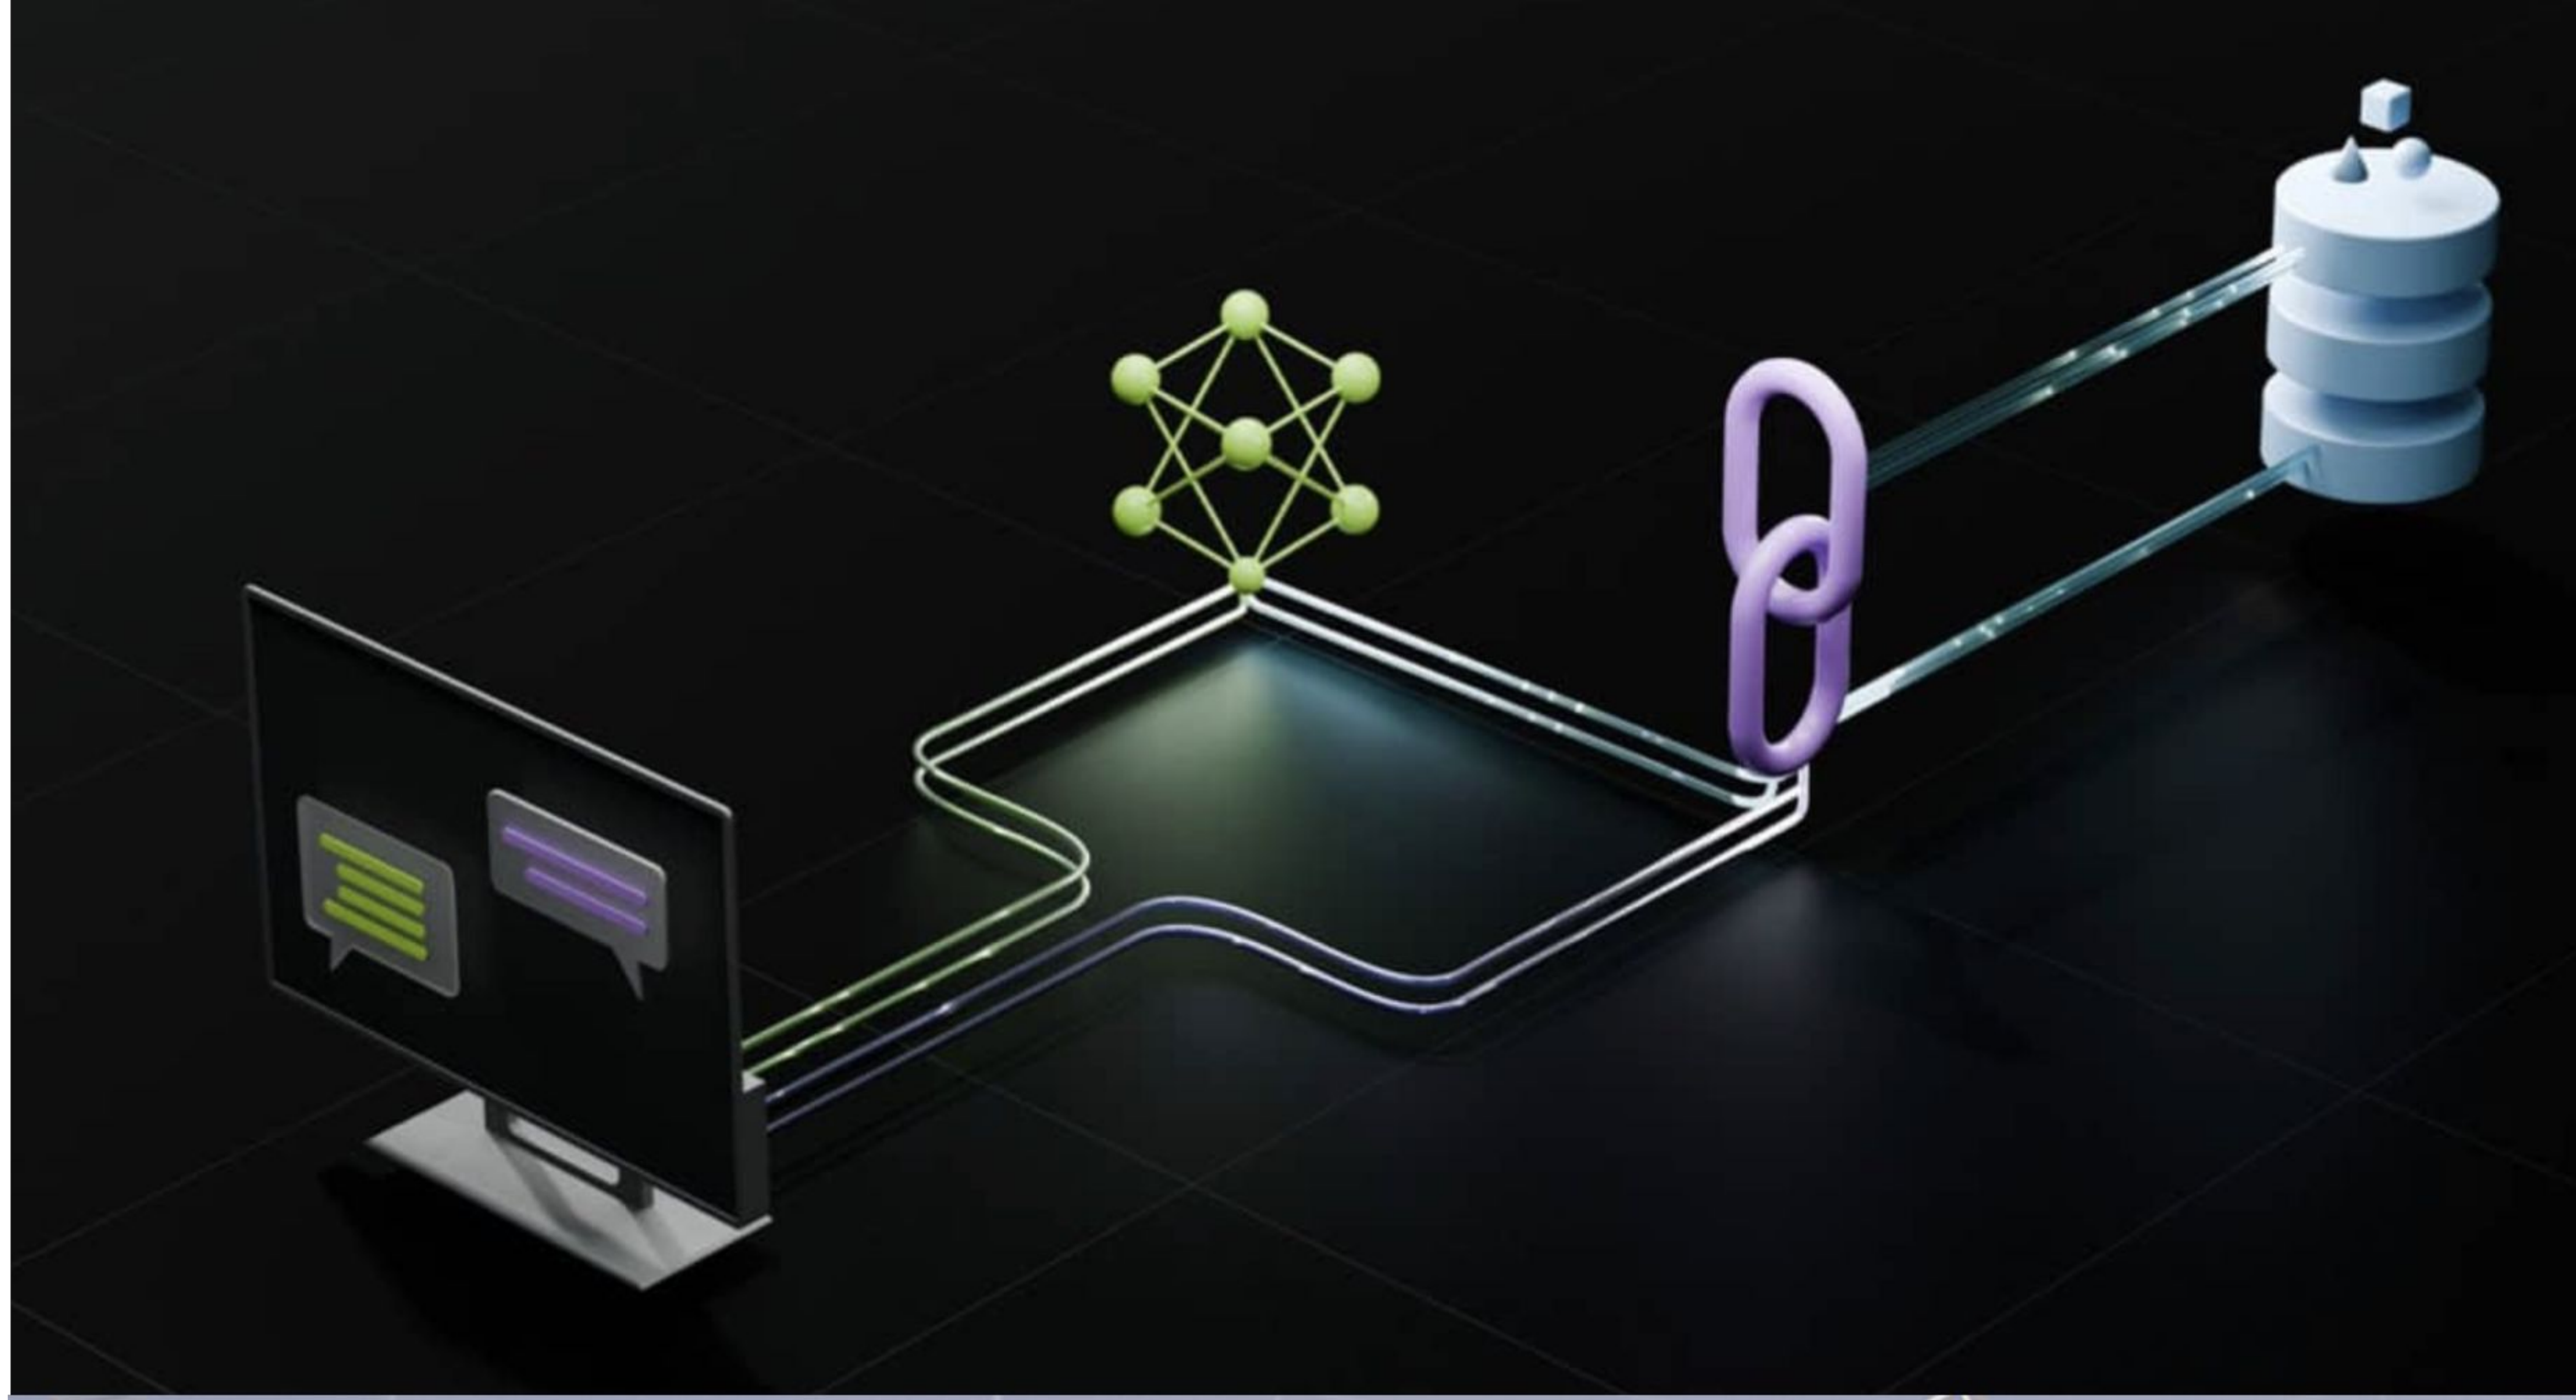

Building RAG Agents with LLMs

Part 1. Introduction to LLM

Introduction to LLM

What is the Large Language Model?

Proprietary(Closed) Models

Global Opensource Models

Korea AI Models

Too Many Models!!

Introduction to LLM

What is the Large Language Model?

User Area

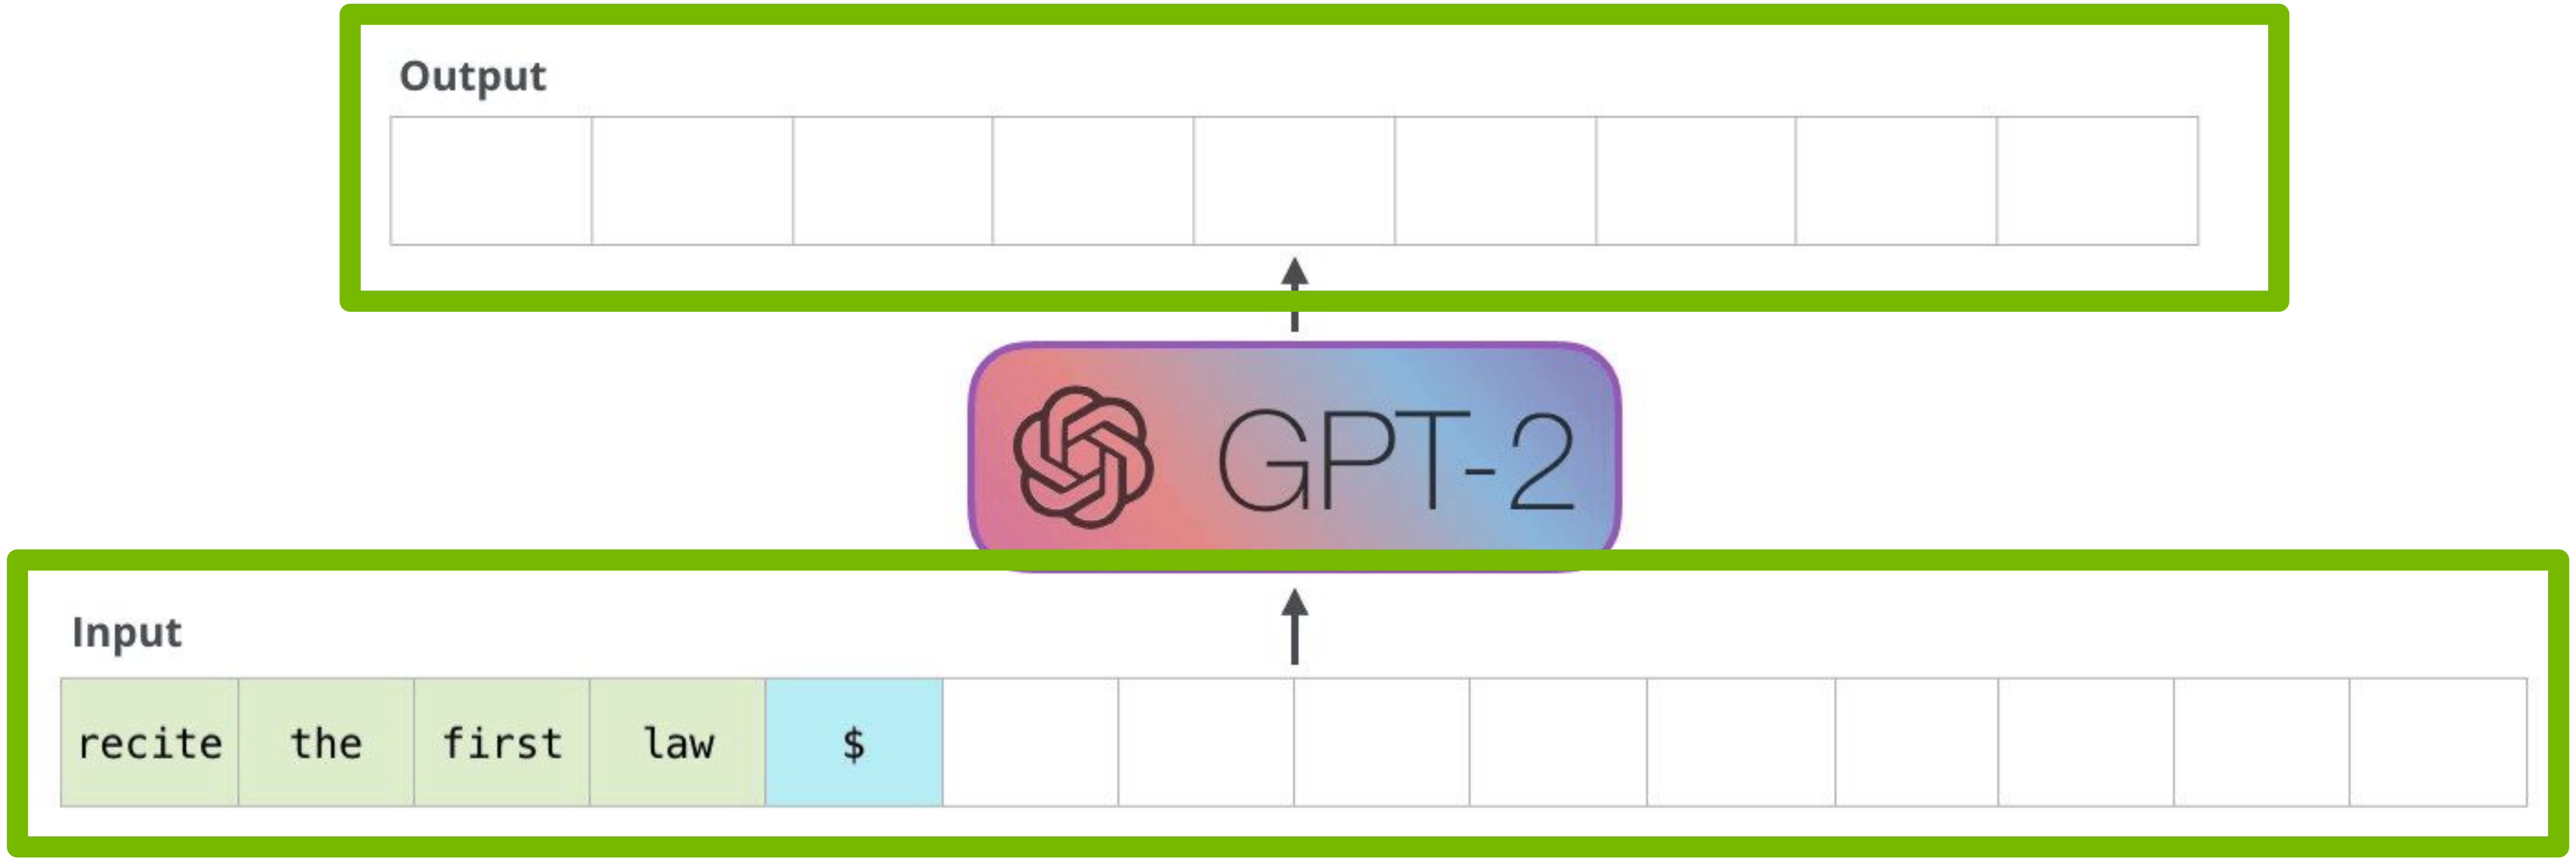

What's under the hood?

Auto Regressive (AR)

https://nlpinkorean.github.io/visualizing-neural-machine-translation-mechanics-of-seq2seq-models-with-attention/

Introduction to LLM

What is the Large Language Model?

New insight of view

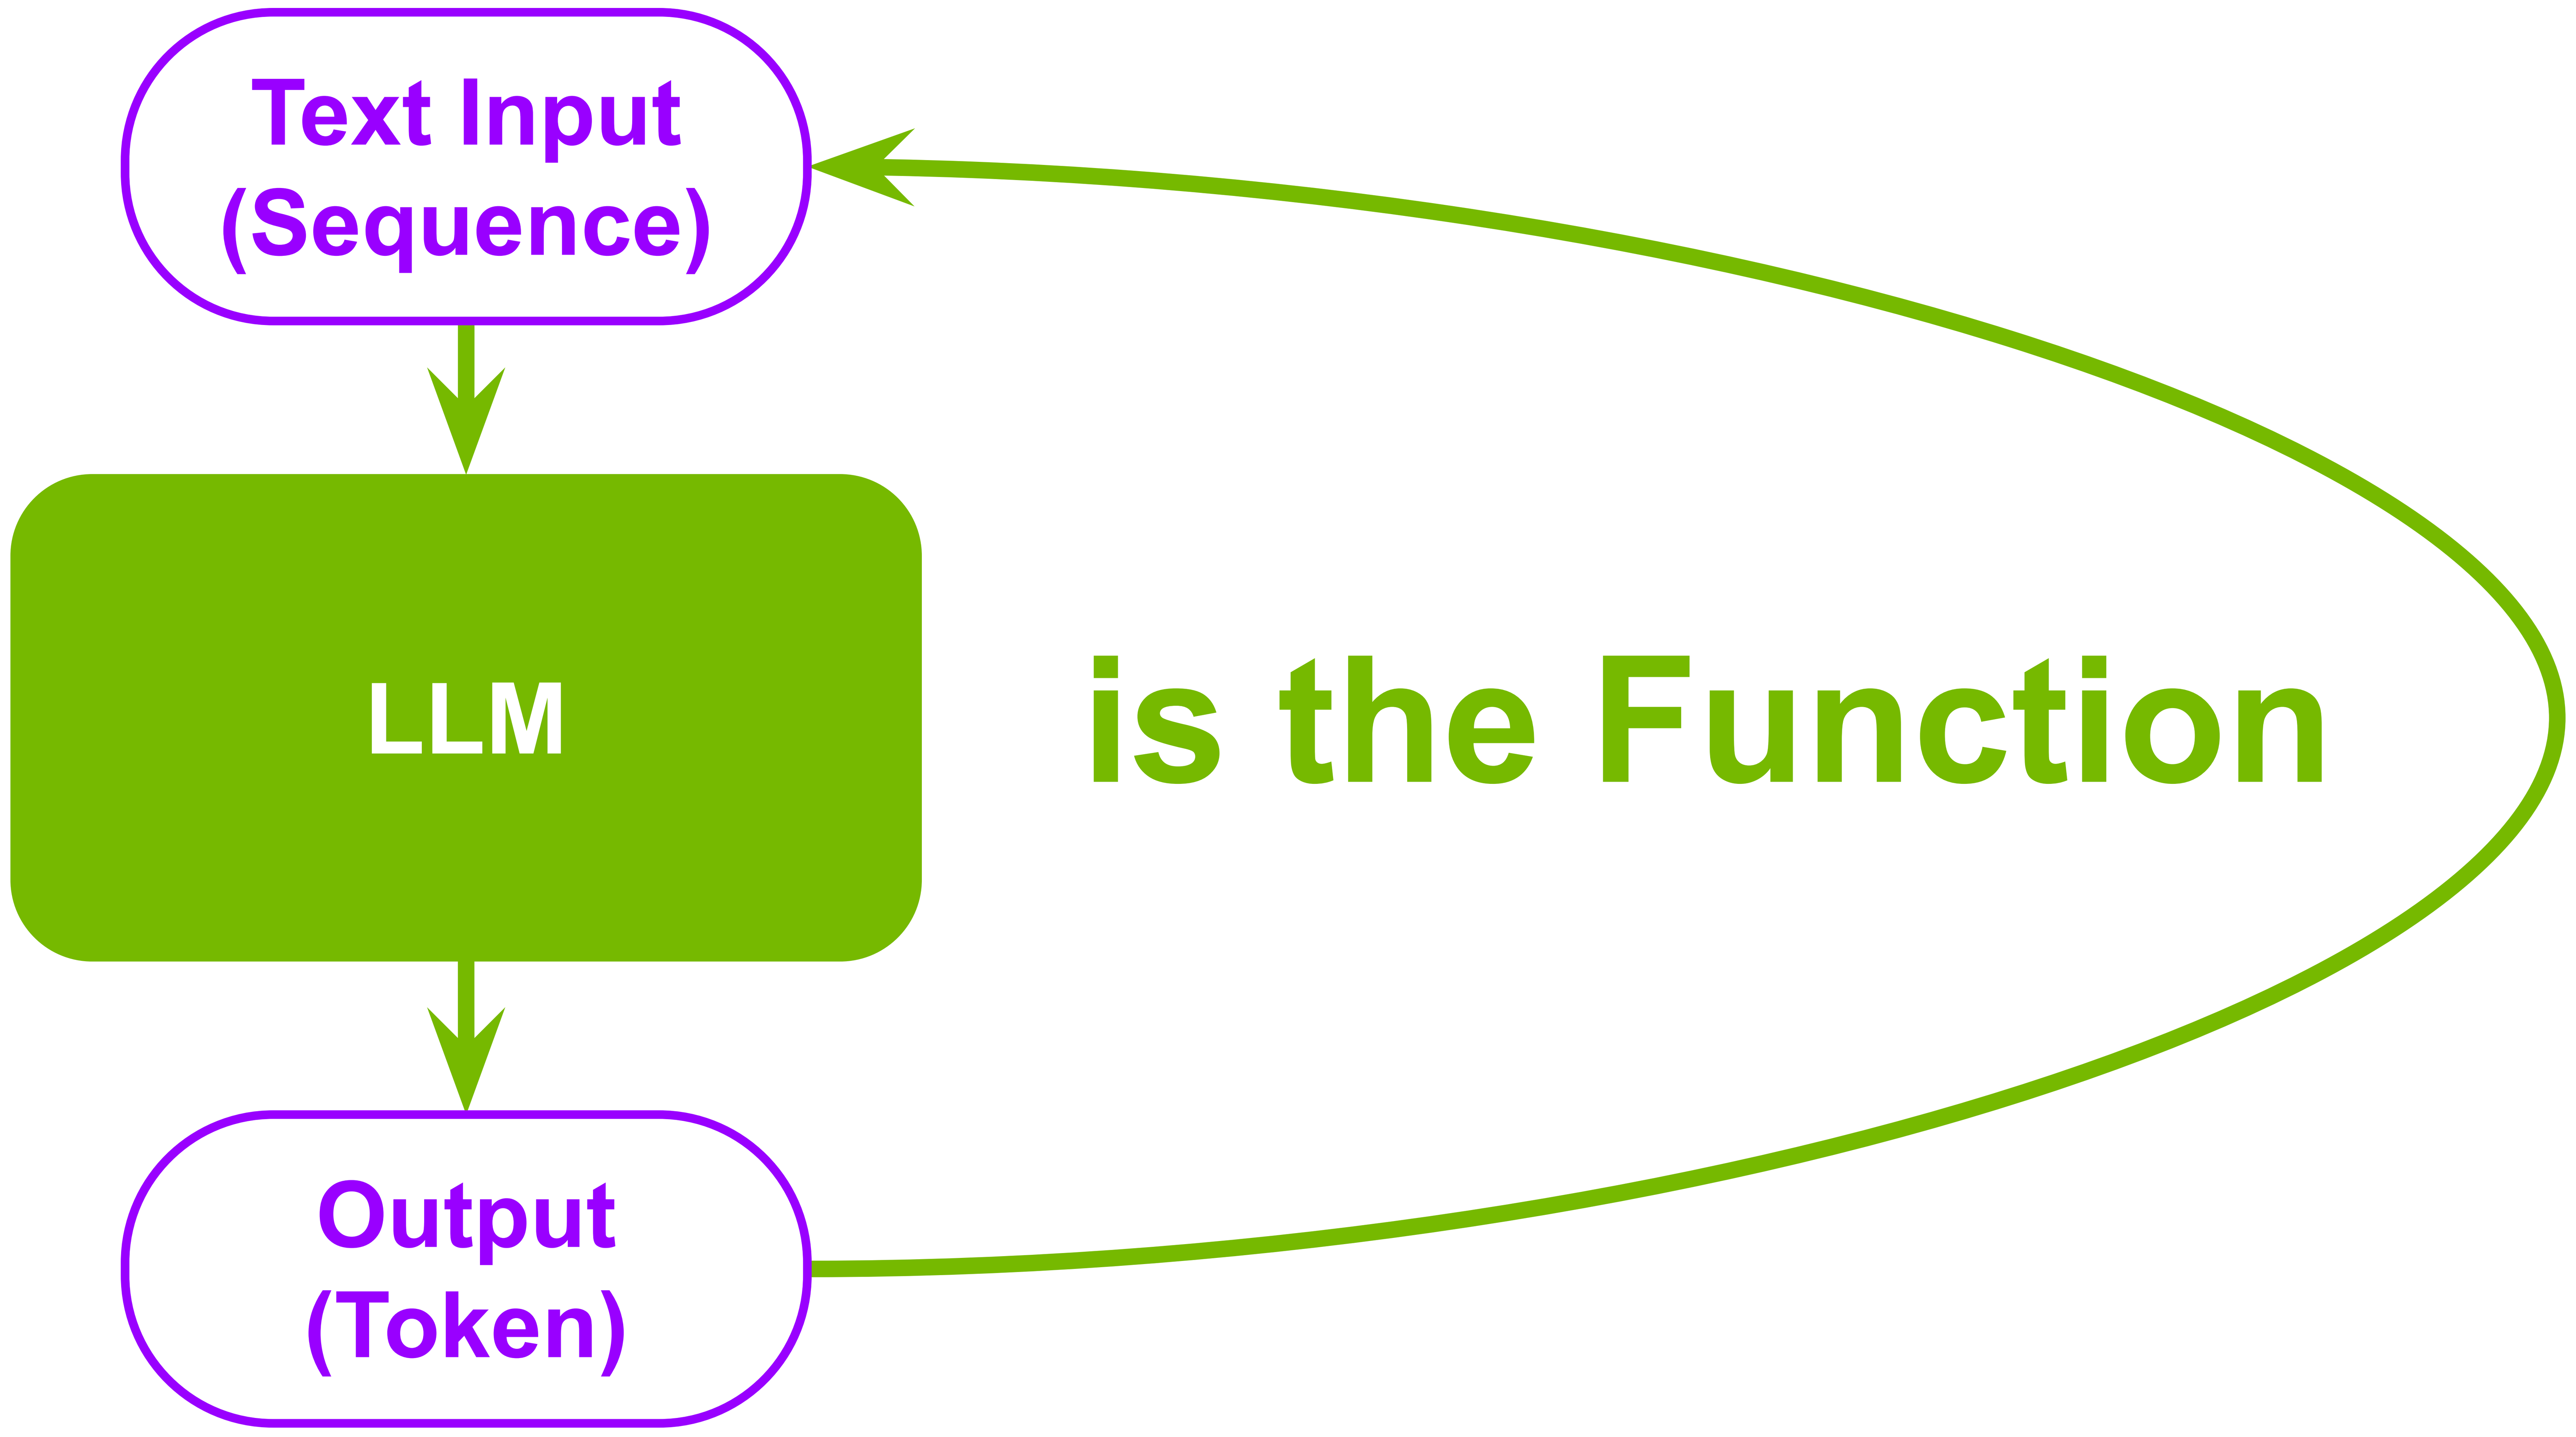

Baseline Architecture
→ Transformer

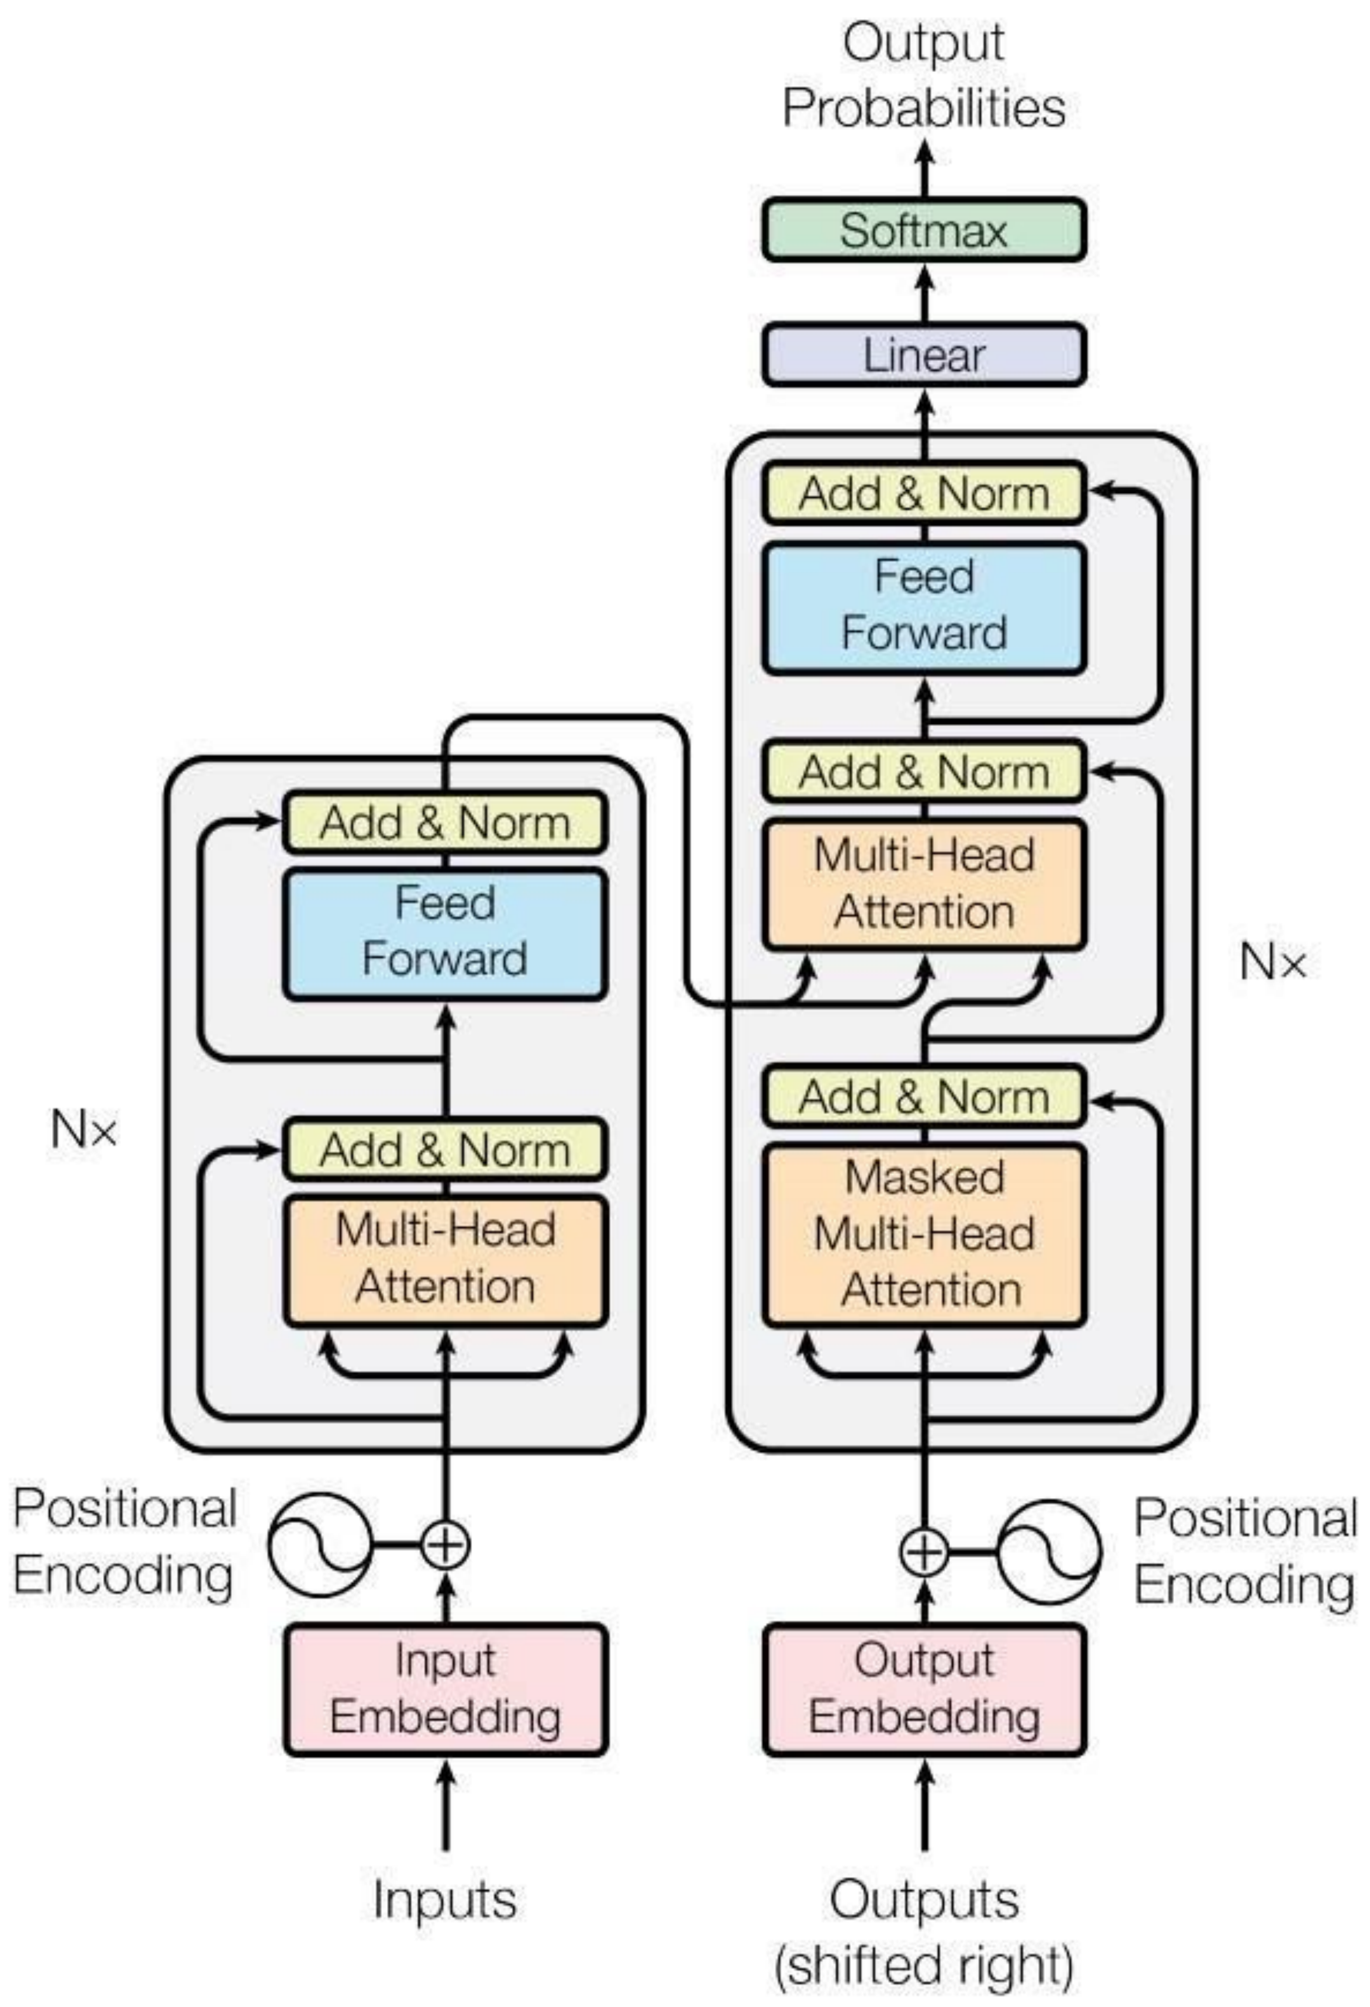

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in neural information processing systems, 30.*

Introduction to LLM

What is the Large Language Model?

Vector Space
High Dimensional Space Where Similar Words are Close and Different Words are Far

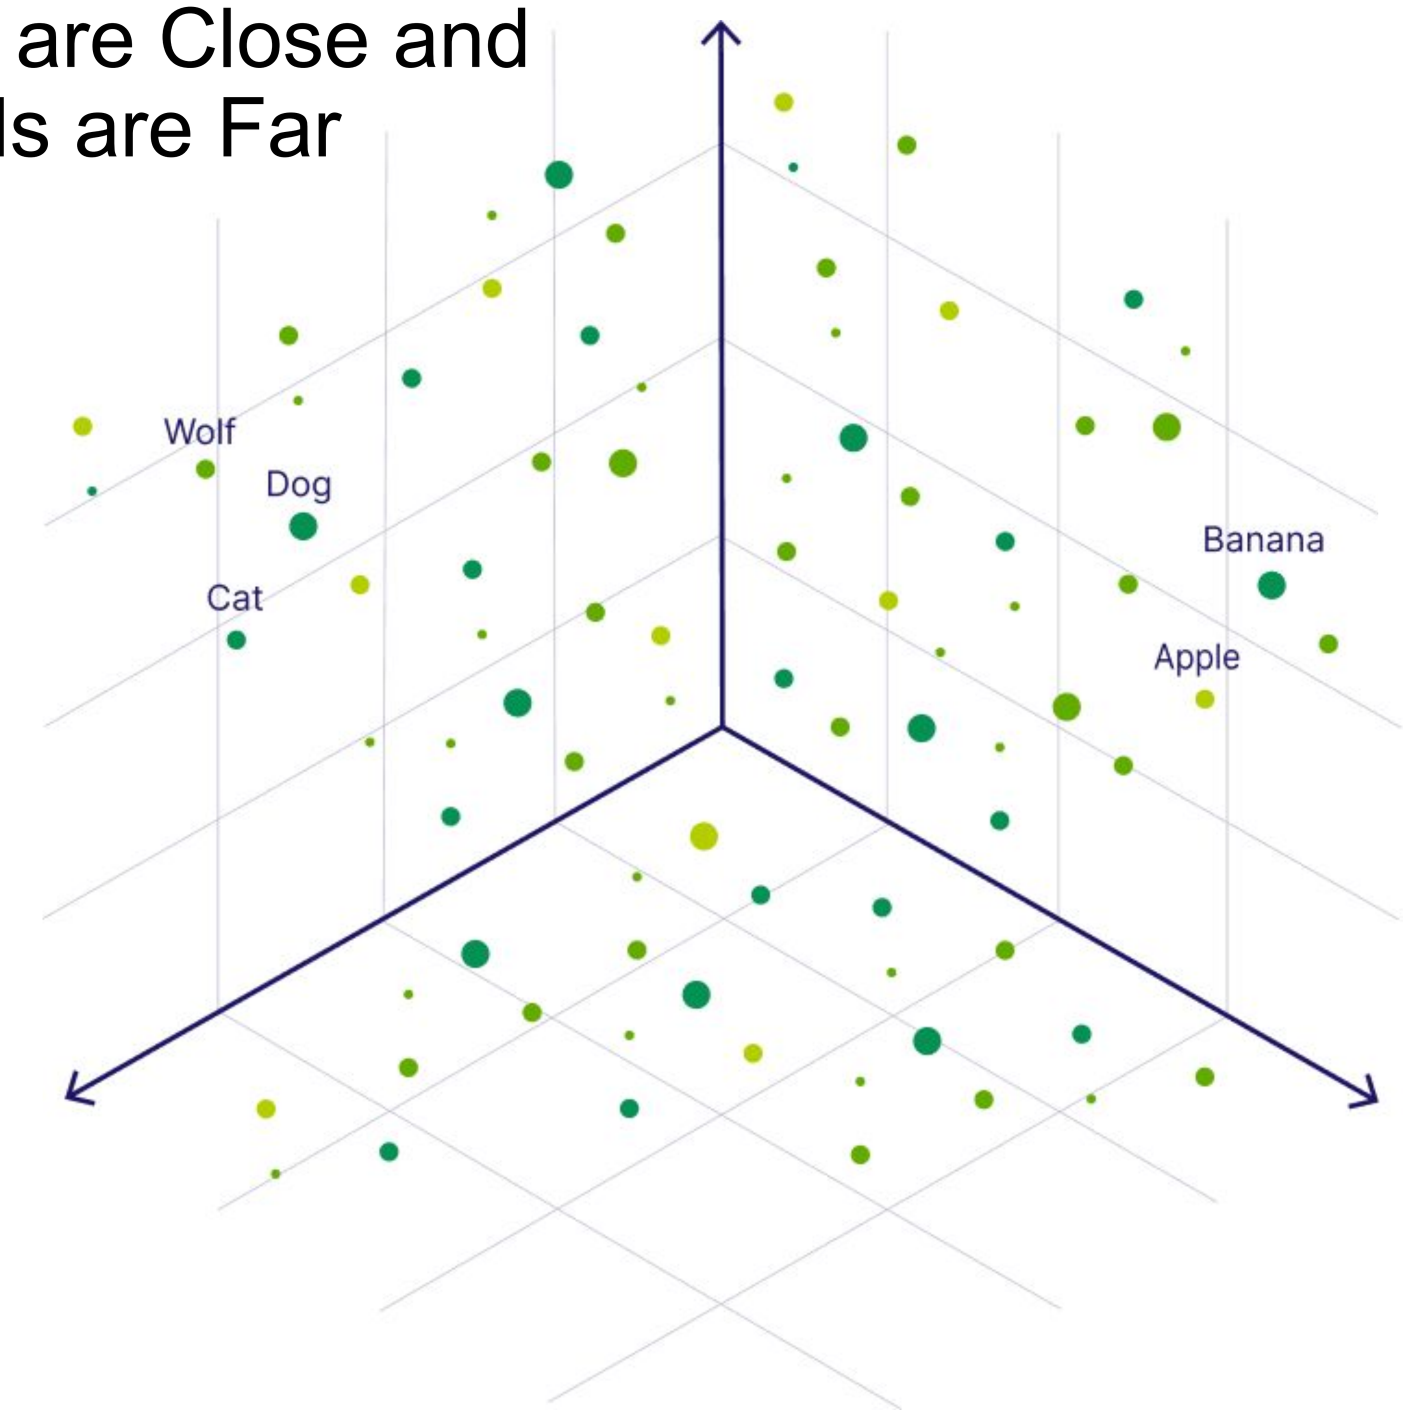

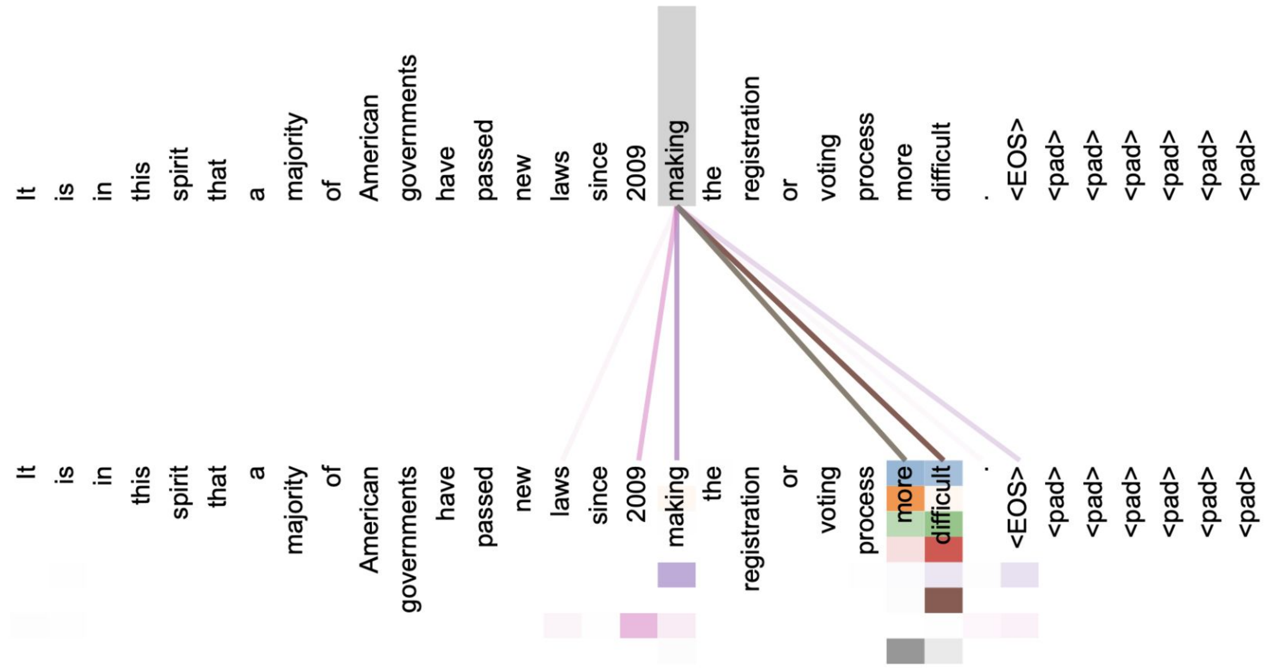

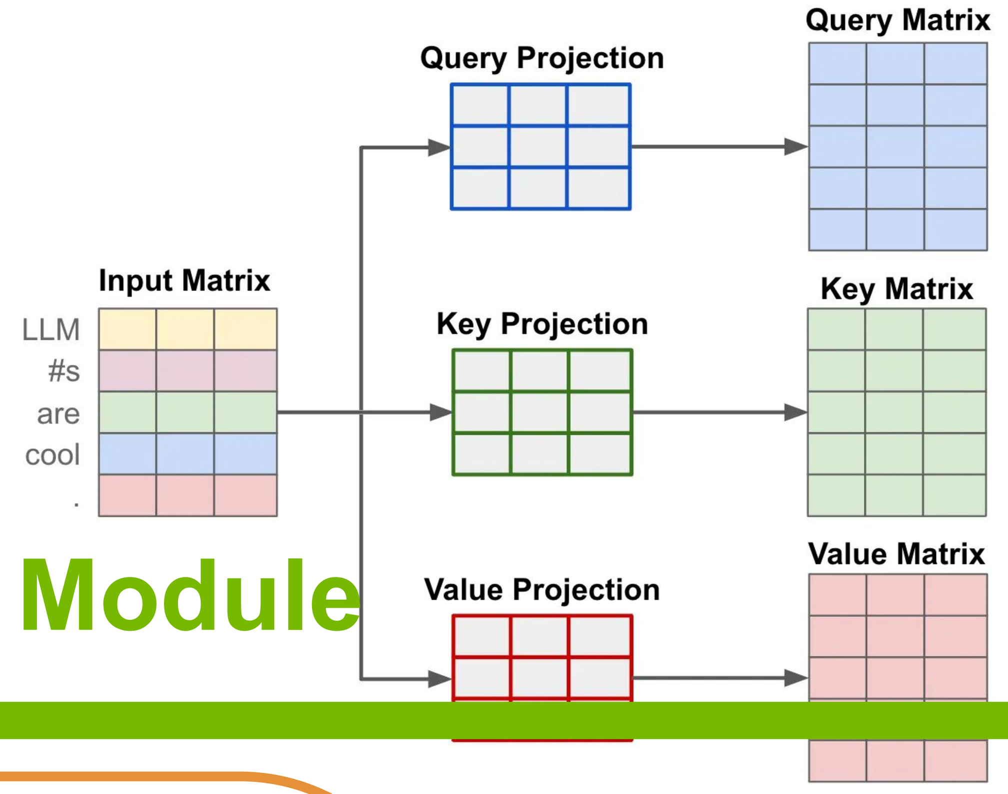

Most of the model parameters are at MLP stage

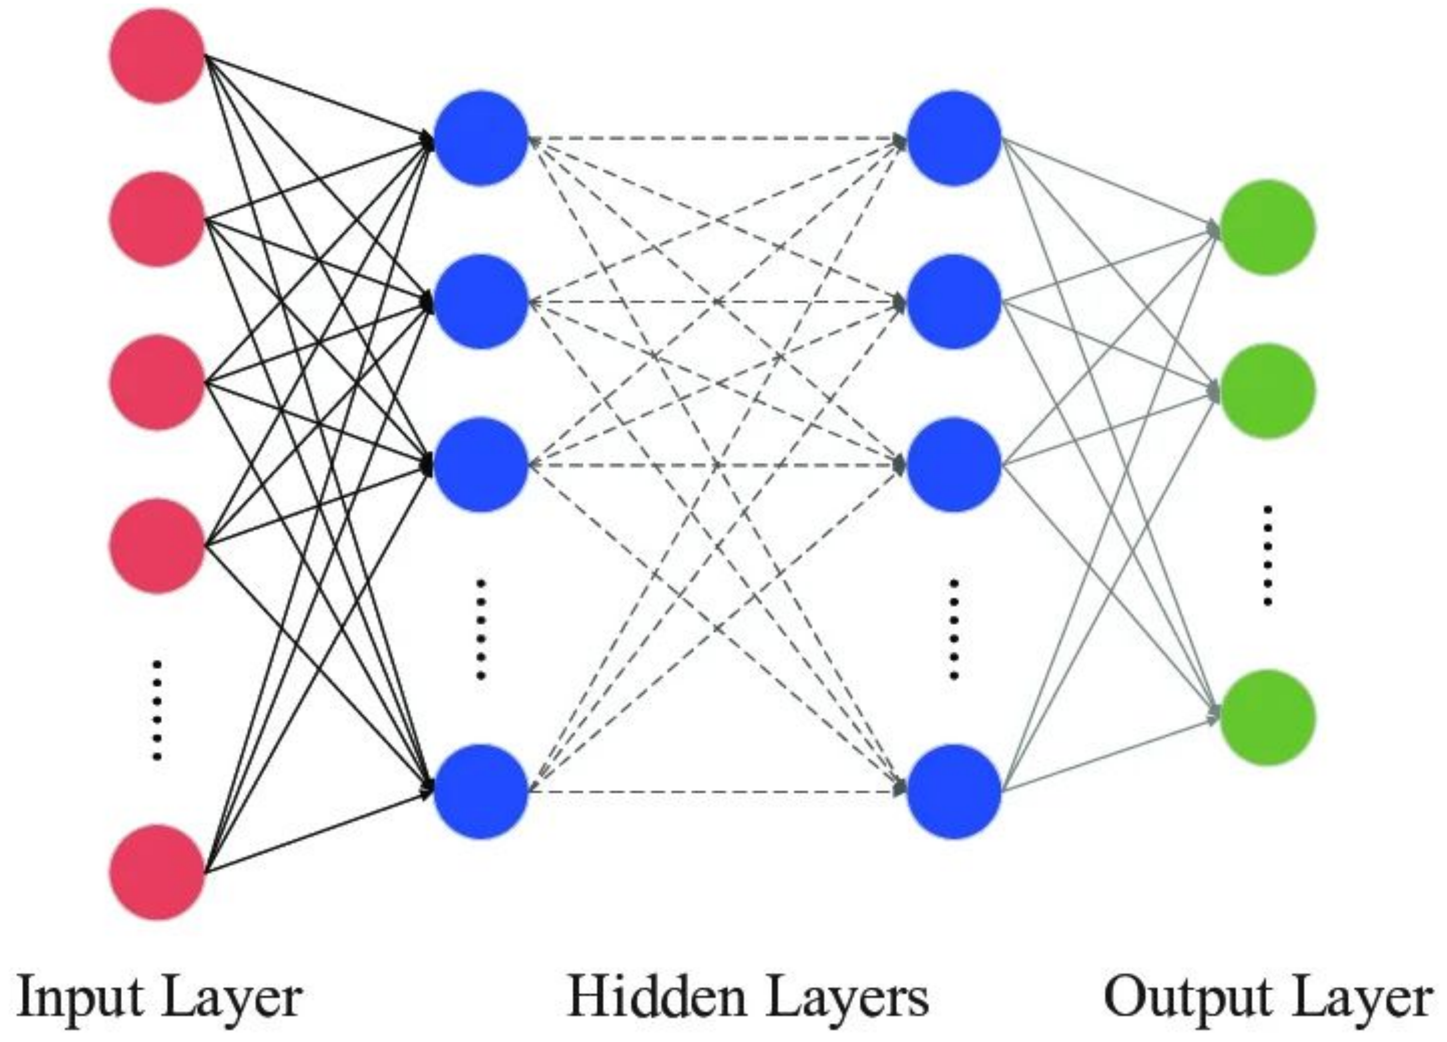

Decoder Module

Text Input (Sequence) → **Embedding** Vector Space → **Attention** Inter-Word Similarity → **MLP** Memory Space → **Output (Token)**

Tens of Decoder Layers are repeated in Decoder Only LLM

NVIDIA

Introduction to LLM

What is the Large Language Model?

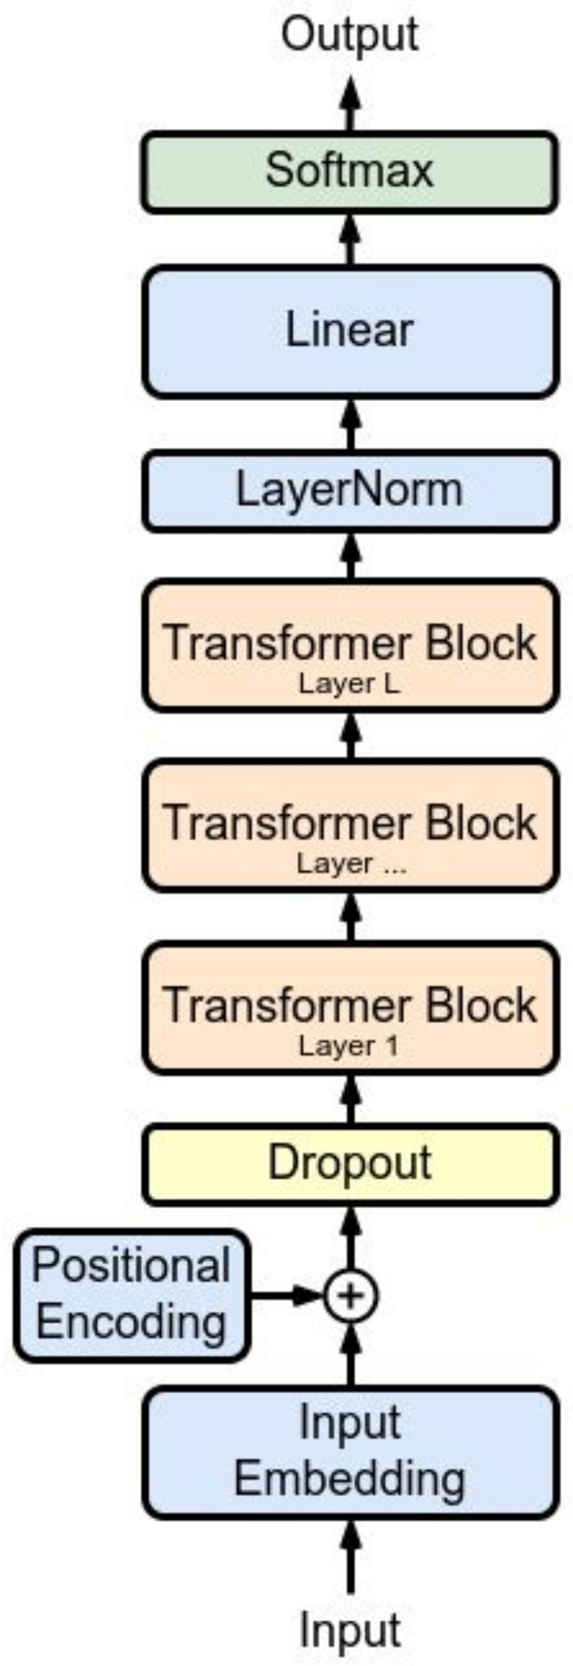

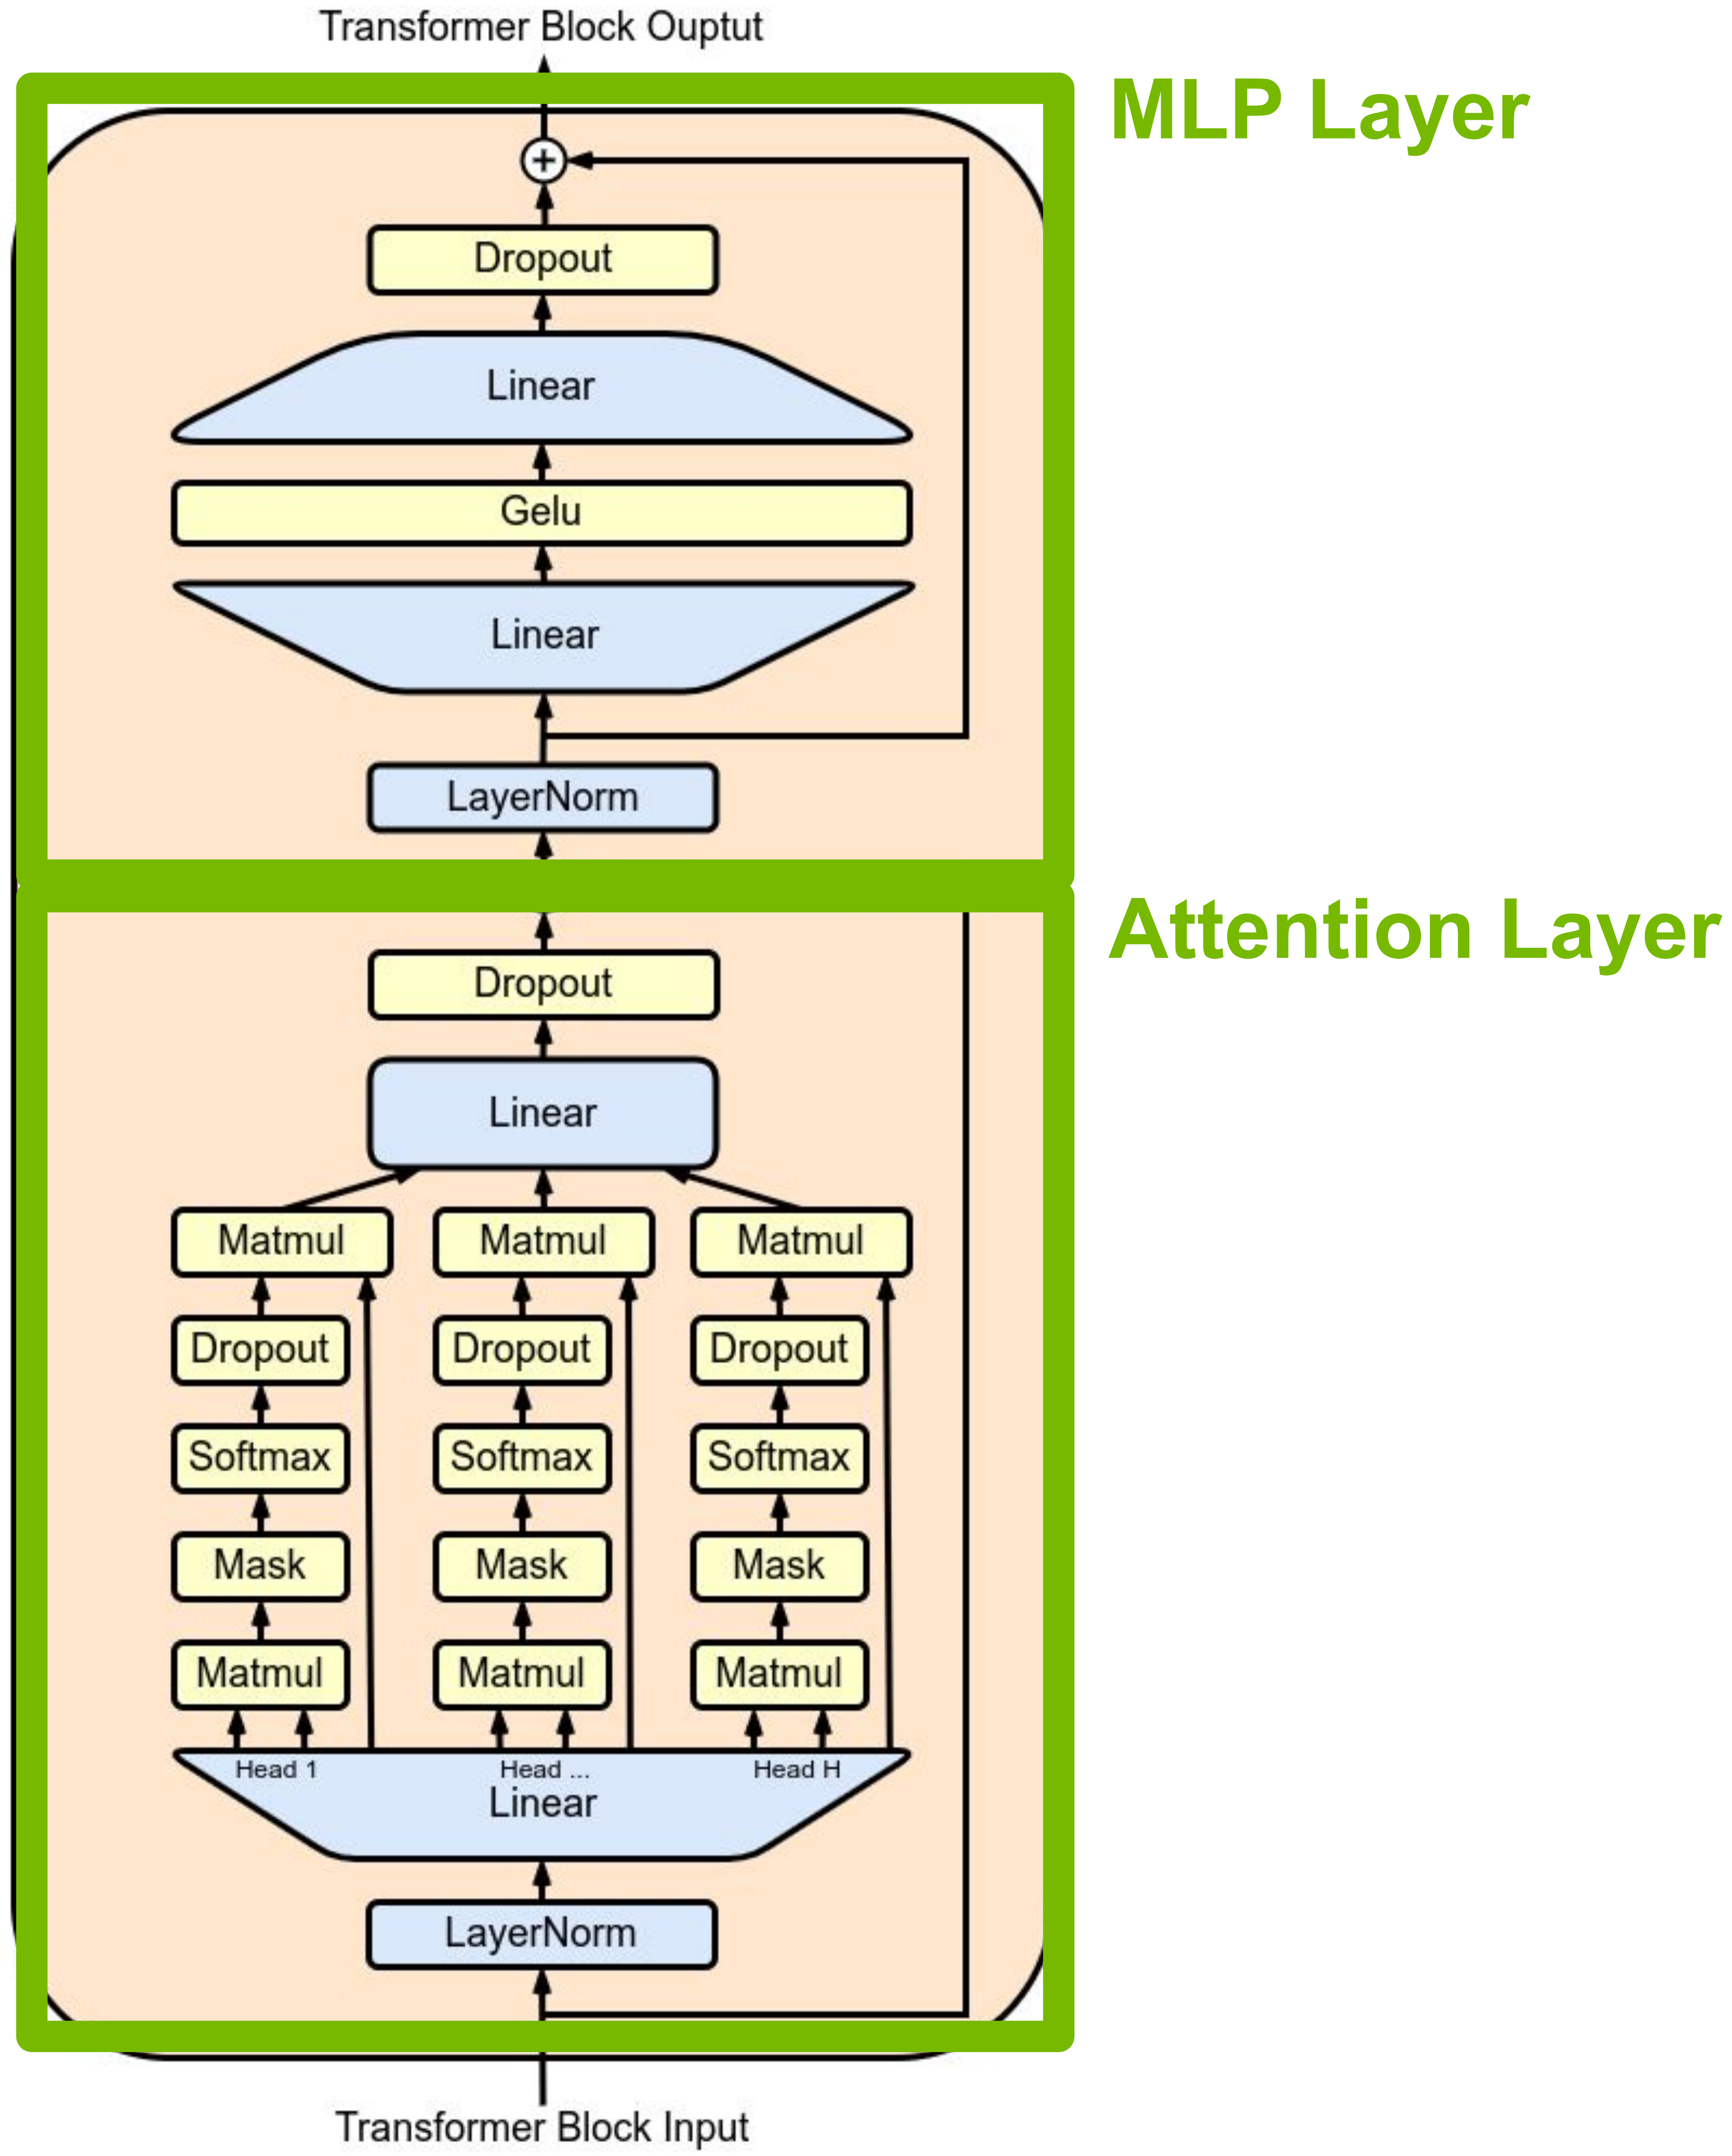

Introduction to LLM

Rapid Advances in LLM Performance

ChatGPT 3.5

Remember the 2022?
The ChatGPT Moment!

You
What is this?

ChatGPT
This is a text-based conversational AI powered by OpenAI's GPT-3.5 model. It is designed to provide responses and engage in conversation on a wide range of topics. You can ask questions, seek information, or have a general conversation with me. Is there something specific you would like to know or discuss?

You
Chatbots have advanced significantly since I last used one

ChatGPT
Indeed, chatbot technology has made significant advancements in recent years. With the introduction of machine learning techniques such as deep learning and neural networks,

Message ChatGPT...

ChatGPT can make mistakes. Consider checking important information.

Introduction to LLM

Rapid Advances in LLM Performance

연합뉴스

오픈AI, 차세대 AI모델 '아스트라' 예고…"수학 난제 10개 해결"

송고 2026-08-03 02:27

AI Models in 2026
→ AI That Goes Beyond Human Capabilities

Claude Mythos

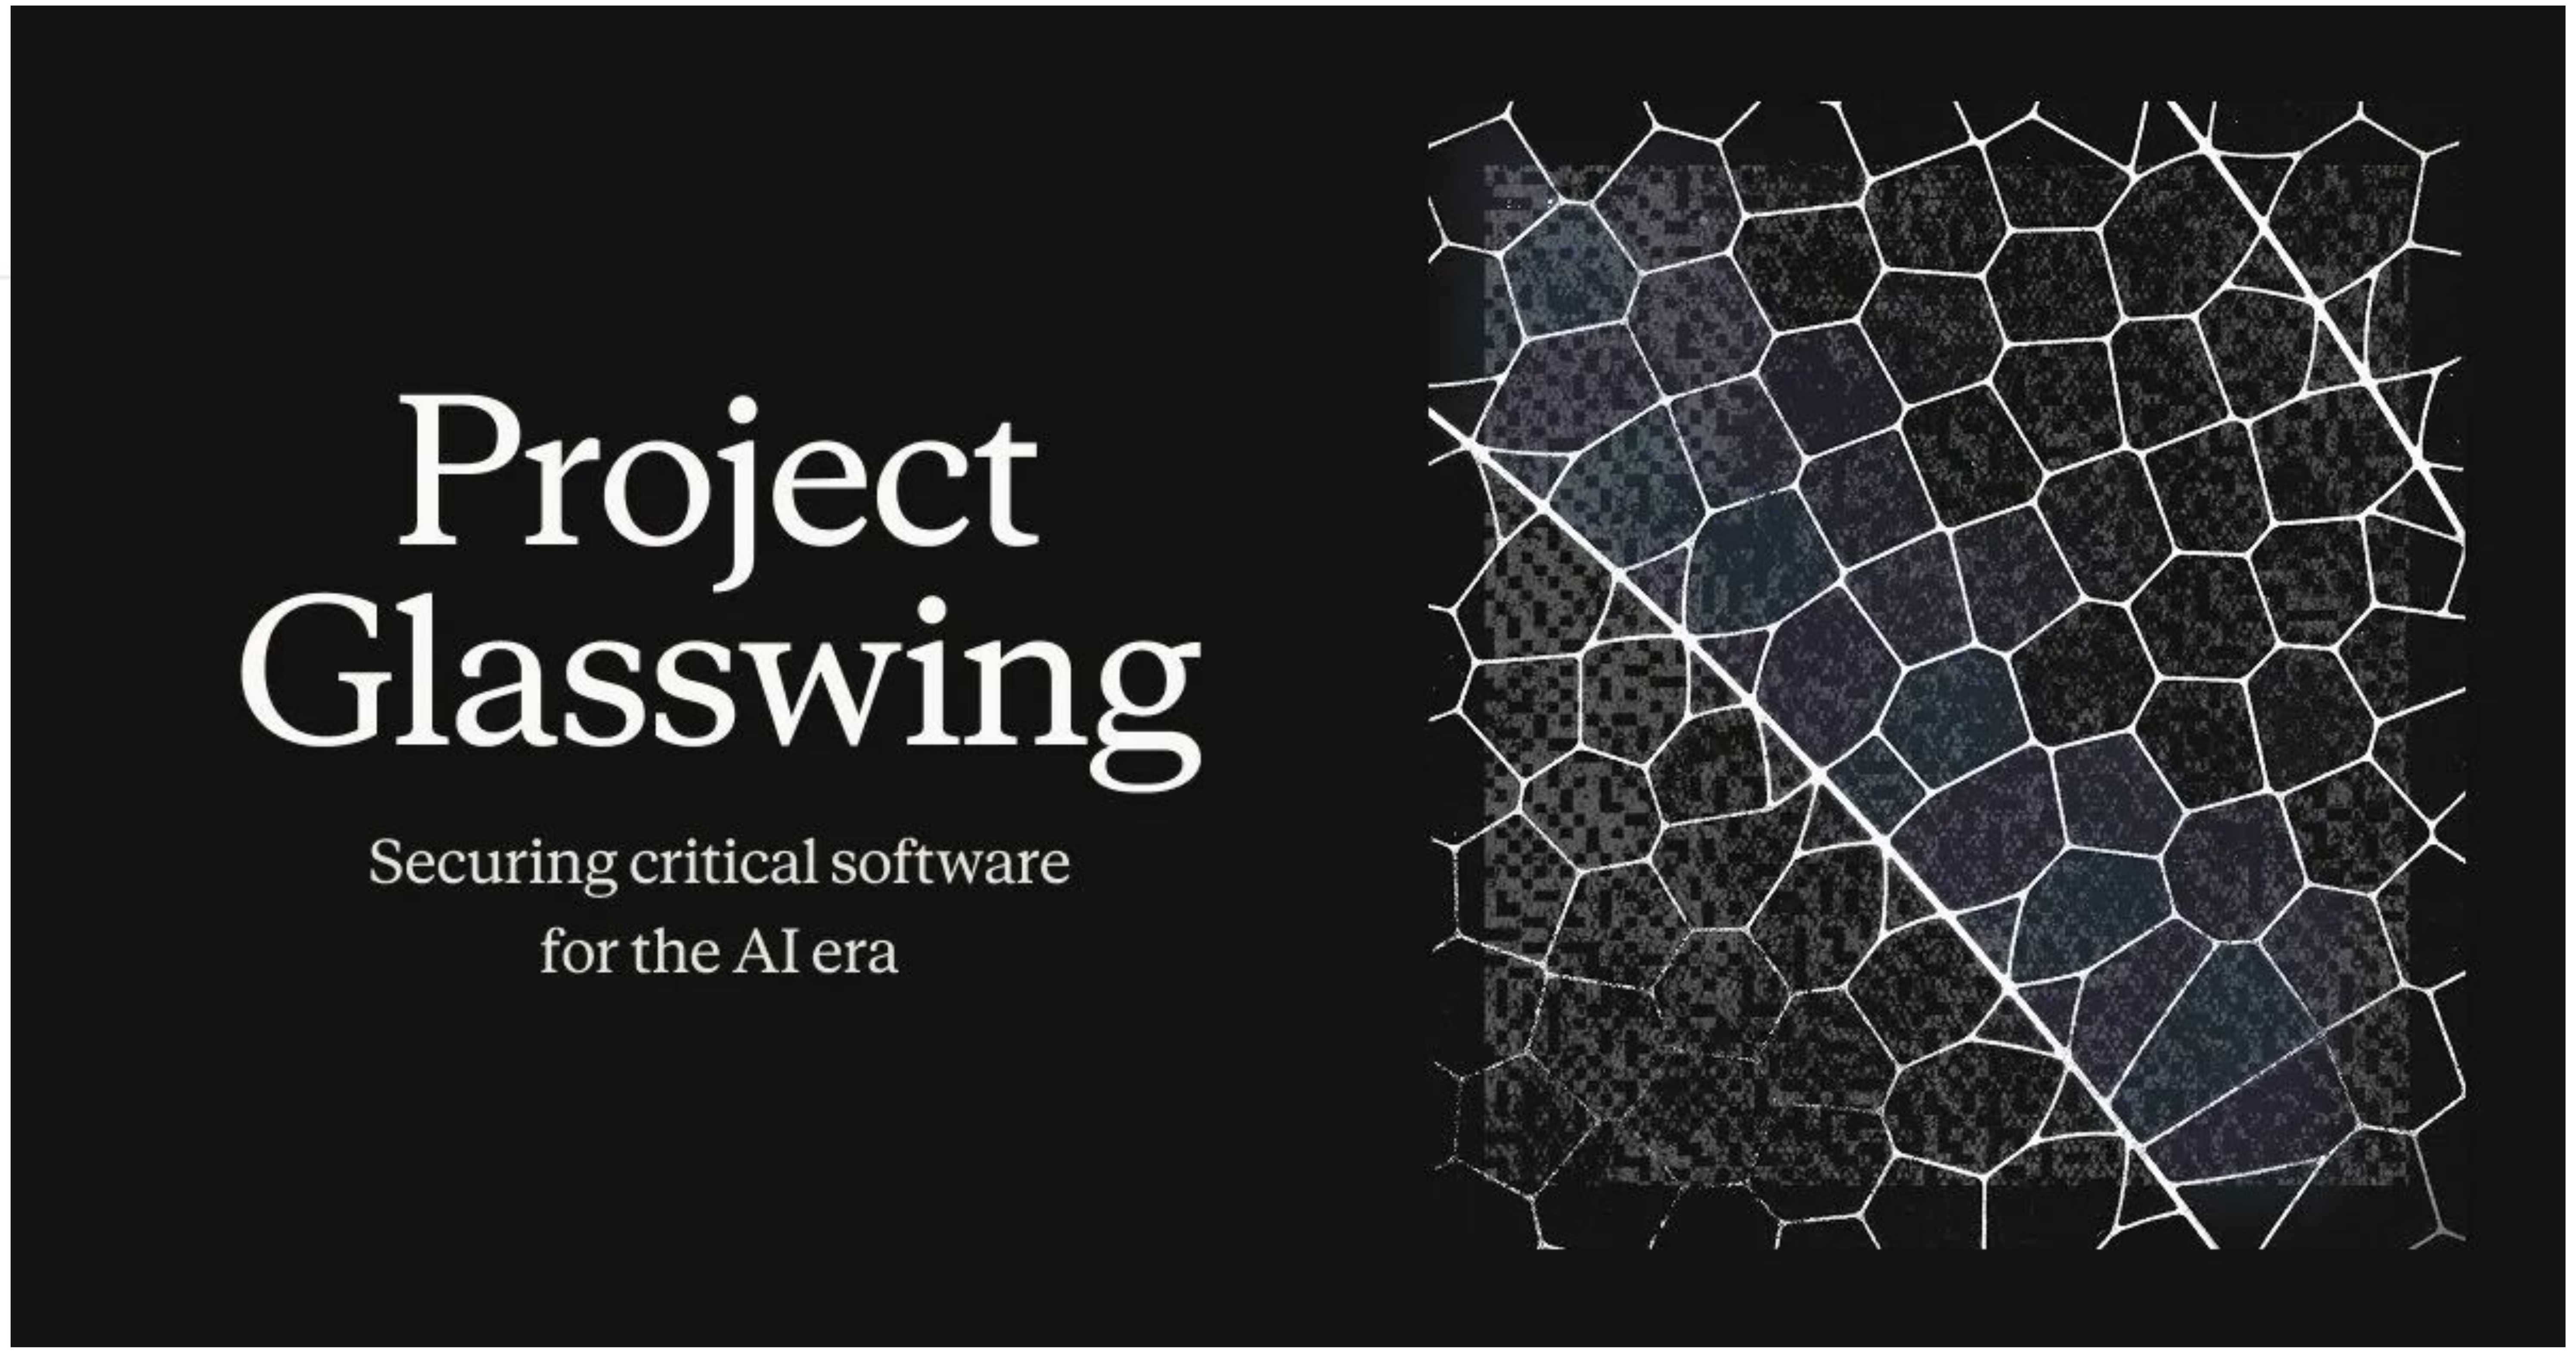

Introduction to LLM

Rapid Advances in LLM Performance

Frontier Language Model Intelligence, Over Time

Artificial Analysis Intelligence Index v4.1.1 incorporates 9 evaluations: GDPval-AA v2, τ^3-Banking, Terminal-Bench v2.1, SciCode, Humanity's Last Exam, GPQA Diamond, CritPt, AA-Omniscience, AA-LCR

https://artificialanalysis.ai/?models=muse-spark-1-2%2Cgemini-3-5-flash-lite%2Cinkling%2Ck-exaone-2-0-0803%2Cglm-5-3-flash%2Cminimax-m3%2Cnemotron-3-5-lightning%2Ccommand-a-plus%2Cgpt-5-6-luna%2Cmuse-glimmer%2Cnvidia-nemotron-3-ultra-550b-a55b%2Csolar-open2-250b%2Ca-x-k2%2Cdeepseek-v4-pro%2Cqwen3-8-2-4t-a95b%2Cclaude-4-5-haiku-reasoning%2Cqwen3-8-27b%2Cmotif-3%2Cgemini-3-7-flash%2Cclaude-opus-5%2Cgpt-5-6-terra%2Cgrok-4-6%2Cclaude-fable-5%2Cglm-5-3%2Cgpt-5-6-sol%2Cmistral-medium-3-5%2Cgpt-5-5-pro%2Cgpt-oss-120b%2Ckimi-k3%2Cgpt-35-turbo

Introduction to LLM

Rapid Advances in LLM Performance

https://artificialanalysis.ai/?models=muse-spark-1-2%2Cgemini-3-5-flash-lite%2Cinkling%2Ck-exaone-2-0-0803%2Cglm-5-3-flash%2Cminimax-m3%2Cnemotron-3-5-lightning%2Ccommand-a-plus%2Cgpt-5-6-luna%2Cmuse-glimmer%2Cnvidia-nemotron-3-ultra-550b-a55b%2Csolar-open2-250b%2Ca-x-k2%2Cdeepseek-v4-pro%2Cqwen3-8-2-4t-a95b%2Cclaude-4-5-haiku-reasoning%2Cqwen3-8-27b%2Cmotif-3%2Cgemini-3-7-flash%2Cclaude-opus-5%2Cgpt-5-6-terra%2Cgrok-4-6%2Cclaude-fable-5%2Cglm-5-3%2Cgpt-5-6-sol%2Cmistral-medium-3-5%2Cgpt-5-5-pro%2Cgpt-oss-120b%2Ckimi-k3%2Cgpt-35-turbo

Building RAG Agents with LLMs

Part 2. LLM Reasoning and Agentic AI

LLM Reasoning and Agentic AI

2025–2026: The Era of AI Agent

AI Models in 2026
→ AI Agent

🦞 OpenClaw — Personal AI Assistant

OPENCLAW

EXFOLIATE! EXFOLIATE!

BUILD PASSING | RELEASE V2026.2.12 | DISCORD 18K ONLINE | LICENS

OpenClaw is a *personal AI assistant* you run on your own devices. It answers you on the channels you already use (WhatsApp, Telegram, Slack, Discord, Google Chat, Signal, iMessage, Microsoft Teams, WebChat), plus extension channels like BlueBubbles, Matrix, Zalo, and Zalo Personal. It can speak and listen on macOS/iOS/Android, and can render a live Canvas you control. The Gateway is just the control plane — the product is the assistant.

If you want a personal, single-user assistant that feels local, fast, and always-on, this is it.

LLM Reasoning and Agentic AI

2025–2026: The Era of AI Agent

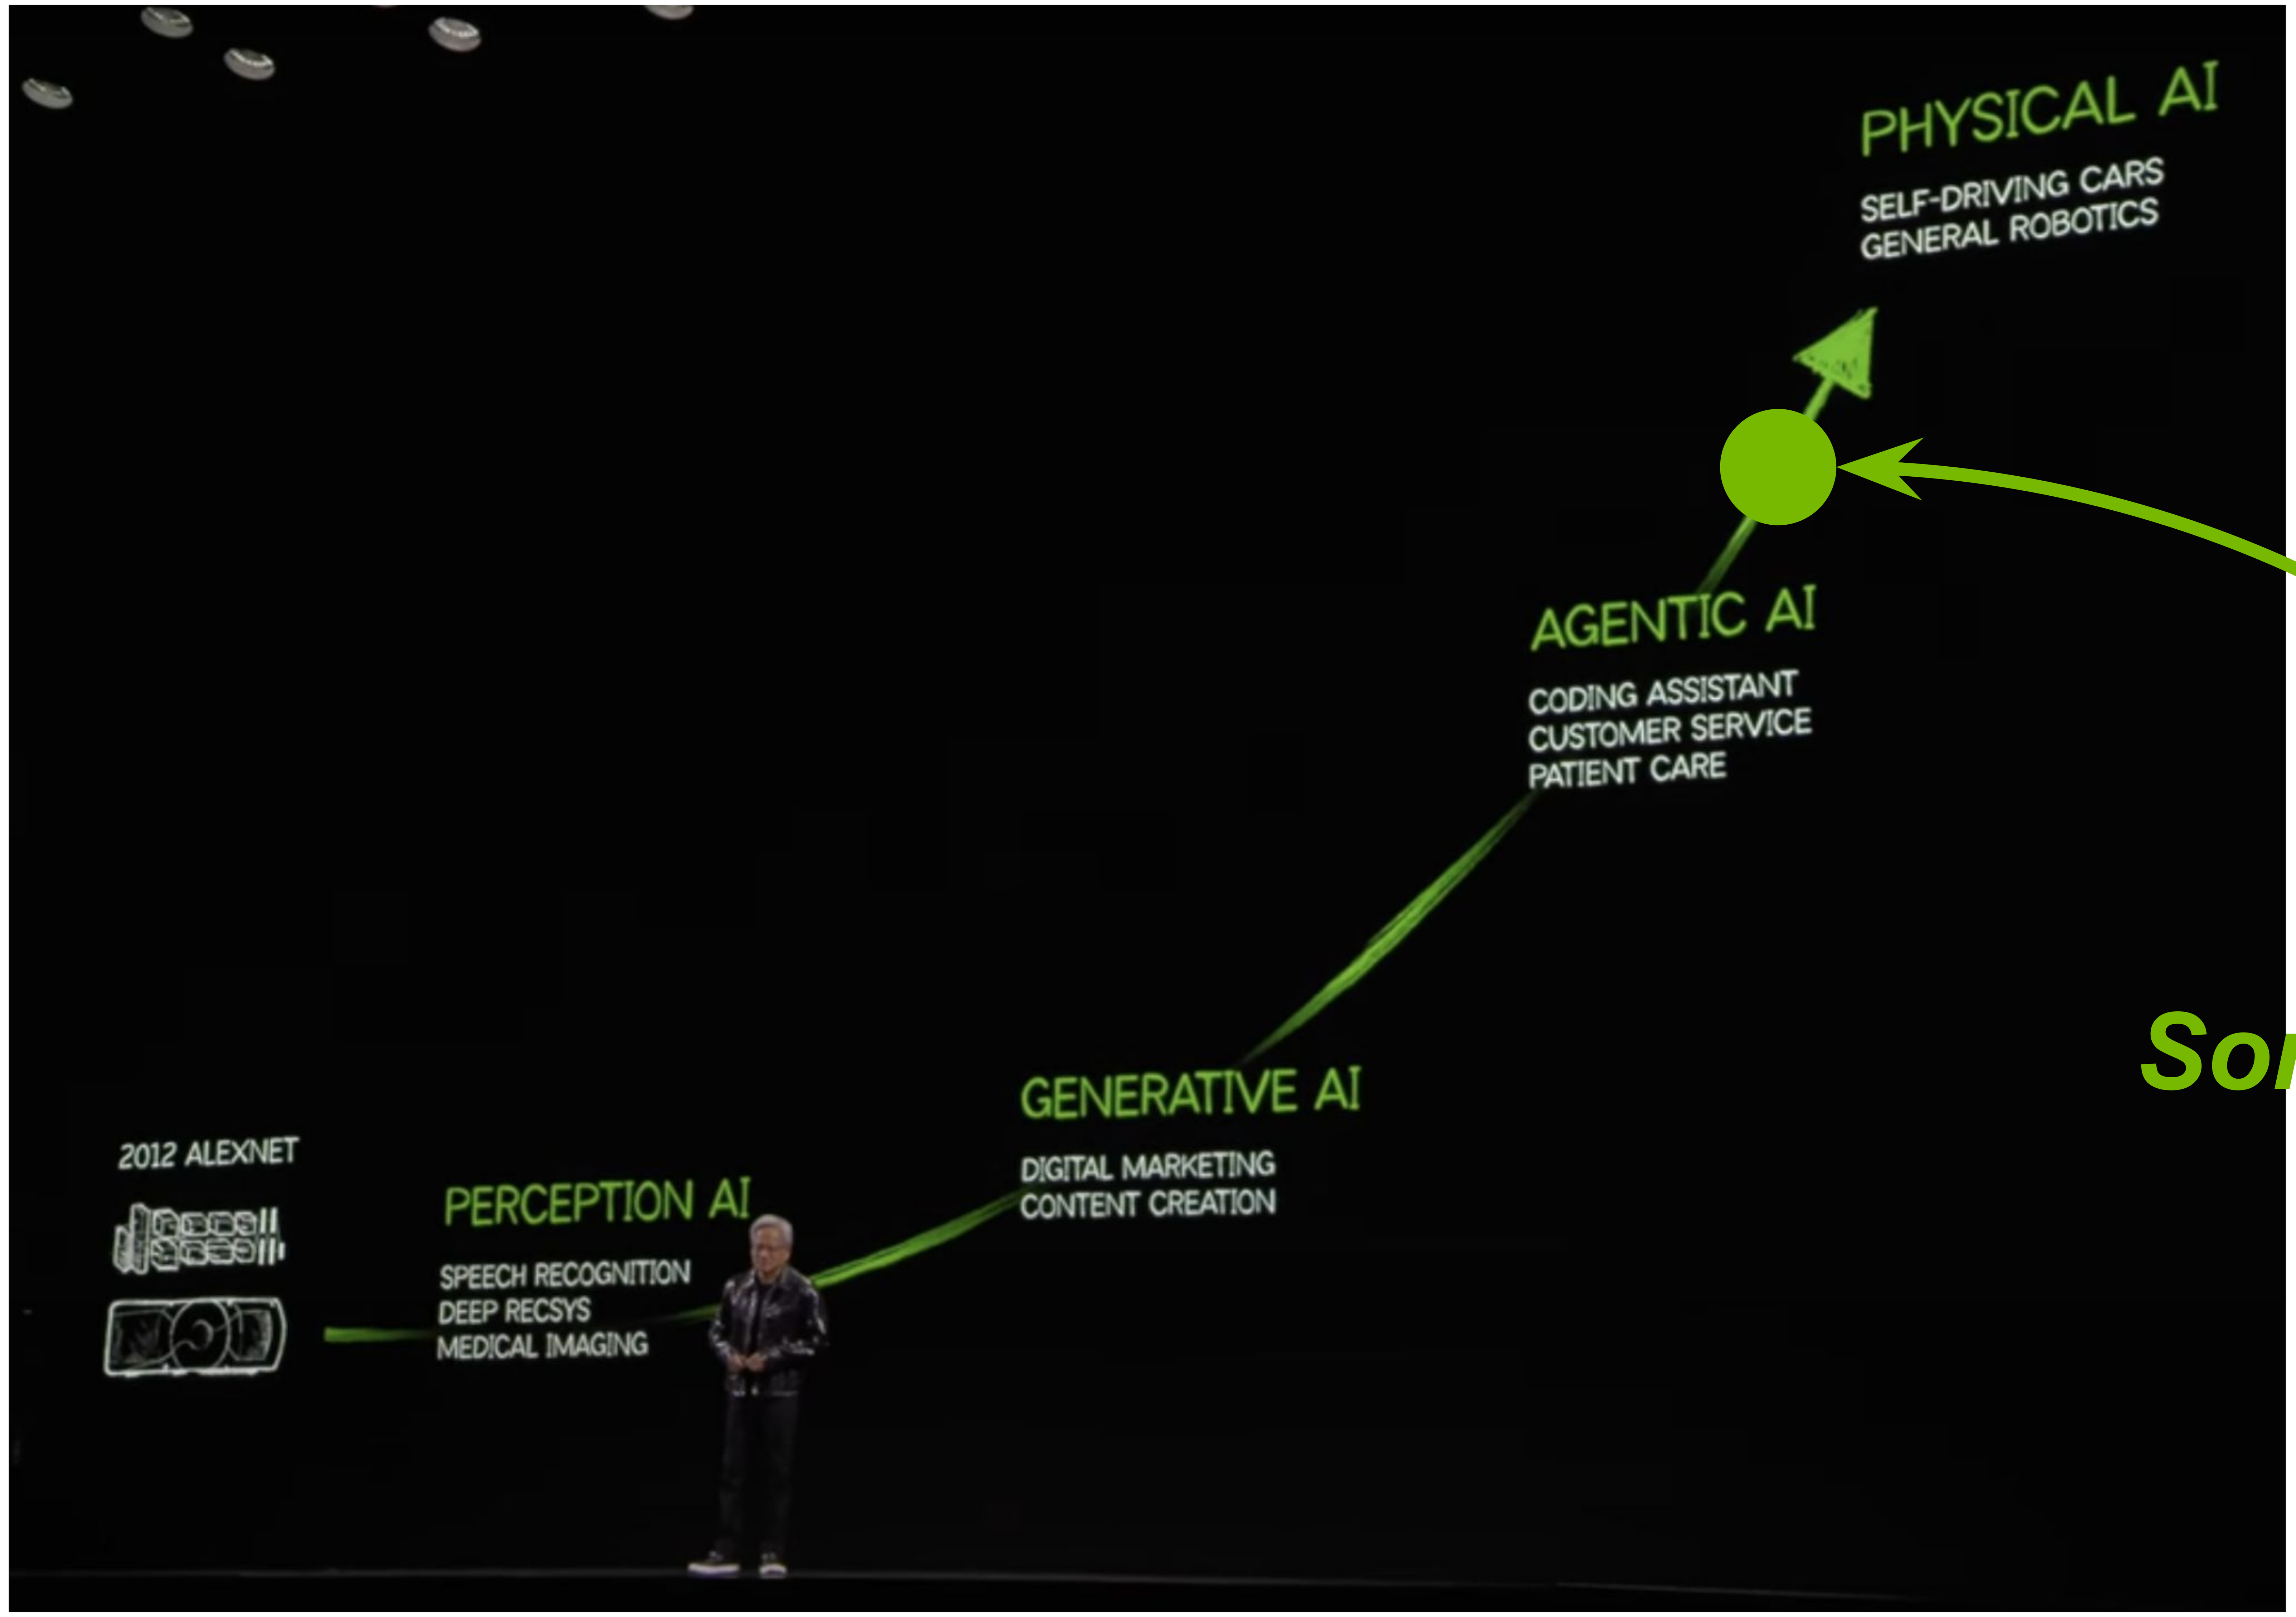

Somewhere in here

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

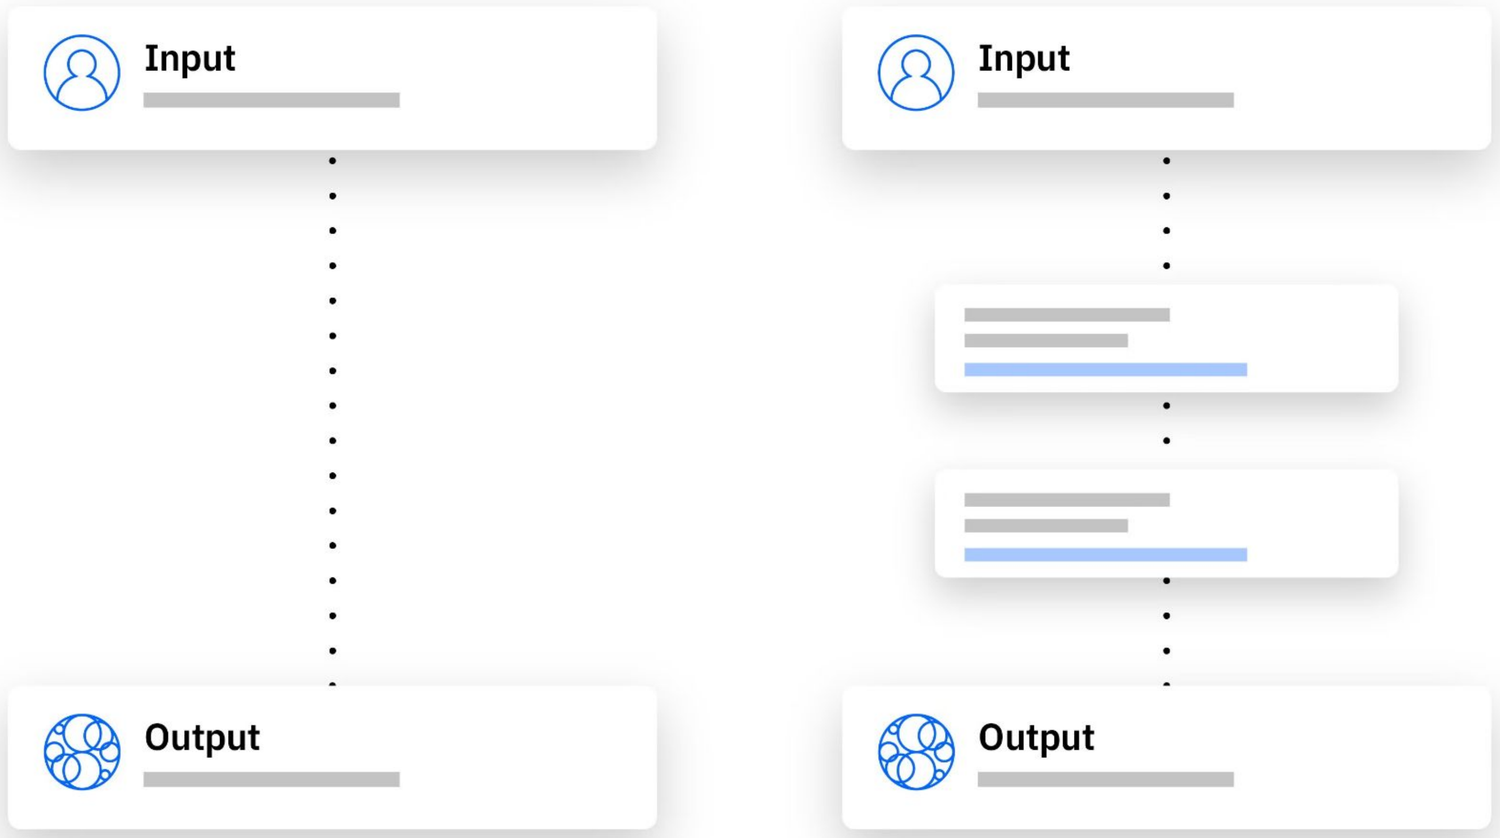

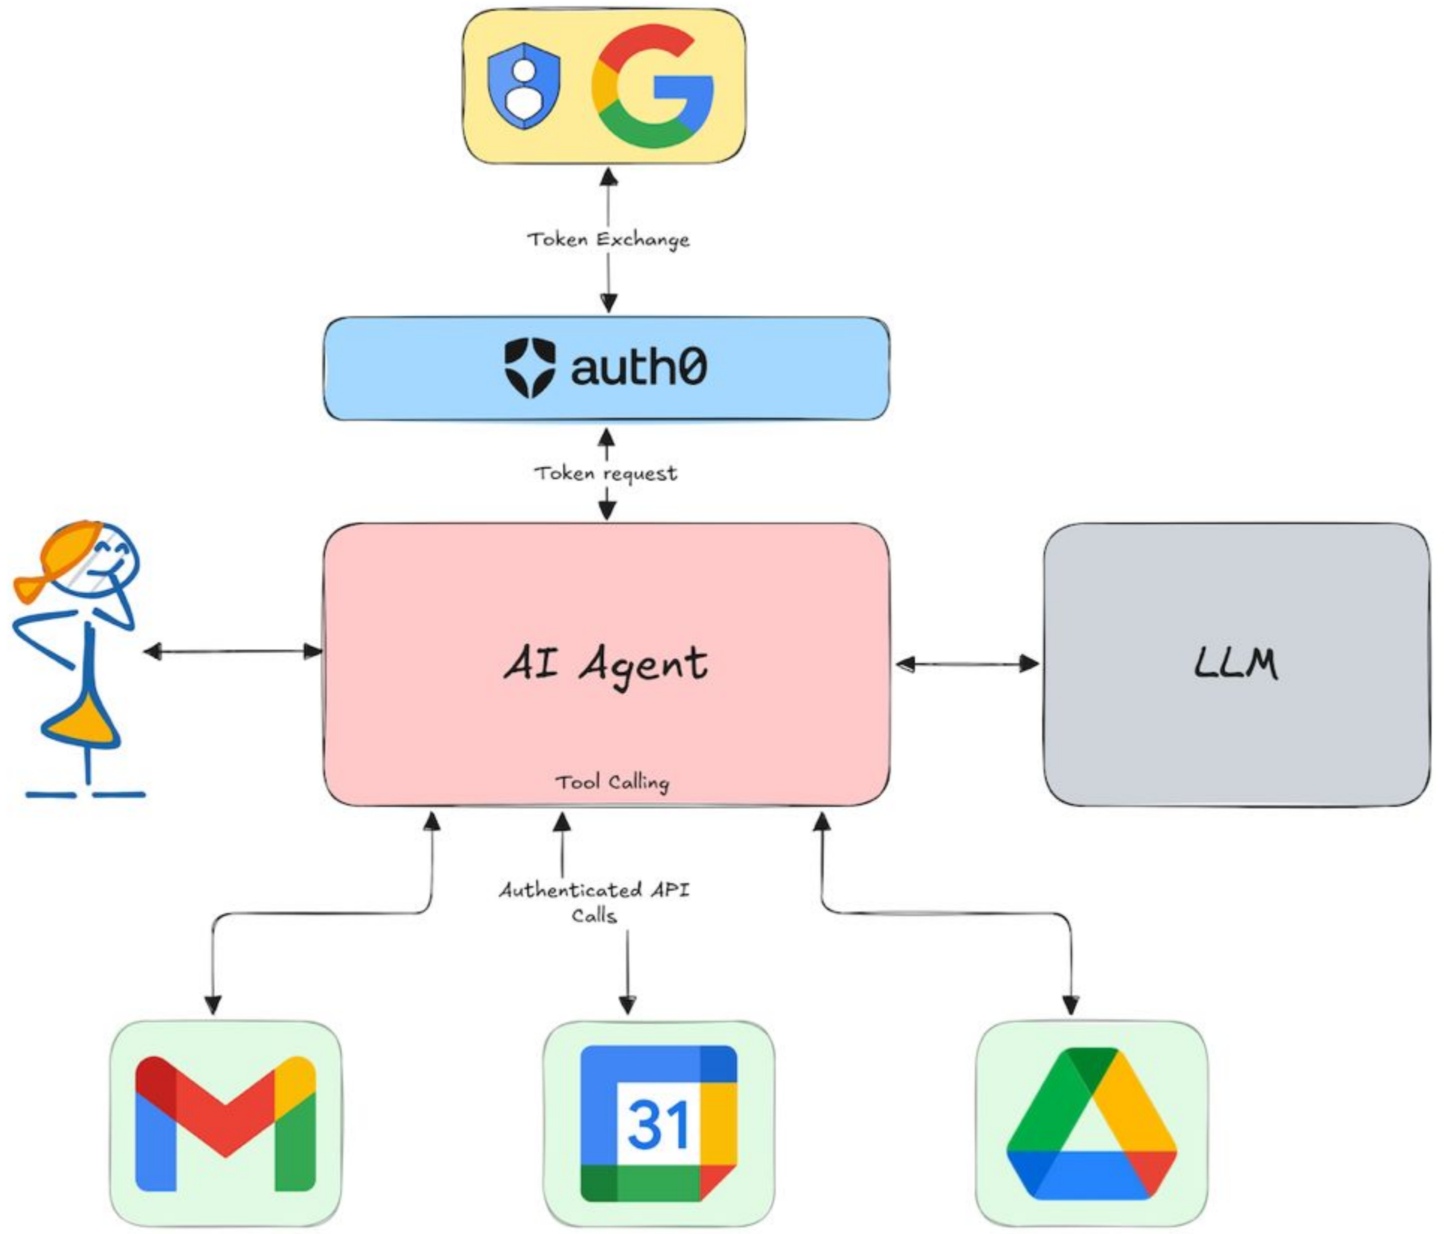

Chain-of-Thought → Function Calling → Agentic AI

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

<|system|> You are a helpful assistant. When solving a problem, think step by step before giving the final answer.
<|user|> 한 상자에 연필이 12자루 있습니다. 이런 상자가 5개 있다면 연필은 모두 몇 자루인가요?

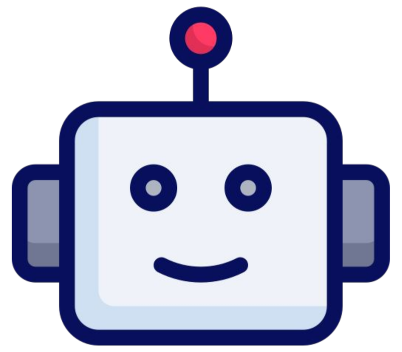

<|assistant|> <think> I need to determine the total number of pencils. Each box contains 12 pencils. There are 5 boxes. Total pencils = 12 × 5 = 60. So there are 60 pencils in total. </think> 총 60자루의 연필이 있습니다.

1st Evolution: Enhancing Reasoning Capabilities through CoT

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

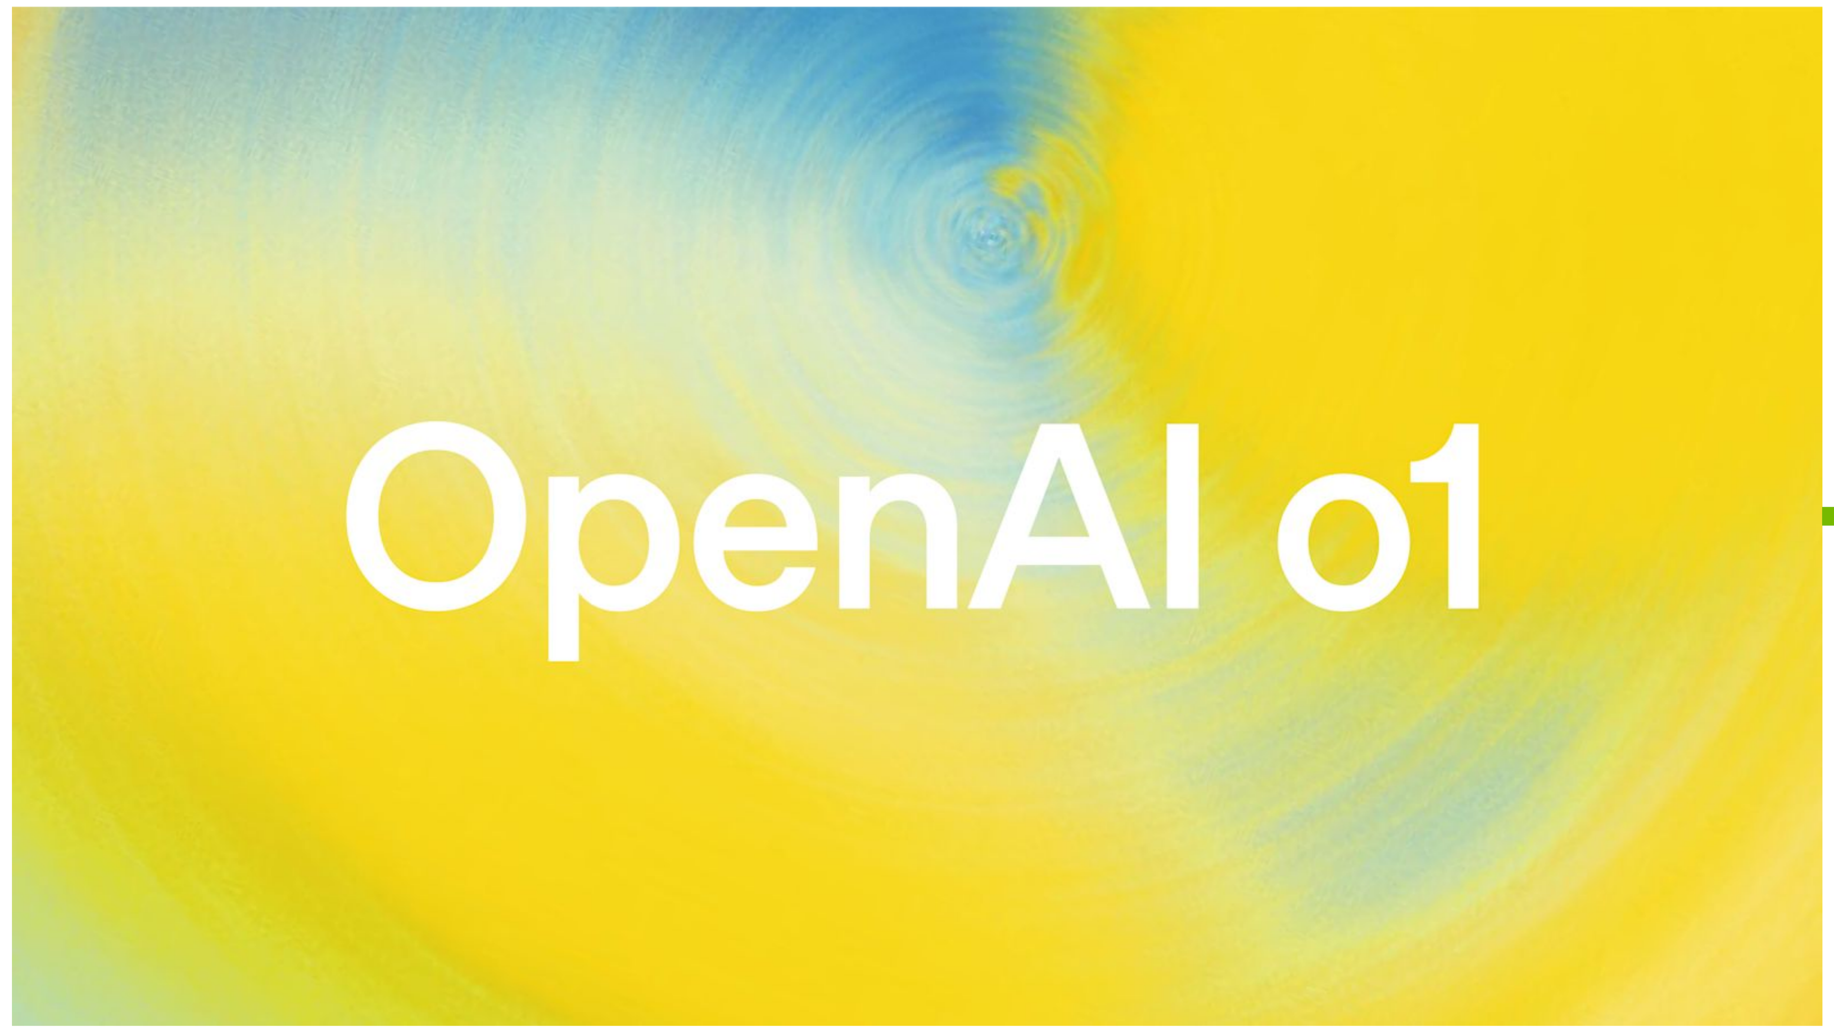

Remember the Deepseek Shock?

→ in the Academia Field, Deepseek Moment

Academic Reproduction of Reasoning via RL

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

deepseek

DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning

DeepSeek-AI

research@deepseek.com

DeepSeek-R1: How CoT Emerges?

Abstract

General reasoning represents a long-standing and formidable challenge in artificial intelligence. Recent breakthroughs, exemplified by large language models (LLMs) (Brown et al., 2020; OpenAI, 2023) and chain-of-thought prompting (Wei et al., 2022b), have achieved considerable success on foundational reasoning tasks. However, this success is heavily contingent upon extensive human-annotated demonstrations, and models' capabilities are still insufficient for more complex problems. Here we show that the reasoning abilities of LLMs can be incentivized through pure reinforcement learning (RL), obviating the need for human-labeled reasoning trajectories. The proposed RL framework facilitates the emergent development of advanced reasoning patterns, such as self-reflection, verification, and dynamic strategy adaptation. Consequently, the trained model achieves superior performance on verifiable tasks such as mathematics, coding competitions, and STEM fields, surpassing its counterparts trained via conventional supervised learning on human demonstrations. Moreover, the emergent reasoning patterns exhibited by these large-scale models can be systematically harnessed to guide and enhance the reasoning capabilities of smaller models.

Guo, D., Yang, D., Zhang, H., Song, J., Wang, P., Zhu, Q., ... & He, Y. (2025). Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948.*

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

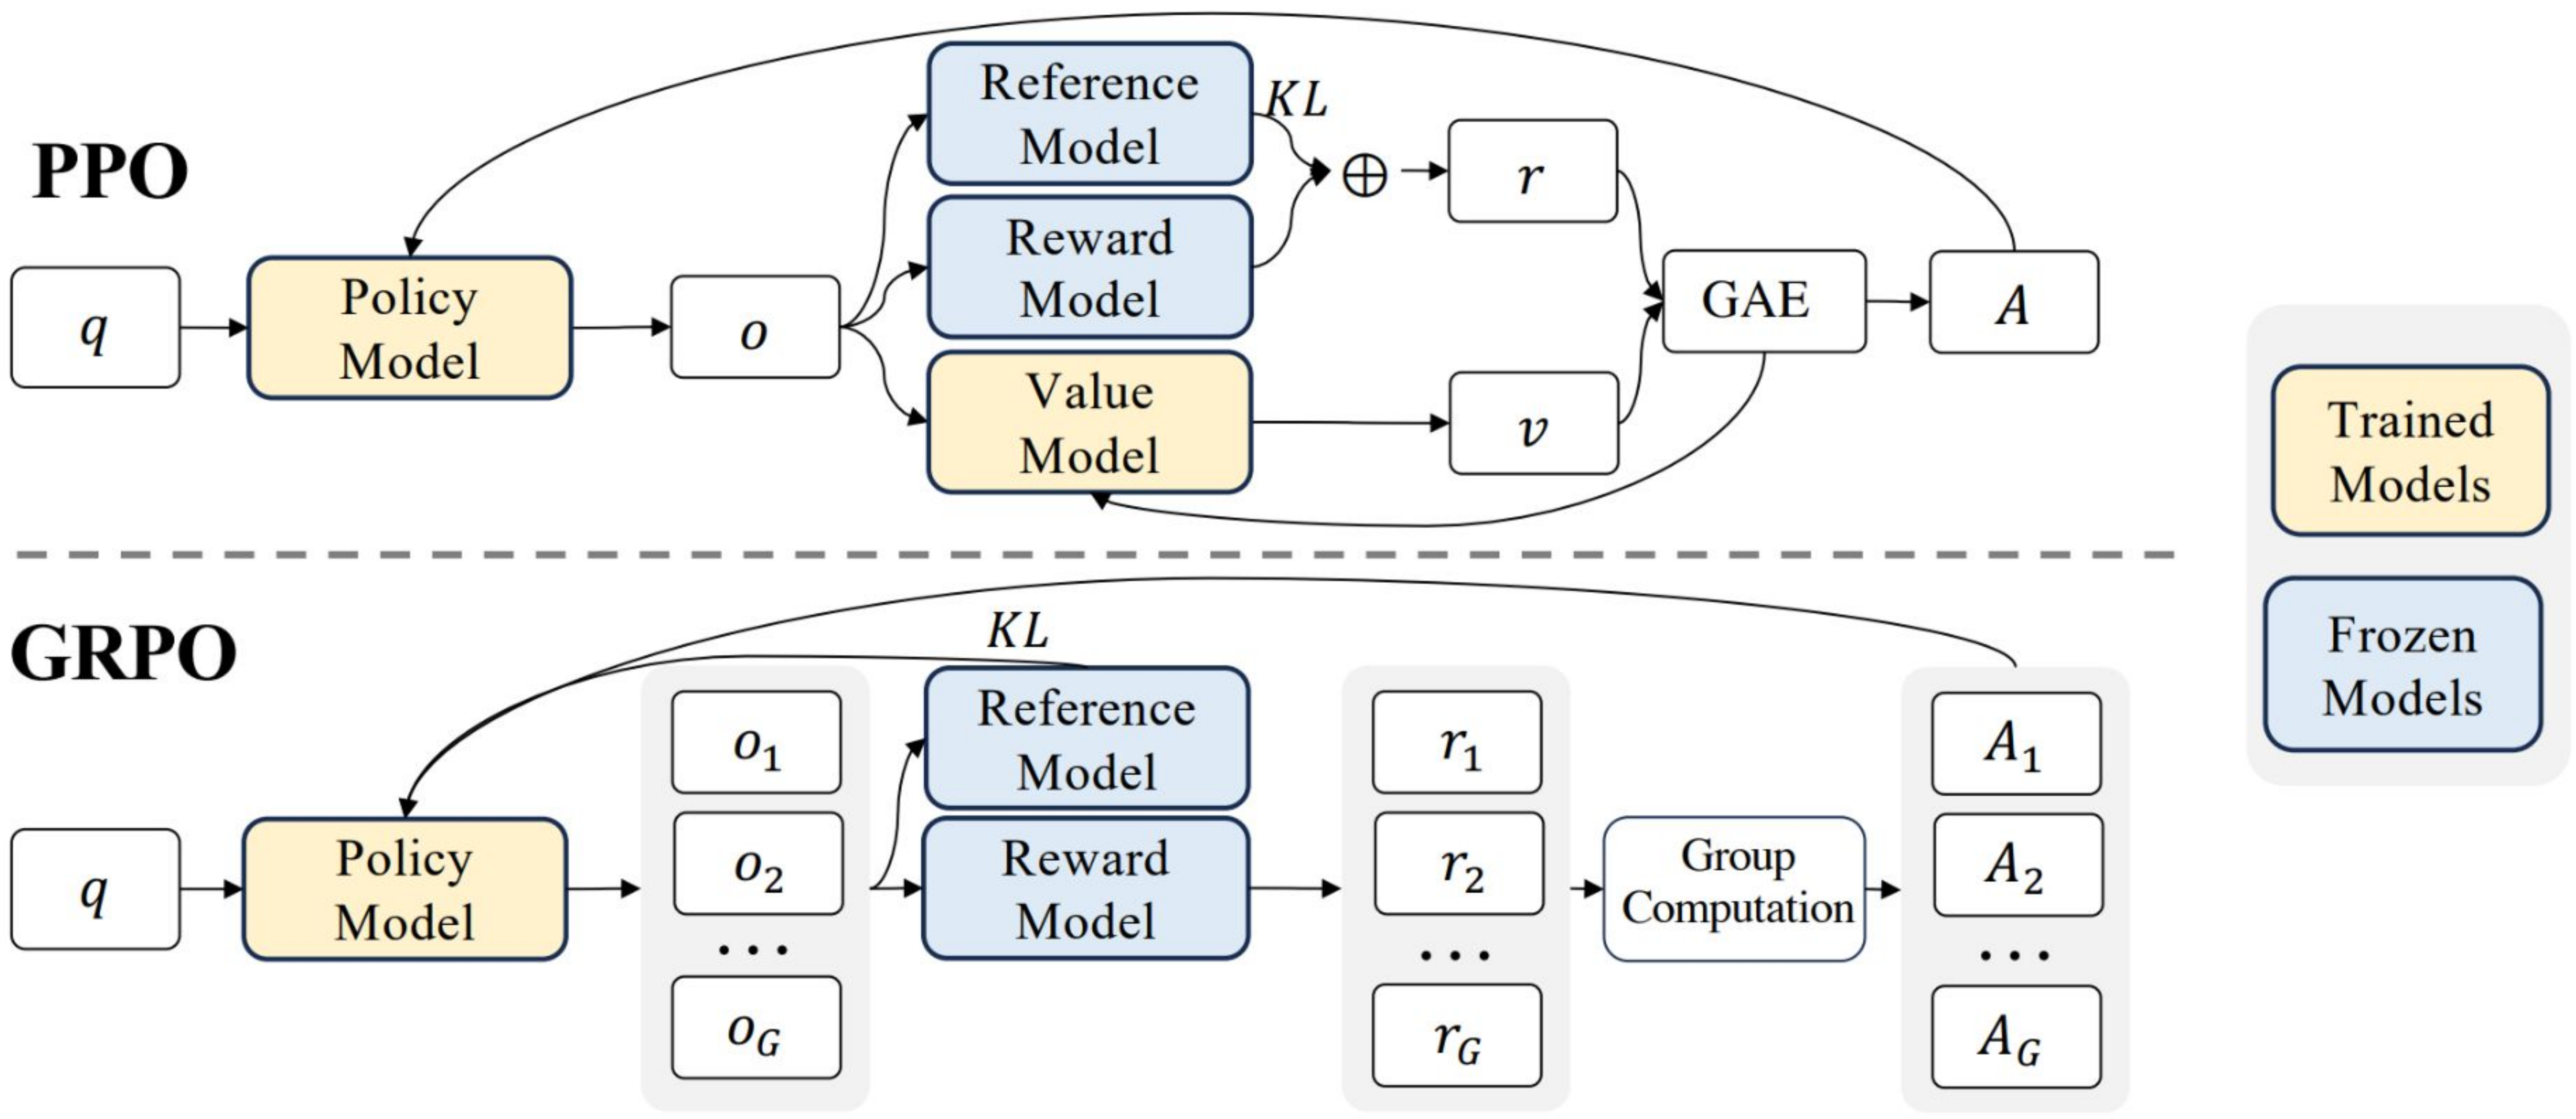

Figure 4 | Demonstration of PPO and our GRPO. GRPO foregoes the value model, instead estimating the baseline from group scores, significantly reducing training resources.

Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., ... & Guo, D. (2024). Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*.

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

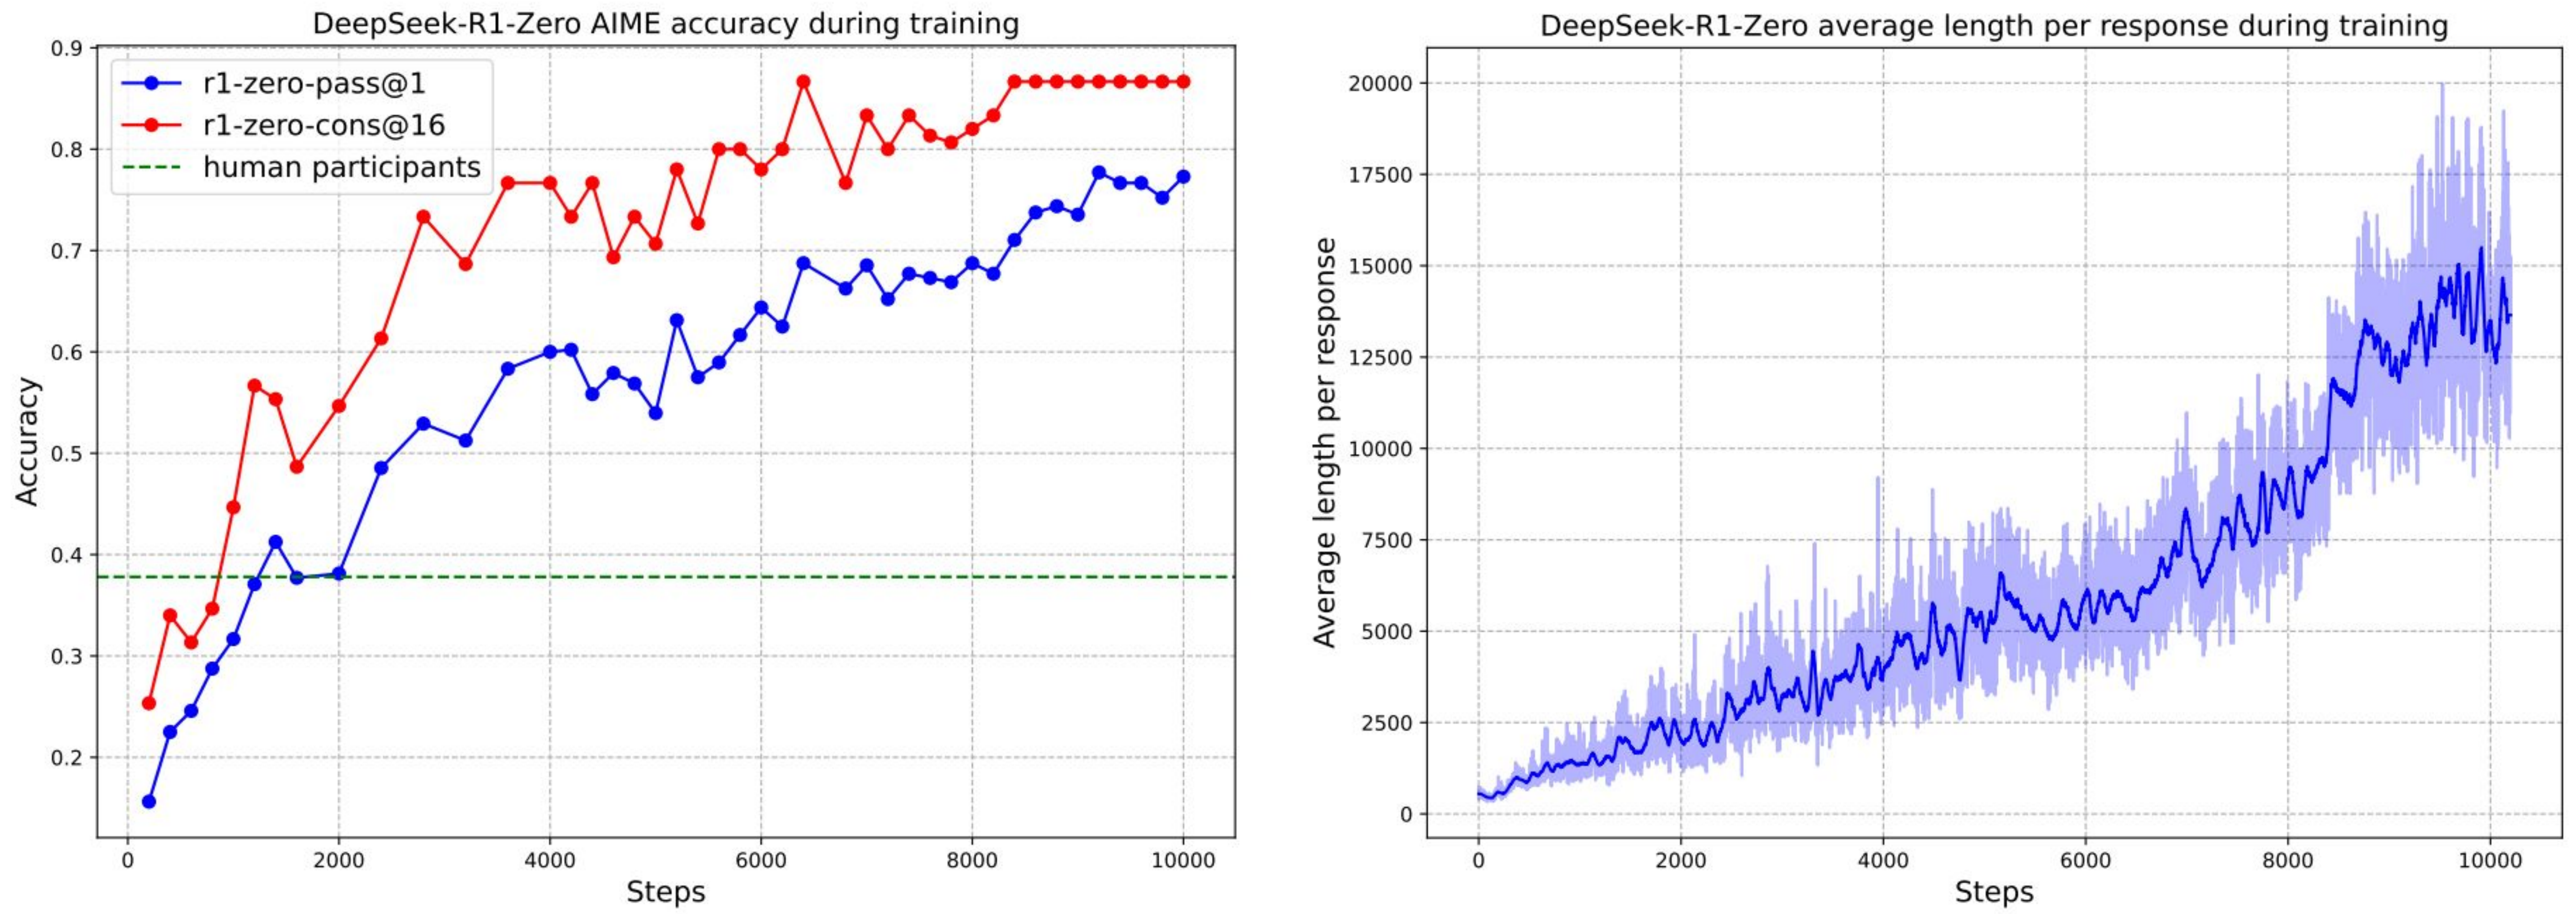

Figure 1 | (a) AIME accuracy of DeepSeek-R1-Zero during training. AIME takes a mathematical problem as input and a number as output, illustrated in Table 32. Pass@1 and Cons@16 are described in Supplementary D.1. The baseline is the average score achieved by human participants in the AIME competition. (b) The average response length of DeepSeek-R1-Zero on the training set during the RL process. DeepSeek-R1-Zero naturally learns to solve reasoning tasks with more thinking time. Note that a training step refers to a single policy update operation.

Guo, D., Yang, D., Zhang, H., Song, J., Wang, P., Zhu, Q., ... & He, Y. (2025). Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948.*

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

Table 2 | An interesting “aha moment” of an intermediate version of DeepSeek-R1-Zero. The model learns to rethink using an anthropomorphic tone. This is also an aha moment for us, allowing us to witness the power and beauty of reinforcement learning.

Question: If $a > 1$, then the sum of the real solutions of $\sqrt{a - \sqrt{a + x}} = x$ is equal to
Response: <think> To solve the equation $\sqrt{a - \sqrt{a + x}} = x$, let’s start by squaring both $\cdots$ $\left(\sqrt{a - \sqrt{a + x}}\right)^2 = x^2 \implies a - \sqrt{a + x} = x^2.$ Rearrange to isolate the inner square root term: $(a - x^2)^2 = a + x \implies a^2 - 2ax^2 + (x^2)^2 = a + x \implies x^4 - 2ax^2 - x + (a^2 - a) = 0$ $\ldots$ Wait, wait. Wait. That’s an aha moment I can flag here. Let’s reevaluate this step-by-step to identify if the correct sum can be $\cdots$ We started with the equation: $\sqrt{a - \sqrt{a + x}} = x$ First, let’s square both sides: $a - \sqrt{a + x} = x^2 \implies \sqrt{a + x} = a - x^2$ Next, I could square both sides again, treating the equation: $\cdots$ $\ldots$

Guo, D., Yang, D., Zhang, H., Song, J., Wang, P., Zhu, Q., ... & He, Y. (2025). Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*.

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

B.4.4. Training Cost

Regarding our research on DeepSeek-R1, we utilized the A100 GPUs to prepare for the experiments with a smaller model (30B parameters). The results from this smaller model have been promising, which has allowed us to confidently scale up to 660B R1-Zero and R1. For the training of DeepSeek-R1-Zero, we employed 64*8 H800 GPUs, and the process required approximately 198 hours. Additionally, during the training phase of DeepSeek-R1, we utilized the same 64*8 H800 GPUs, completing the process in about 4 days, or roughly 80 hours. To create the SFT datasets, we use 5K GPU hours. The details are shown in Table 7.

Training a GPT-level AI for approximately $5.6 million (~₩8 billion)

Model	**AIME 2024**		**MATH-500**	**GPQA Diamond**	**LiveCode Bench**	**CodeForces**
	pass@1	cons@64	pass@1	pass@1	pass@1	rating
OpenAI-o1-mini	63.6	80.0	90.0	60.0	53.8	1820
OpenAI-o1-0912	74.4	83.3	94.8	77.3	63.4	1843
DeepSeek-R1-Zero	71.0	86.7	95.9	73.3	50.0	1444

Table 2 | Comparison of DeepSeek-R1-Zero and OpenAI o1 models on reasoning-related benchmarks.

Guo, D., Yang, D., Zhang, H., Song, J., Wang, P., Zhu, Q., ... & He, Y. (2025). Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948.*

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

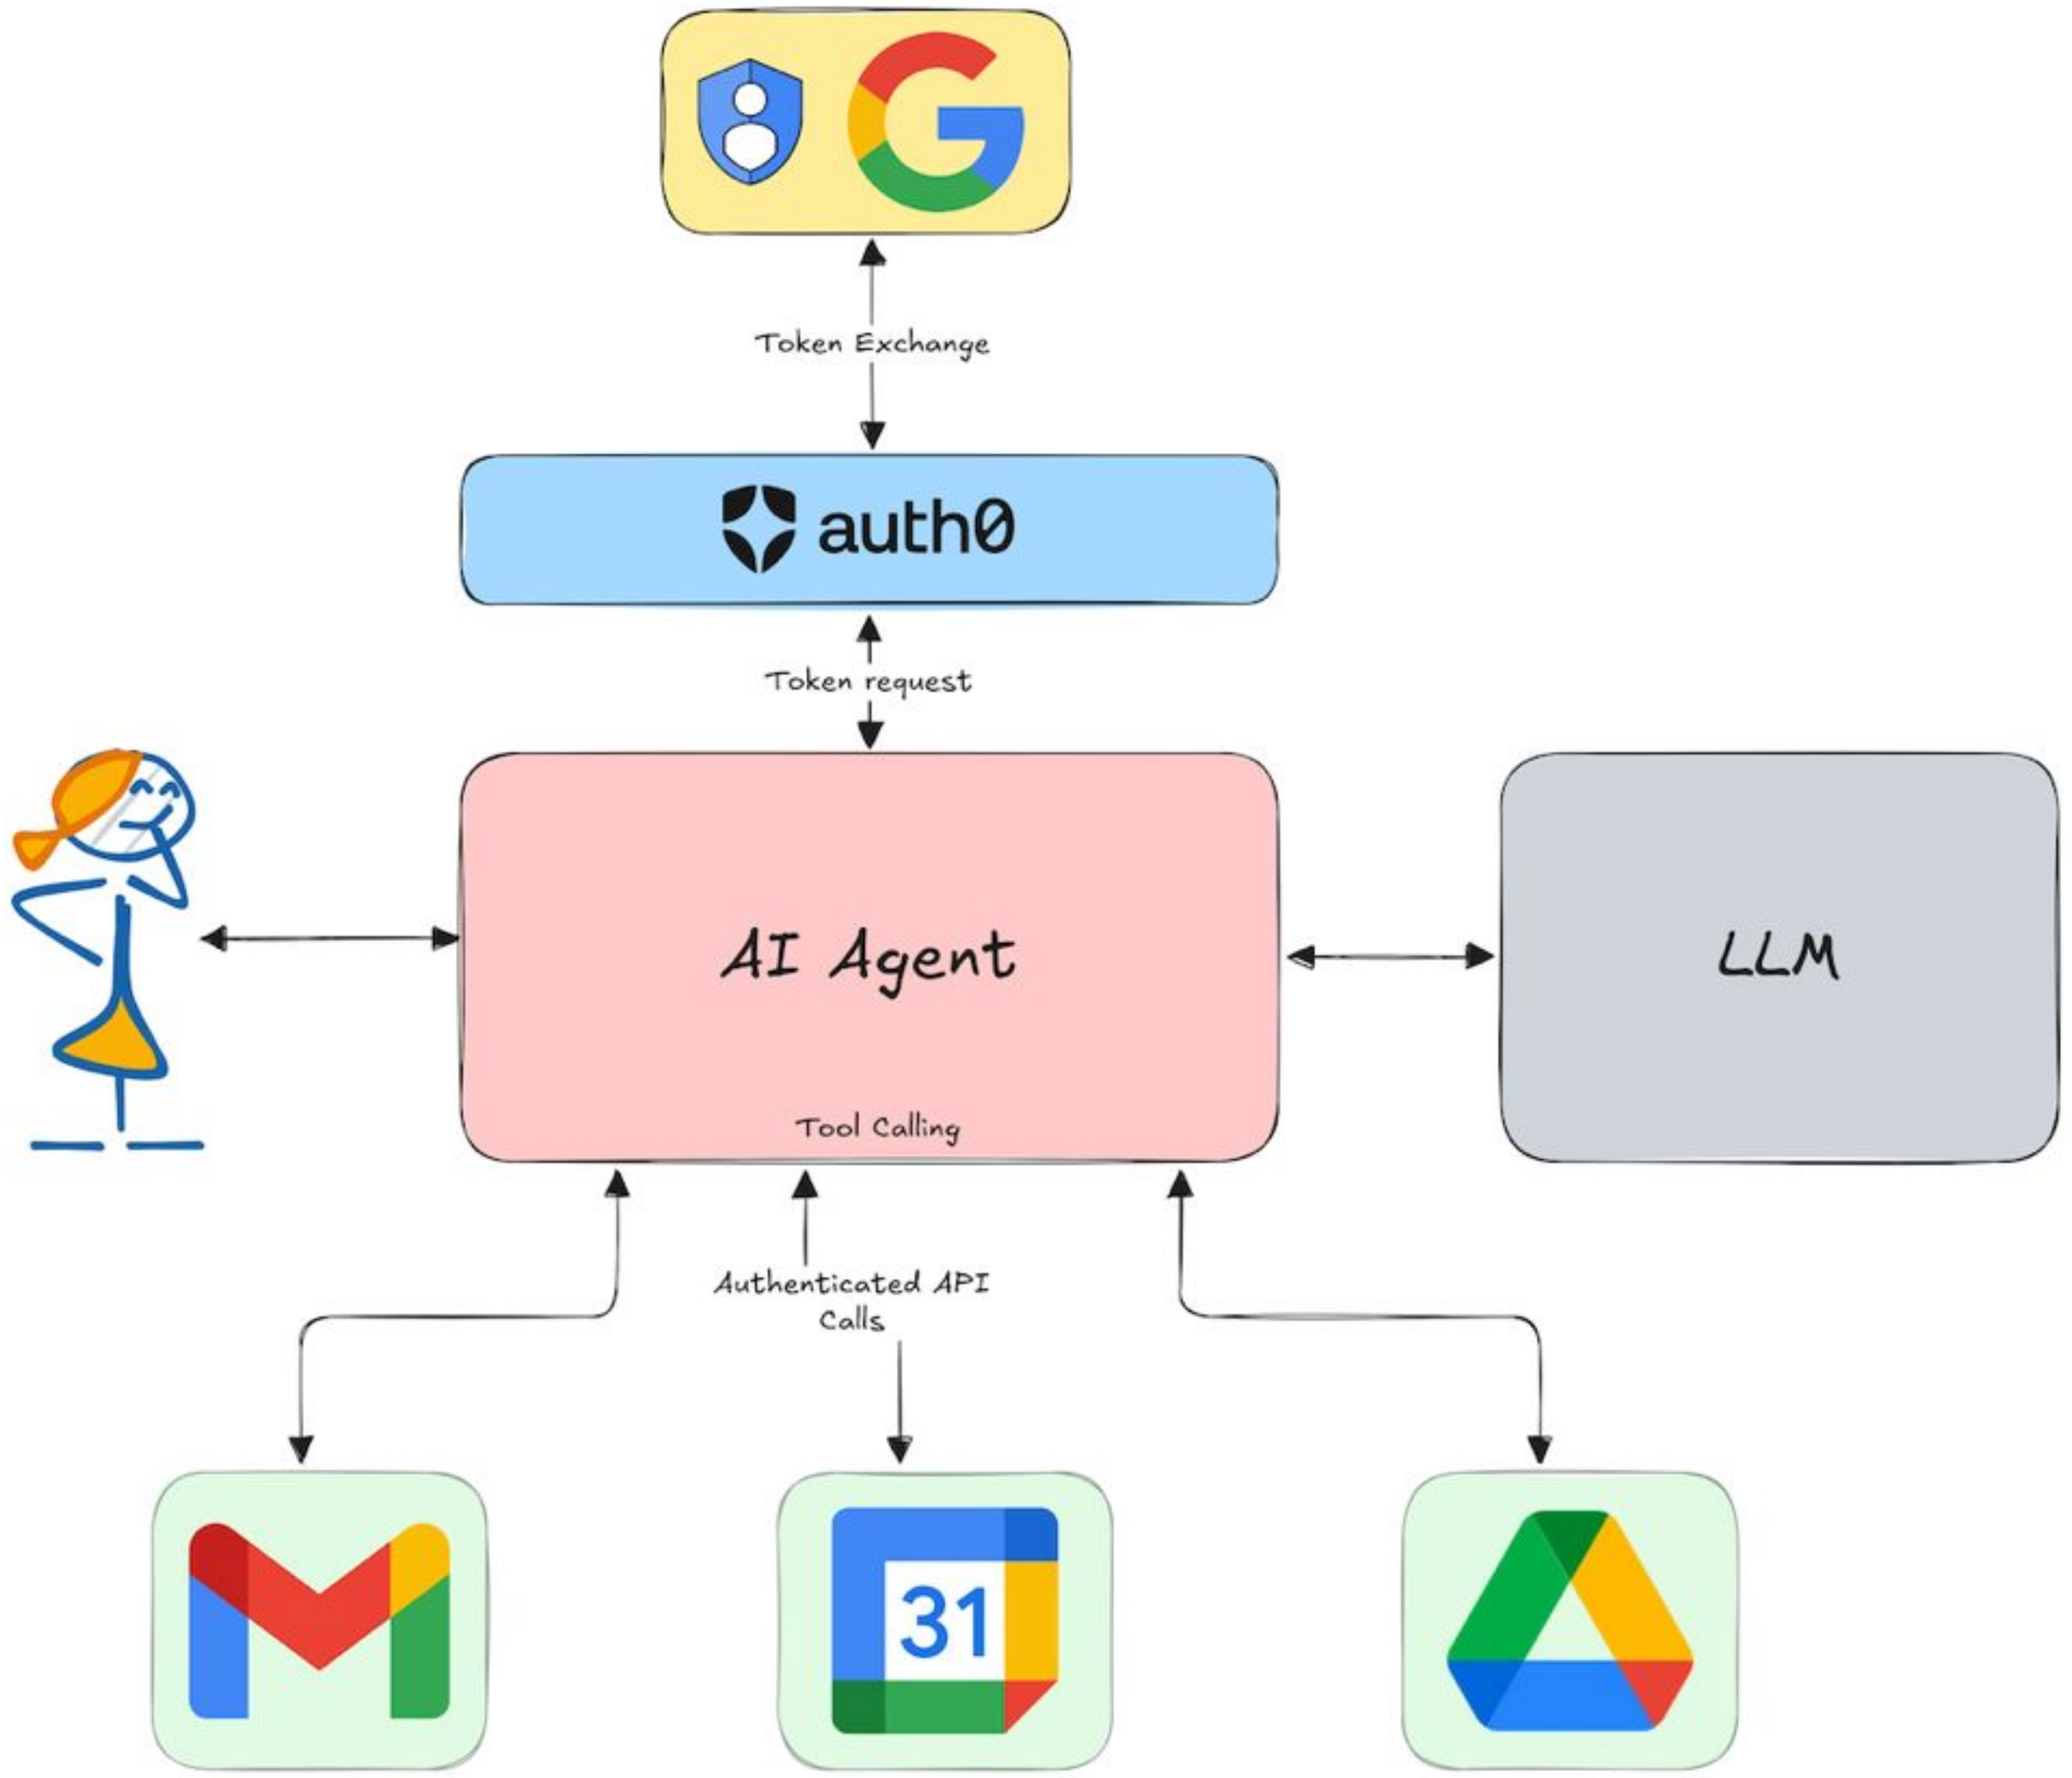

2nd Evolution: Extending Capabilities through Tool Use

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

<|start|>

Tools
You have access to the following tool:

get_weather(location: string)
→ Returns the current weather for a given location.

Tool Call Format
<|tool_call|>
{"name": "<function_name>", "arguments": {...}}
<|end|>

User
What's the weather in Seoul today?

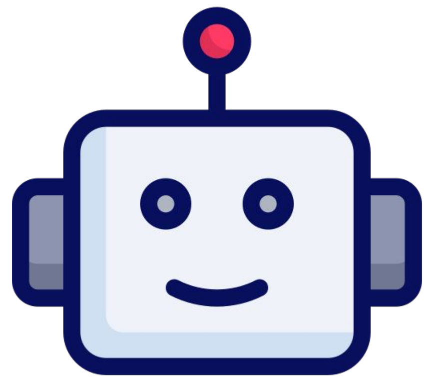

<|assistant|>
<|tool_call|>
{"name": "get_weather", "arguments": {"location": "Seoul"}}
<|end|>

<|tool_result|>
{"temperature": 27, "condition": "Sunny"}
<|end|>

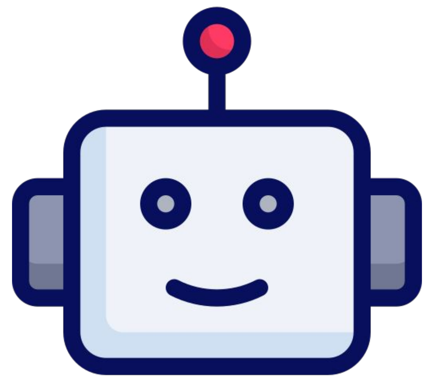

<|assistant|>
It's 27°C and sunny in Seoul today.
<|end|>

2nd Evolution: Extending Capabilities through Tool Use

LLM Reasoning and Agentic AI

Technical Foundations of AI Agents: Reasoning & External Interaction

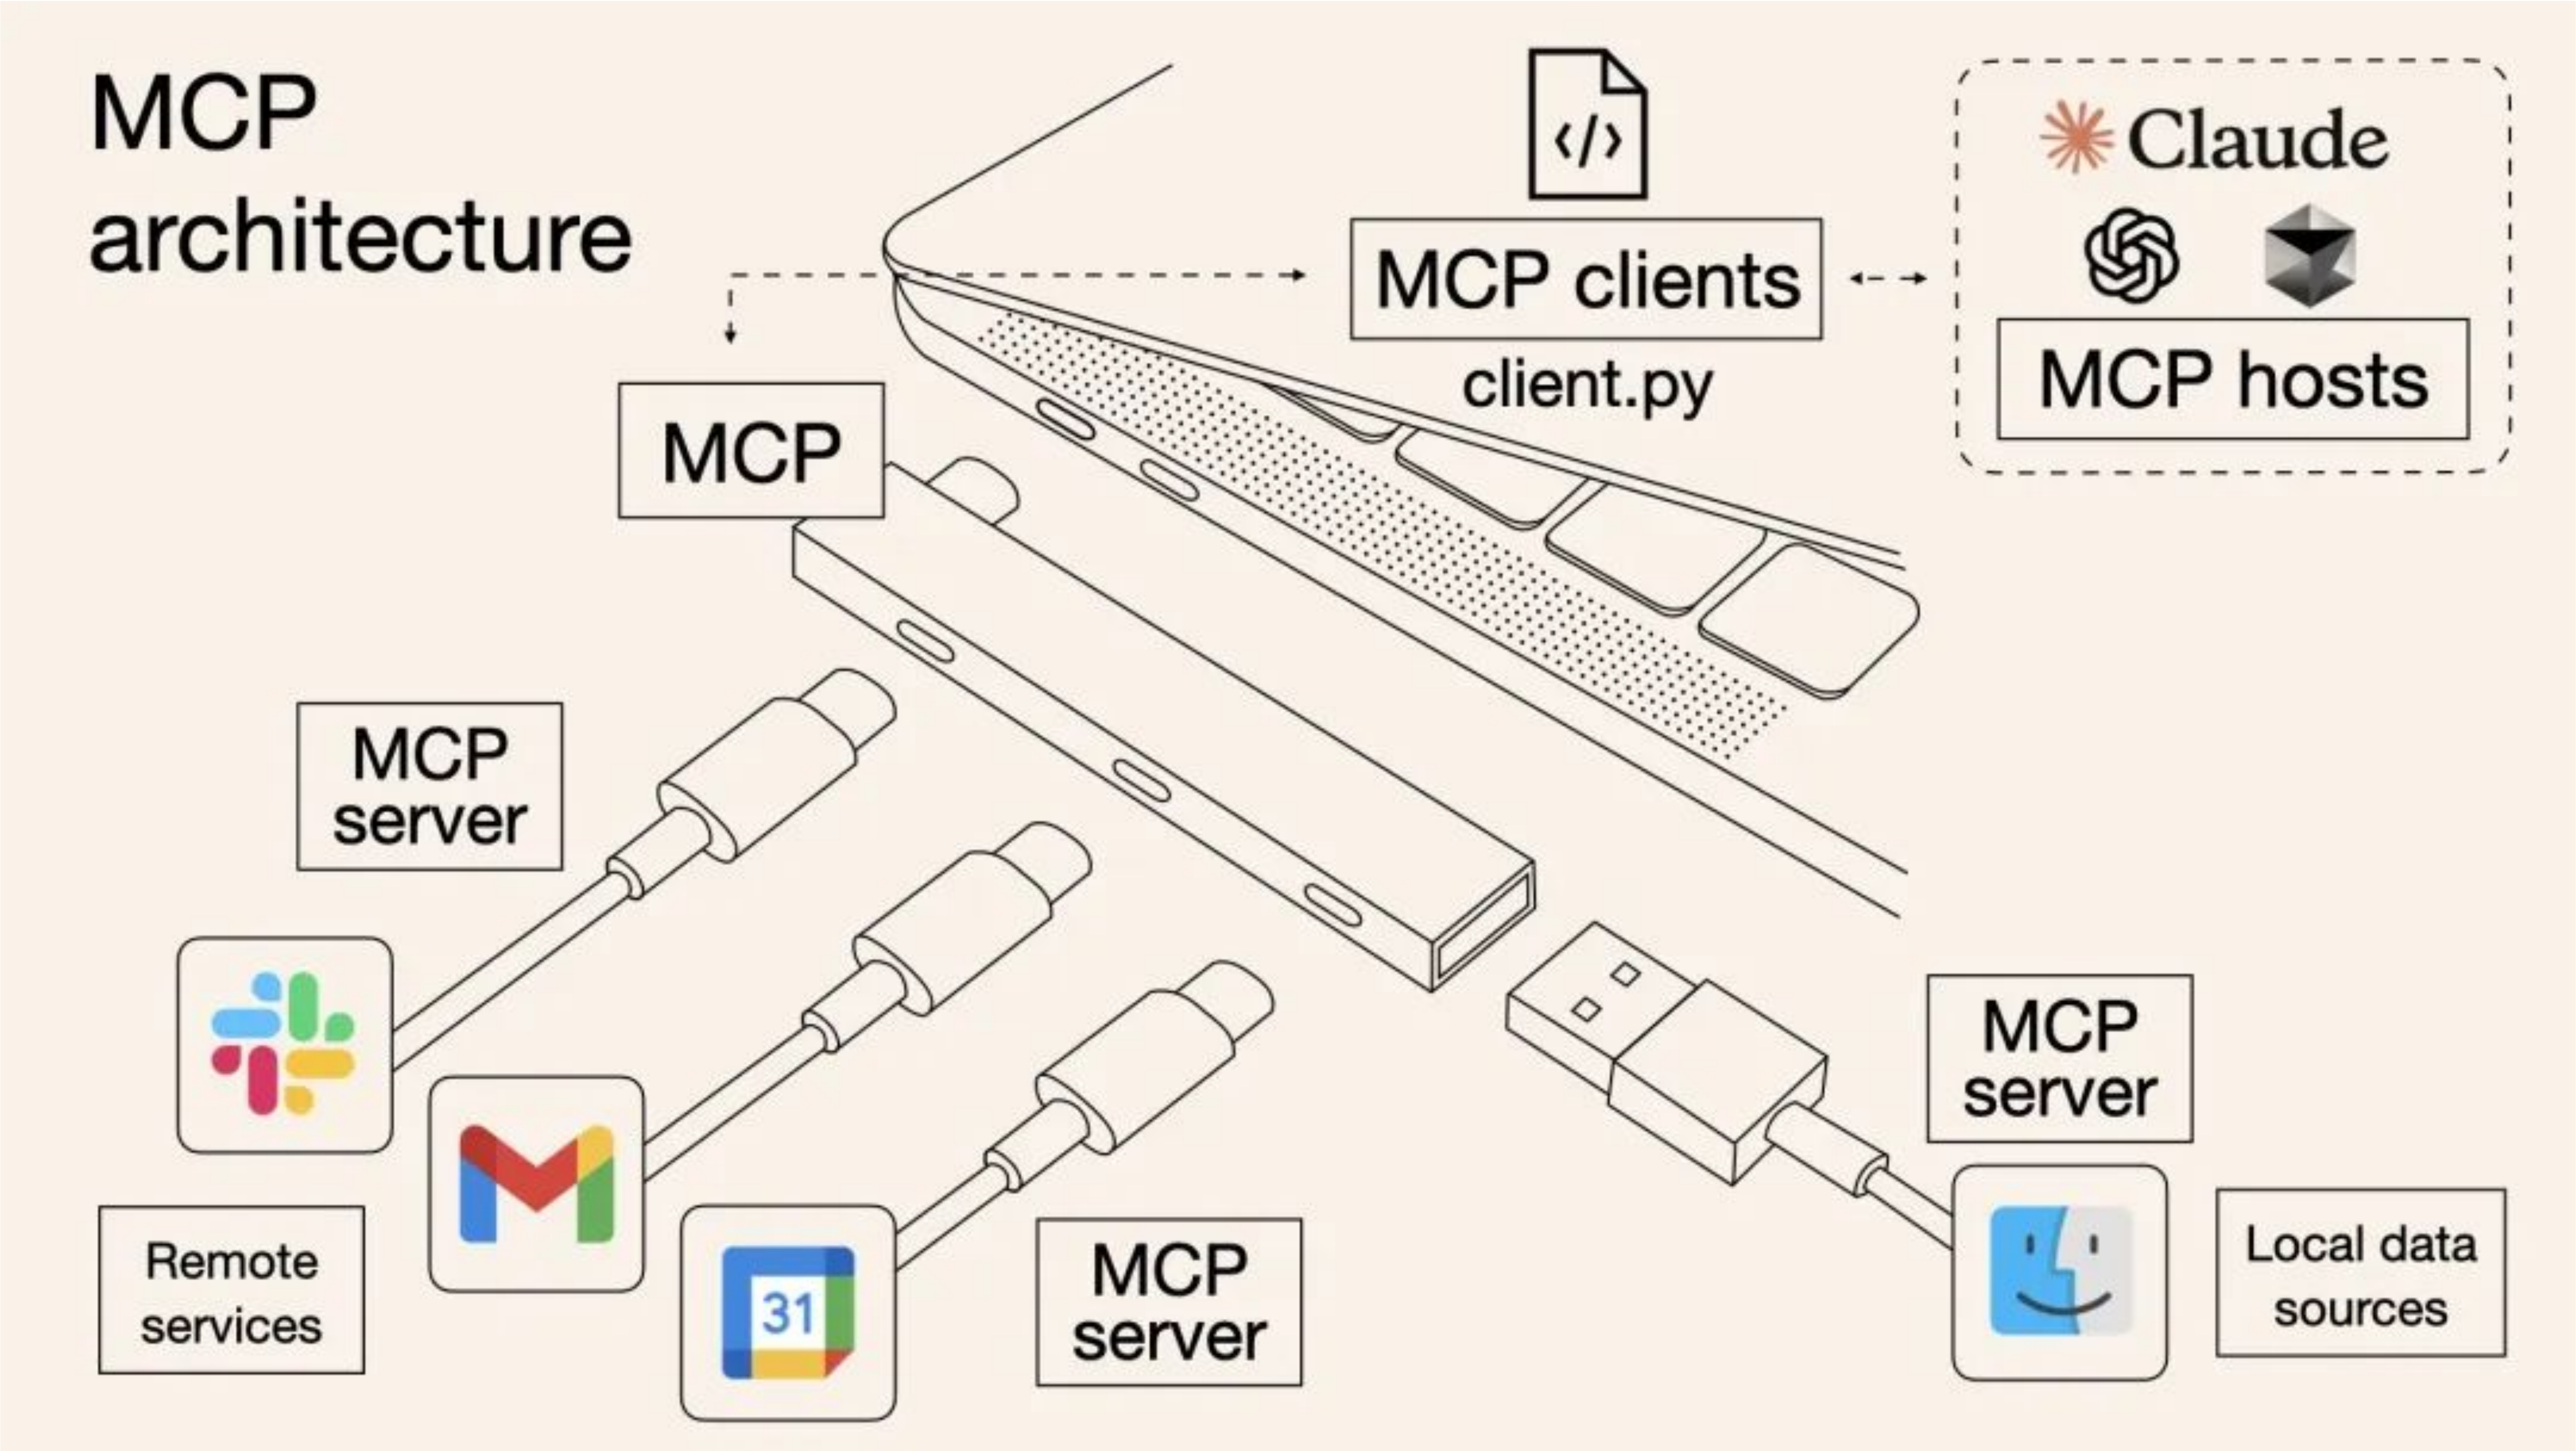

Building RAG Agents with LLMs

Part 3. LLM Training

LLM Training

Why Transformer is Powerful Architecture?

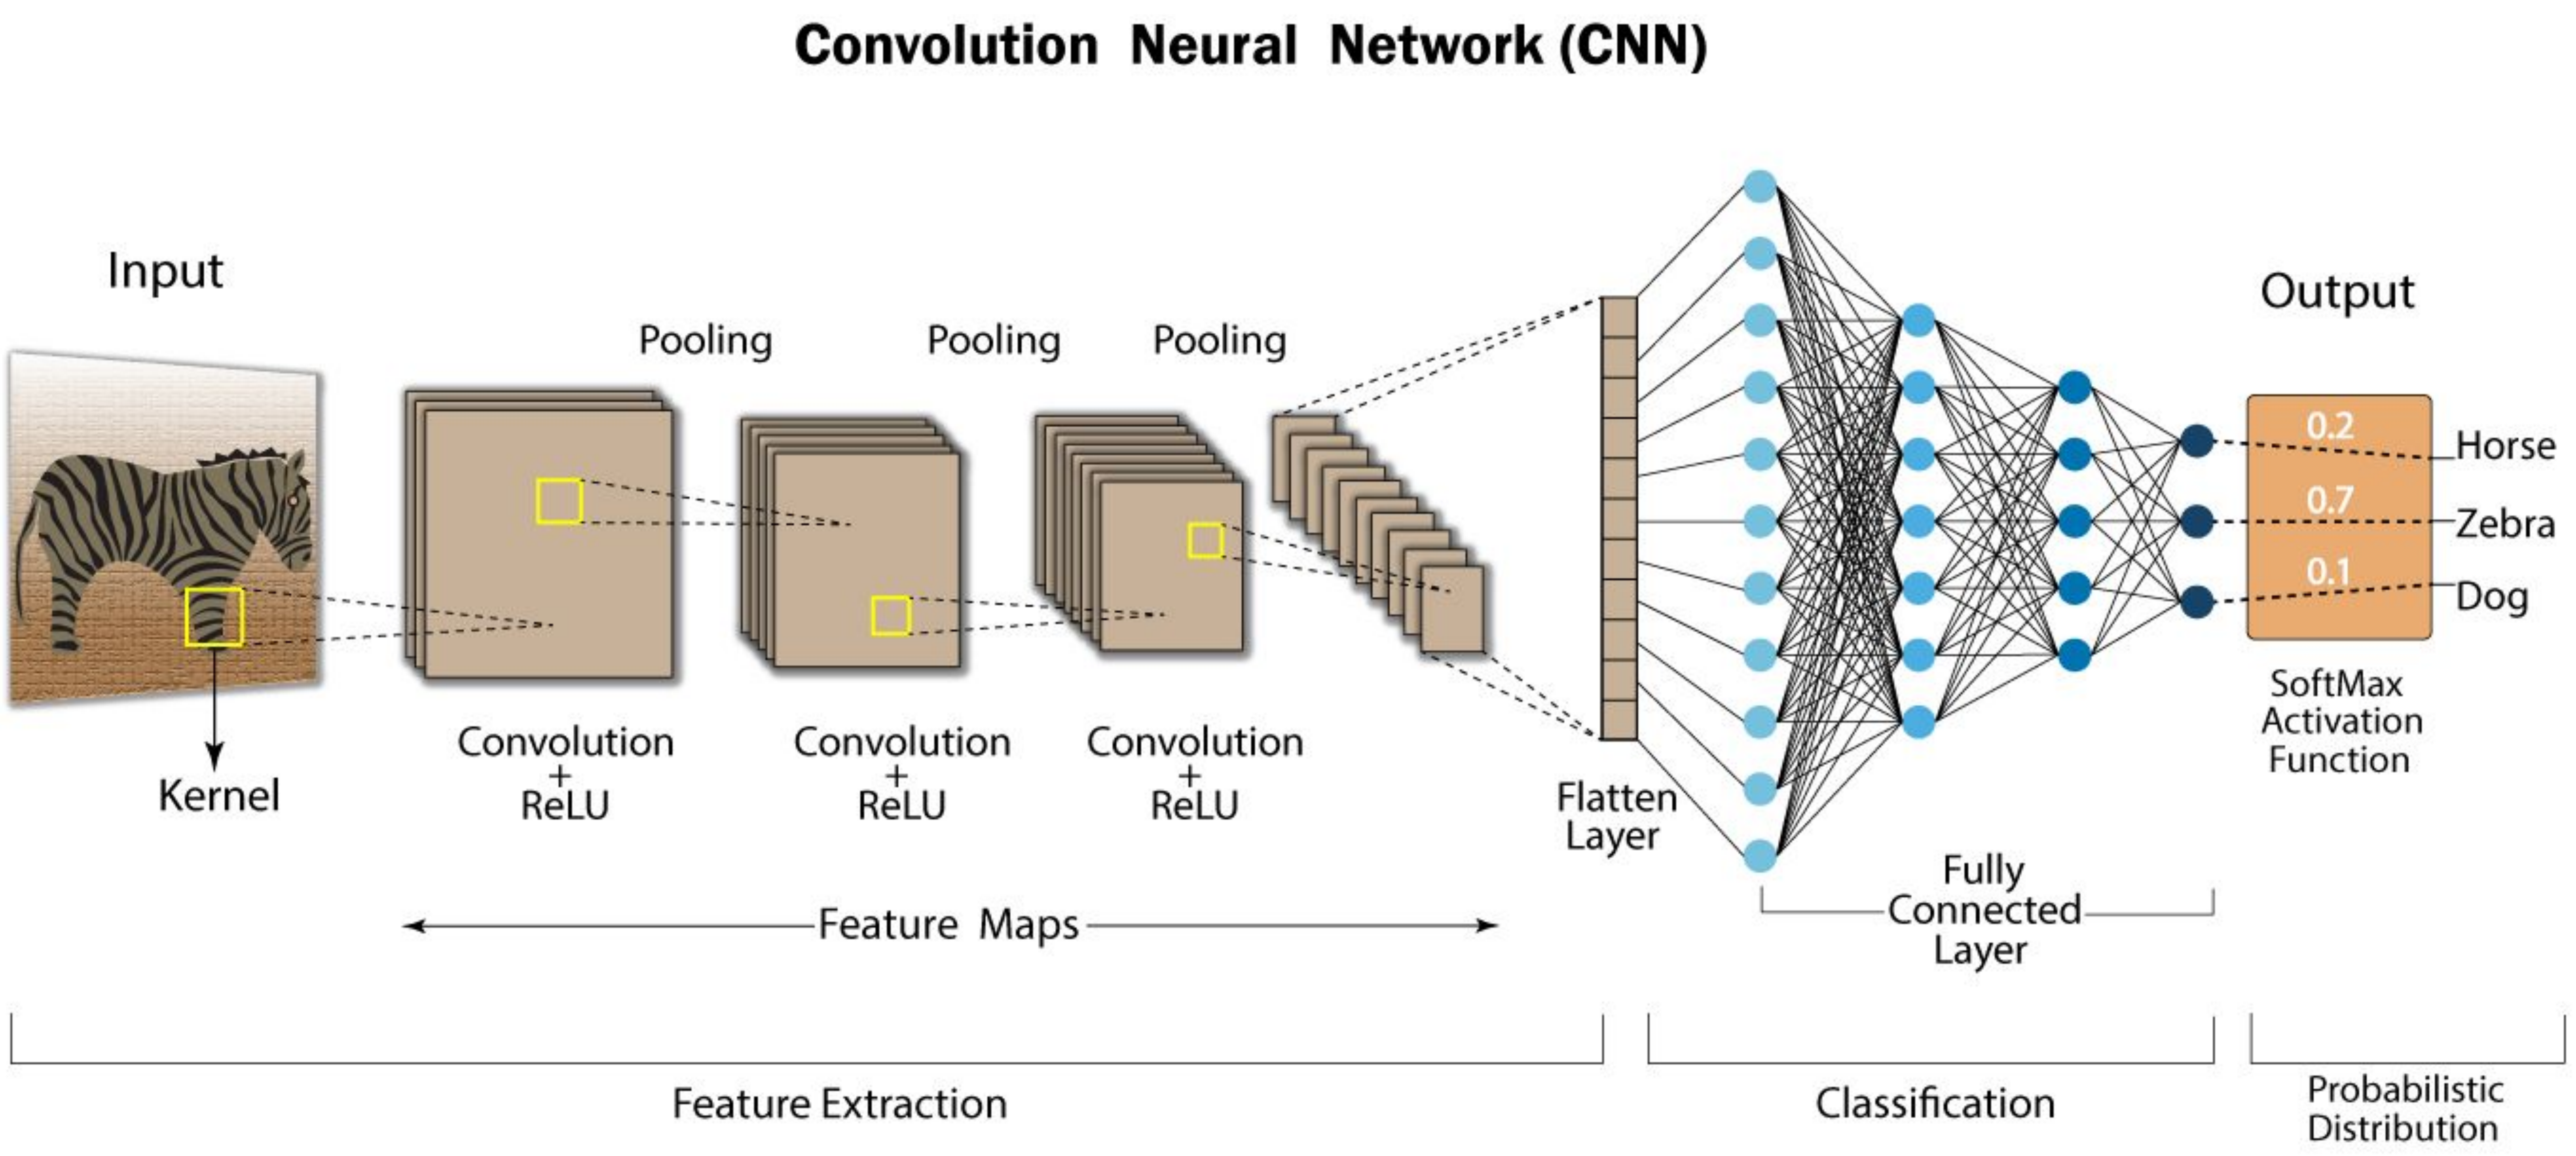

Inductive Bias : Nearby pixels are more likely to be related.
→ Image Prediction Task

LLM Training

Why Transformer is Powerful Architecture?

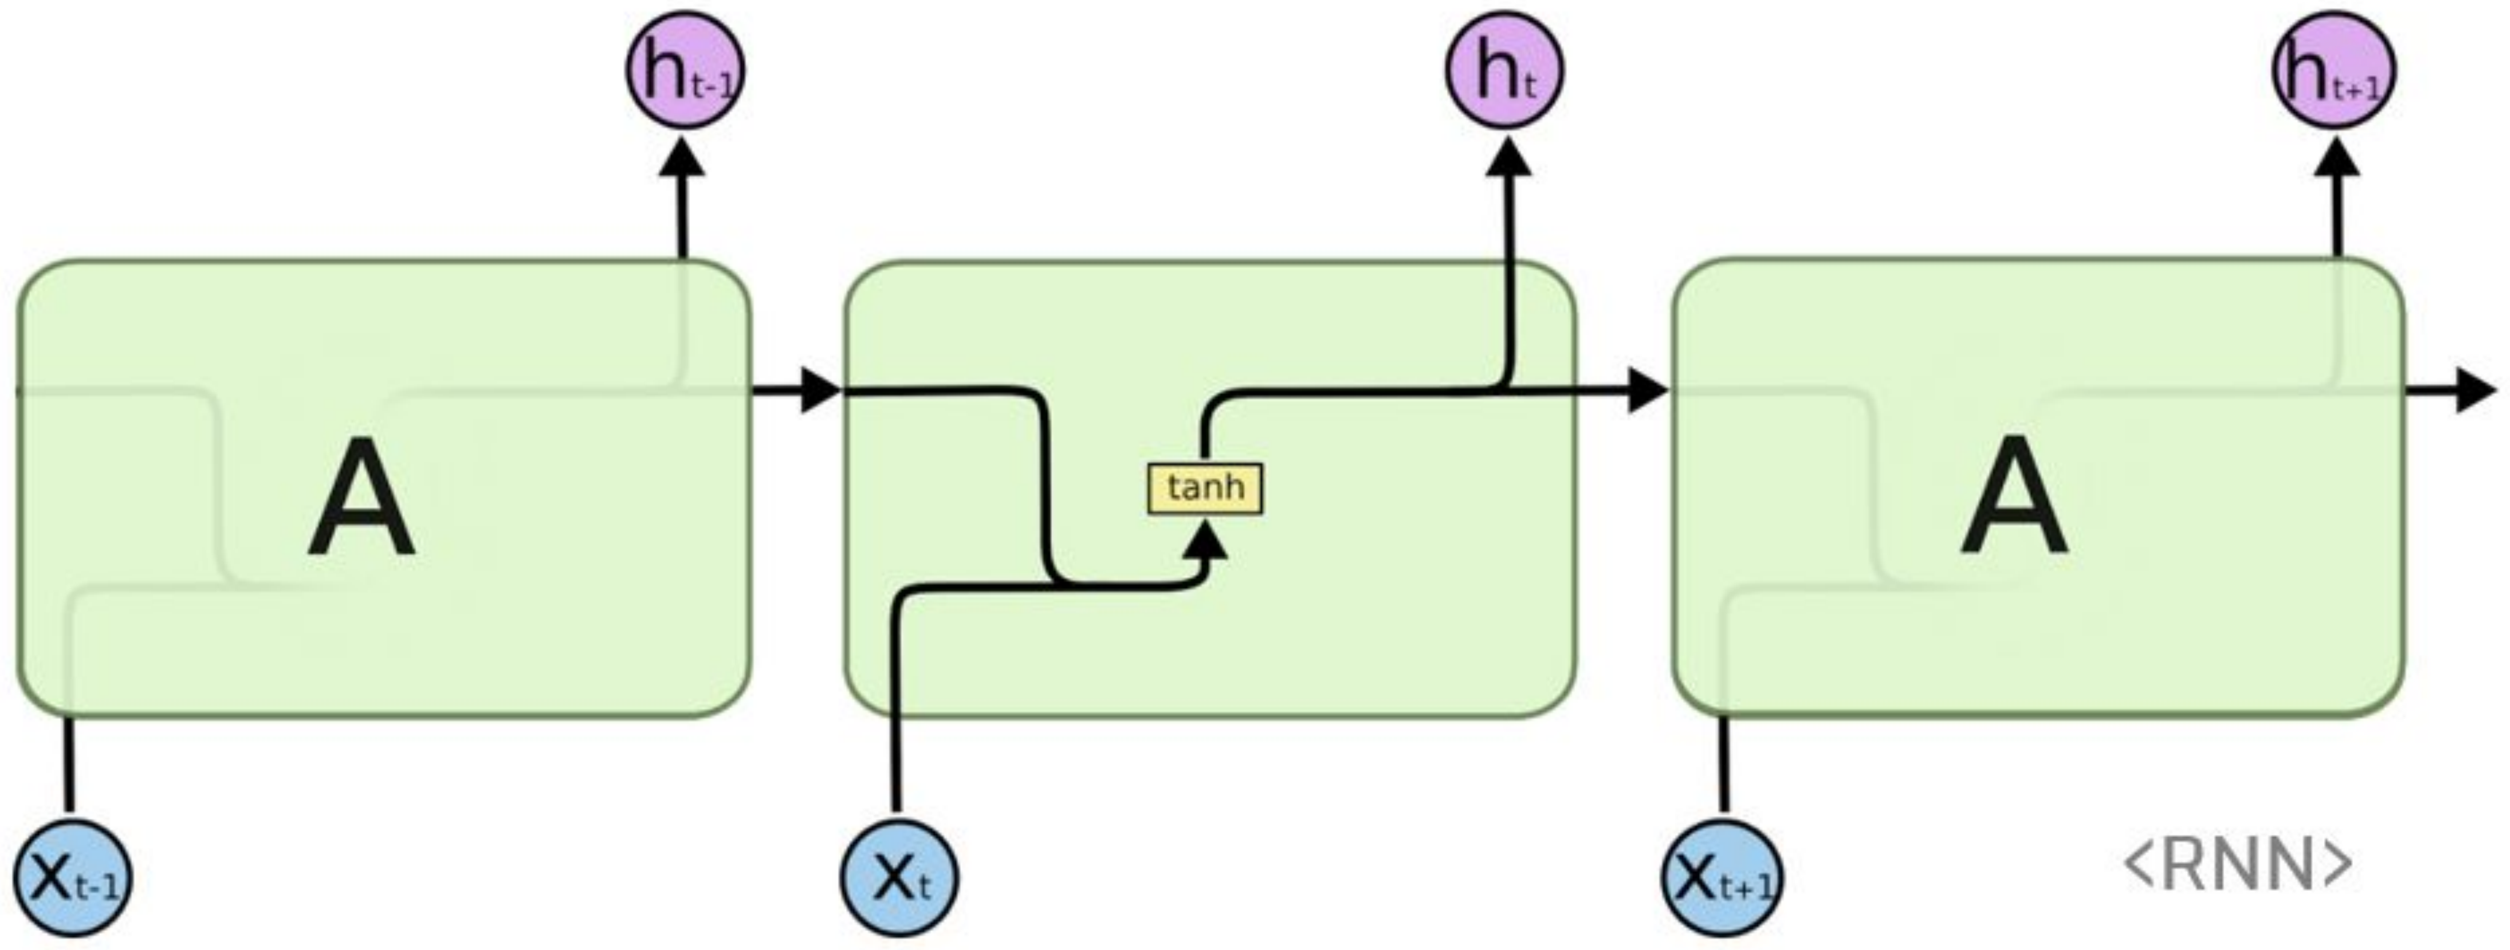

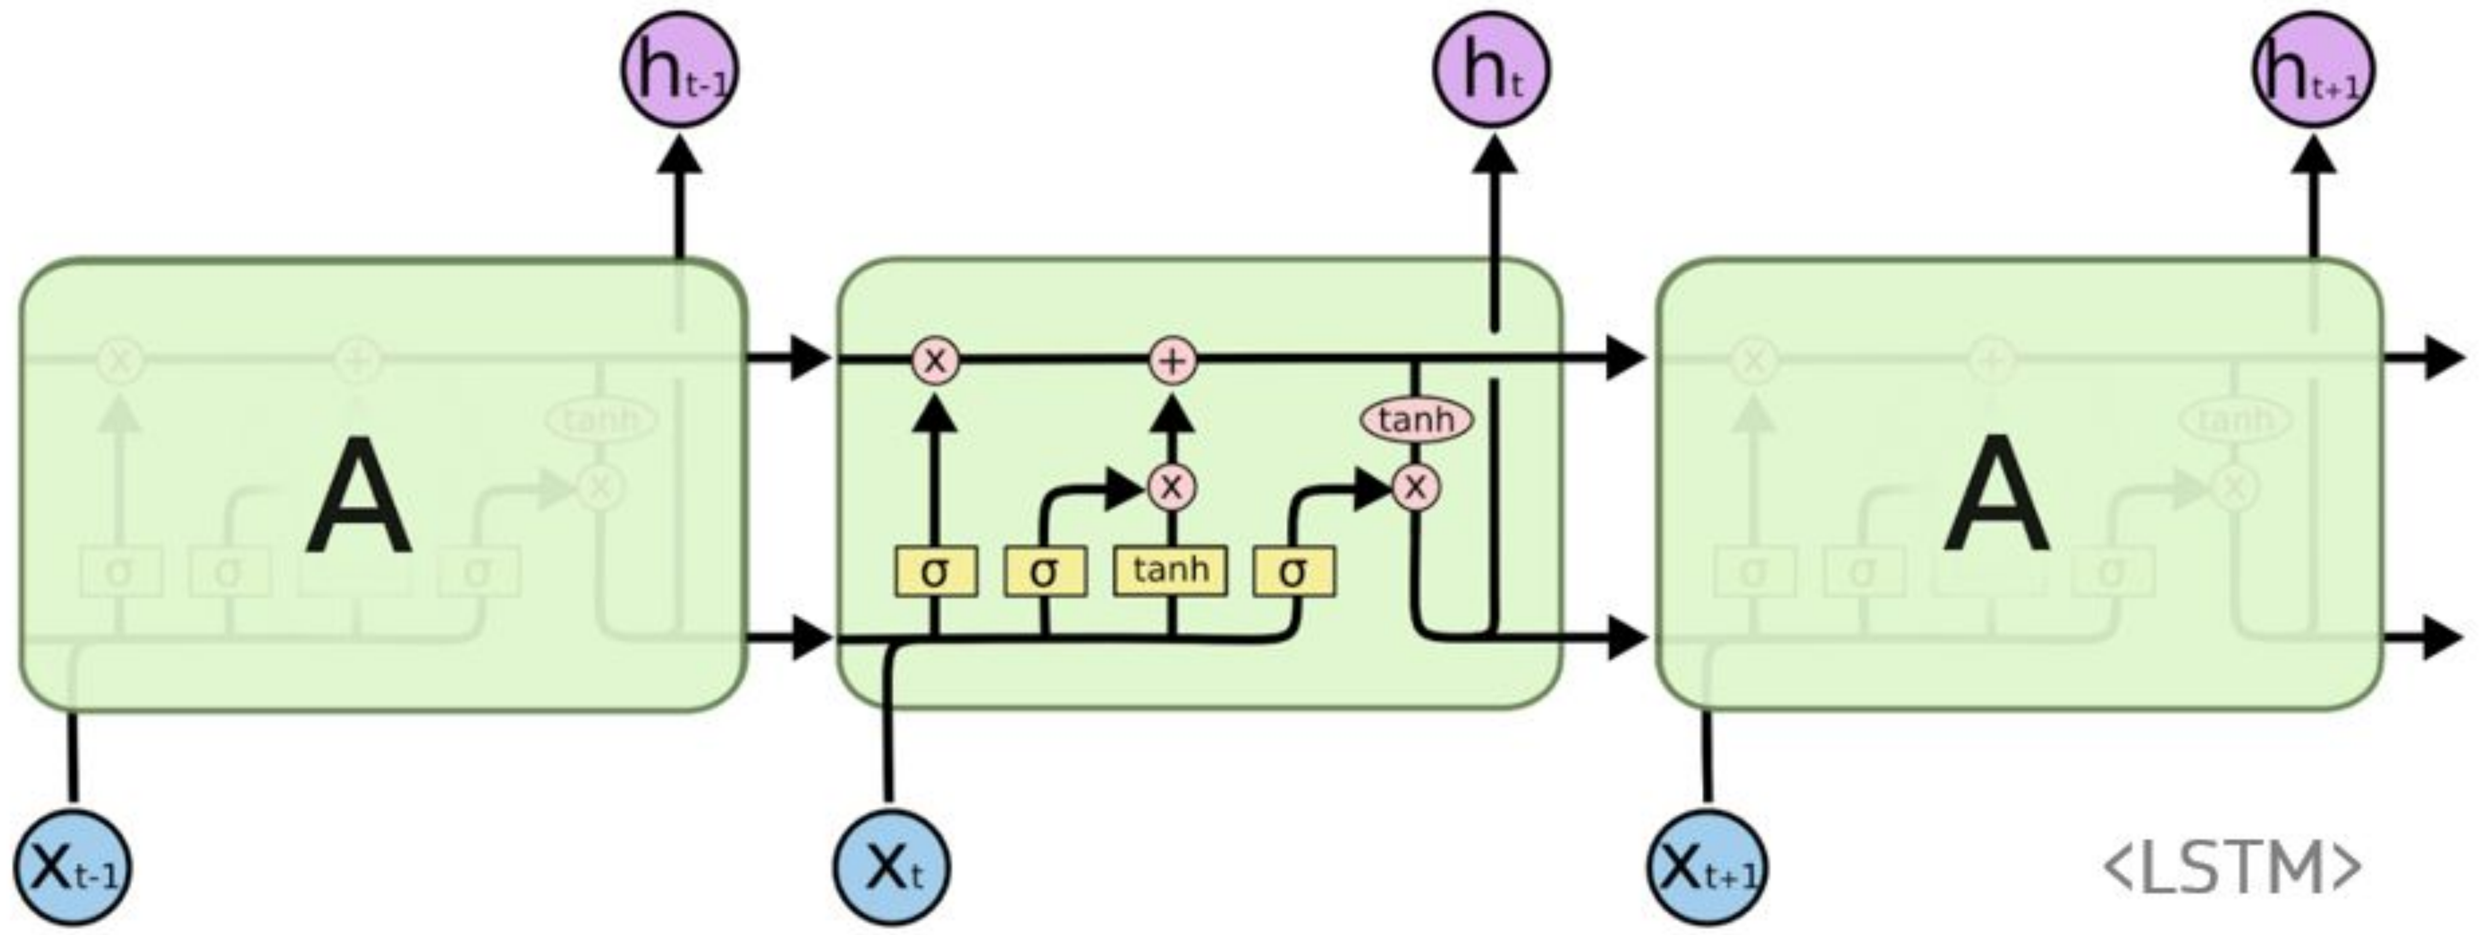

Inductive Bias : The current state depends on previous states.
→ Sequence Prediction Task

LLM Training

Why Transformer is Powerful Architecture?

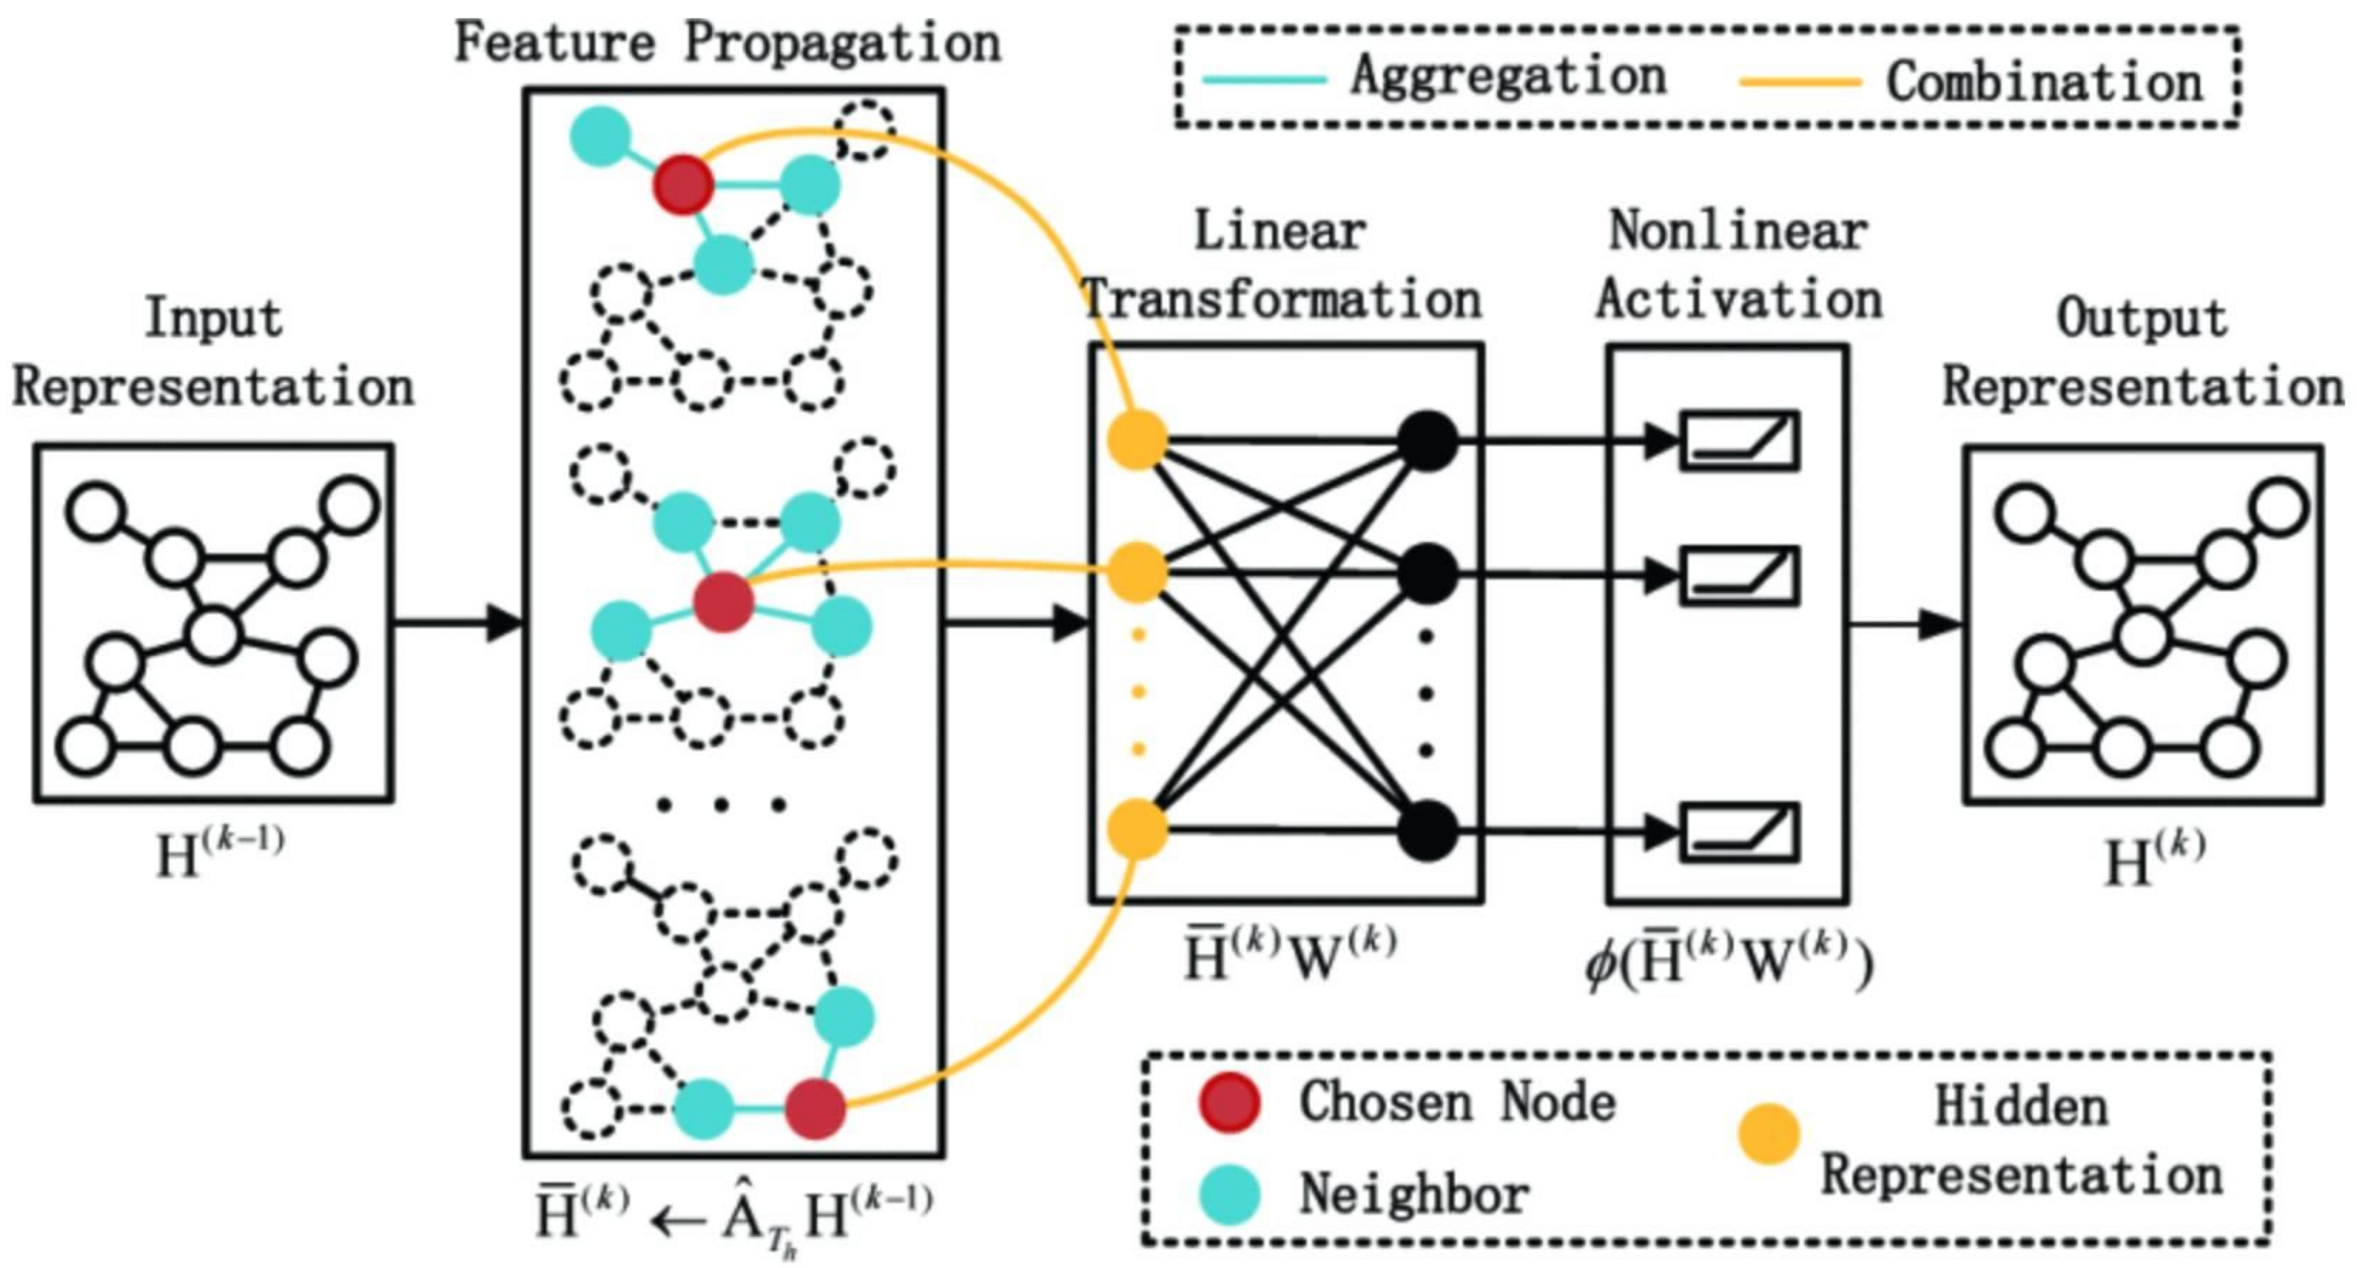

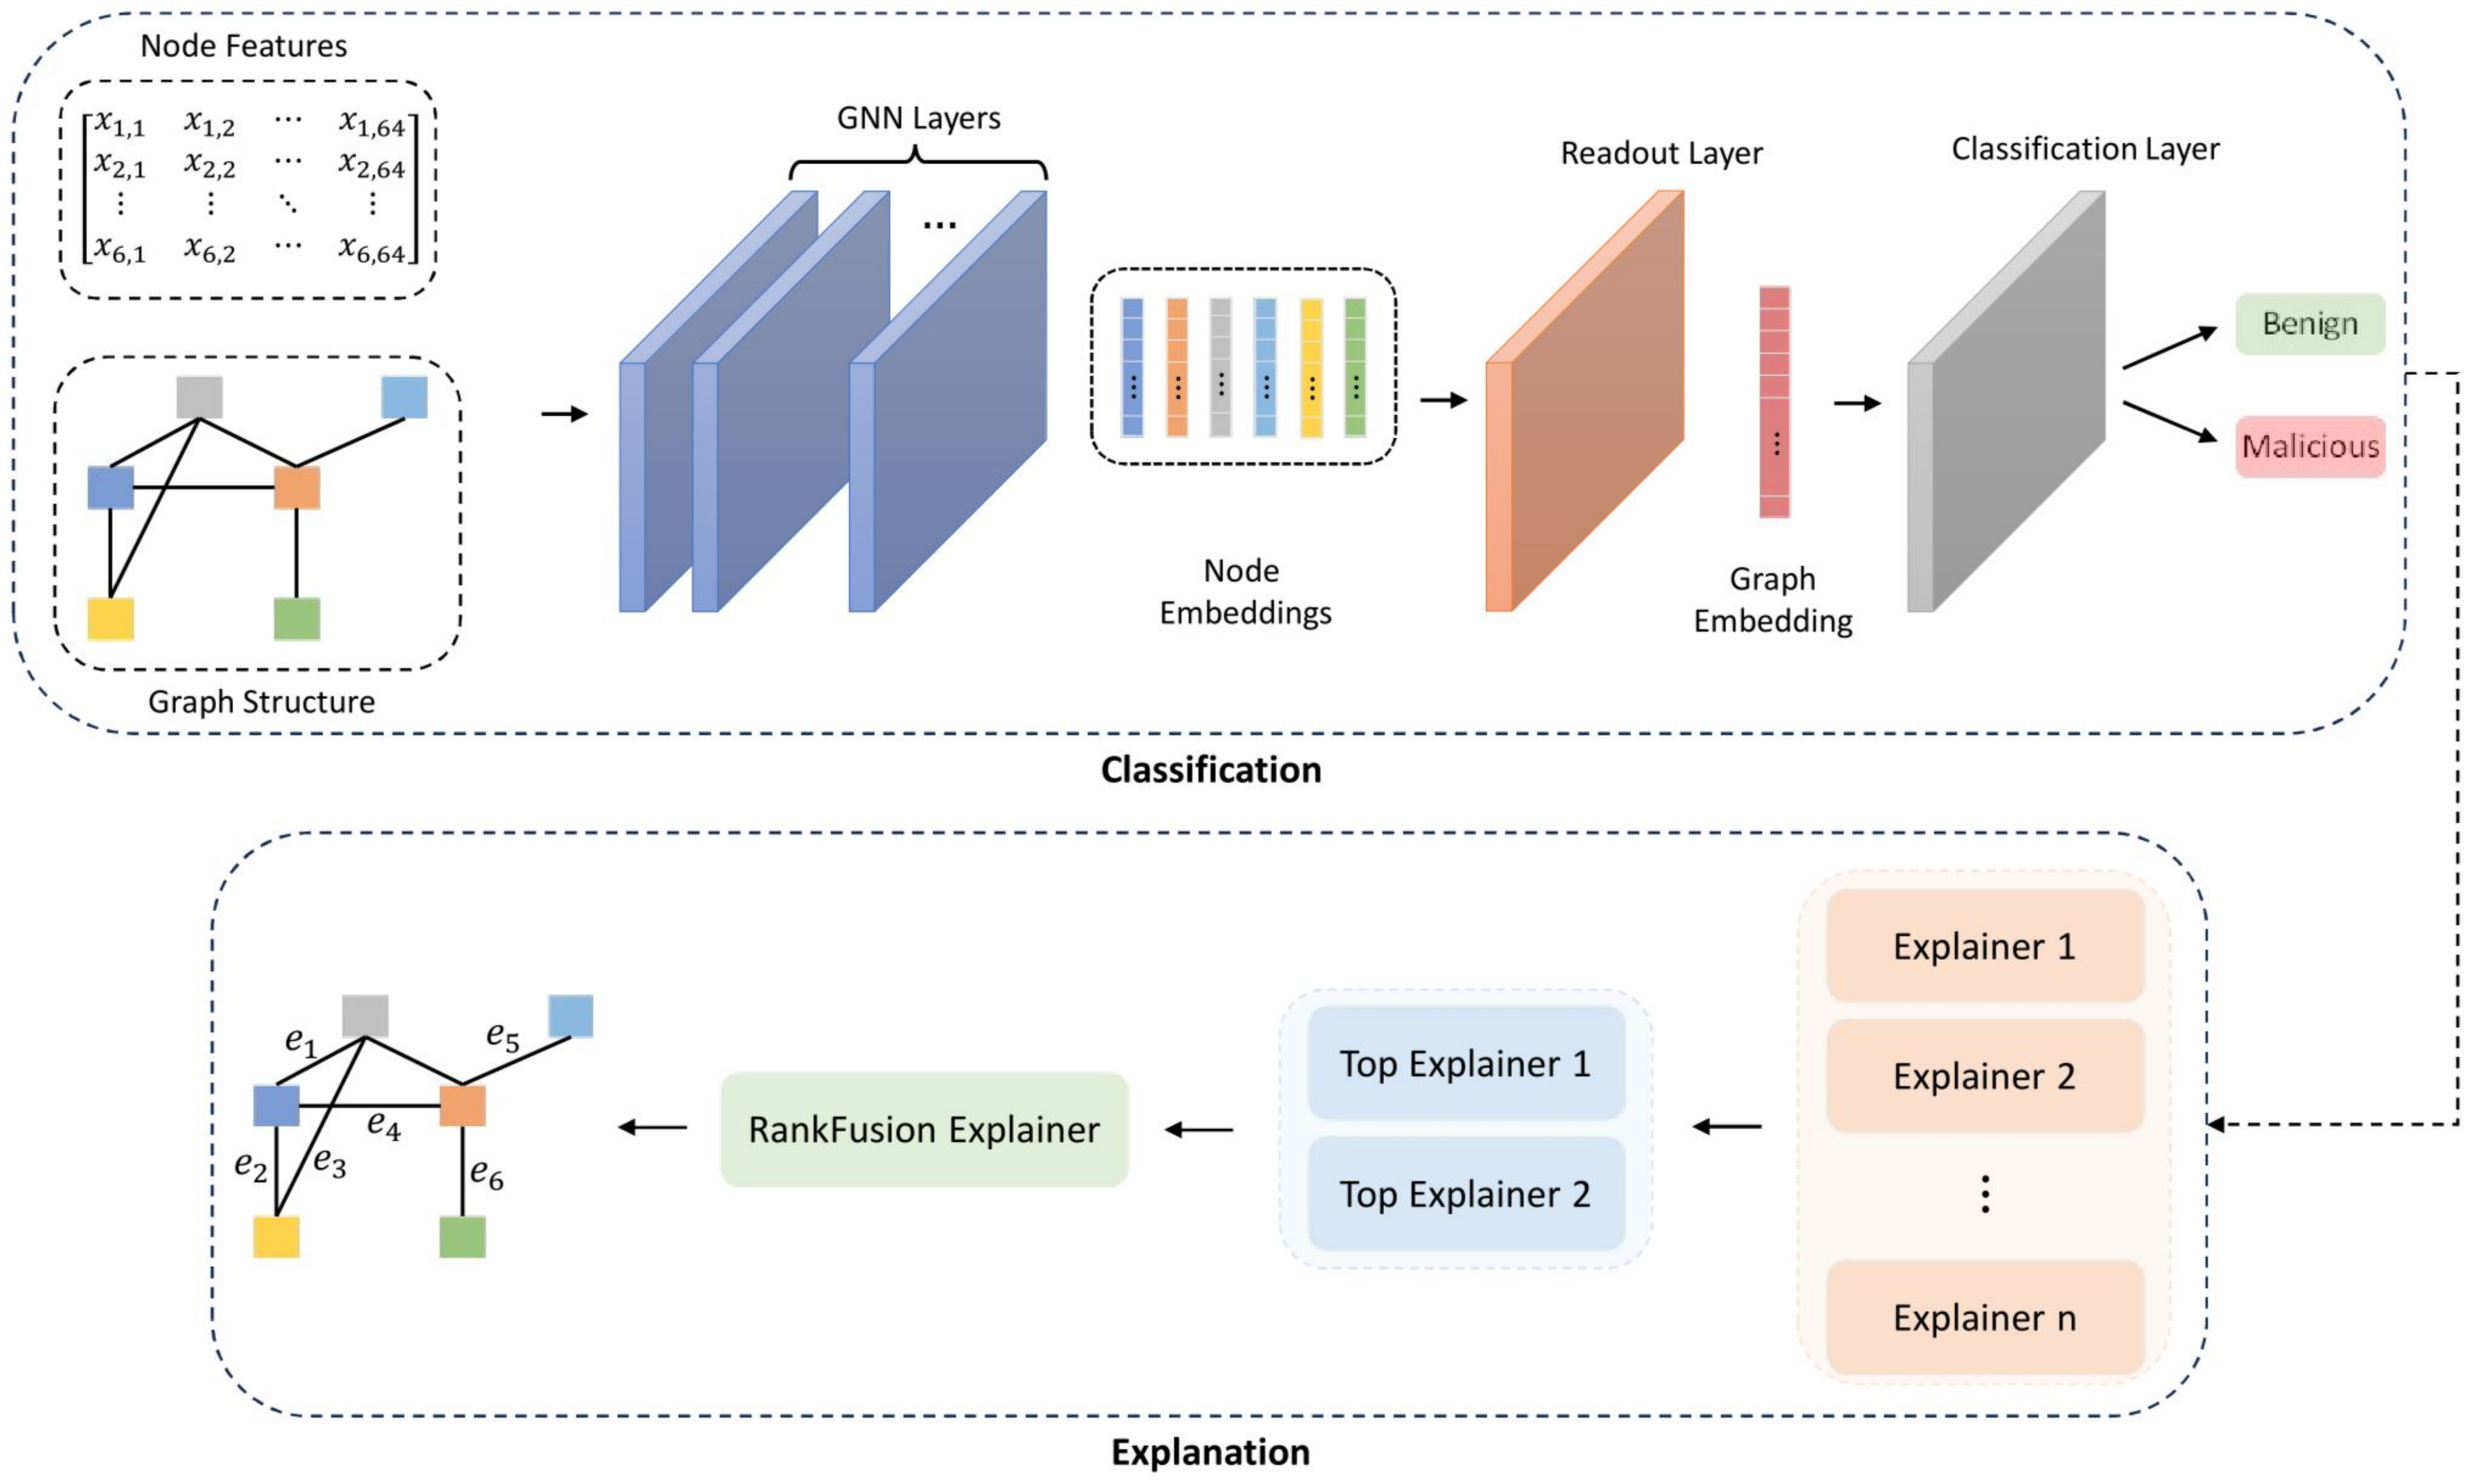

Inductive Bias : Connected nodes are more likely to interact.
→ Graph Prediction Task

LLM Training

Why Transformer is Powerful Architecture?

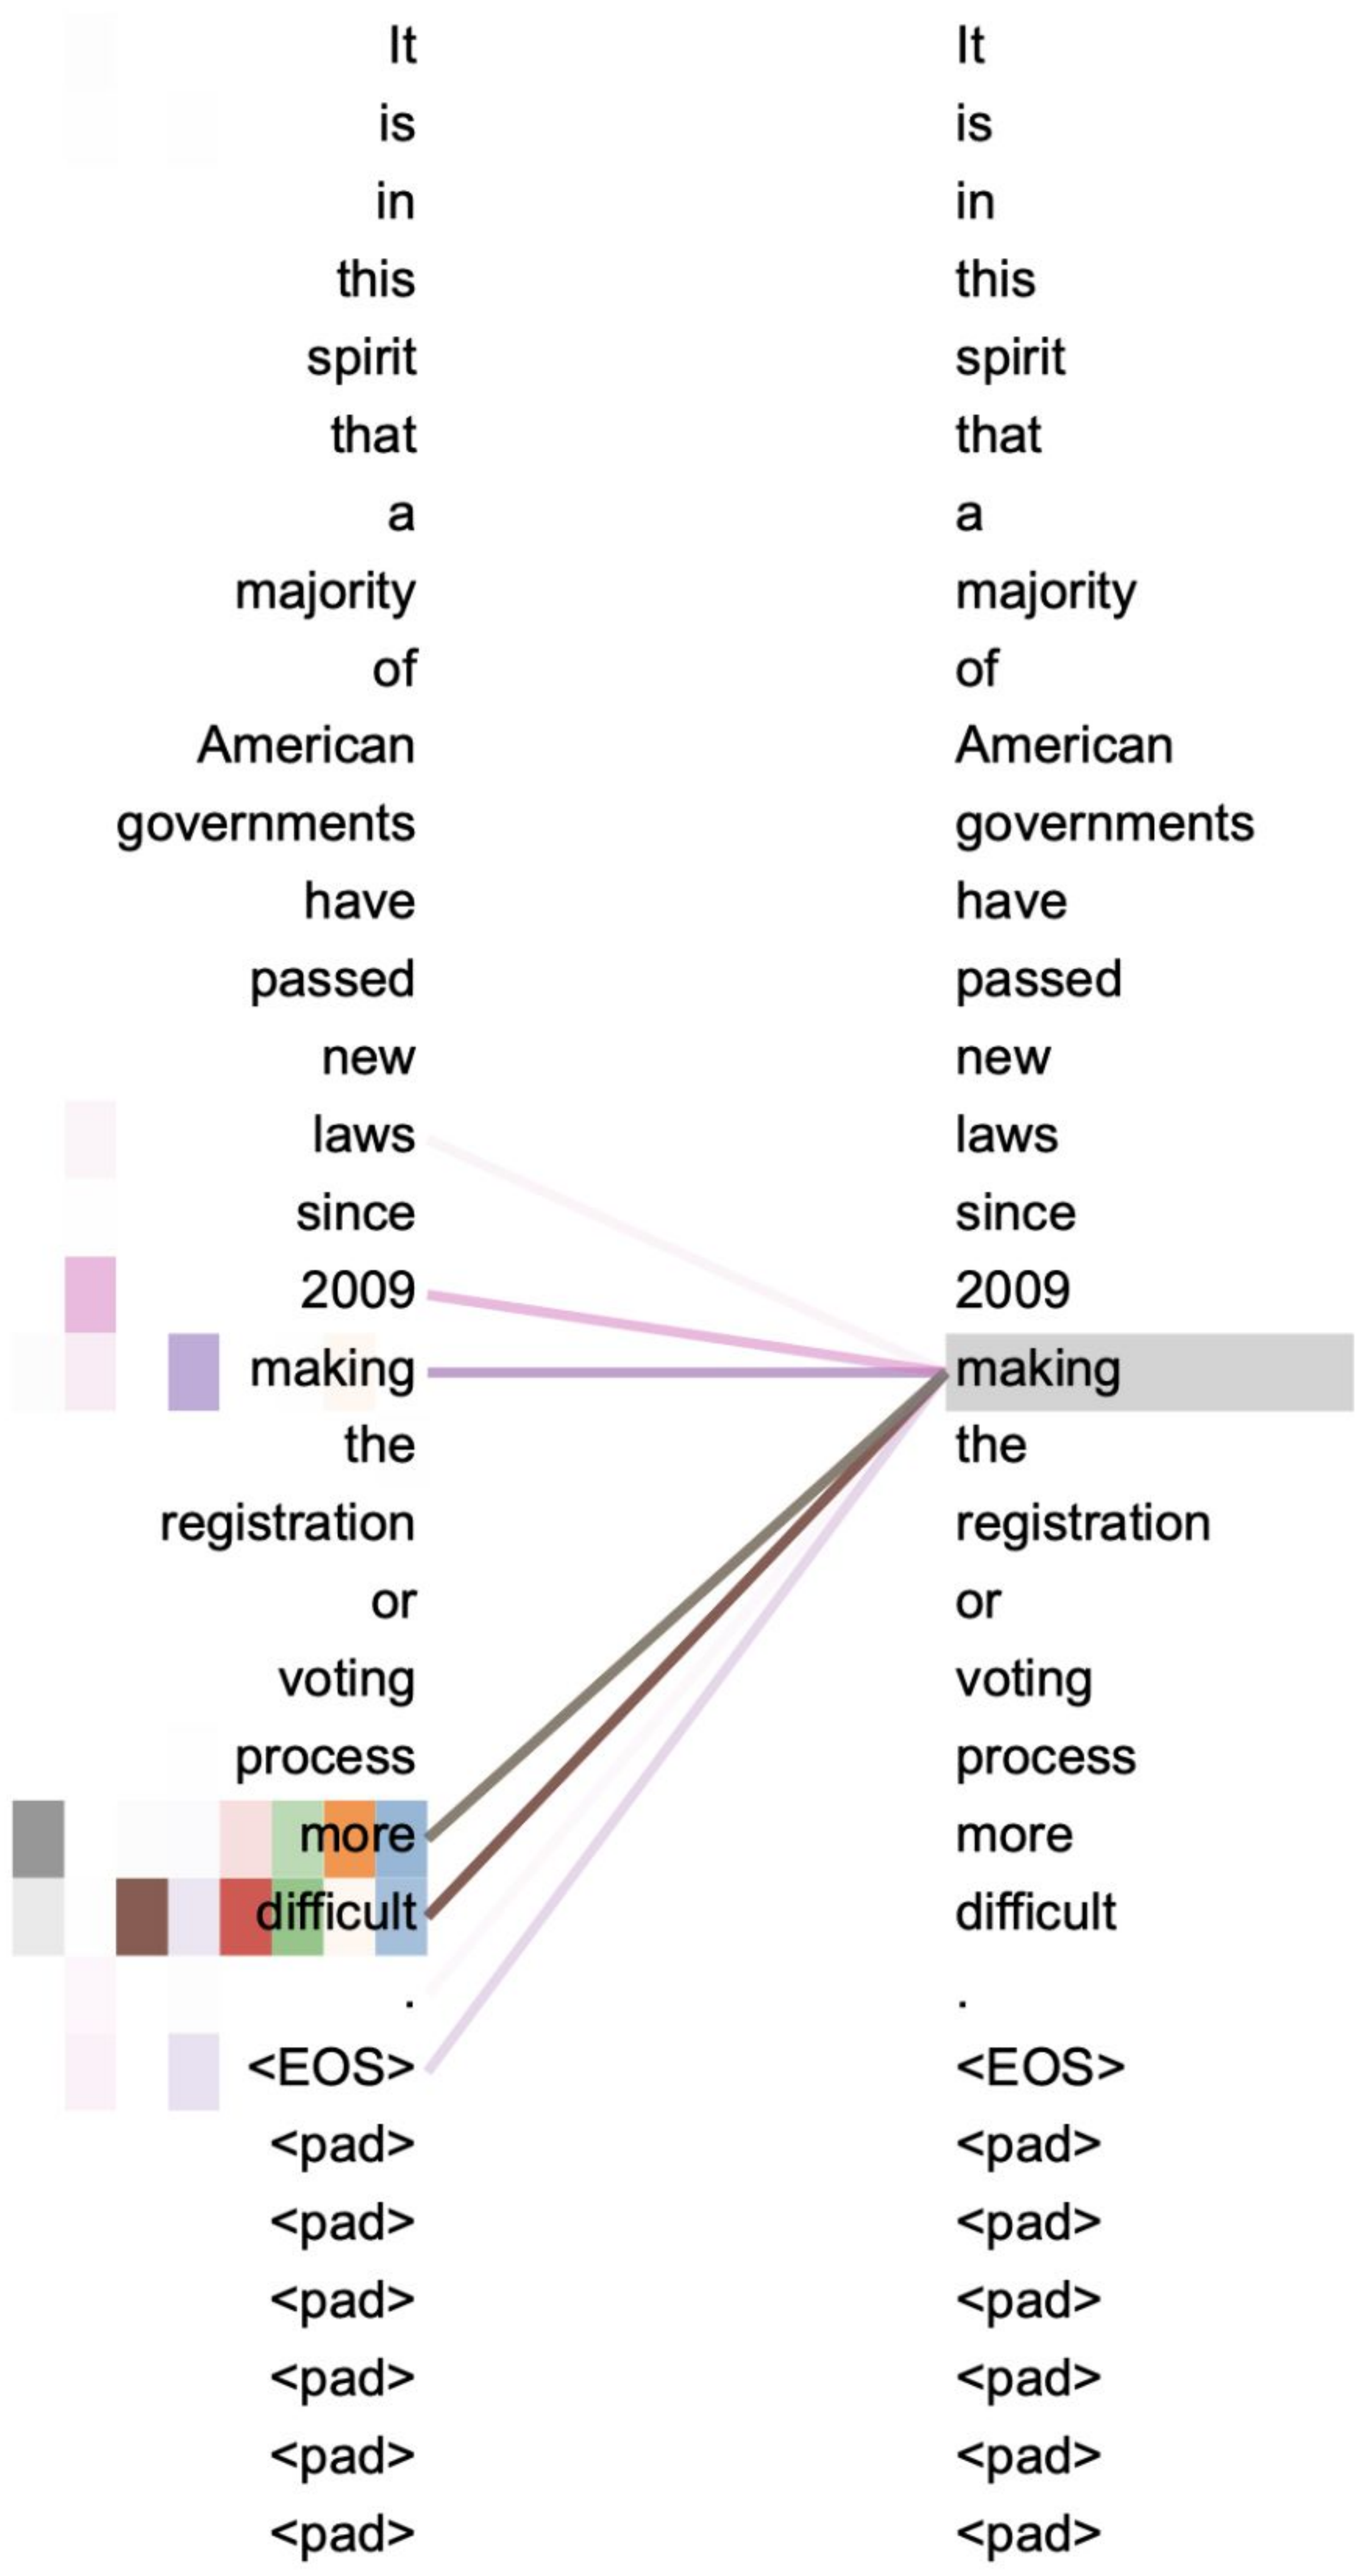

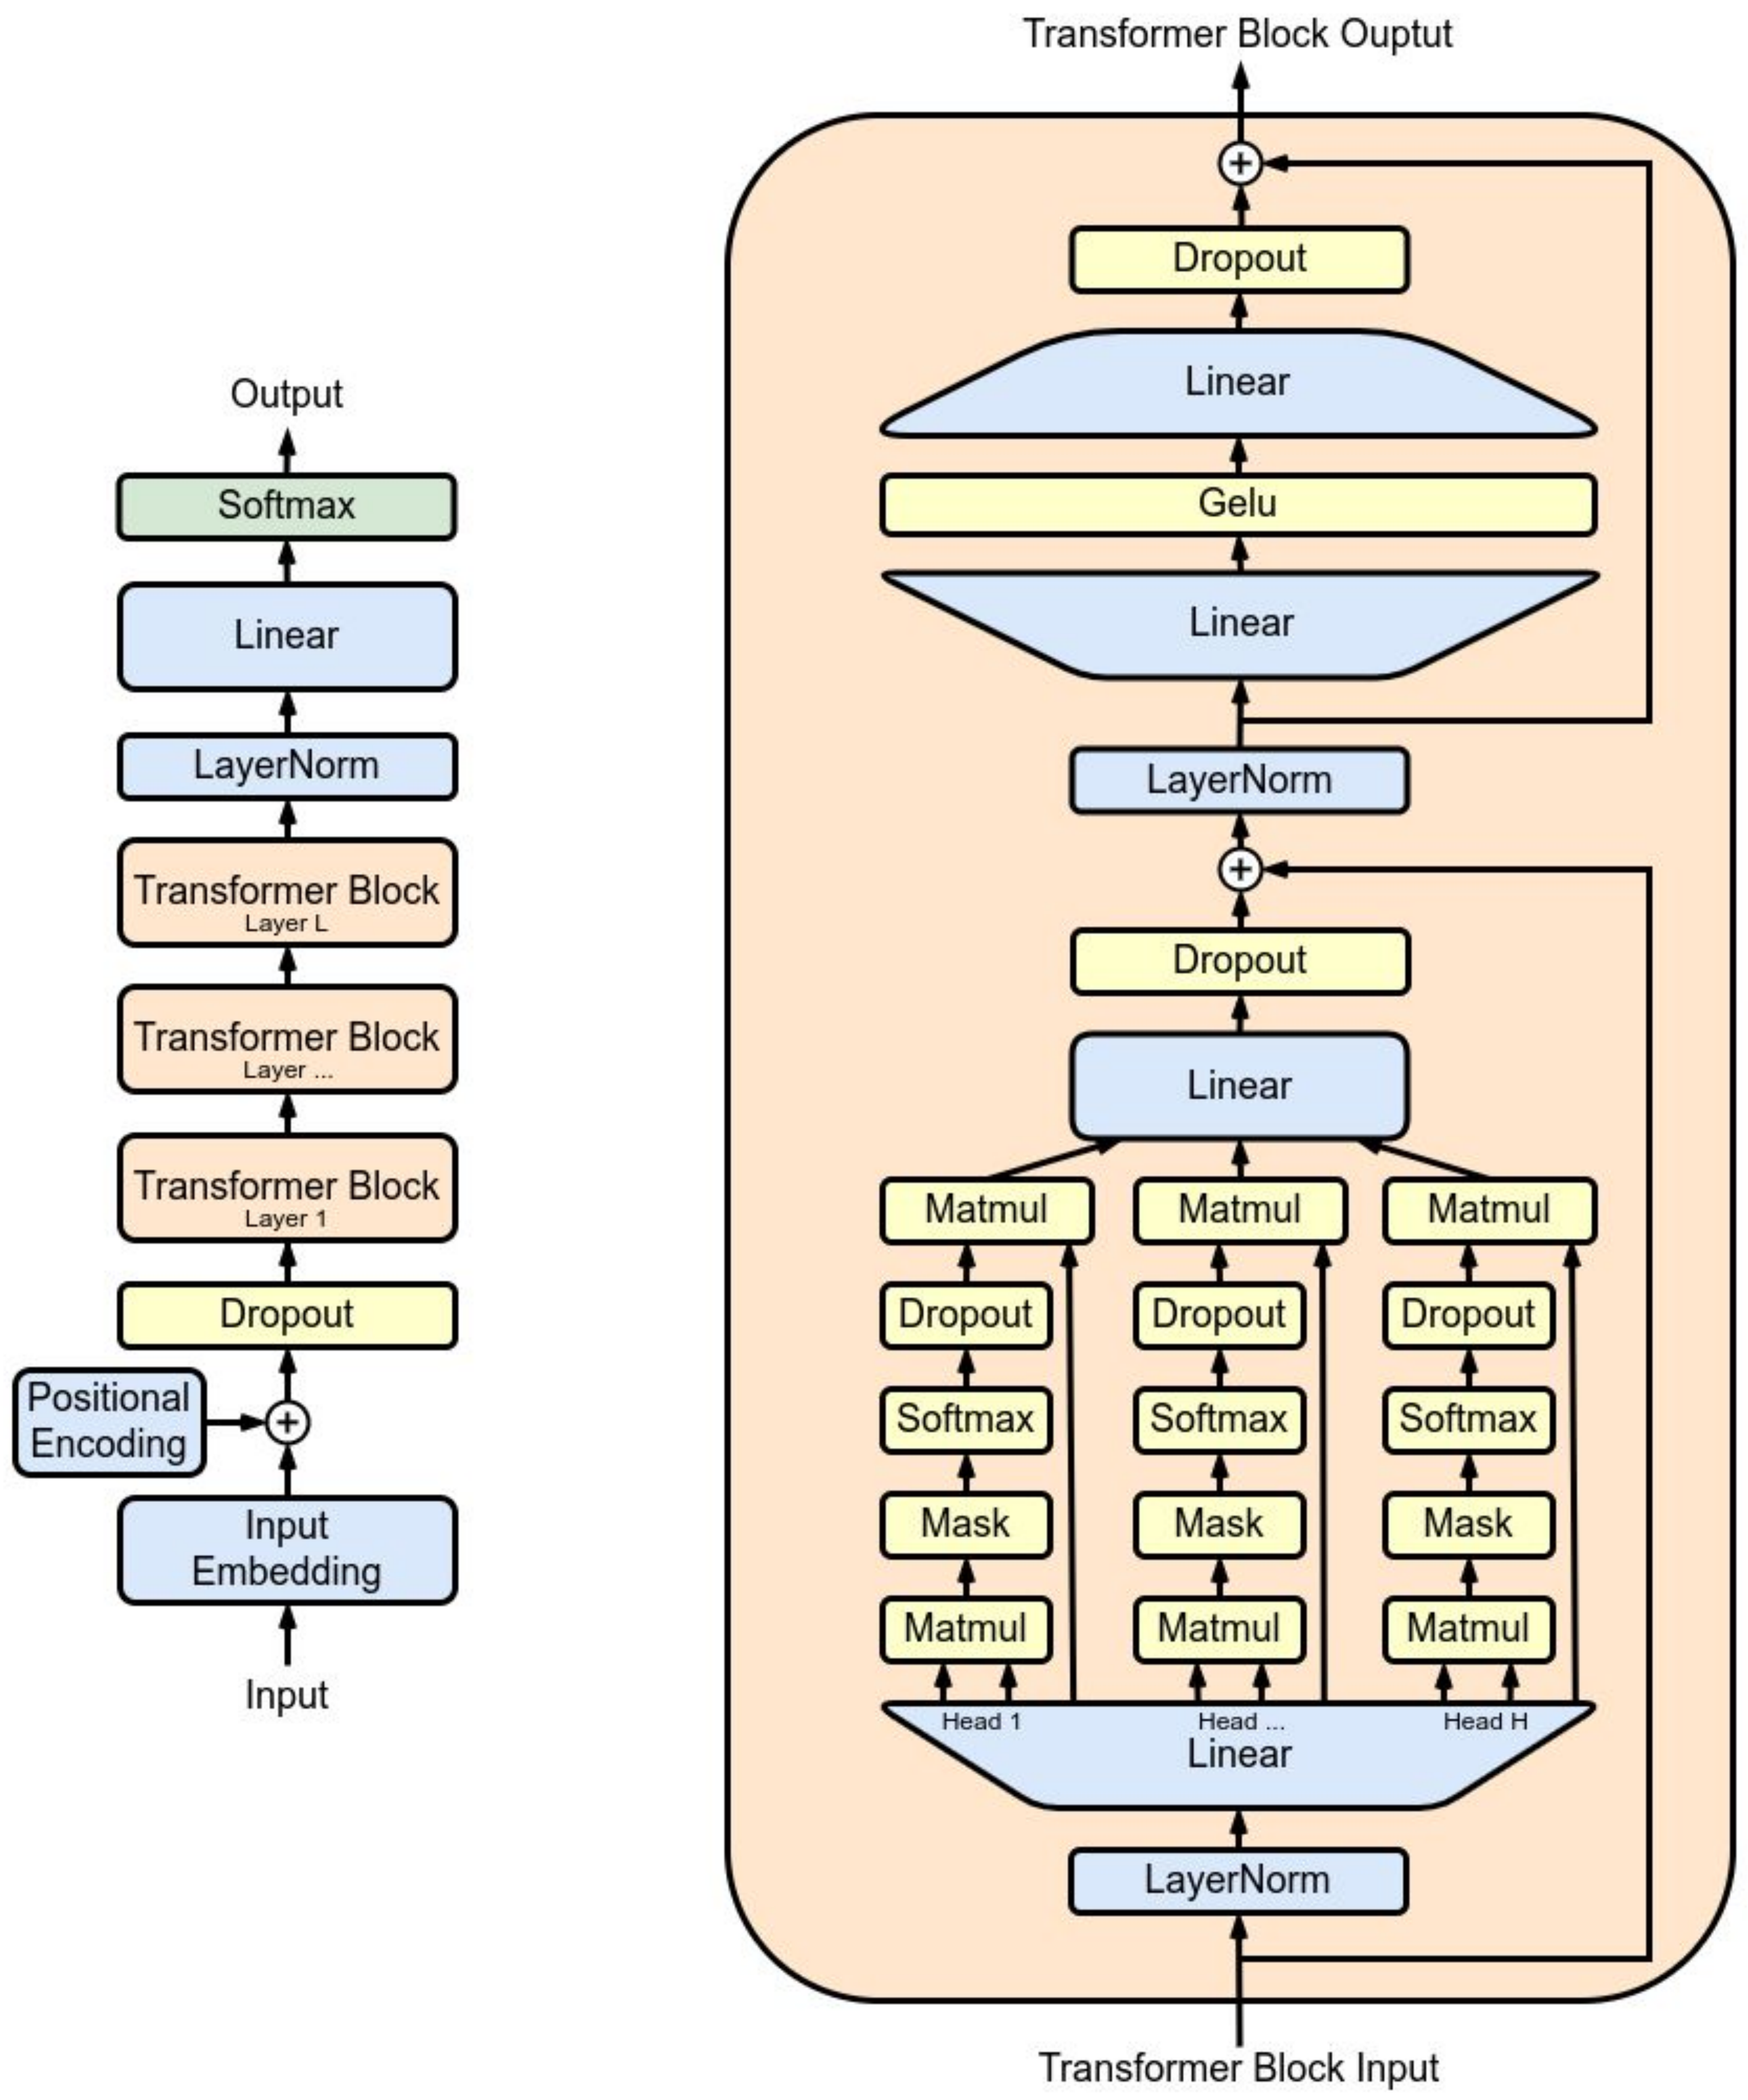

Any token can interact with any other token

CNN → Locality
RNN → Sequentiality
GNN → Connectivity
Transformer
→ Global Token Relationships

LLM Training

Massive Computational and Data Resources for Model Training

LLM Training

Massive Computational and Data Resources for Model Training

Inside an AI start-up's plan to scan and dispose of millions of books

Court filings reveal how AI companies raced to obtain more books to feed chatbots, including by buying, scanning and disposing of millions of titles.

January 27, 2026

Massive Data
to Model Training

LLM Training

Massive Computational and Data Resources for Model Training

LLM Training

How are LLM models trained?

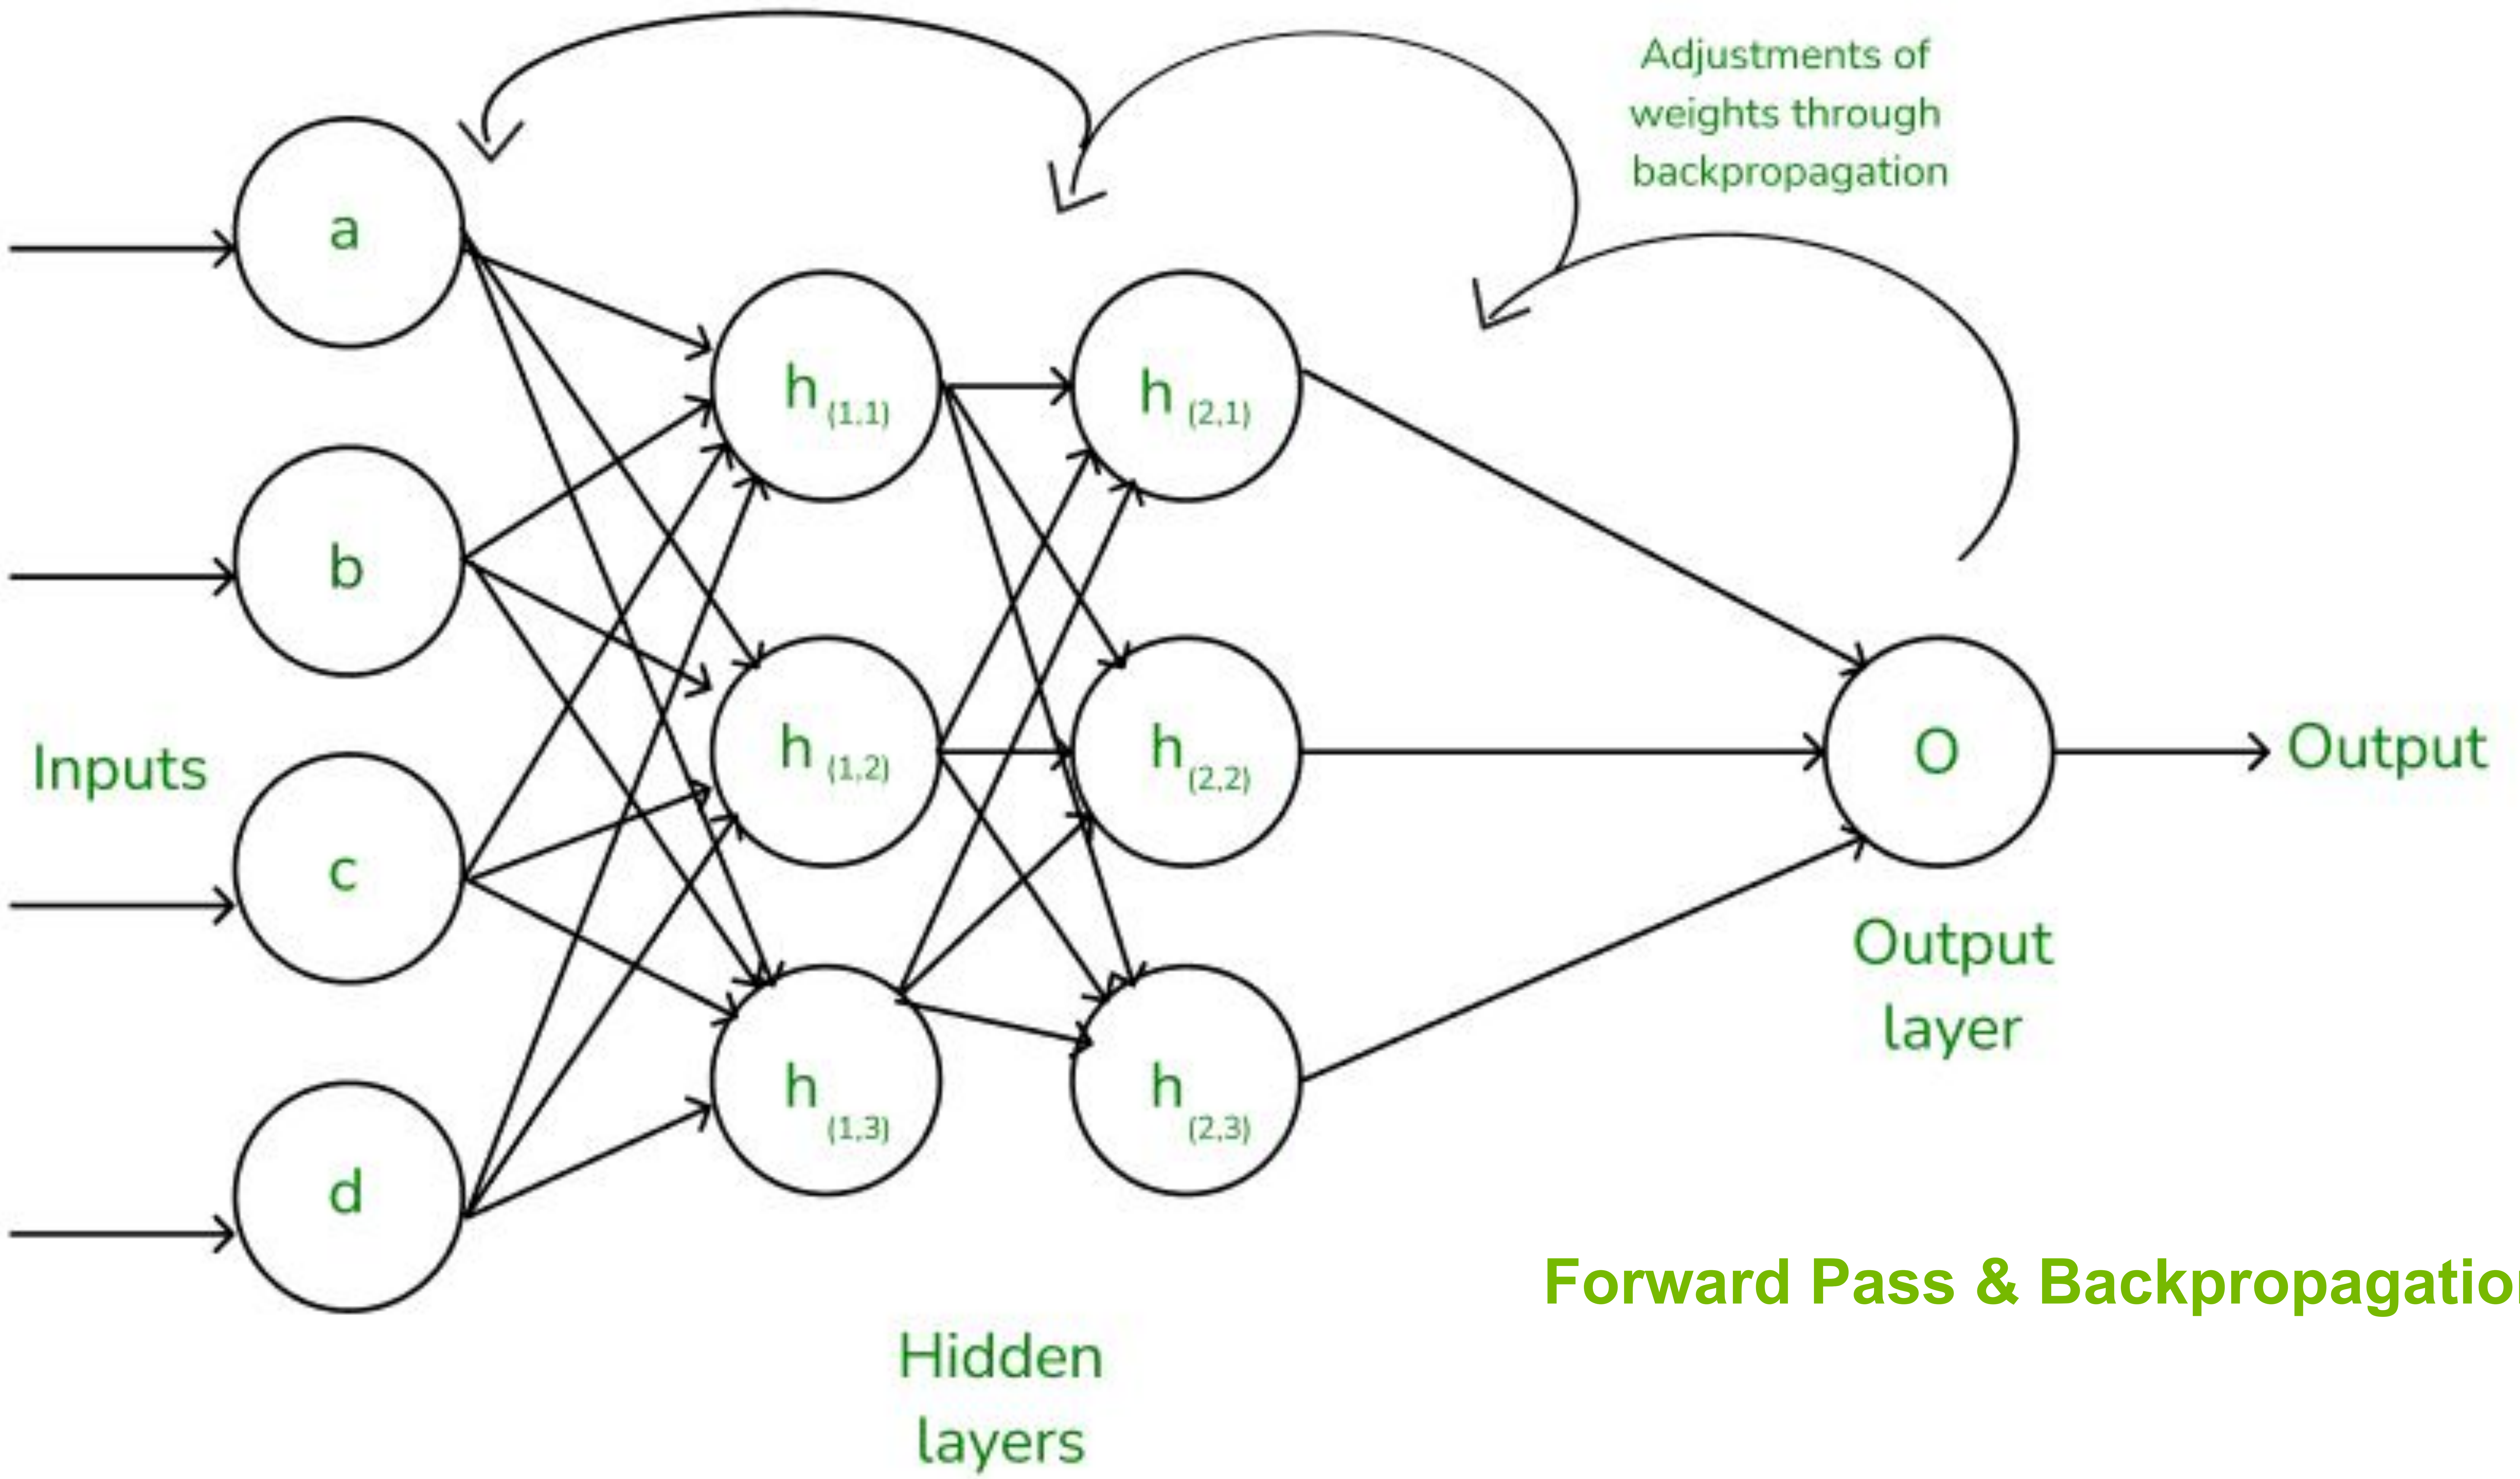

Forward Pass & Backpropagation

LLM Training

How are LLM models trained?

Next-**T**oken **P**rediction Problem
with ~ Trillion Corpus

Code, Math…
Task Specific Knowledge

Guardrail, Human Preference …

Random Initialization → **Pre-Training** → **Post-Training** → **Alignment** → **Production**

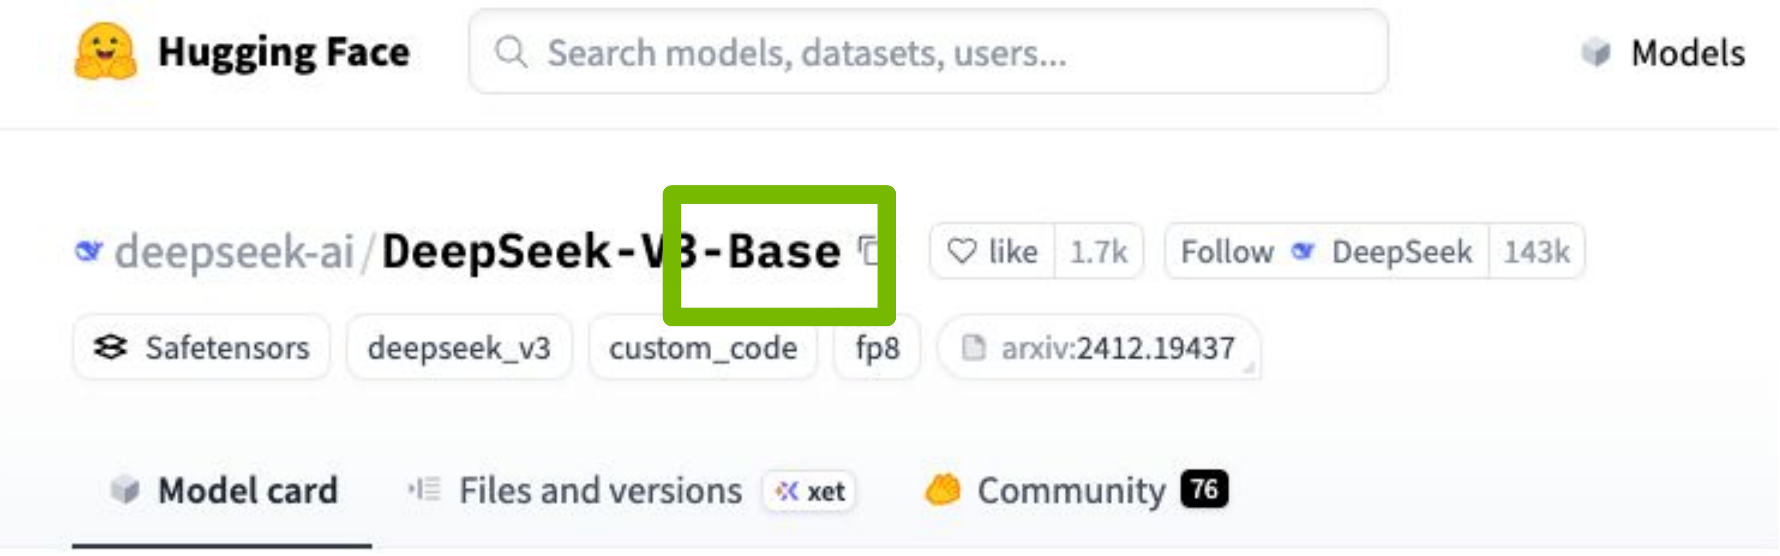

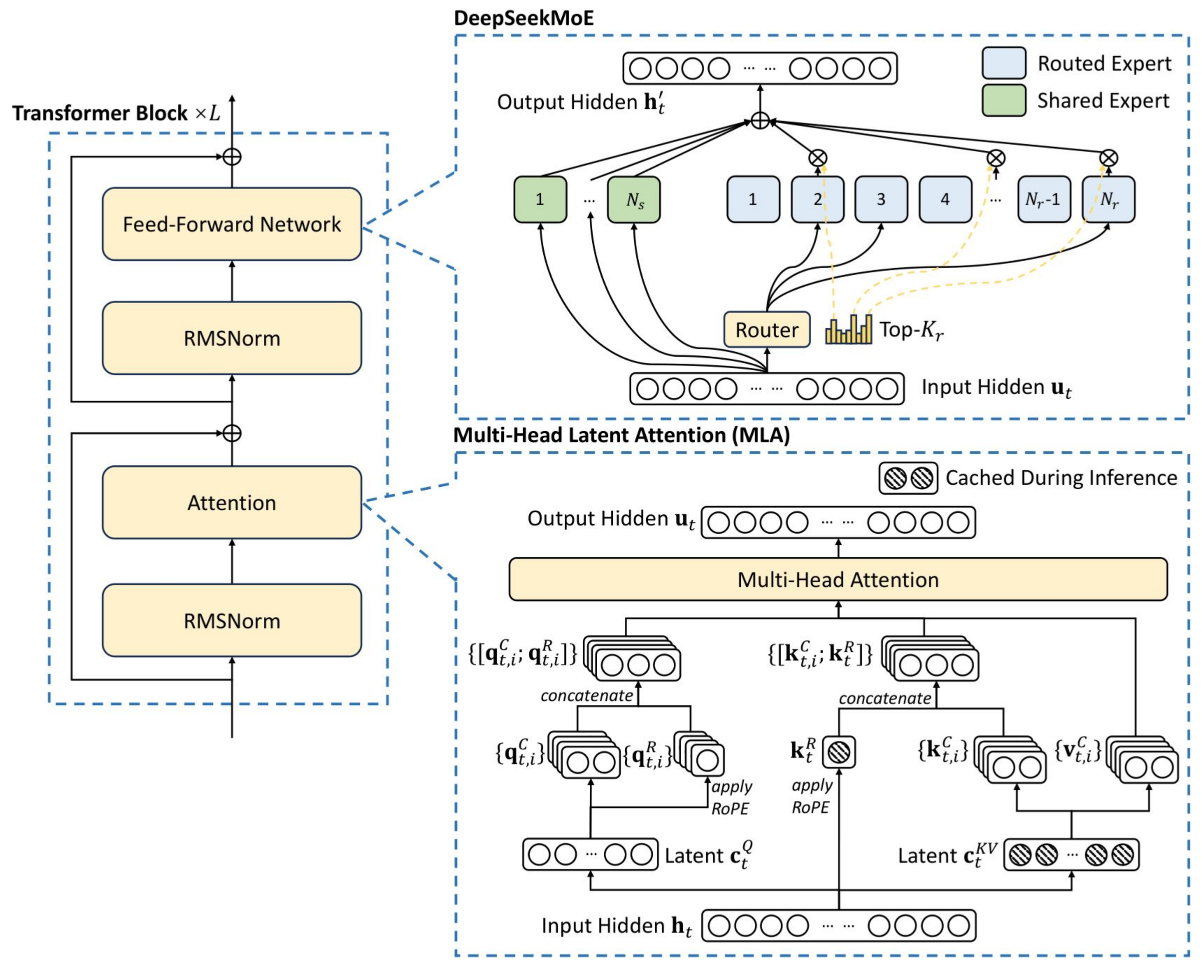

Strong Backbone
Architecture with
randomized parameters

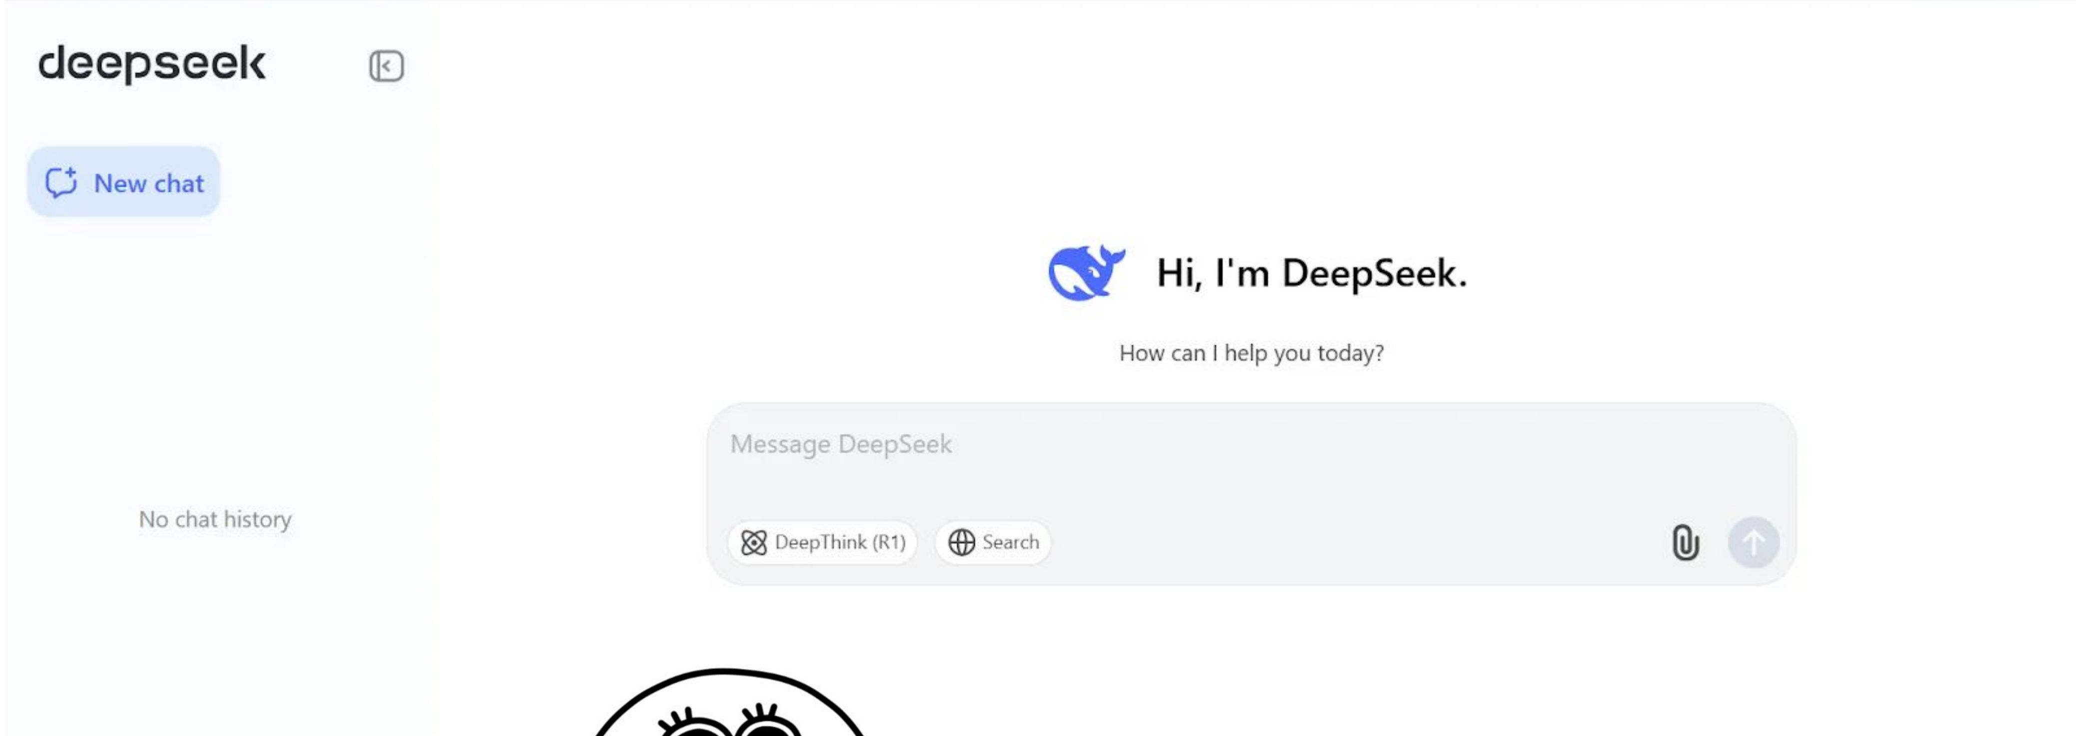

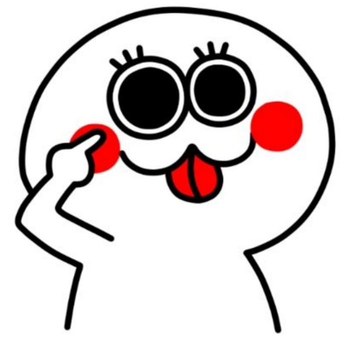

LLM Training

How are LLM models trained?

kanana

Kanana(Kakao) model training recipe from if(kakao)25

https://www.youtube.com/watch?v=B7sc-OlnkCE

LLM Training

How are LLM models trained?

https://www.youtube.com/watch?v=B7sc-OlnkCE

LLM Training

How are LLM models trained?

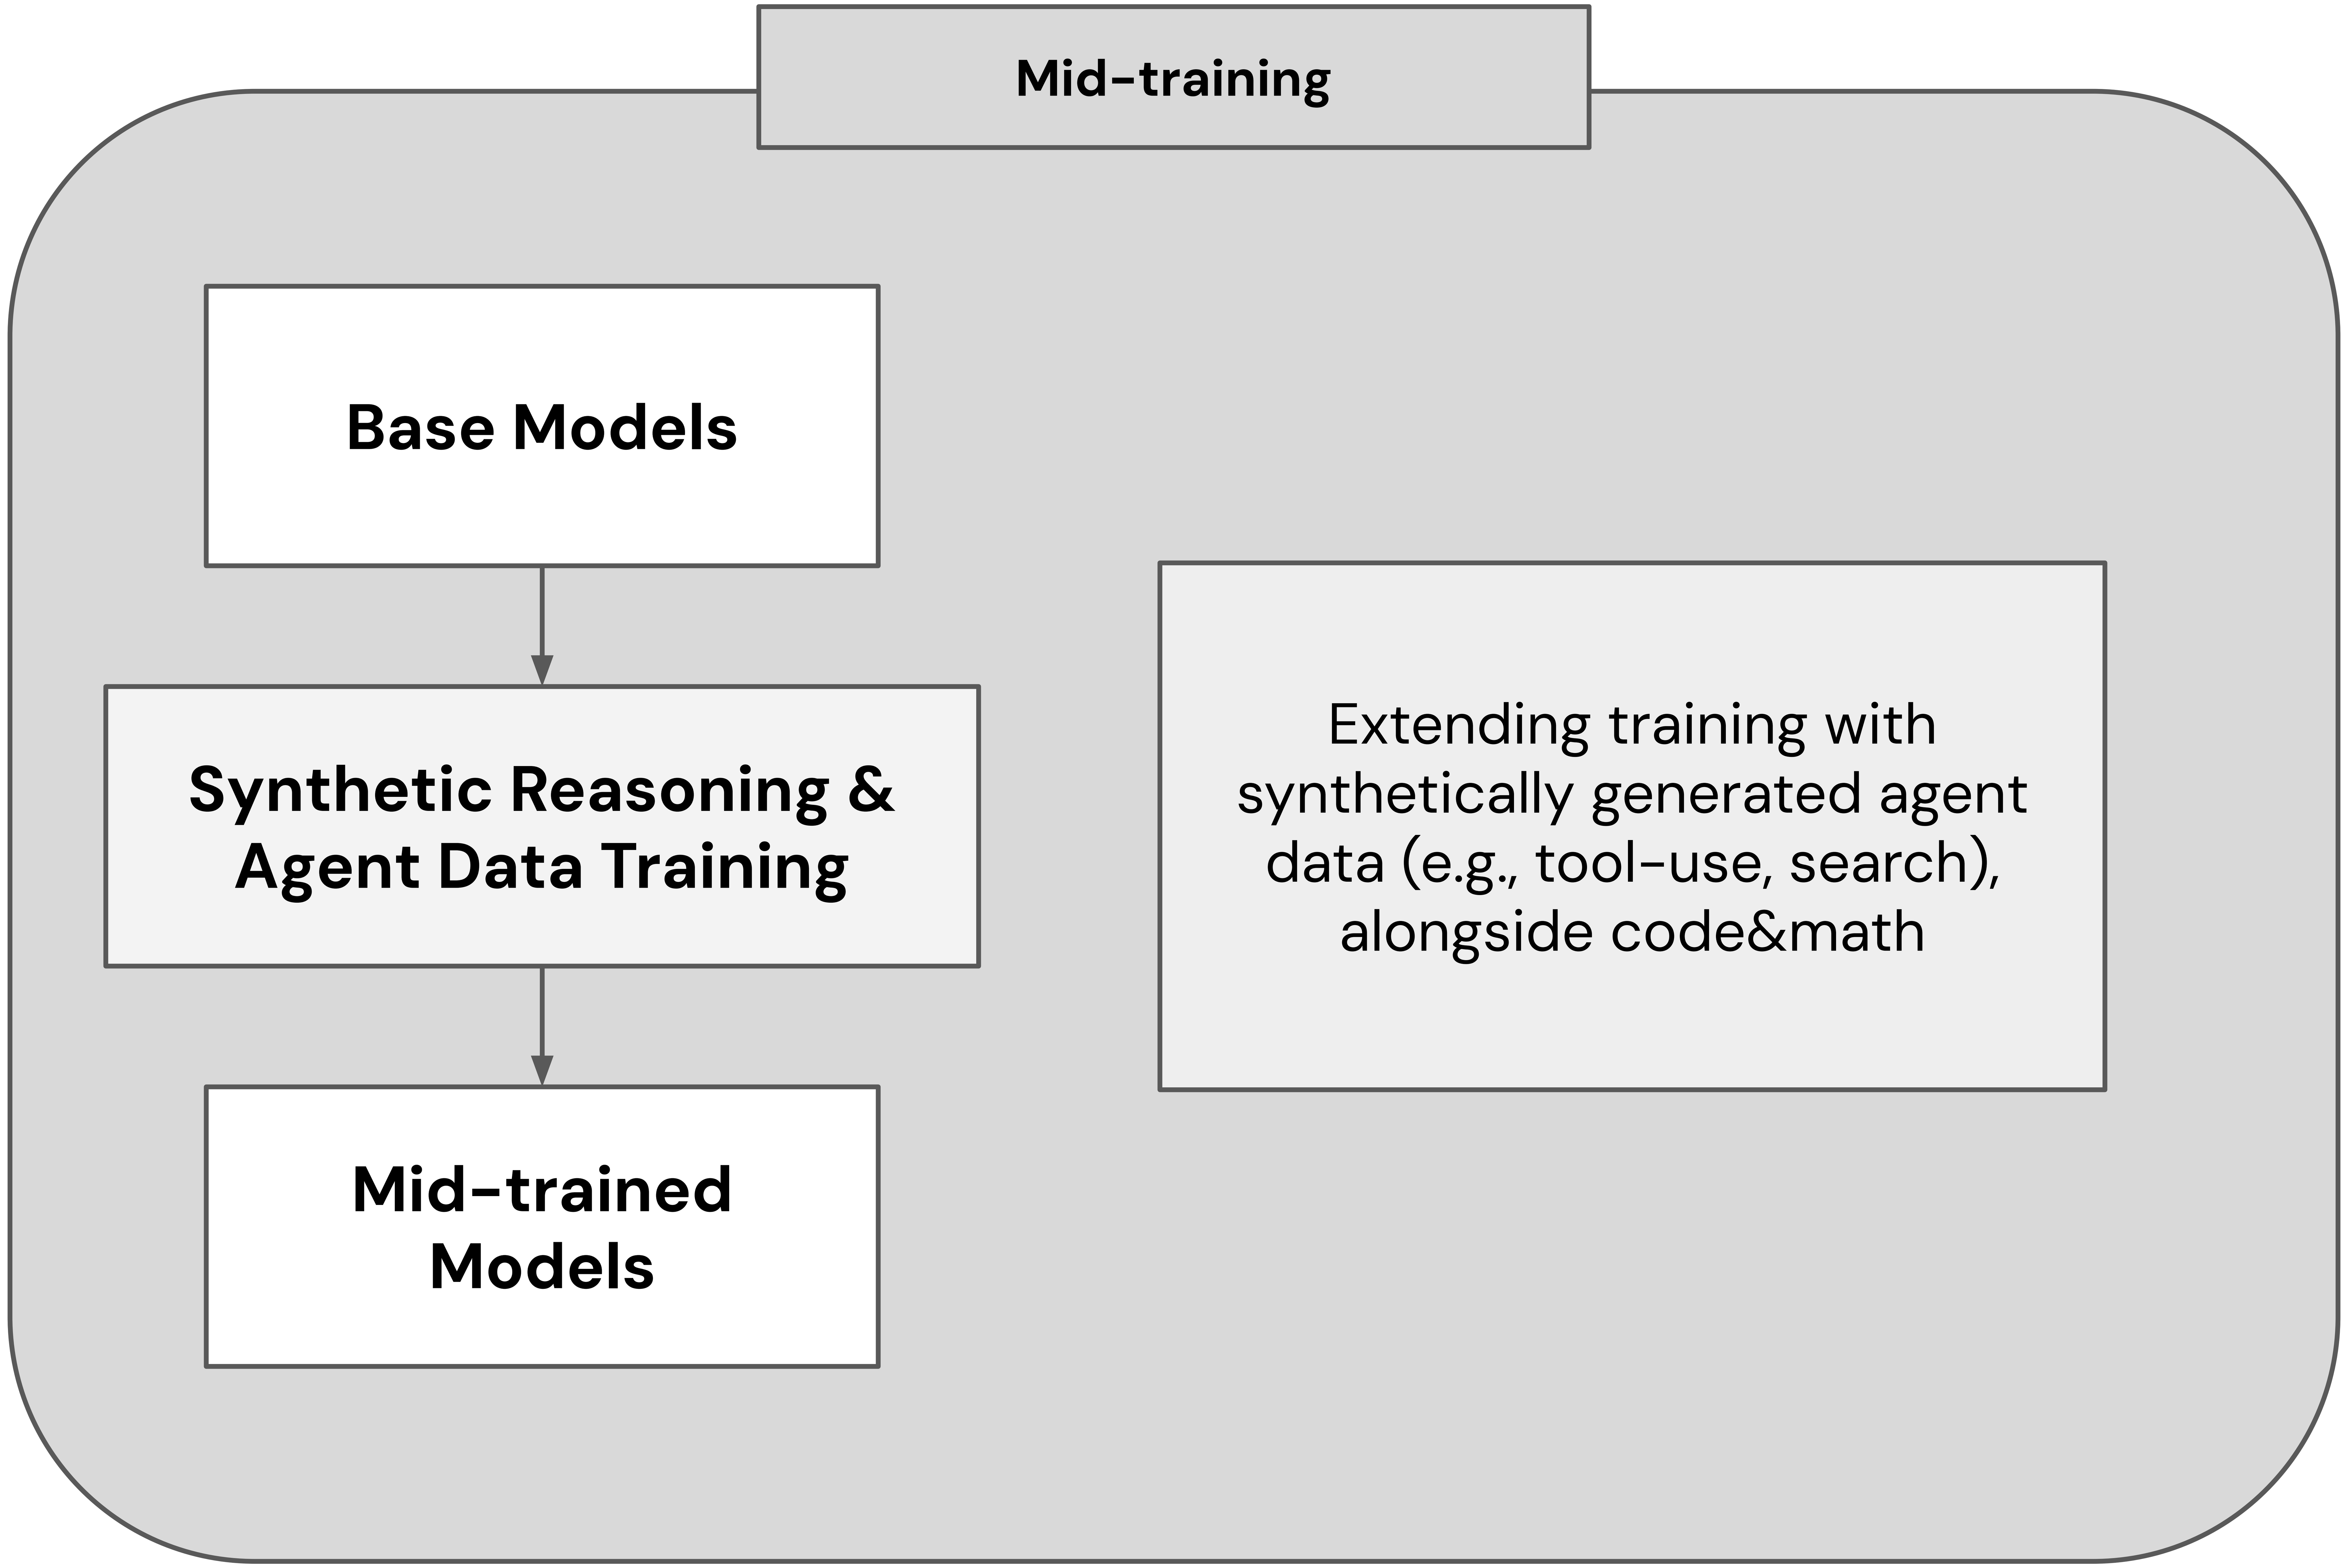

Building RAG Agents with LLMs

Part 4. Evolution of AI Models

Evolution of AI Models

Expansion into Multimodality and Omni-Modality

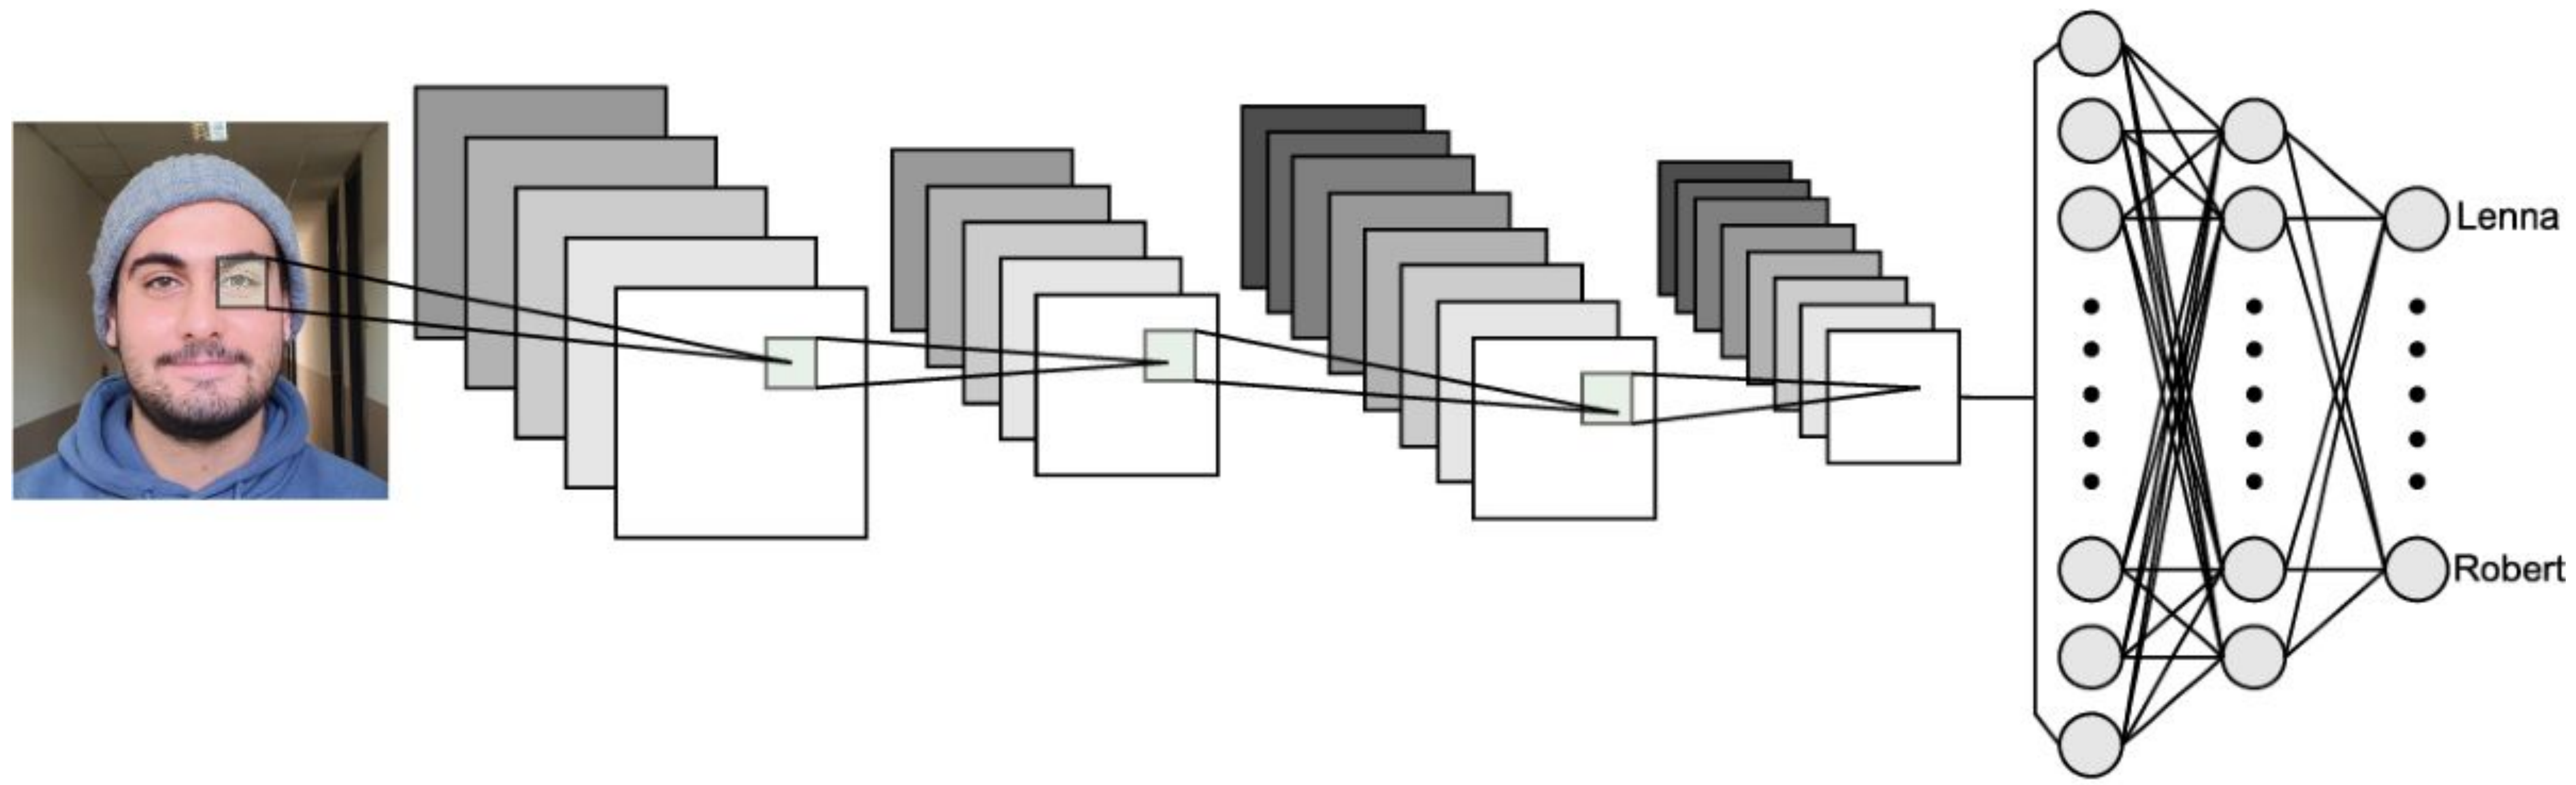

(a) Common CNN architecture

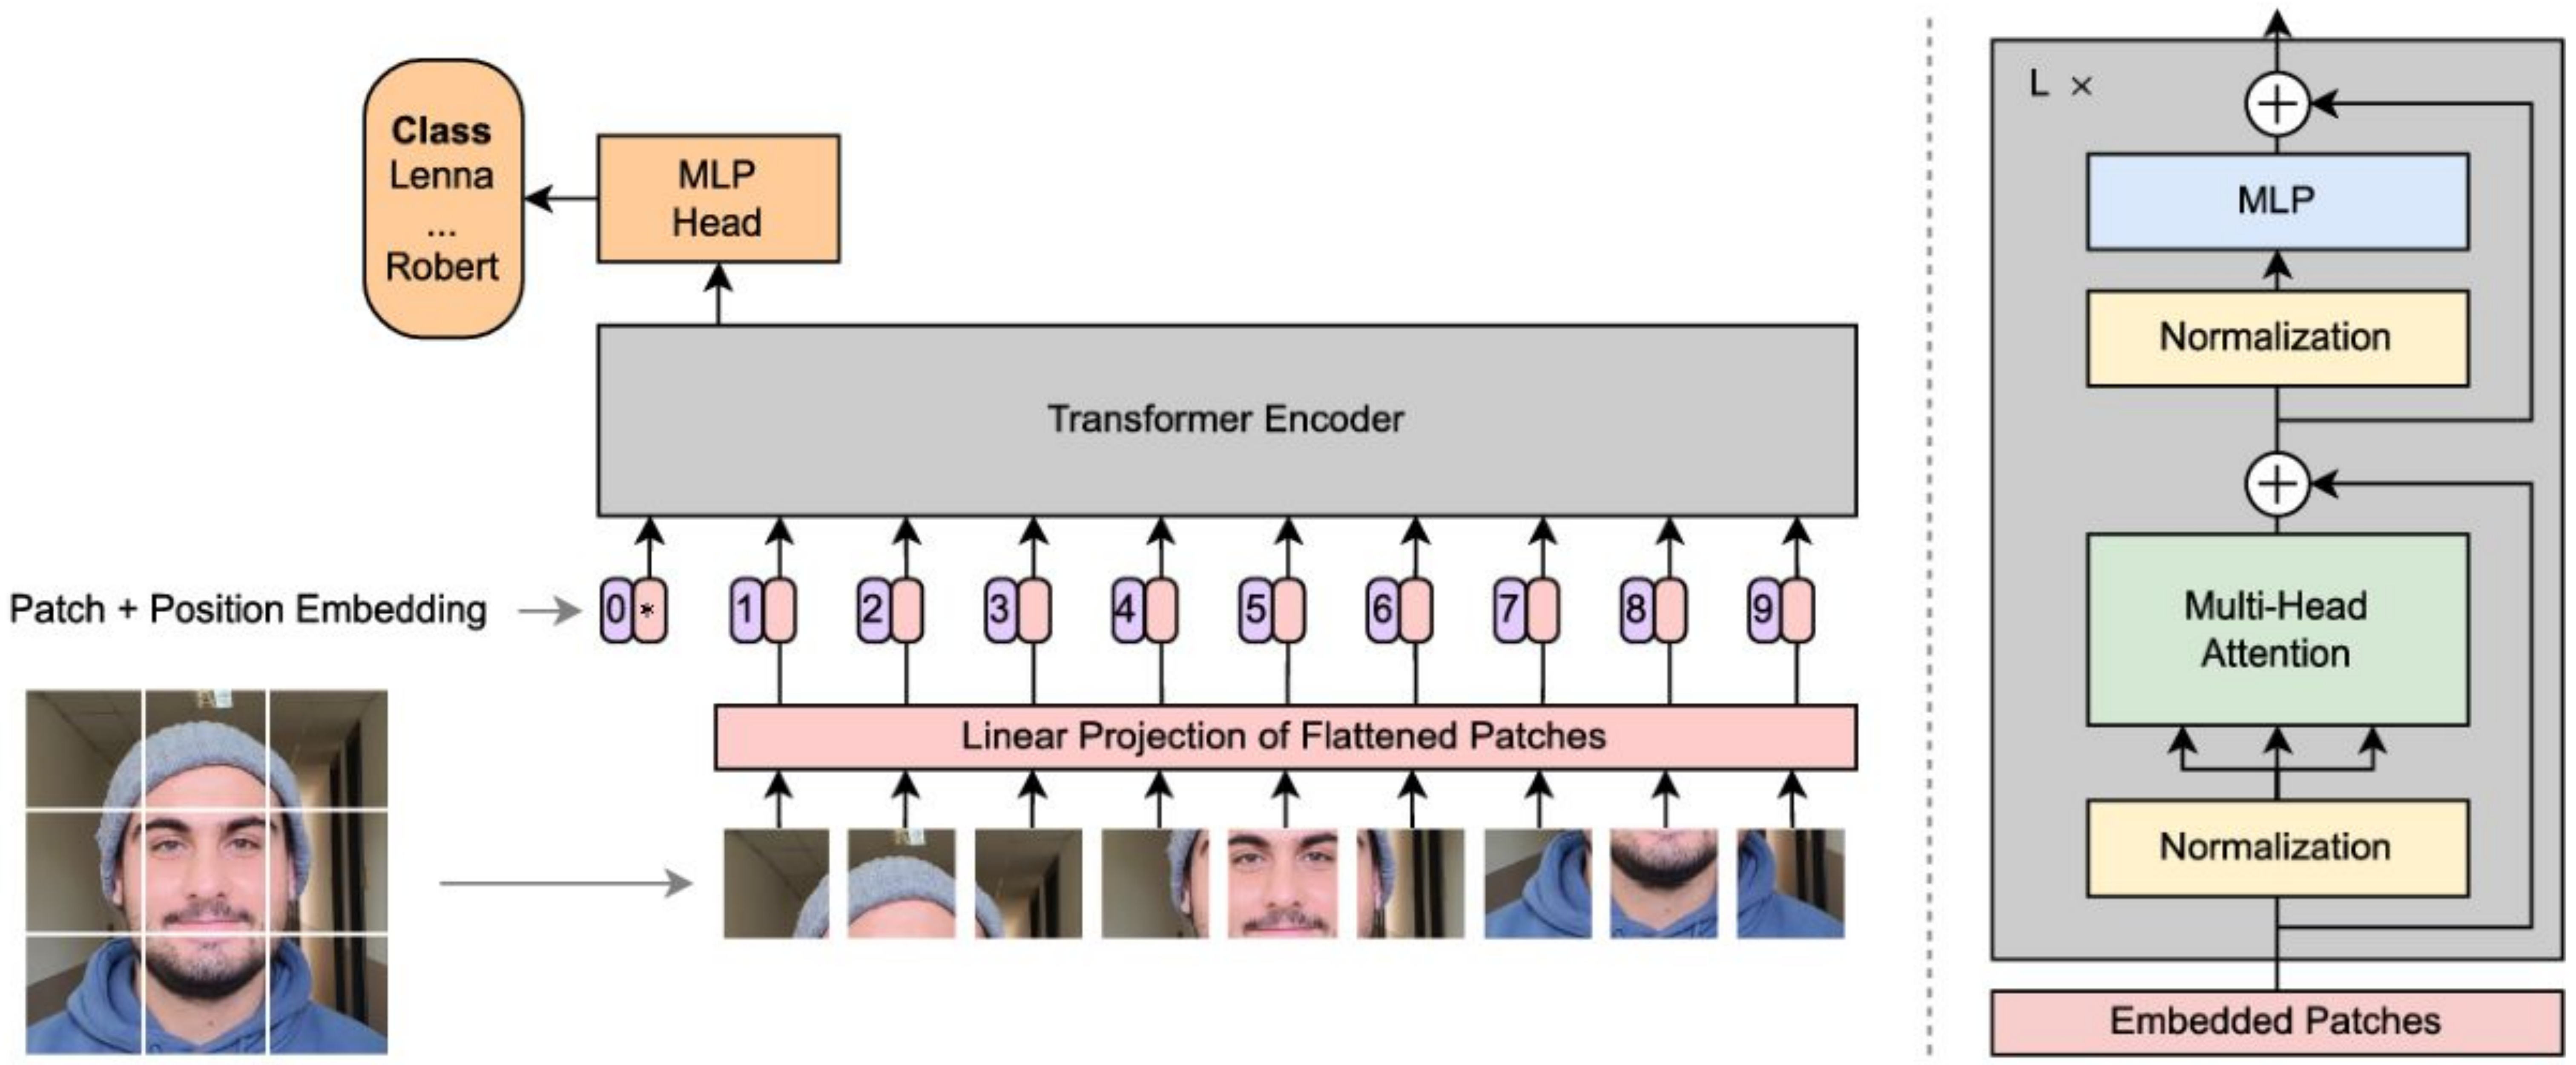

(b) Vision Transformer architecture

Evolution of AI Models

Expansion into Multimodality and Omni-Modality

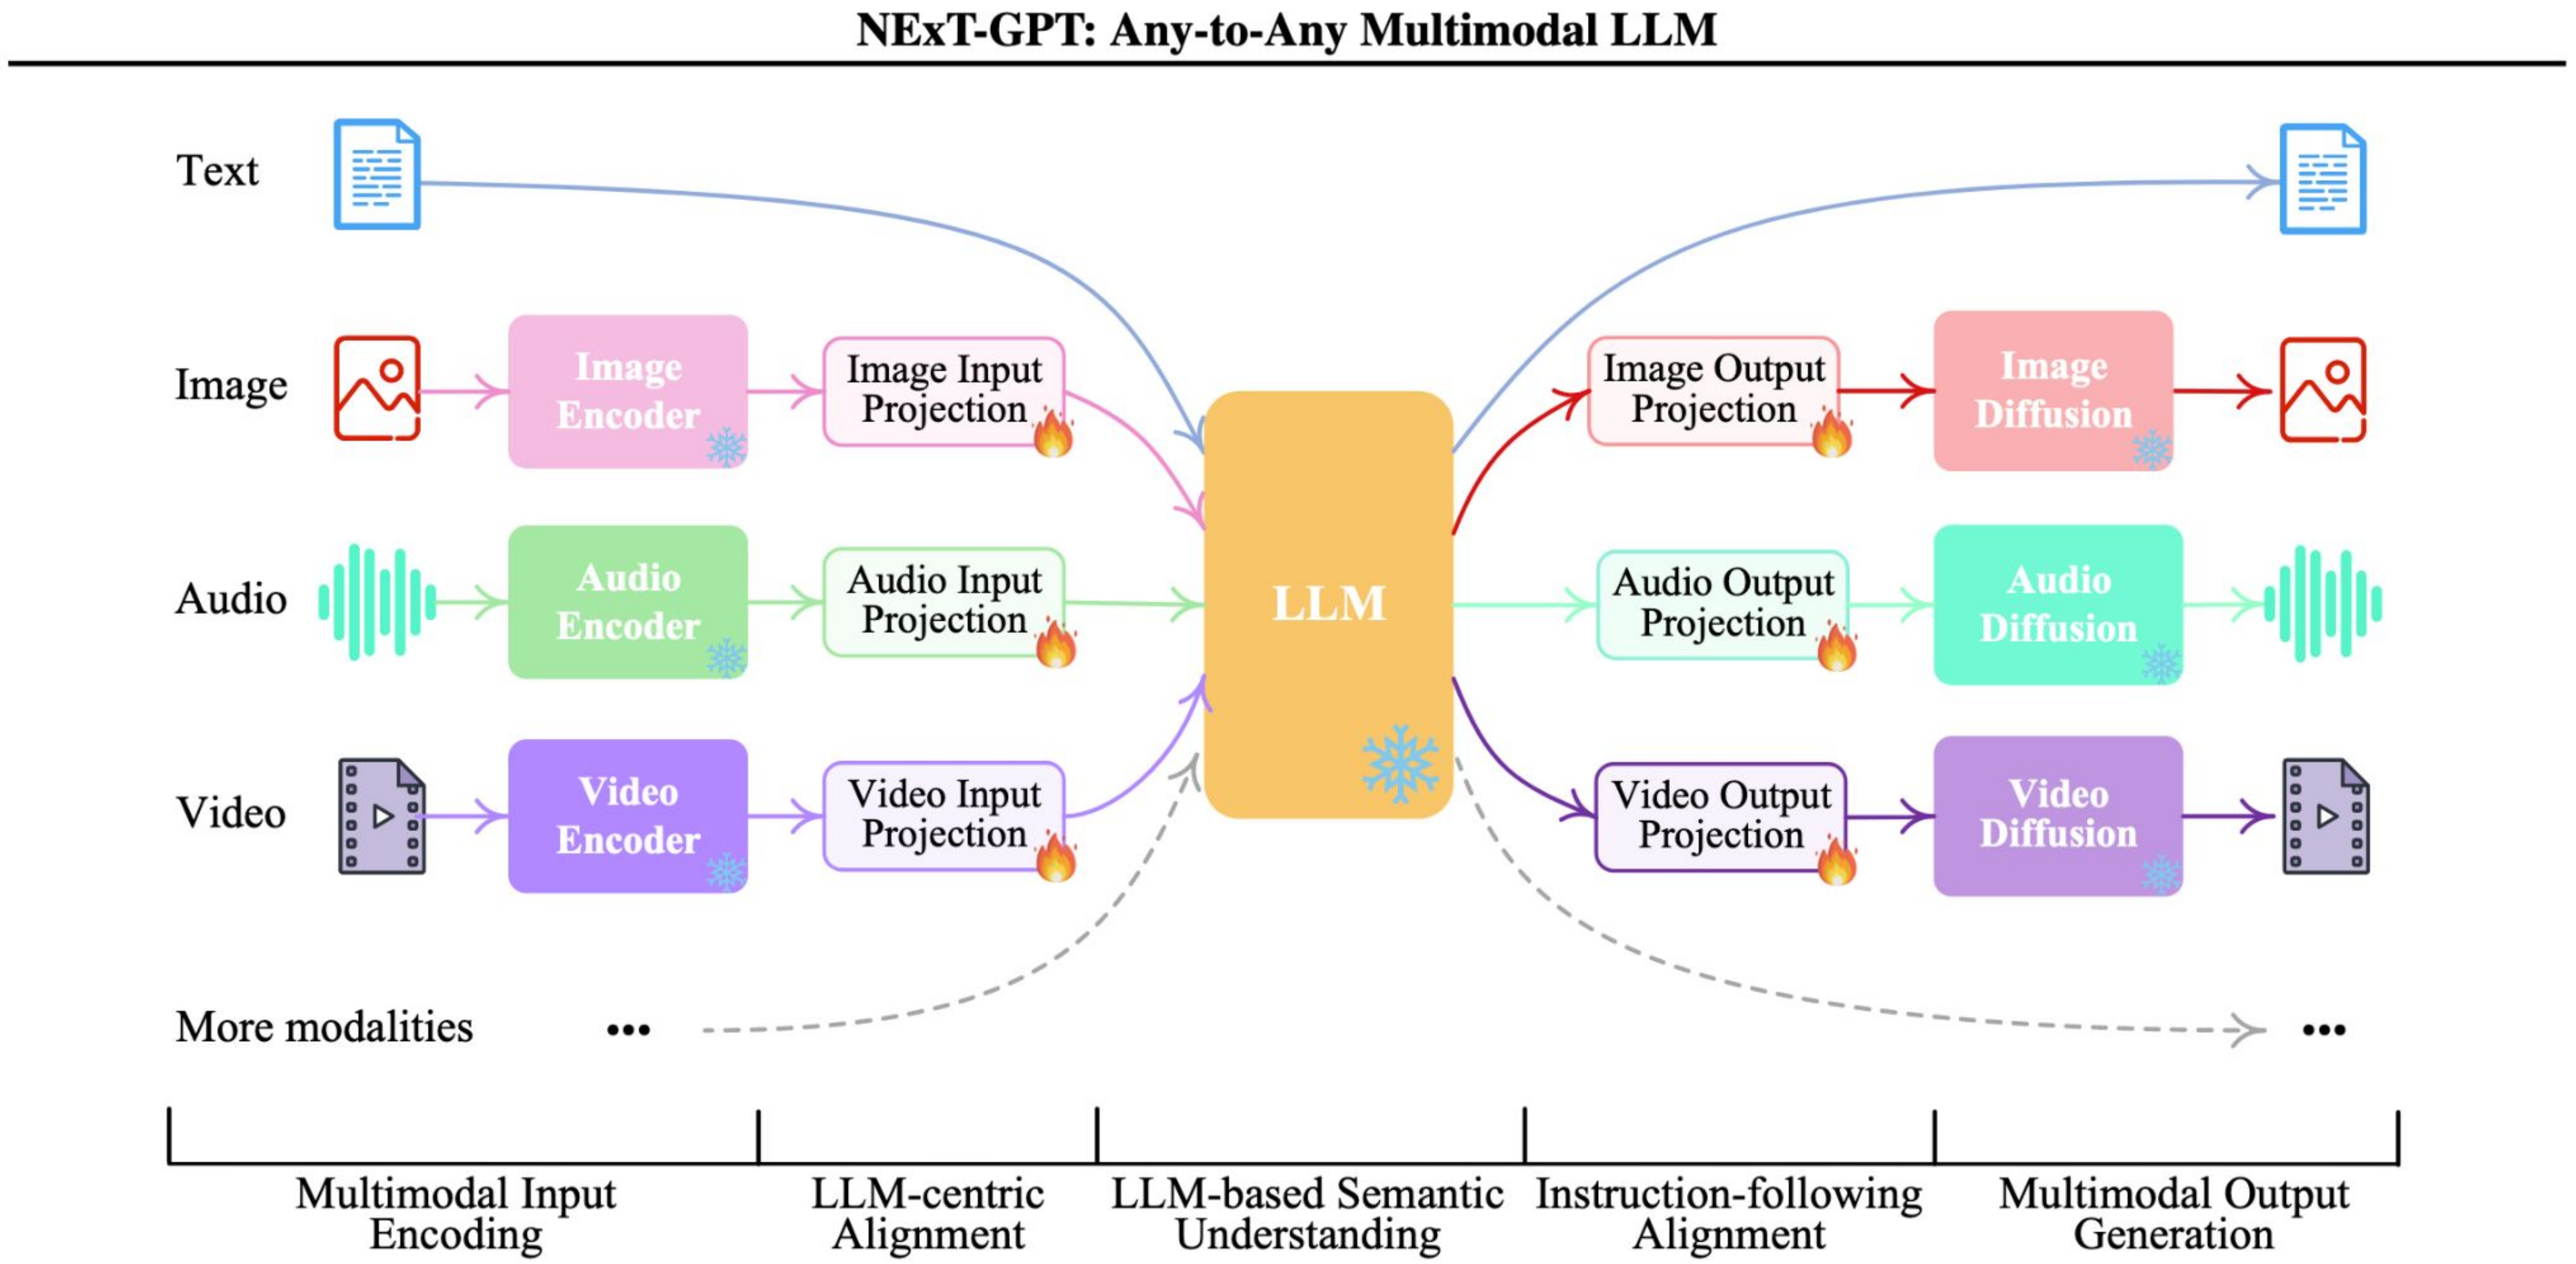

Figure 1. By connecting LLM with multimodal adaptors and diffusion decoders, NExT-GPT achieves universal multimodal understanding and any-to-any modality input and output. ❄ and 🔥 represent the frozen and trainable modules, respectively.

https://arxiv.org/abs/2309.05519

Evolution of AI Models

Expansion into Multimodality and Omni-Modality

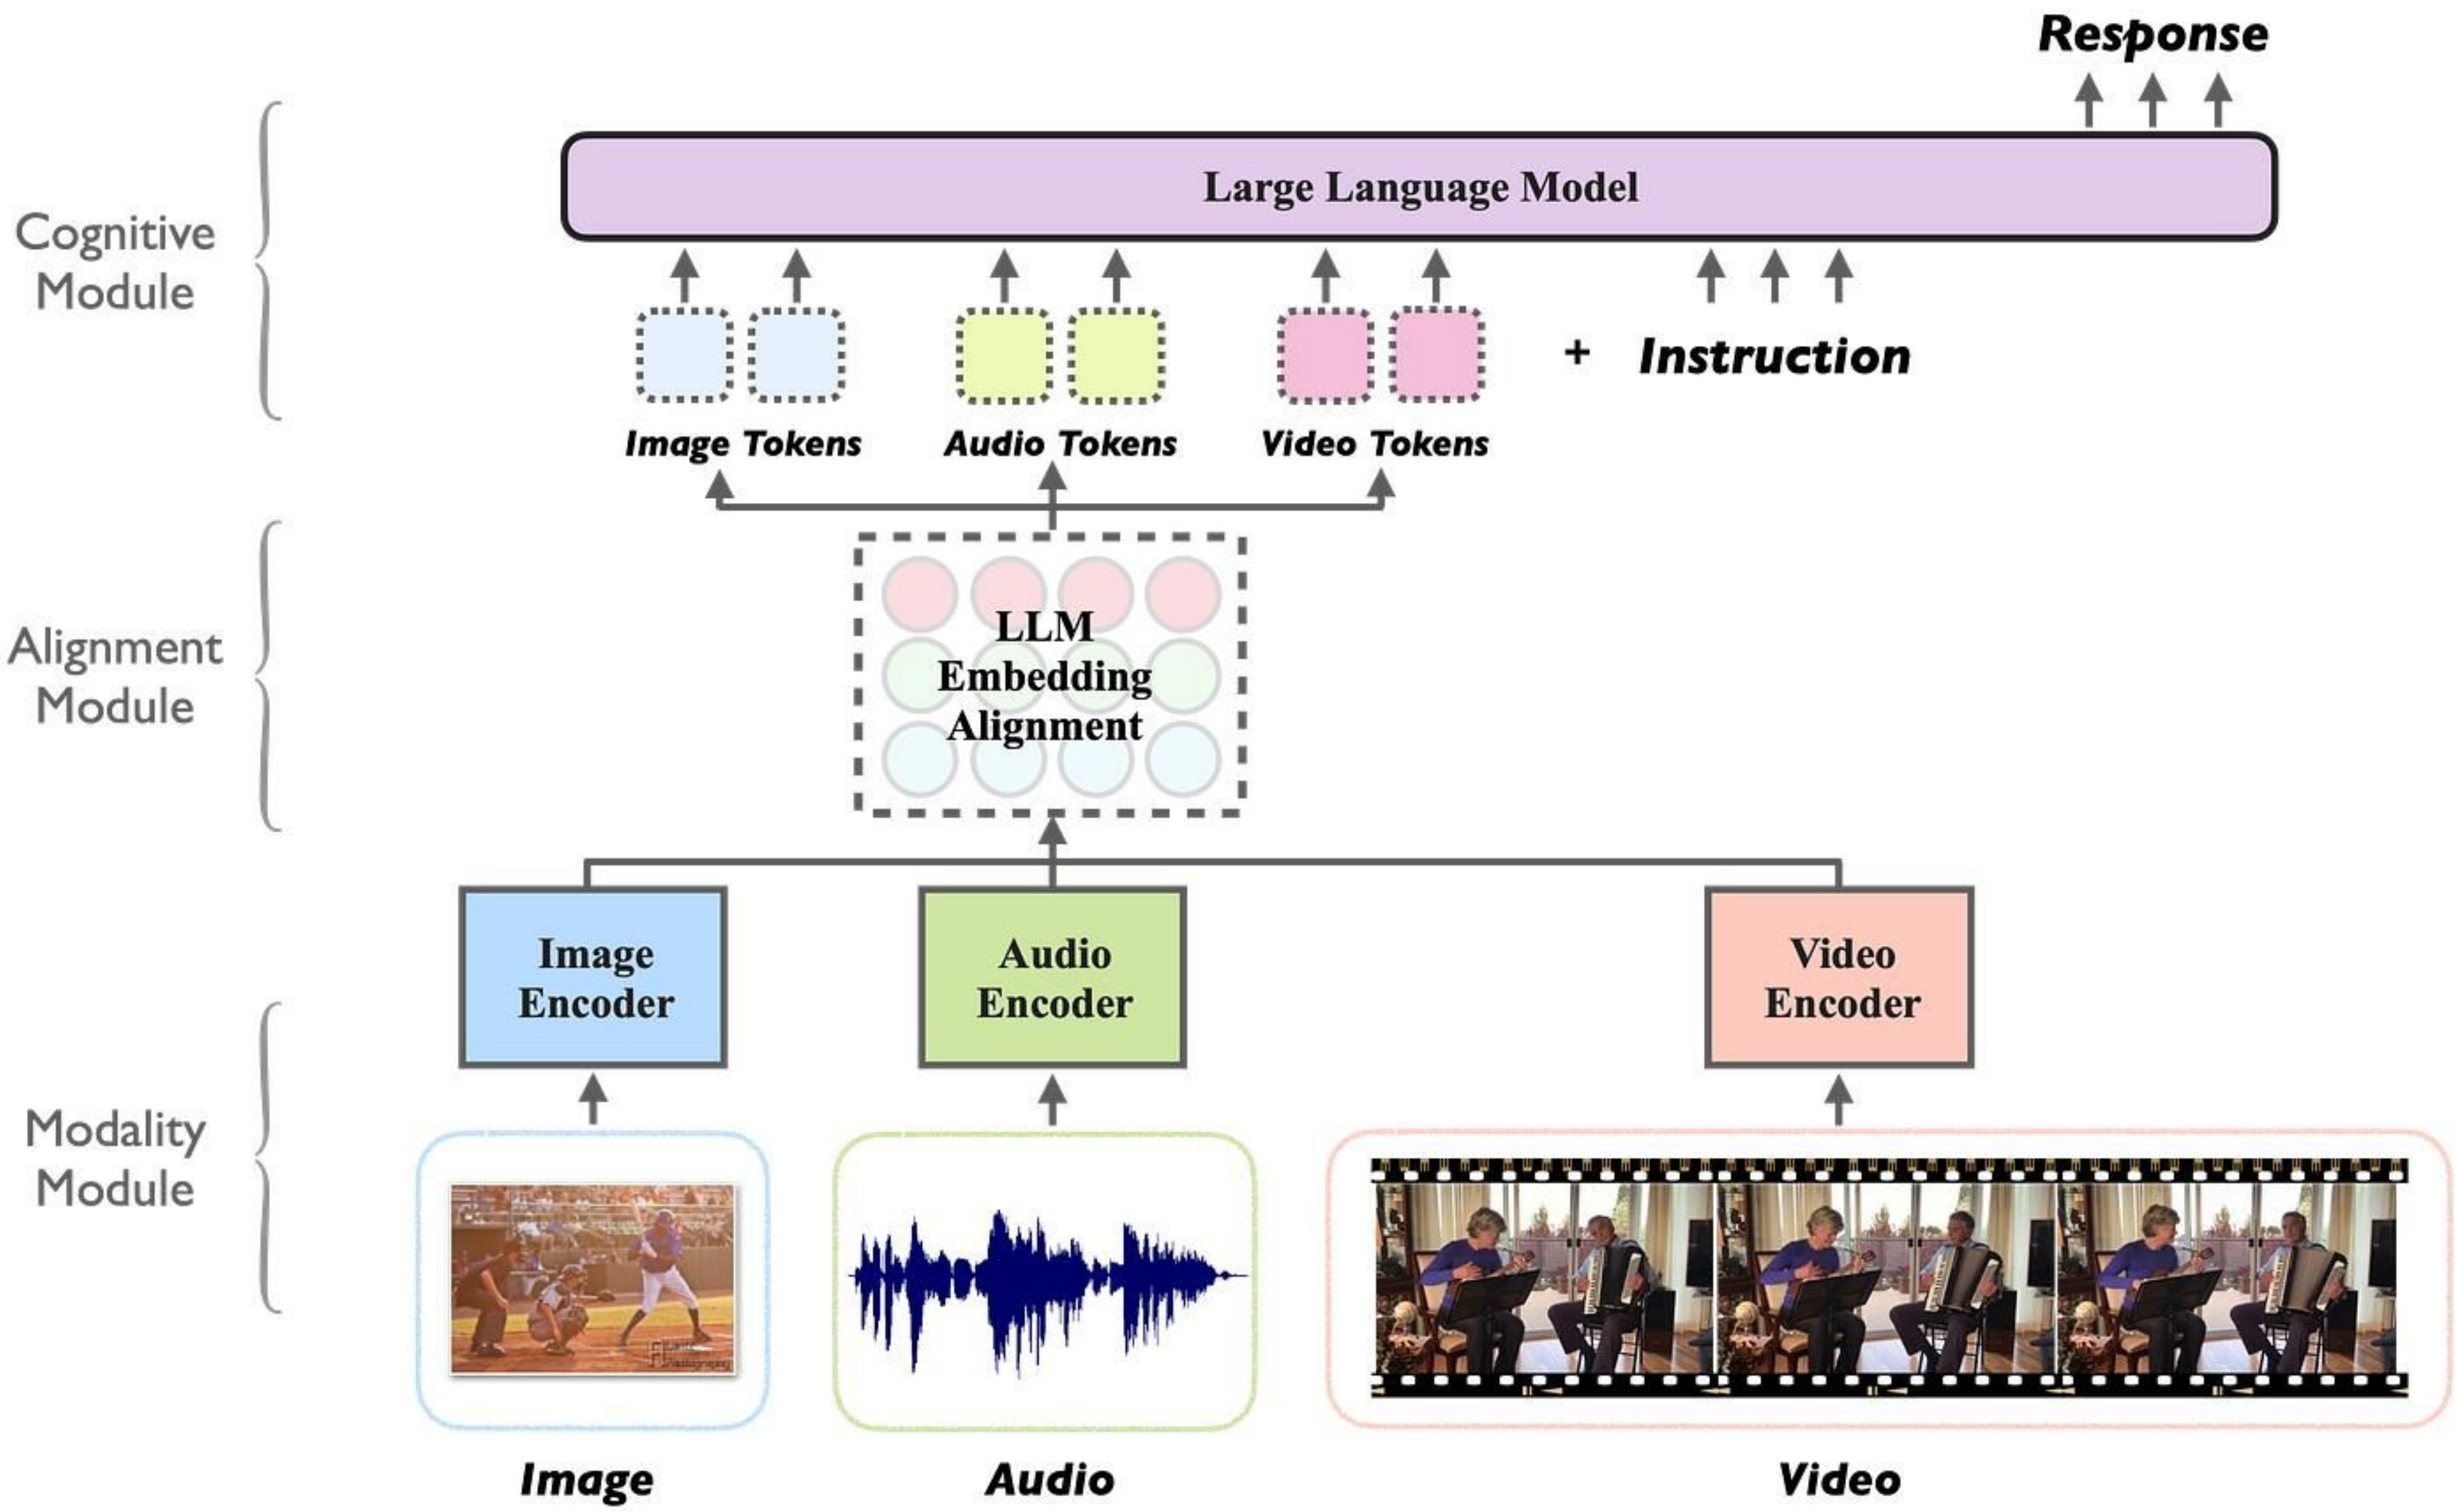

https://medium.com/%40cout.shubham/exploring-multimodal-large-language-models-a-step-forward-in-ai-626918c6a3ec?utm_source=chatgpt.com

Evolution of AI Models

Expansion into Multimodality and Omni-Modality

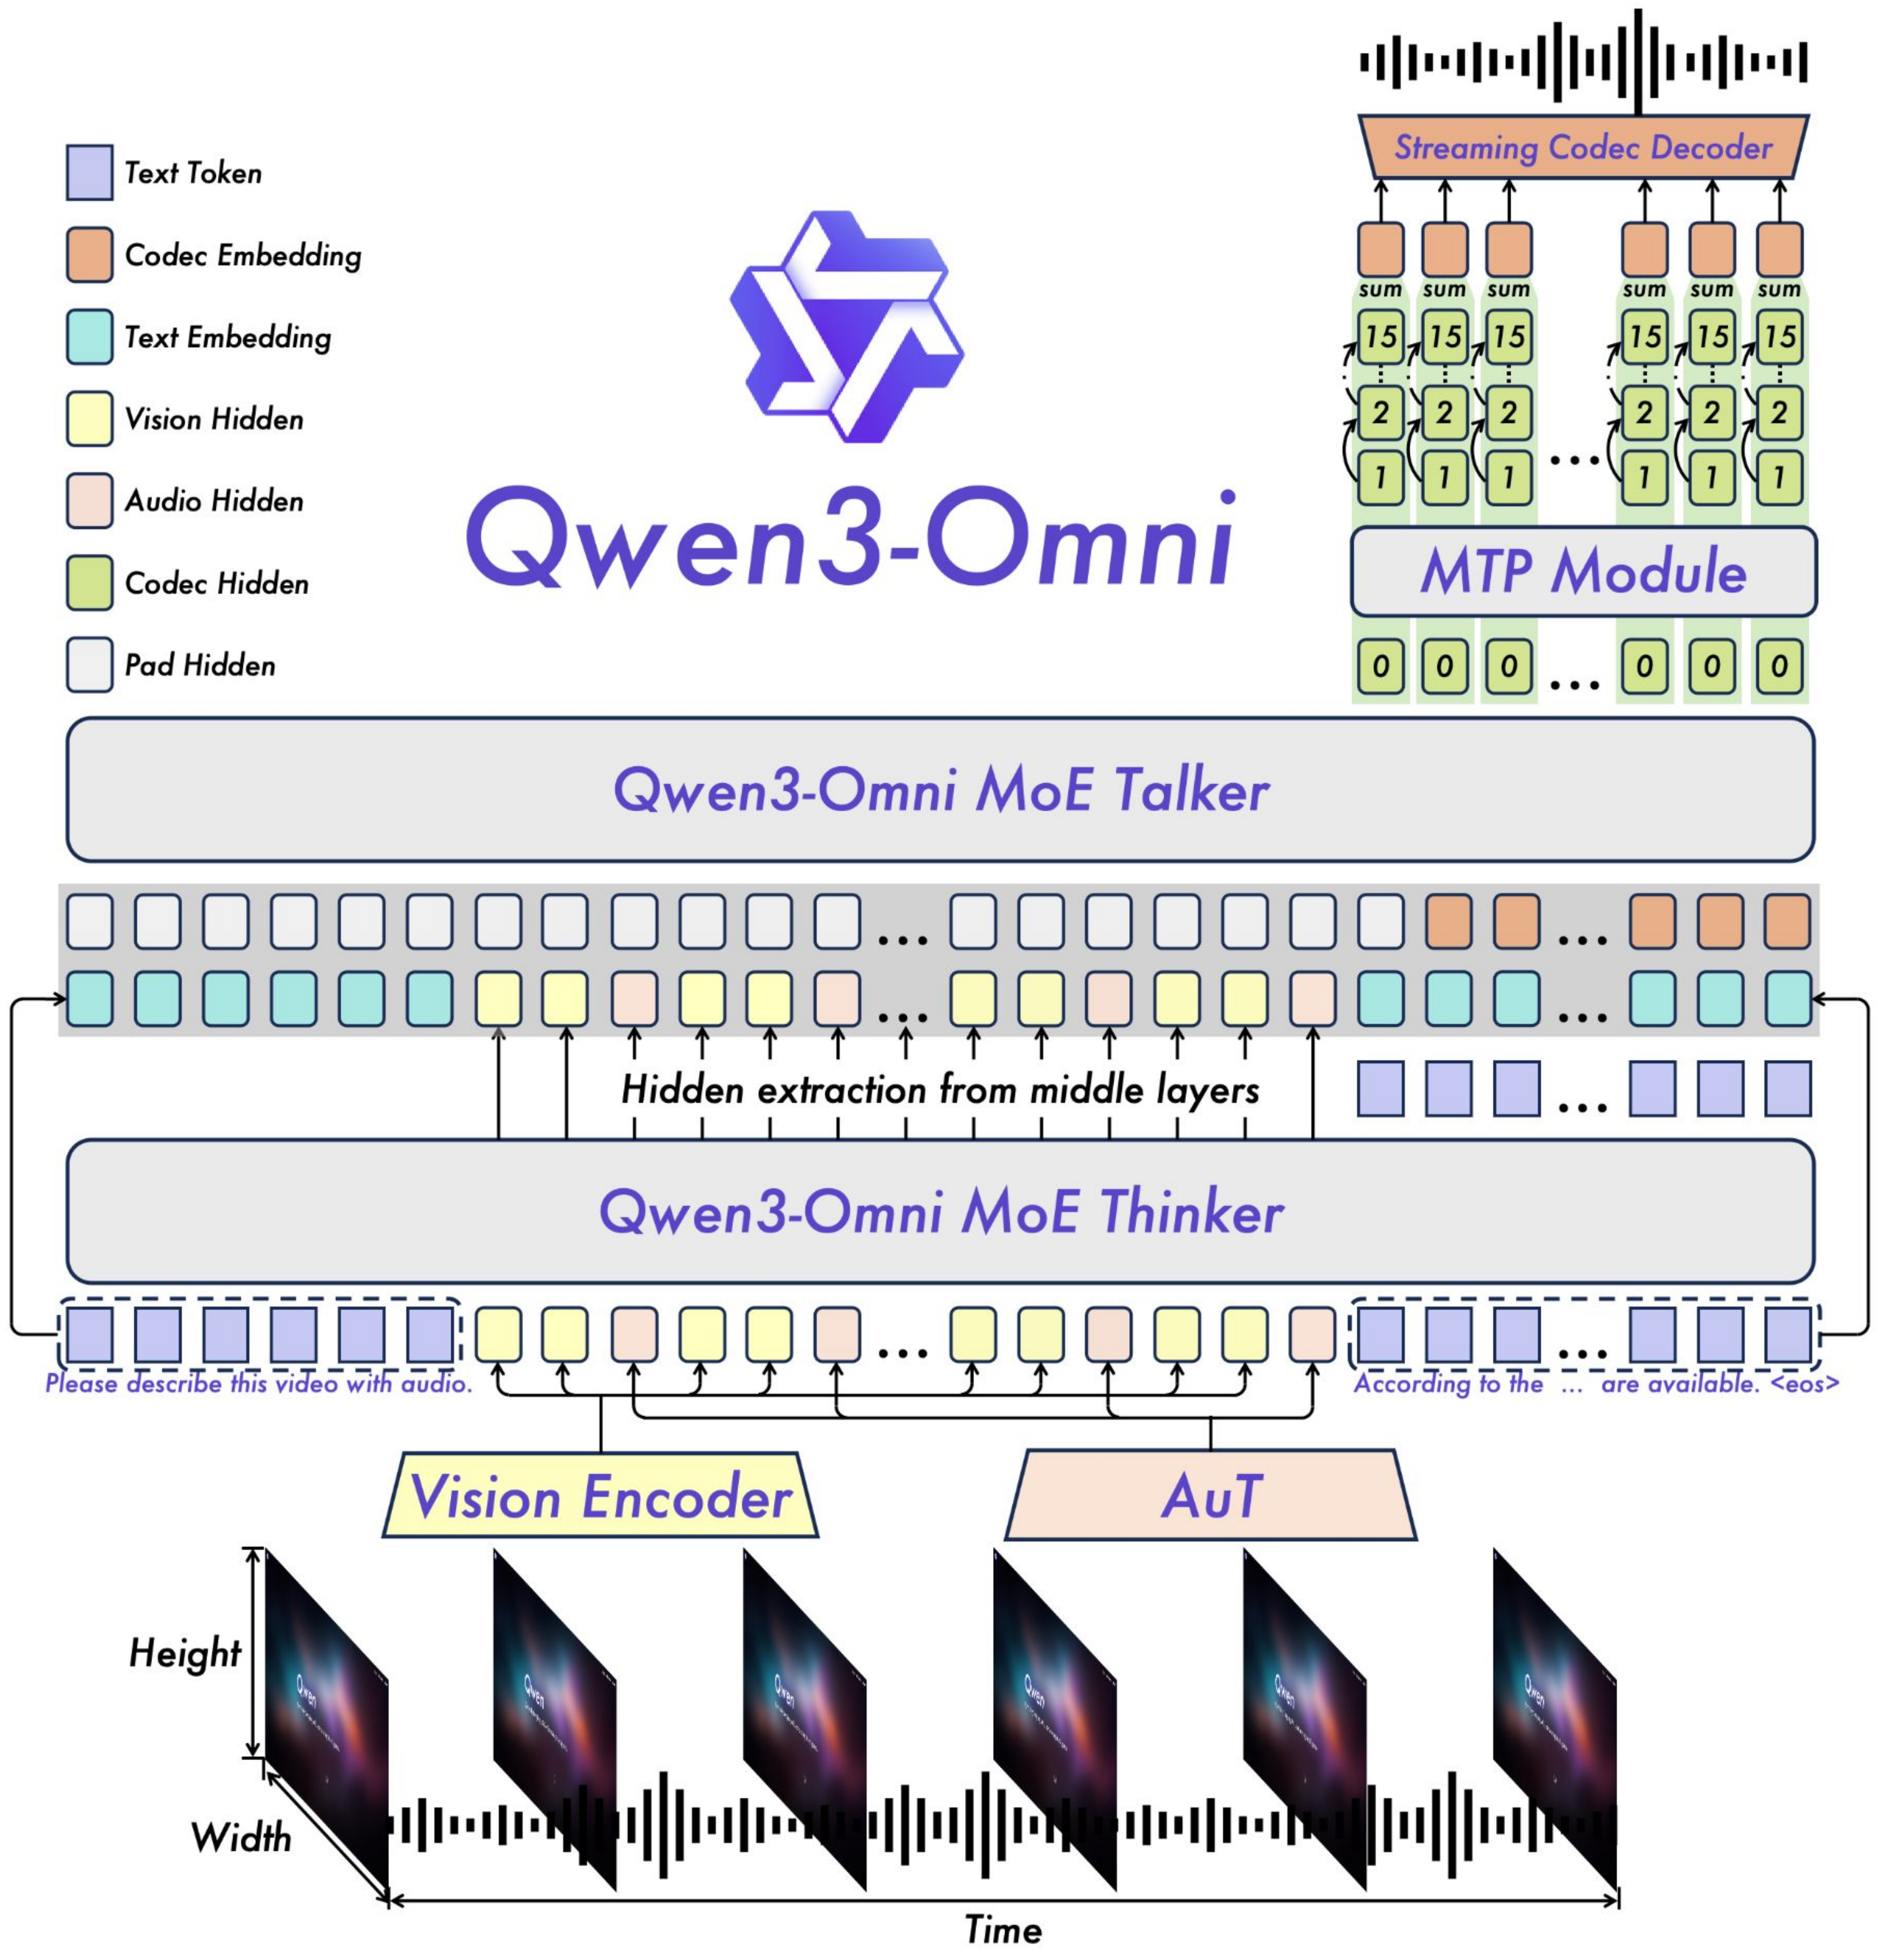

https://qwen.ai/blog?id=fdfbaf2907a36b7659a470c77fb135e381302028&from=research.research-list

Evolution of AI Models

Expansion into Multimodality and Omni-Modality

https://www.youtube.com/watch?v=PAiMe4vftOc

Evolution of AI Models

Expansion into Multimodality and Omni-Modality

https://www.youtube.com/watch?v=PAiMe4vftOc

Building RAG Agents with LLMs

Part 5. GPU Infrastructure

GPU Infrastructure

Recap - Massive Model Parameter

https://huggingface.co/nvidia/NVIDIA-Nemotron-3-Nano-30B-A3B-FP8

GPU Infrastructure

Recap - Massive Model Parameter

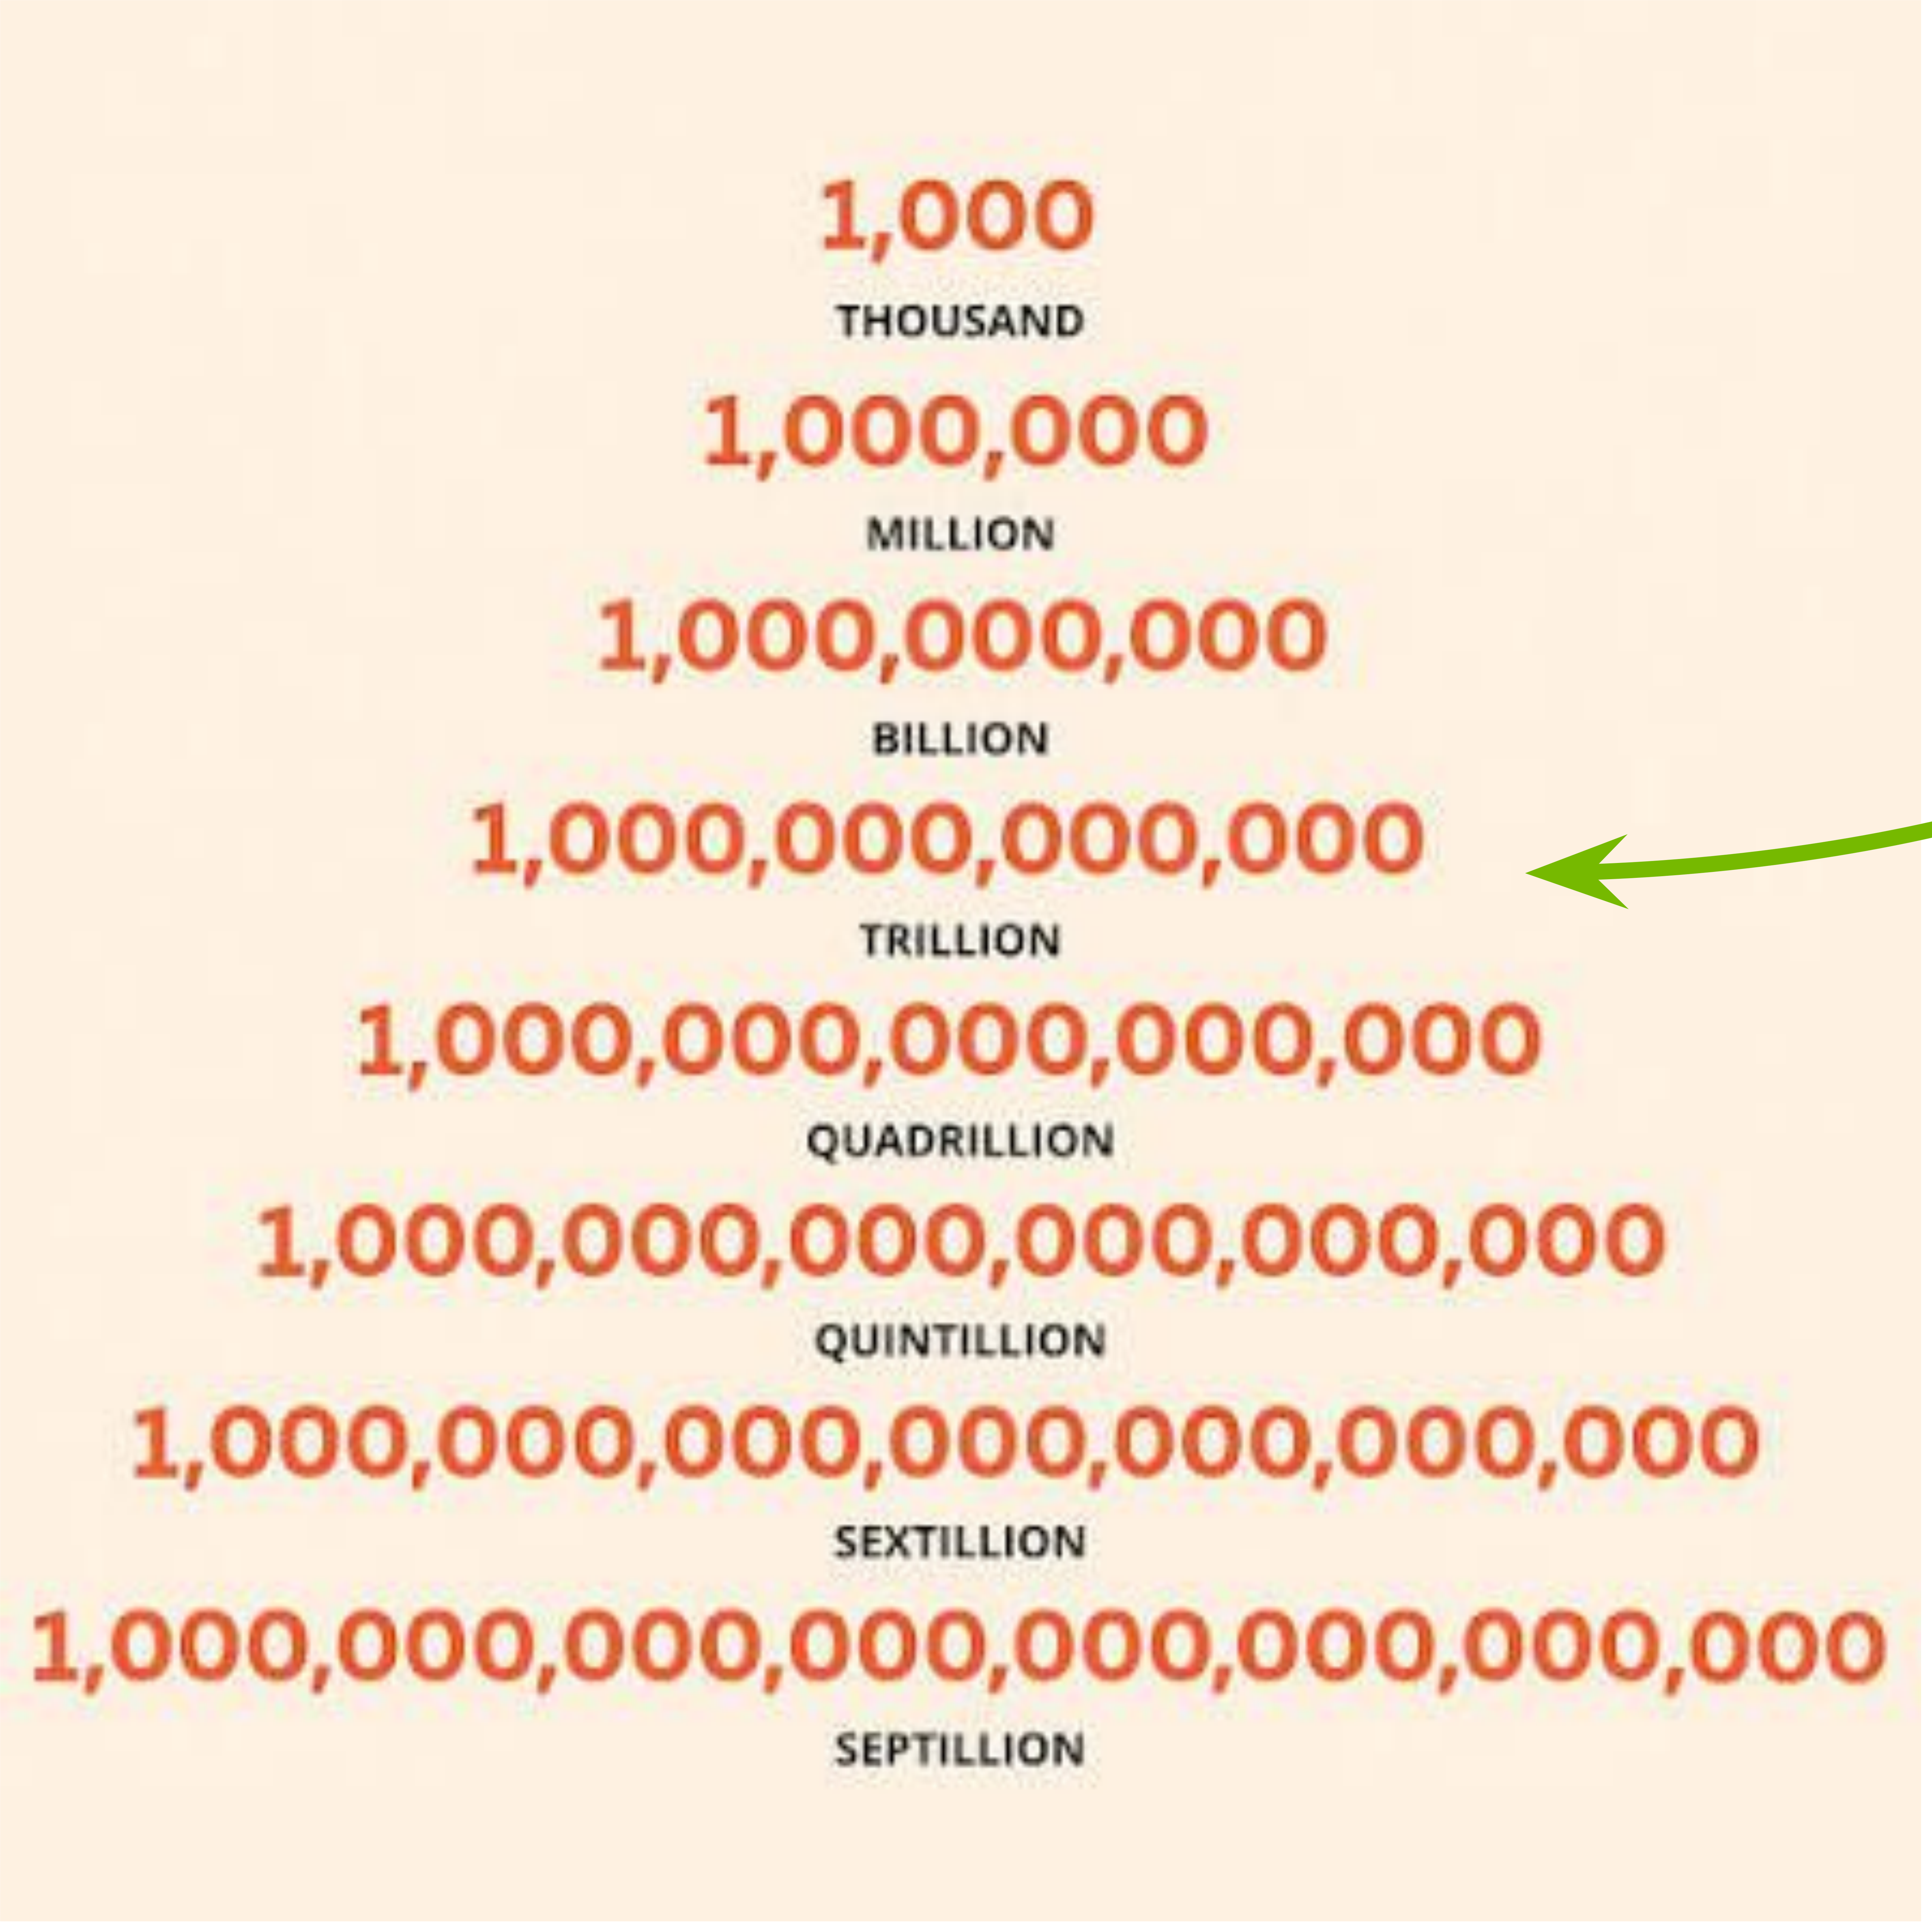

We Are in the Trillion-Parameter Era

How Much Compute Is Required for Training and Inference?

GPU Infrastructure

Modern Architecture - MoE(Mixture of Experts)

Activate Only a Subset of MLP Parameters

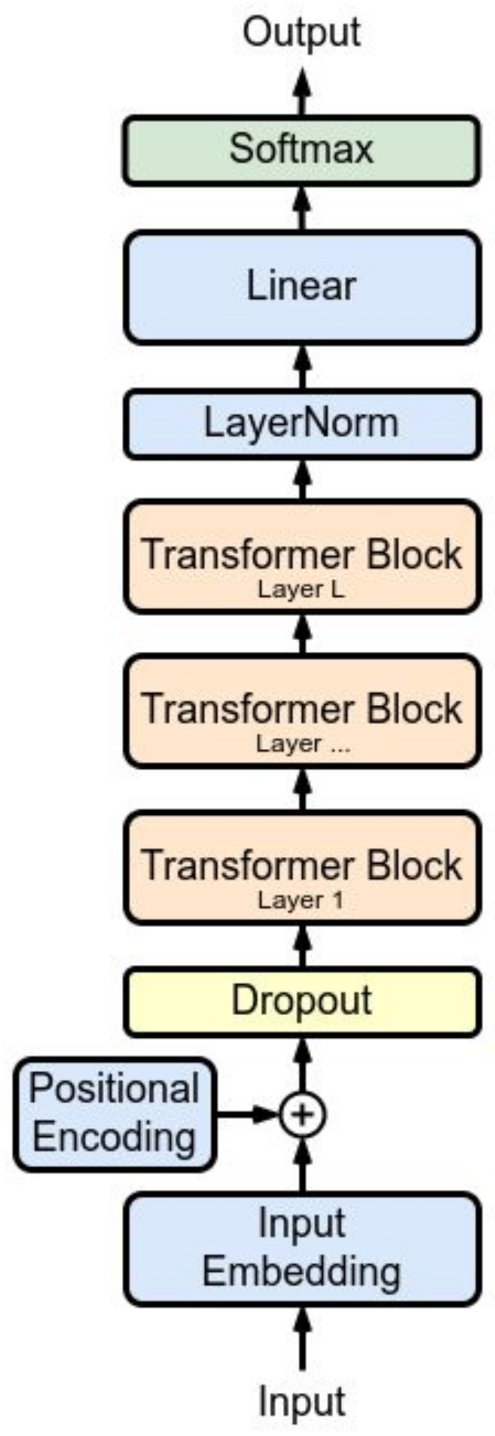

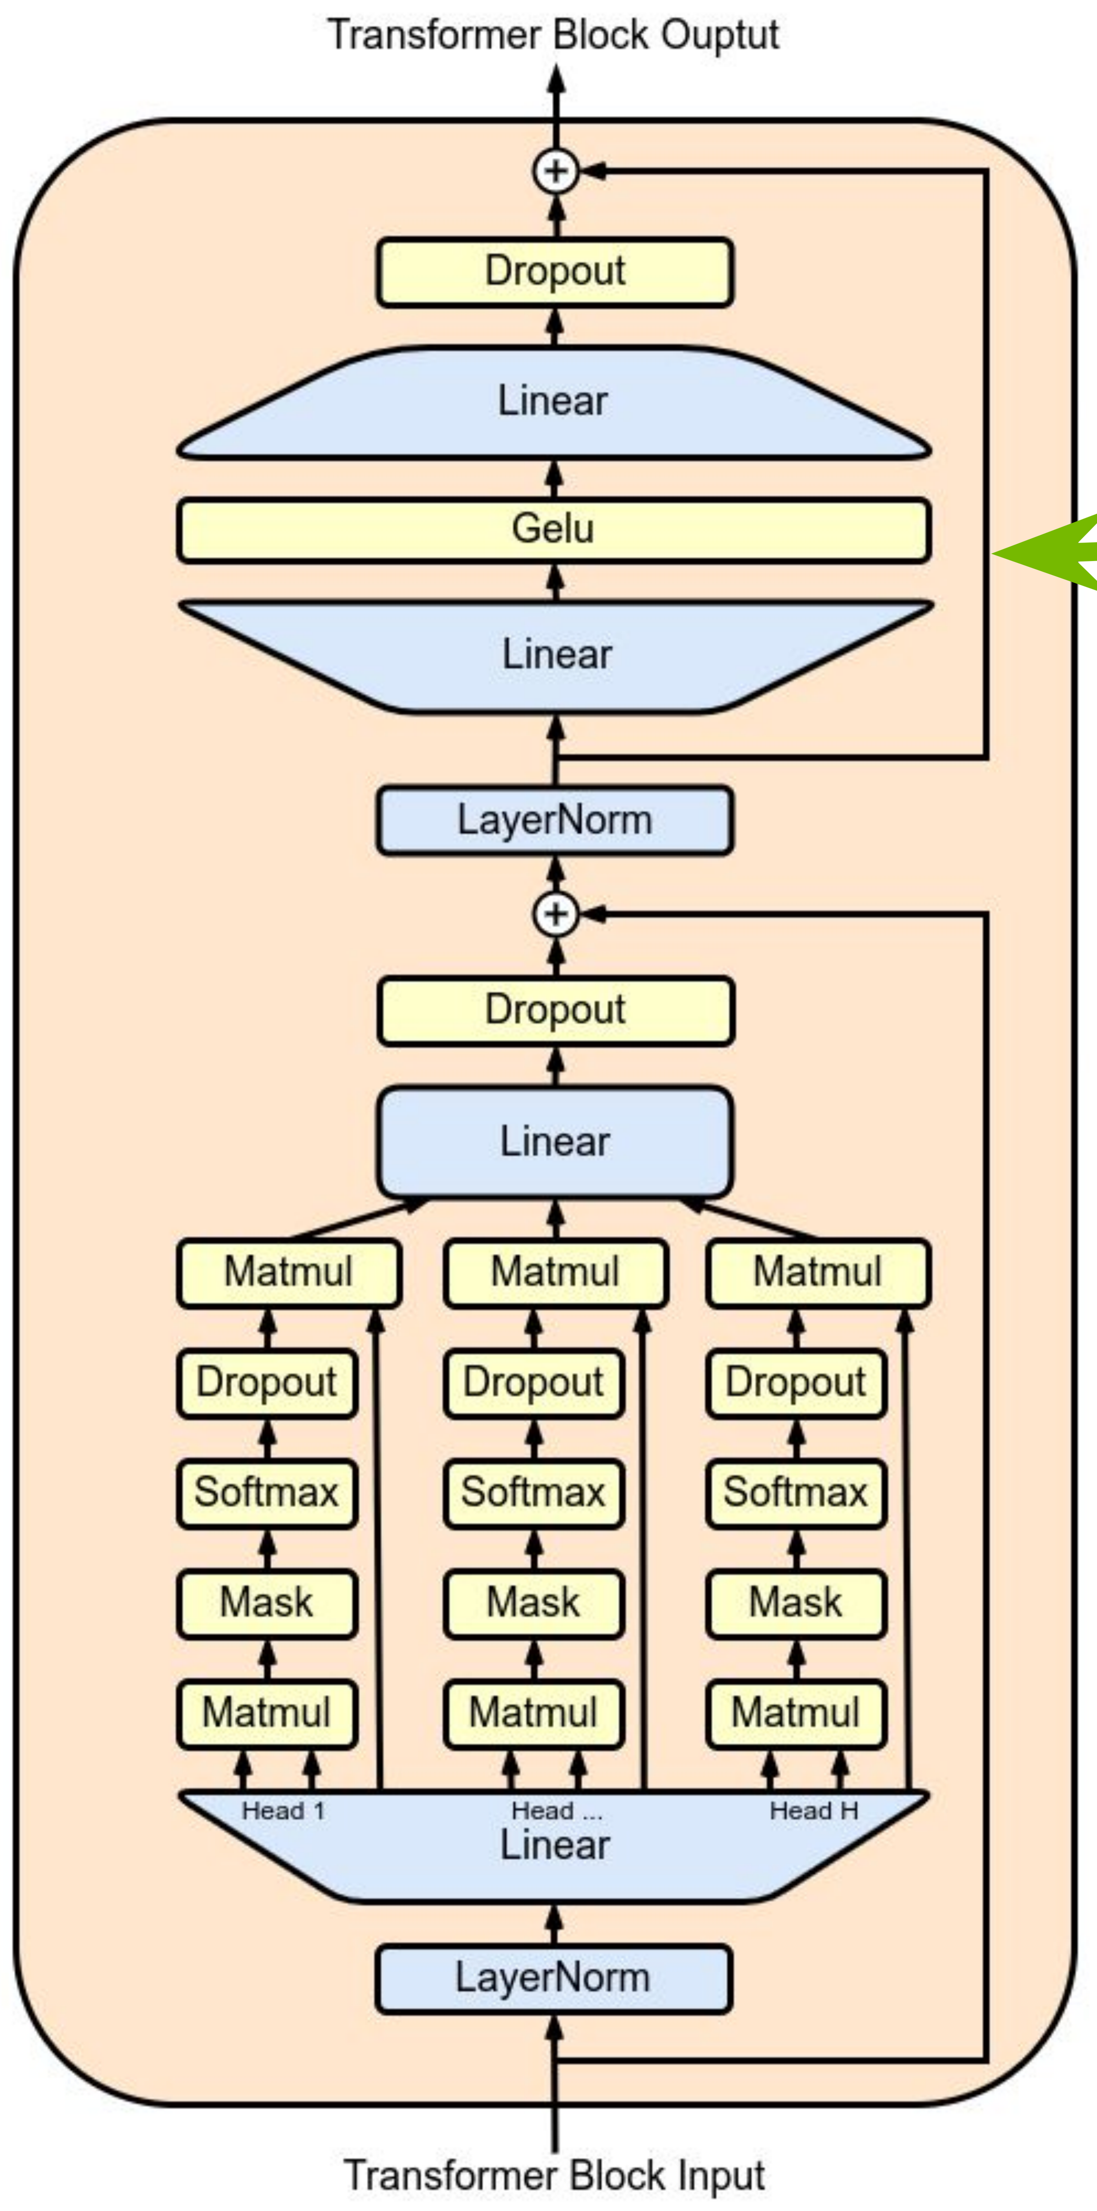

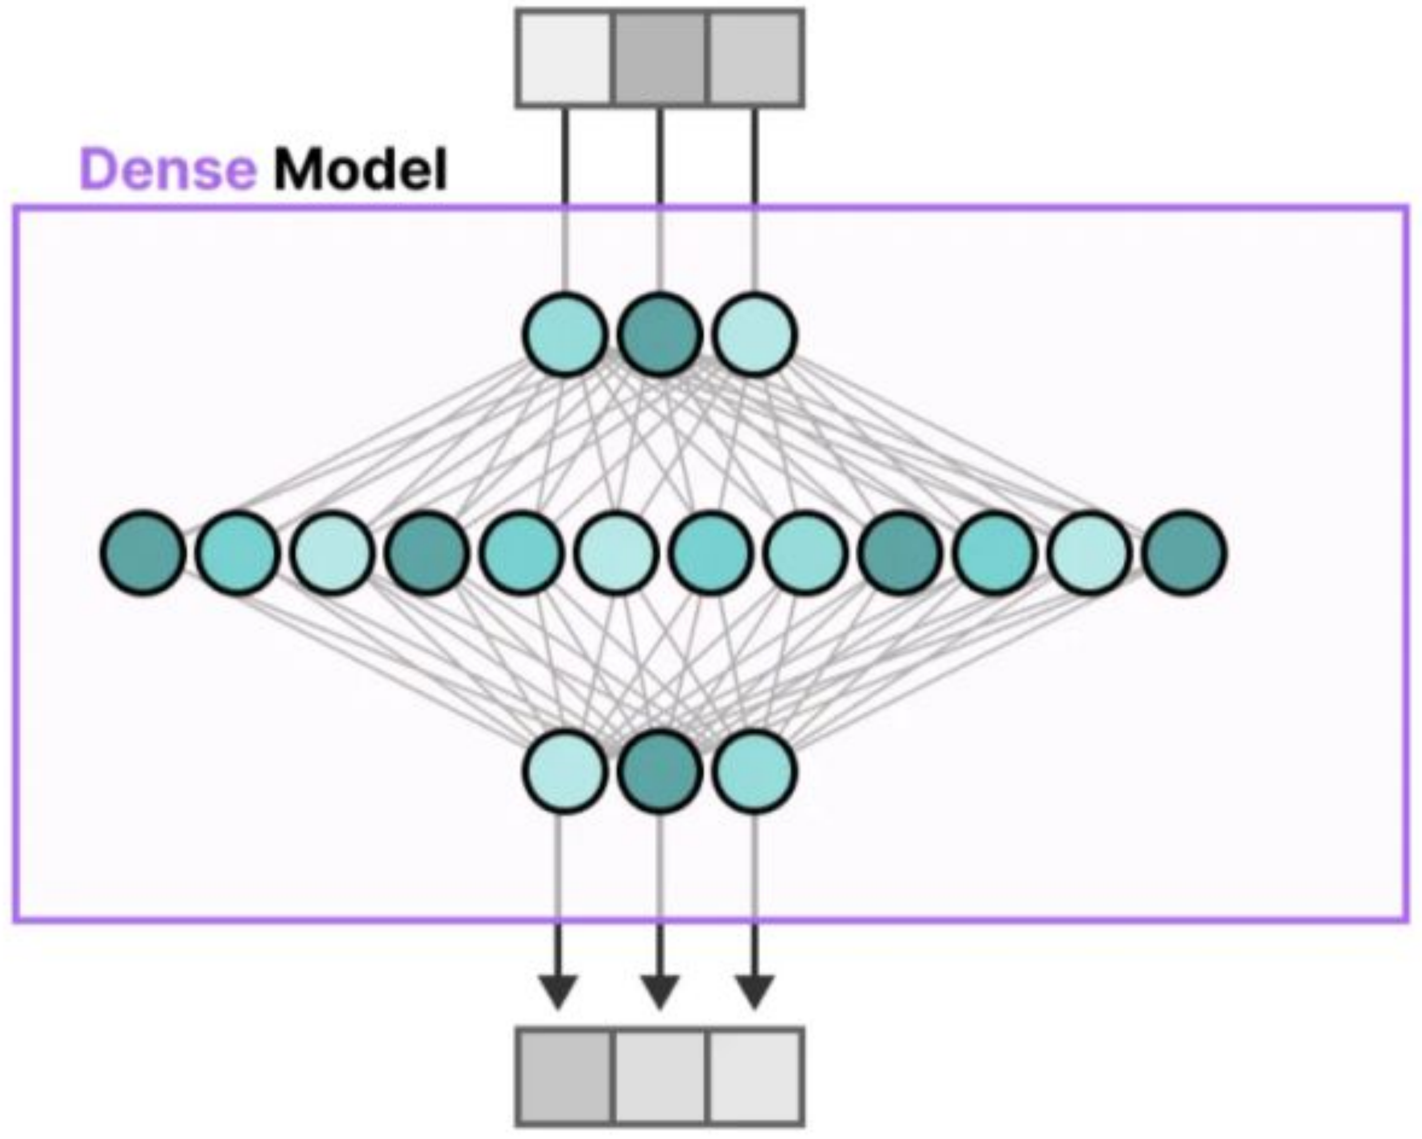

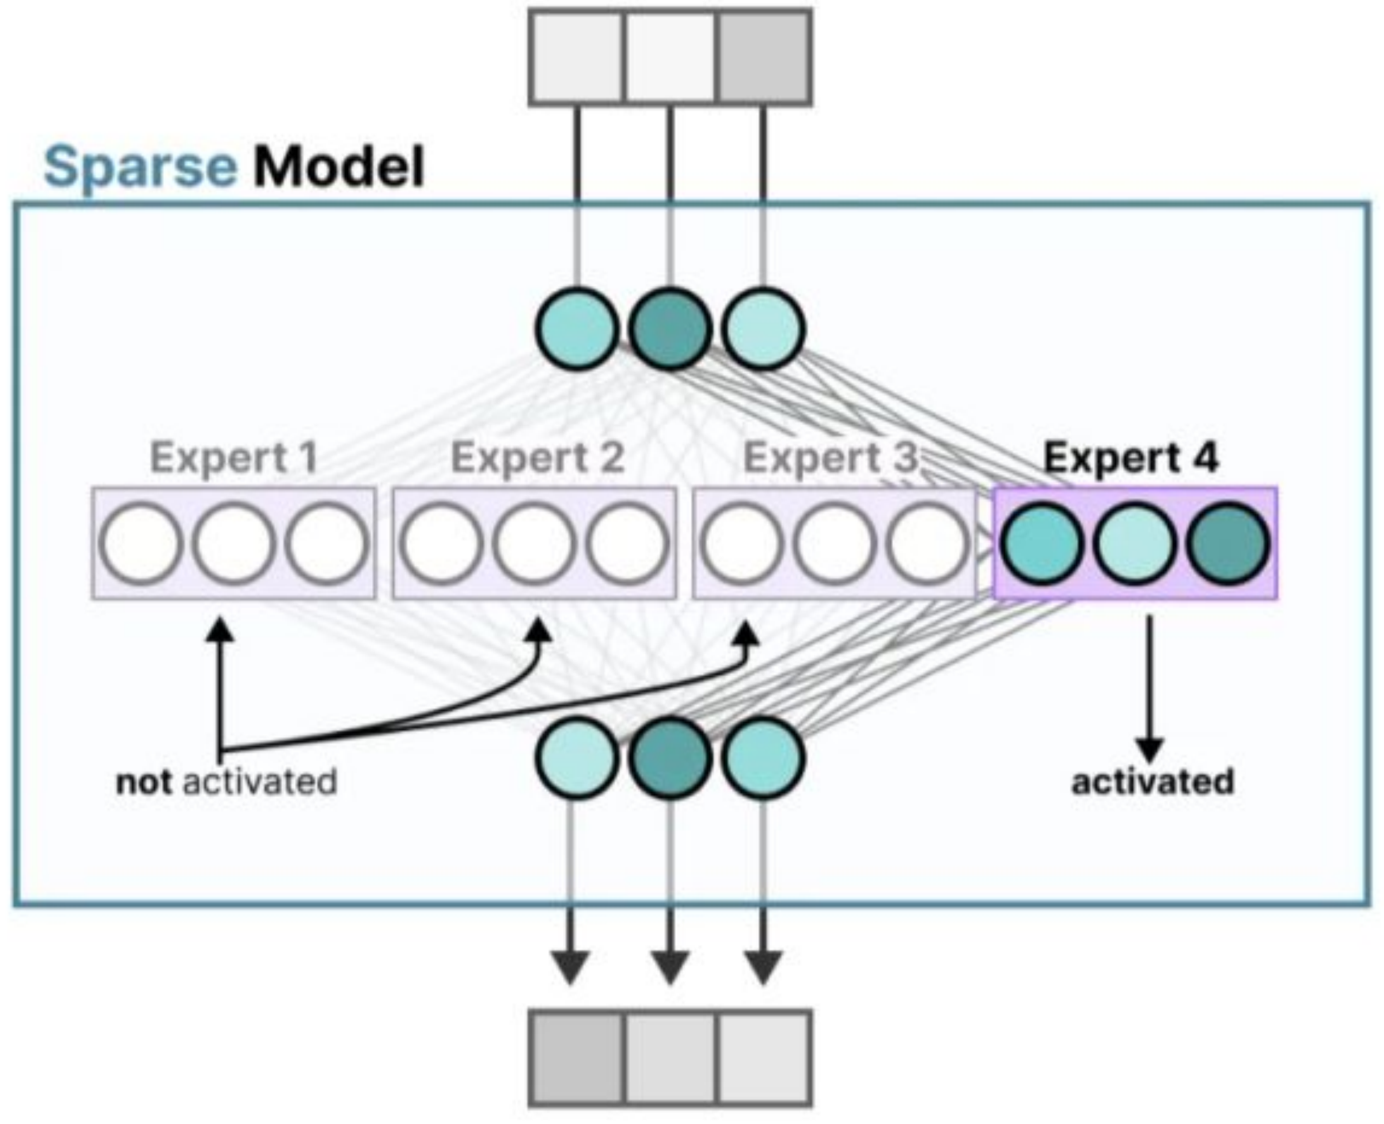

GPU Infrastructure

Modern Architecture - MoE(Mixture of Experts)

Gemma 4

updated 13 days ago

Our most intelligent open models to date

- **google/gemma-4-31B-it**
  Image-Text-to-Text • 31B • Updated Jul 21 • 8.28M • 3.69k

- **google/gemma-4-31B**
  Image-Text-to-Text • 33B • Updated Jul 16 • 717k • 517

- **google/gemma-4-26B-A4B-it**
  Image-Text-to-Text • 26B • Updated Jul 21 • 8.19M • 1.47k

- **google/gemma-4-26B-A4B** → **26B Total Parameters, 4B Active Parameters**
  Image-Text-to-Text • 27B • Updated Jul 16 • 80.5k • 392

GPU Infrastructure

How to Calculate the Computational Cost of a Model

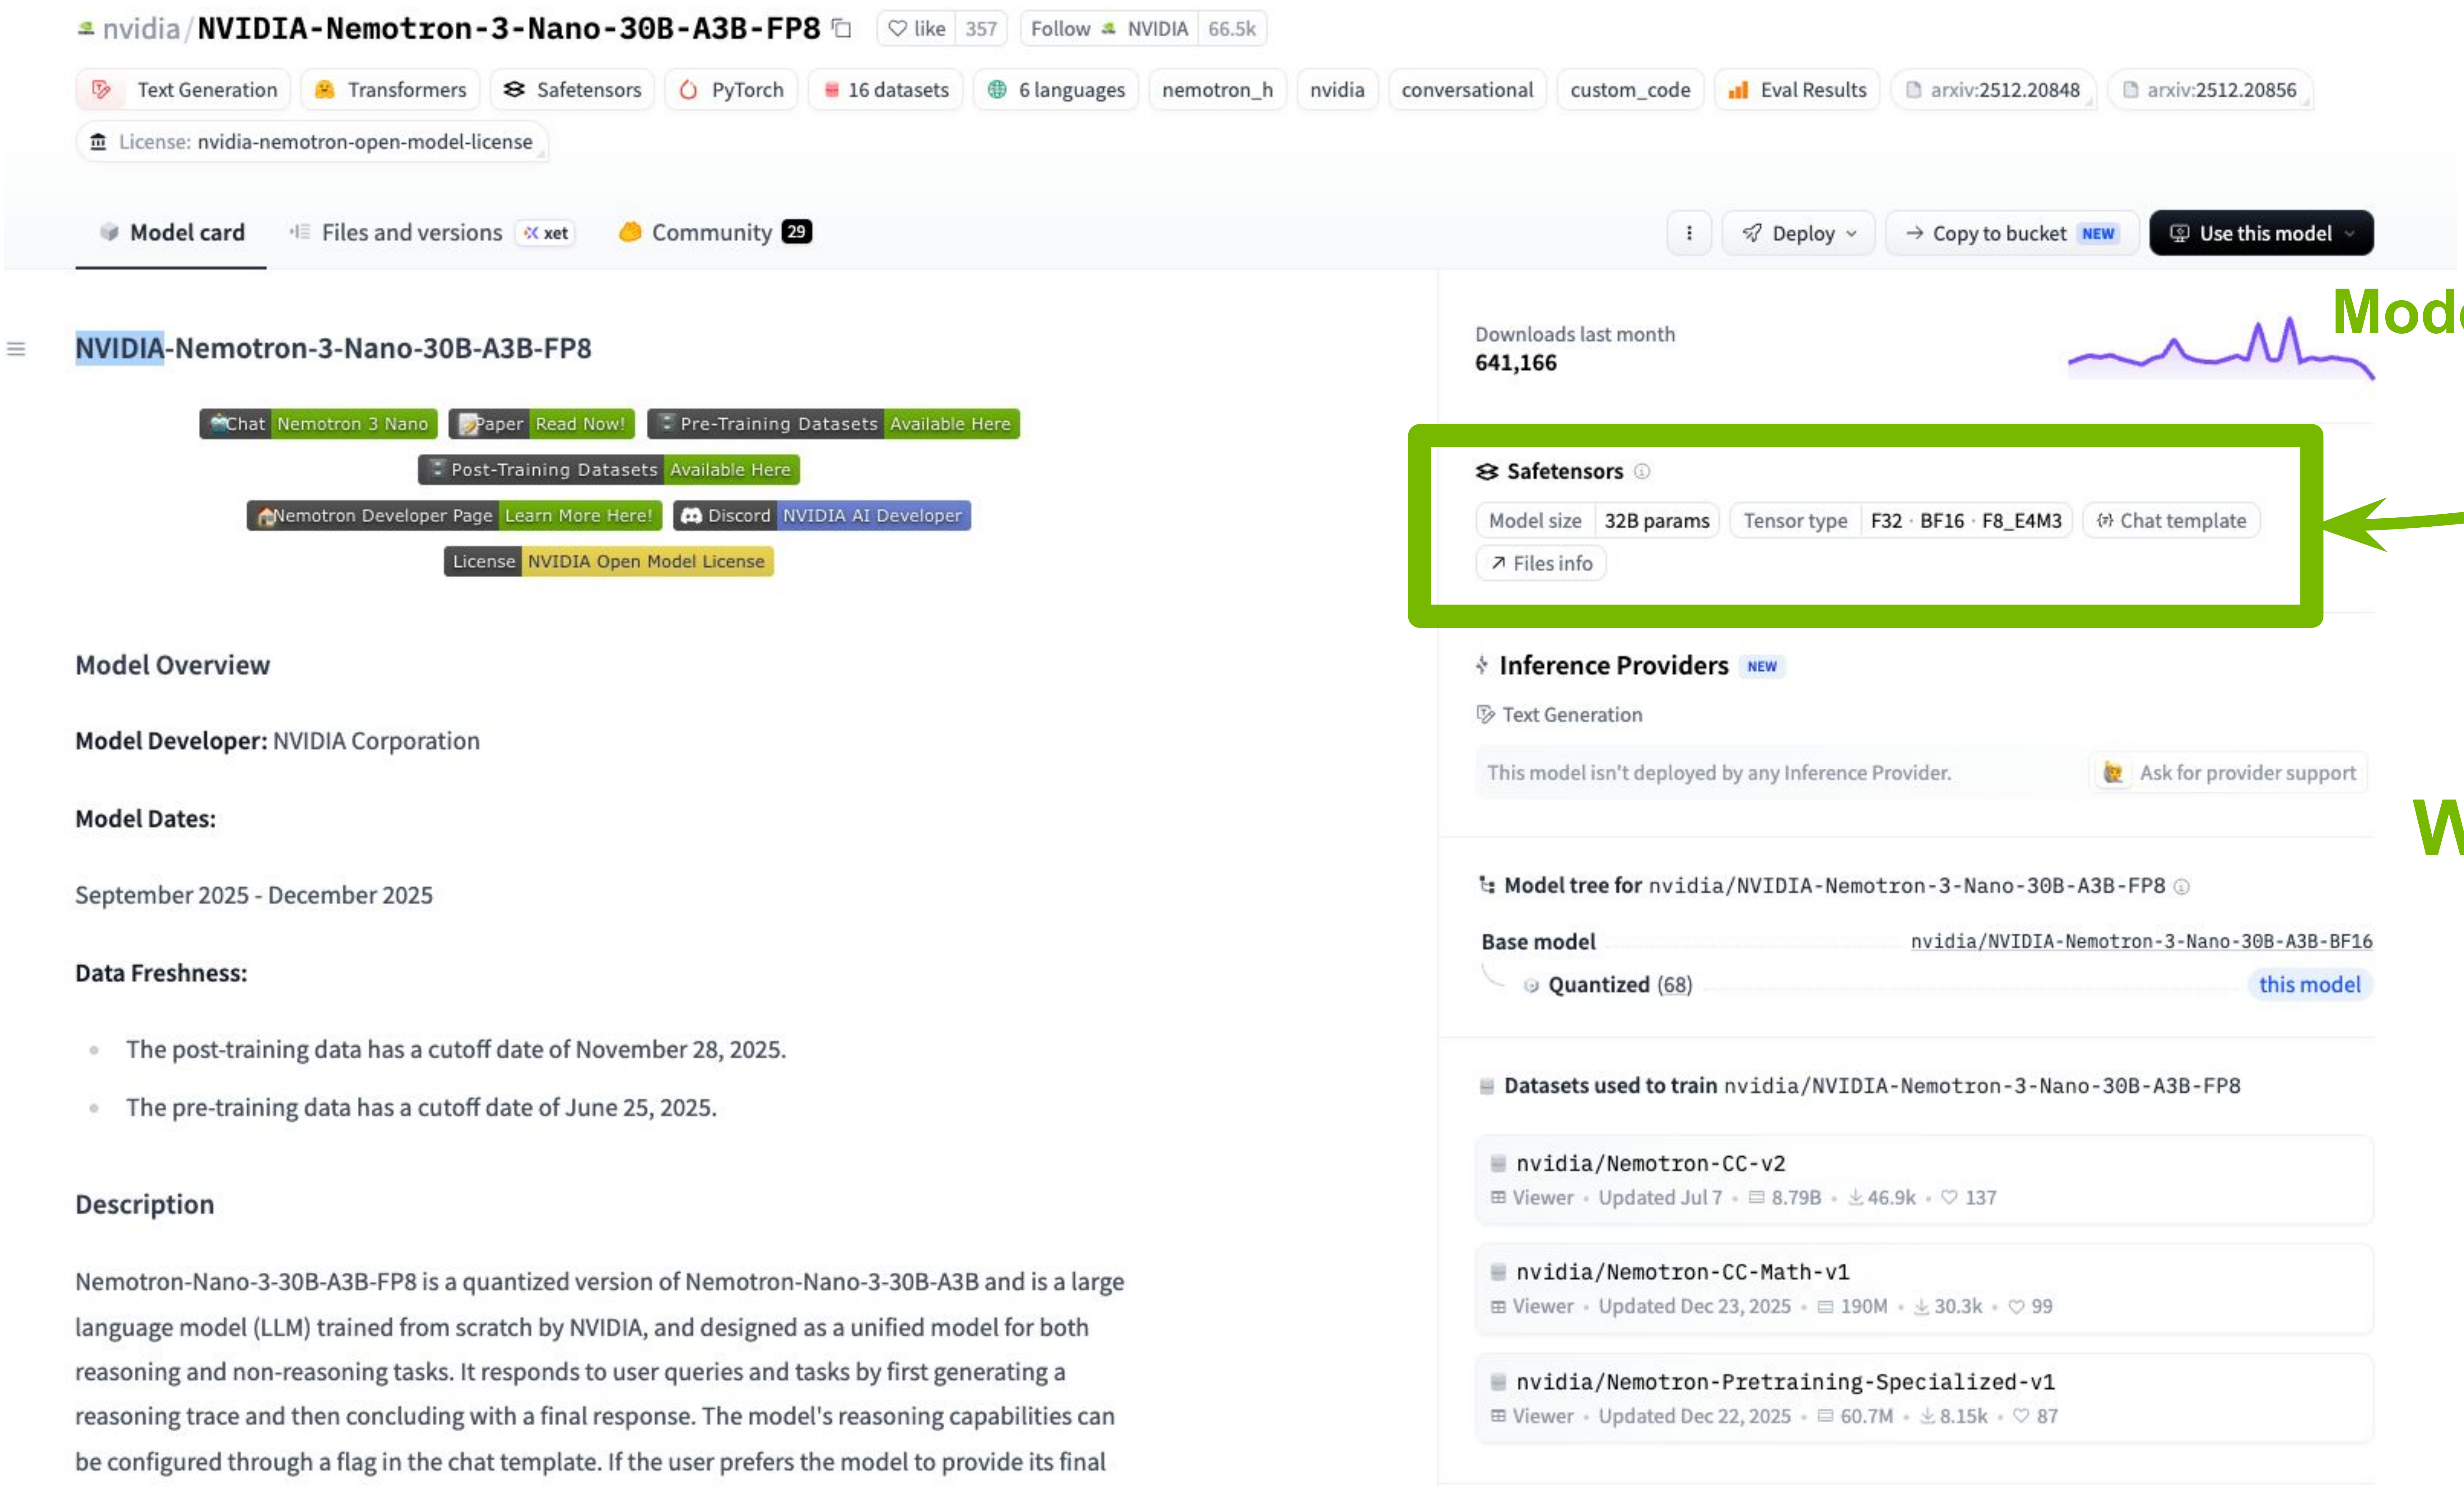

Model Parameter Size : B (Billion)

What is the Tensor Type?

https://huggingface.co/nvidia/NVIDIA-Nemotron-3-Nano-30B-A3B-FP8

GPU Infrastructure

Tensor Type : Bytes per Parameter

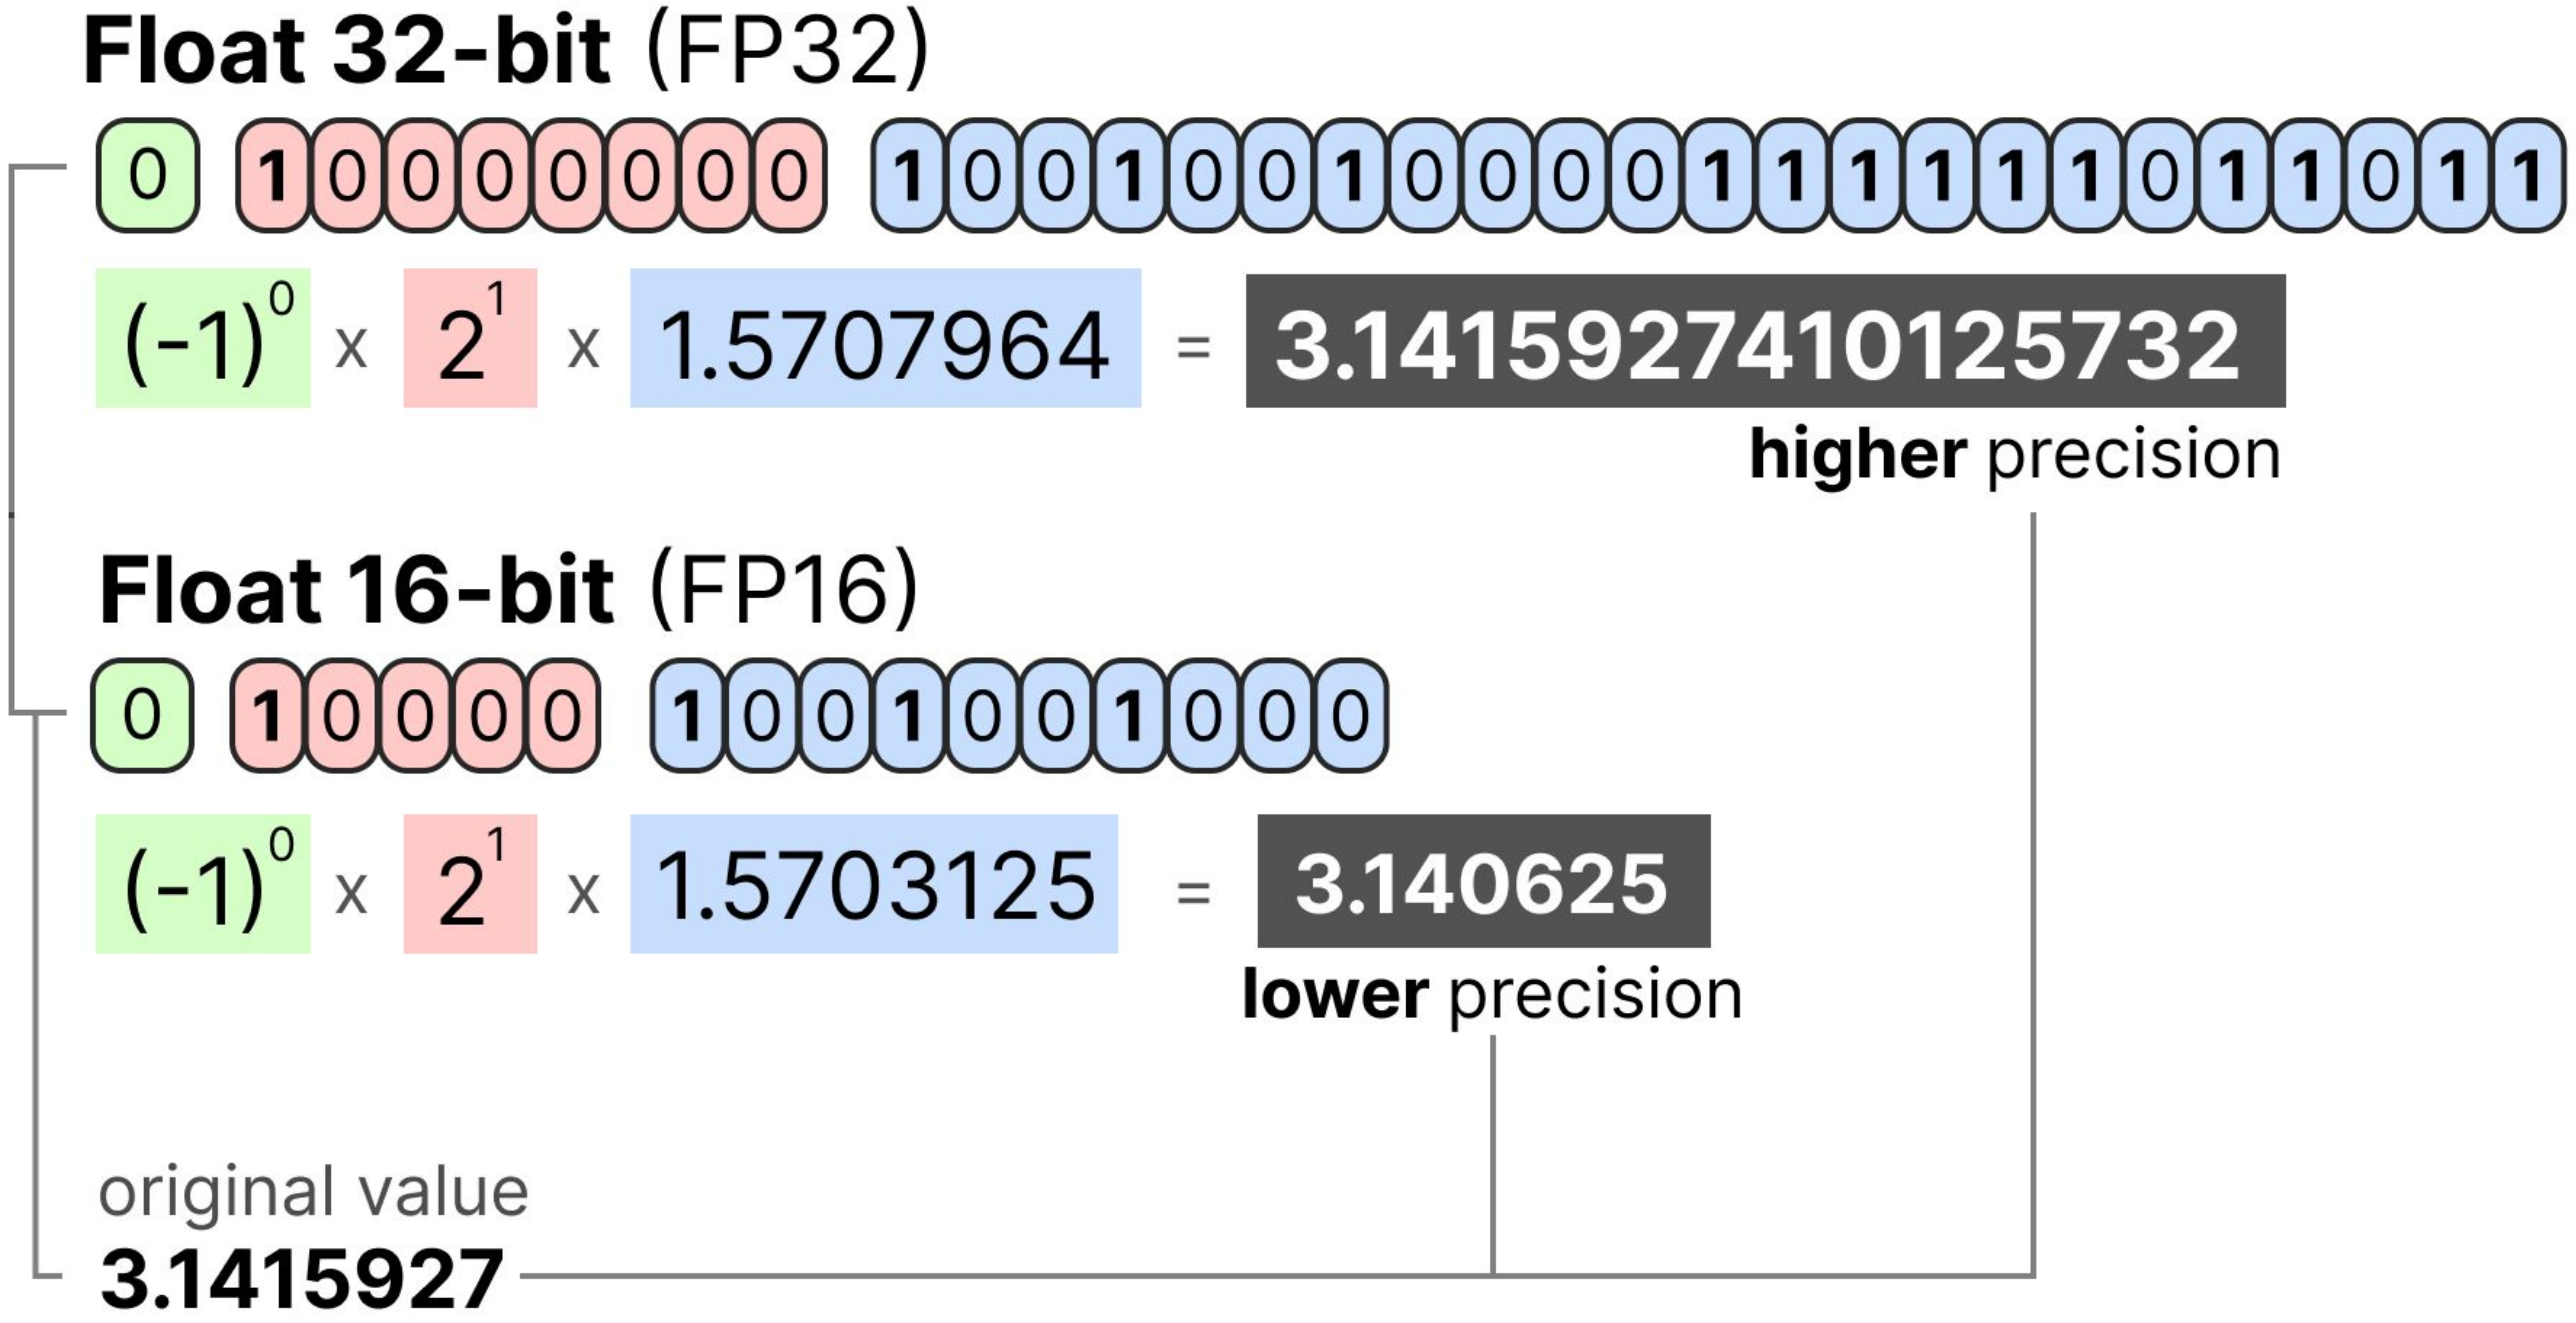

https://newsletter.maartengrootendorst.com/p/a-visual-guide-to-quantization

GPU Infrastructure

Tensor Type : Bytes per Parameter

Data Type	Total Bits	Representable Range
FP32	32	$\pm 3.4 \times 10^{38}$
FP16	16	±65,504
BF16	16	$\pm 3.4 \times 10^{38}$
FP8	8	±448
FP4	4	-6 to +6
INT8	8	-128 to +127
INT4	4	-8 to +7

Table 1. Summary of data types with bit width, representable ranges, and format descriptions (floating point vs. integer)

https://developer.nvidia.com/blog/optimizing-llms-for-performance-and-accuracy-with-post-training-quantization/

GPU Infrastructure

Tensor Type : Bytes per Parameter

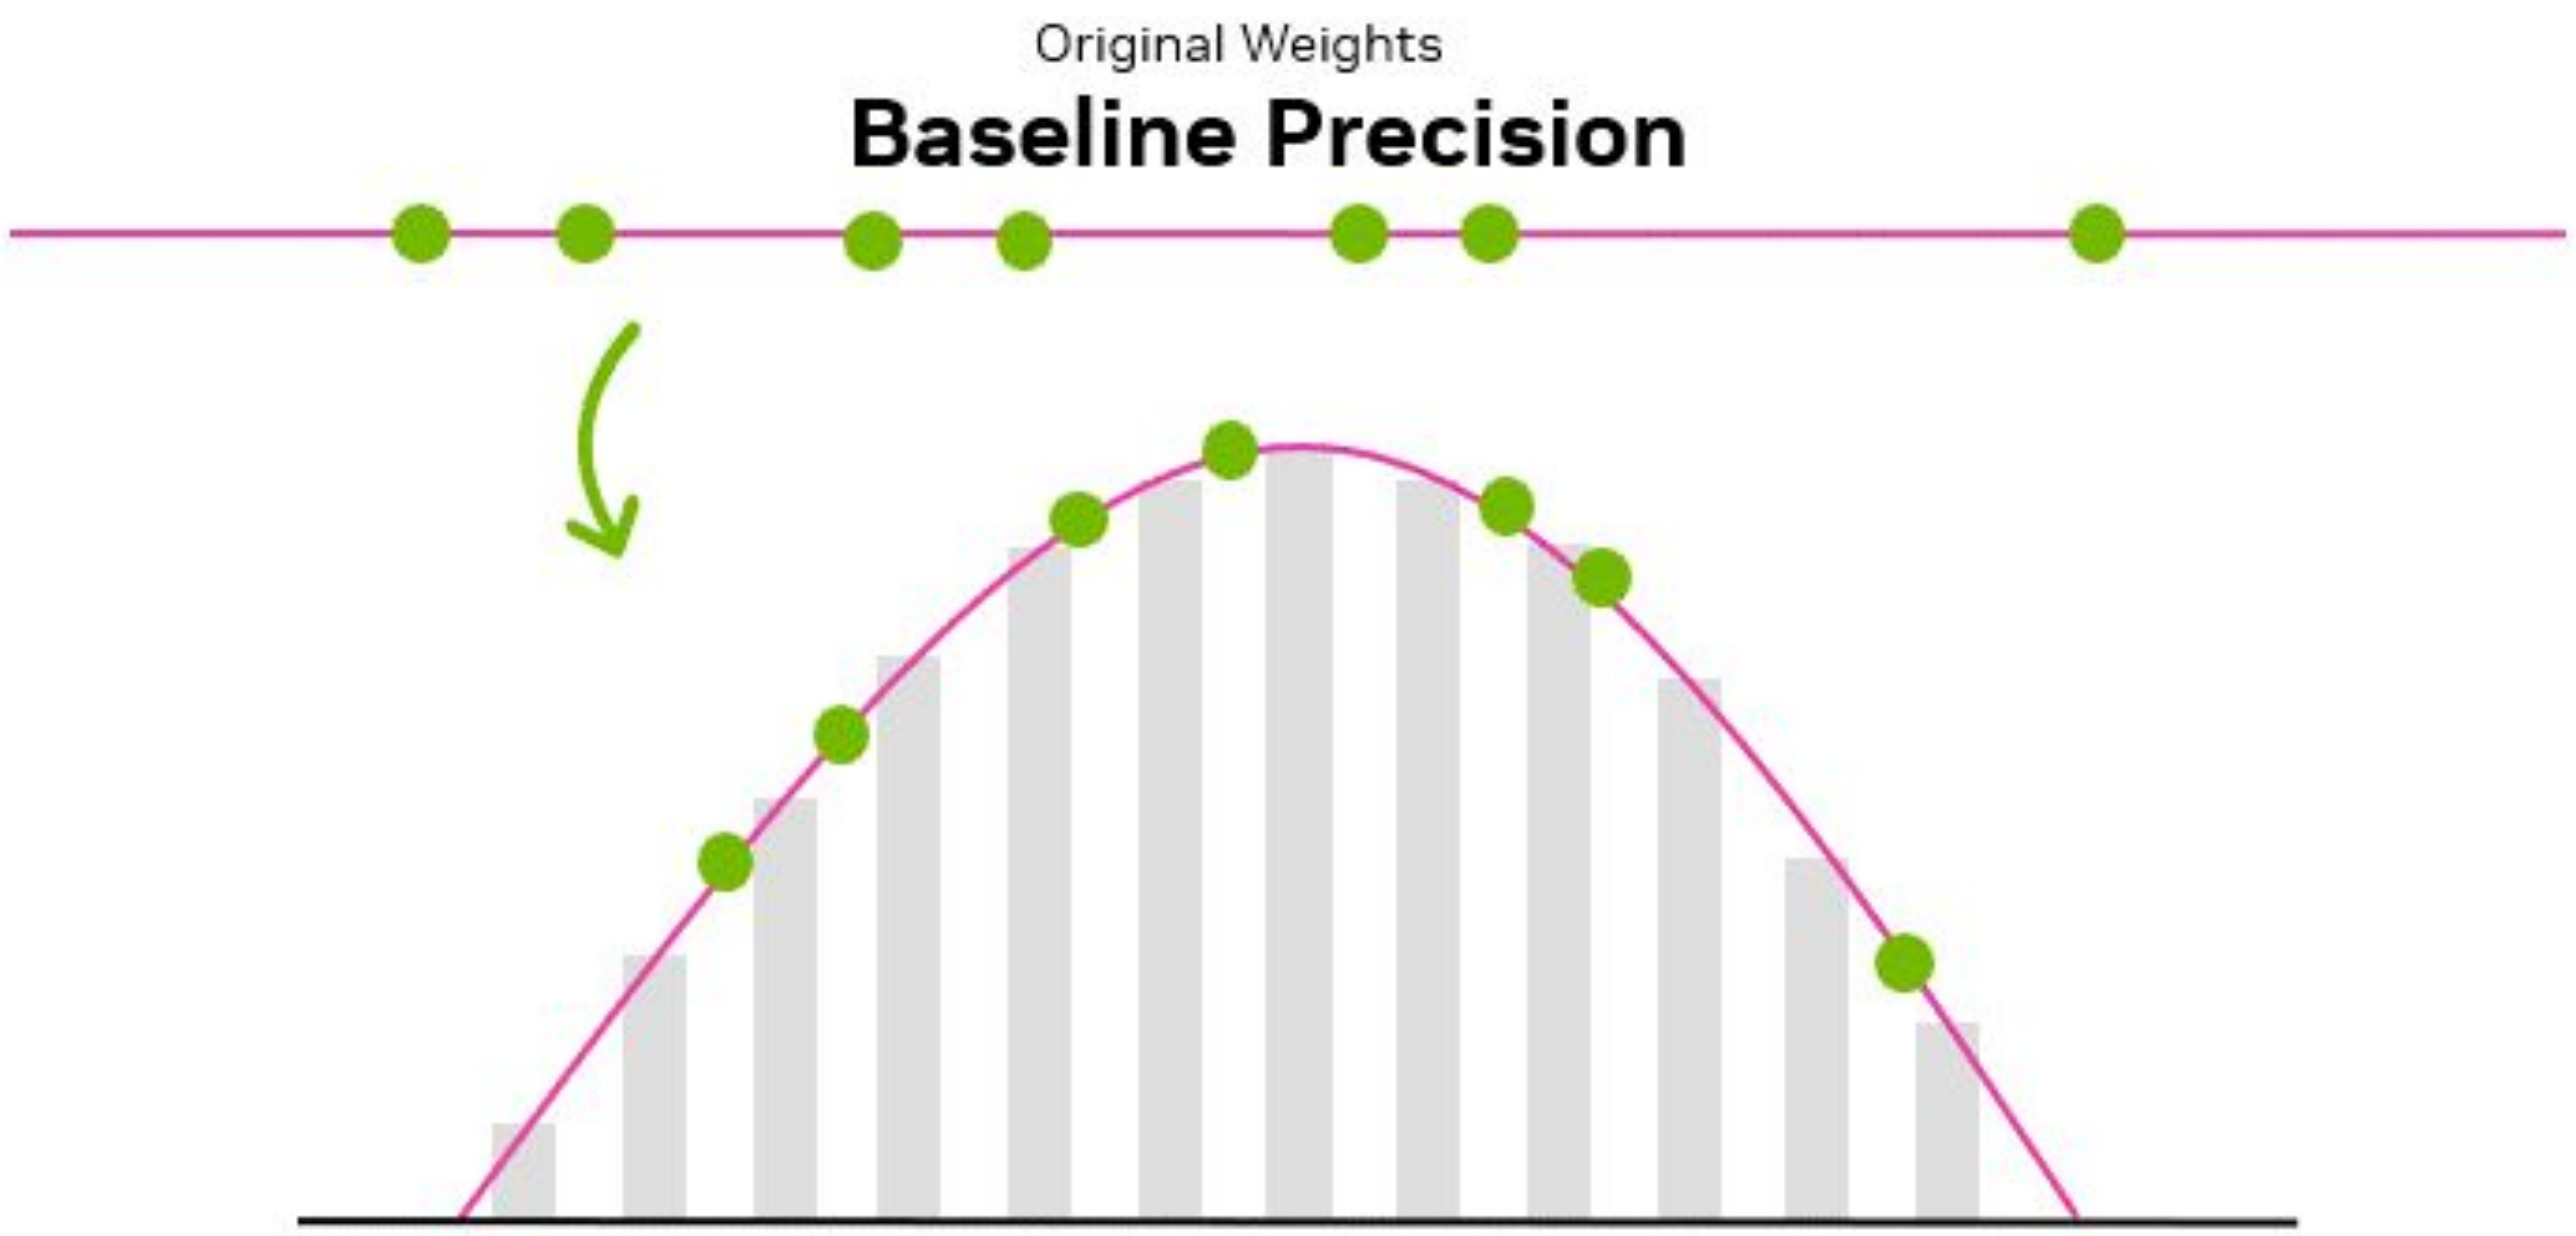

GPU Infrastructure

Tensor Type : Bytes per Parameter

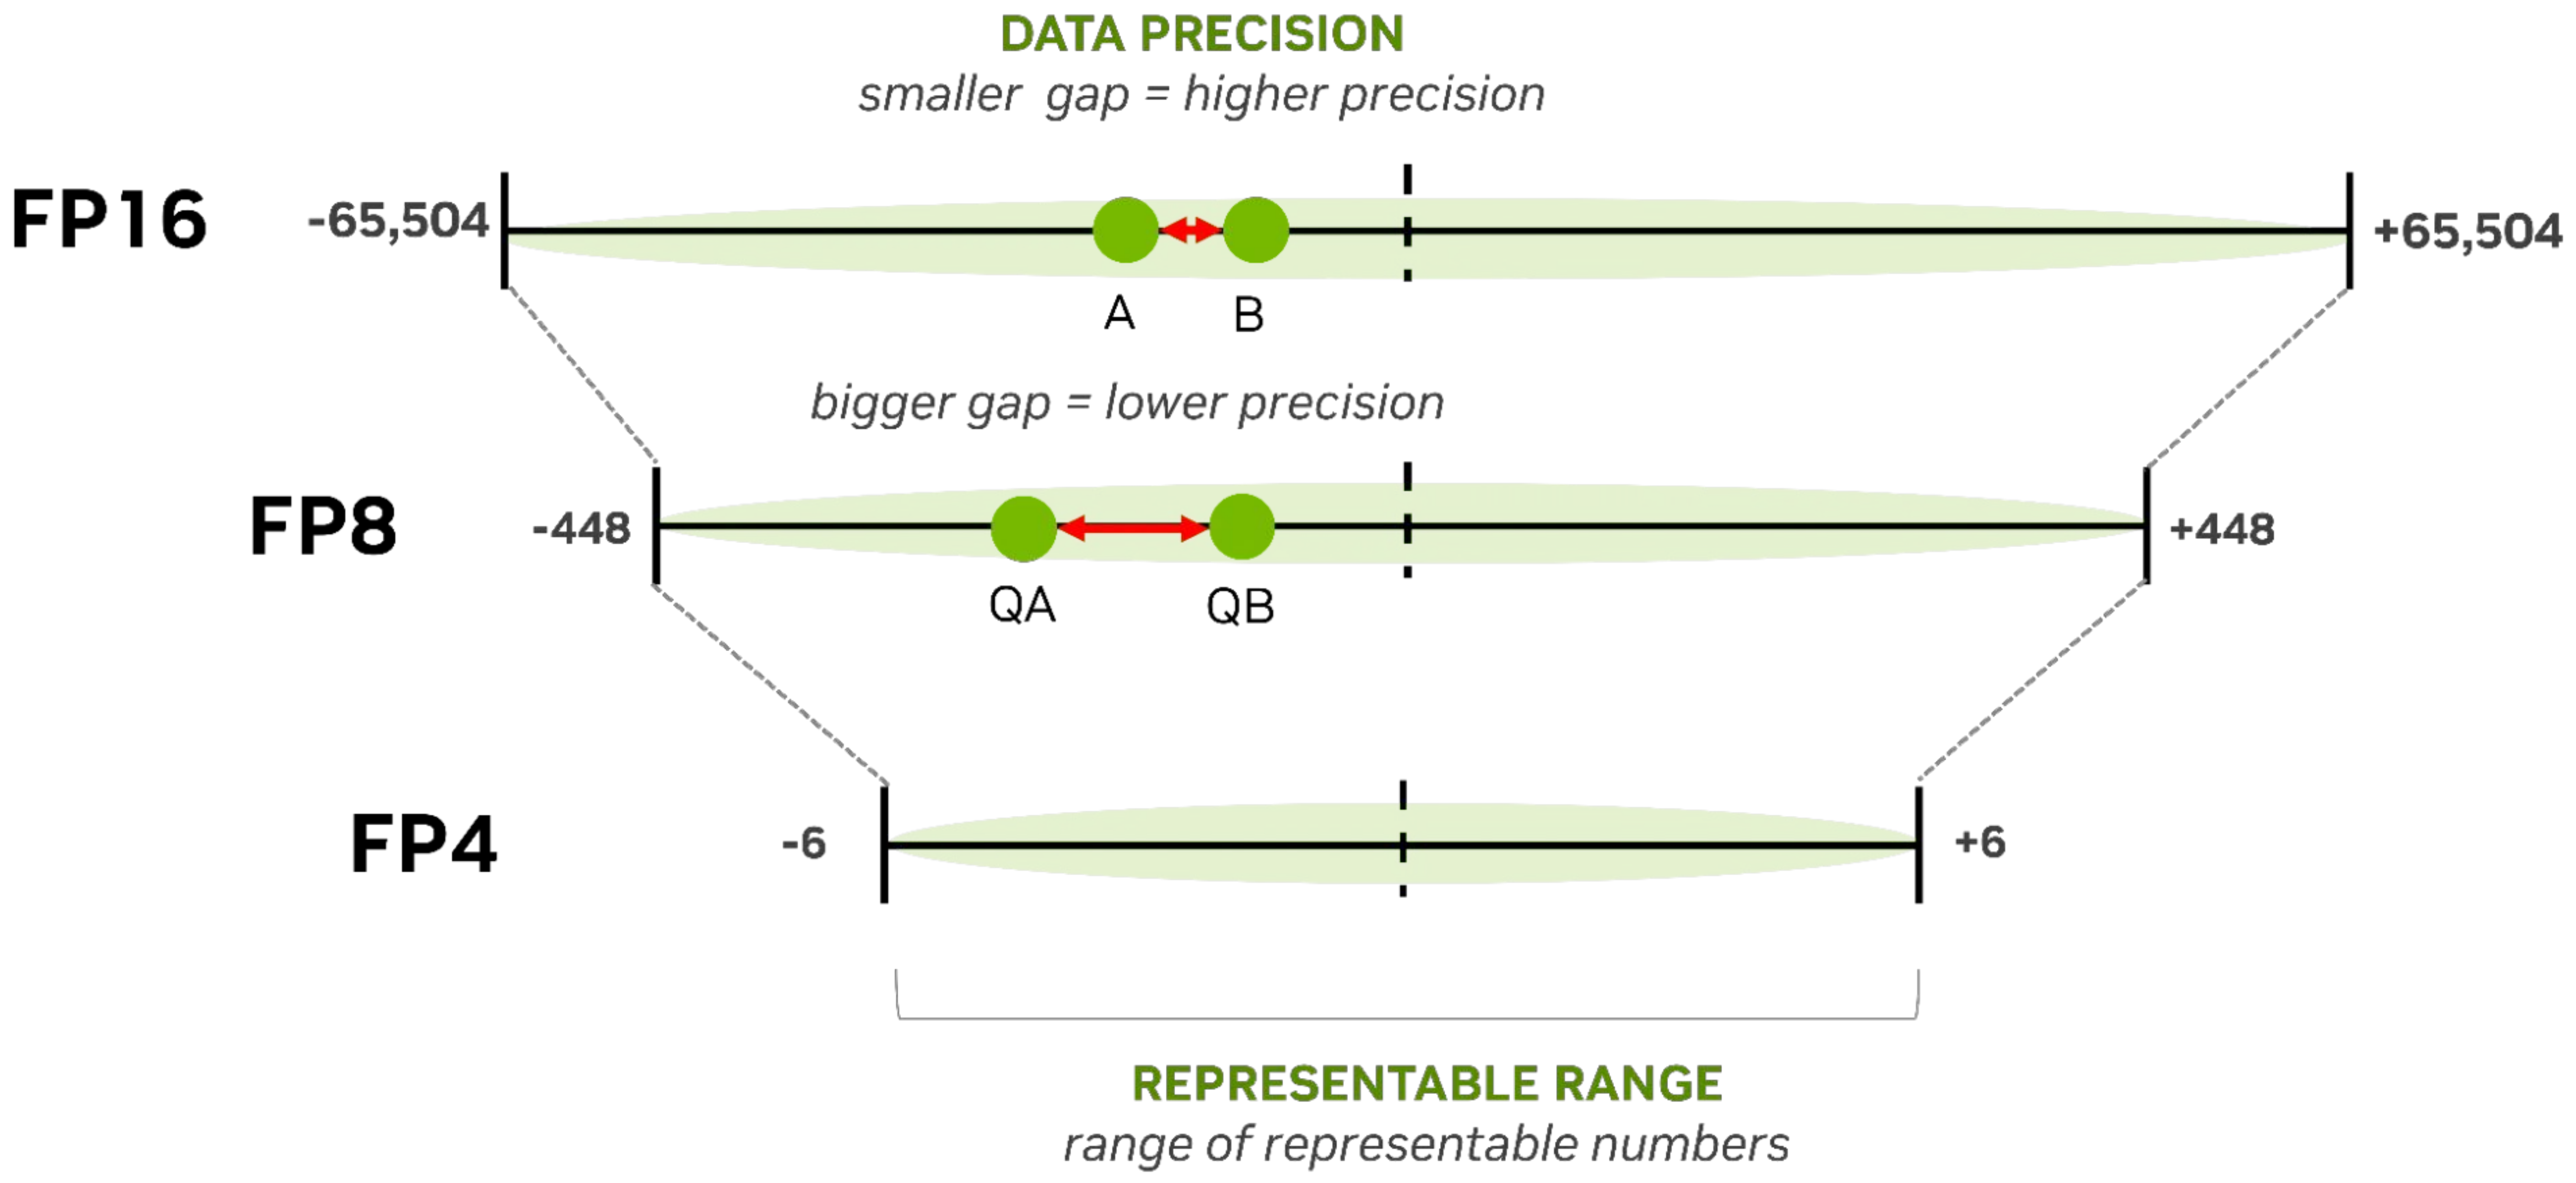

https://developer.nvidia.com/blog/optimizing-llms-for-performance-and-accuracy-with-post-training-quantization/

GPU Infrastructure

Tensor Type : Bytes per Parameter

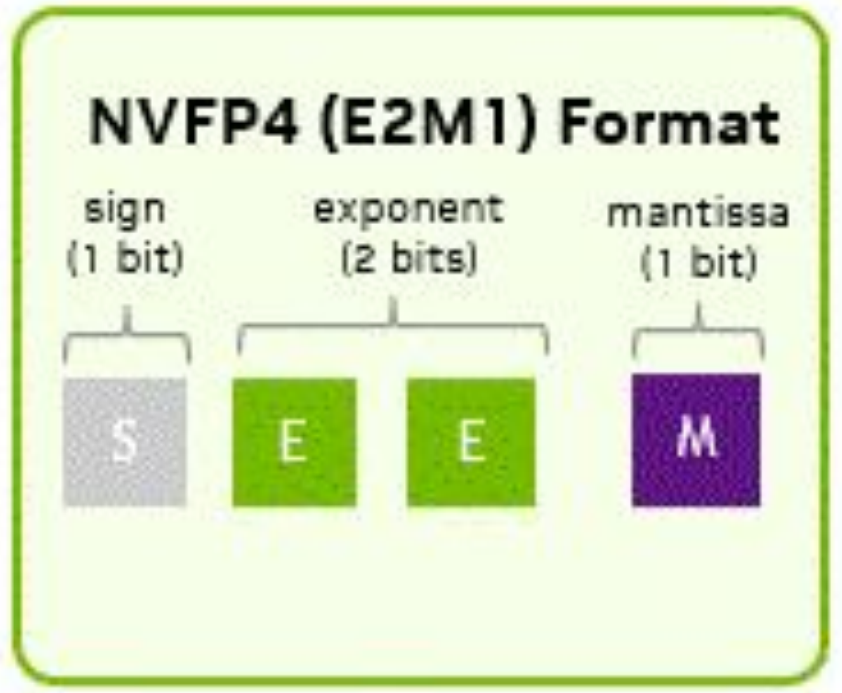

https://developer.nvidia.com/blog/optimizing-llms-for-performance-and-accuracy-with-post-training-quantization/

GPU Infrastructure

How to Calculate the Computational Cost of a Model

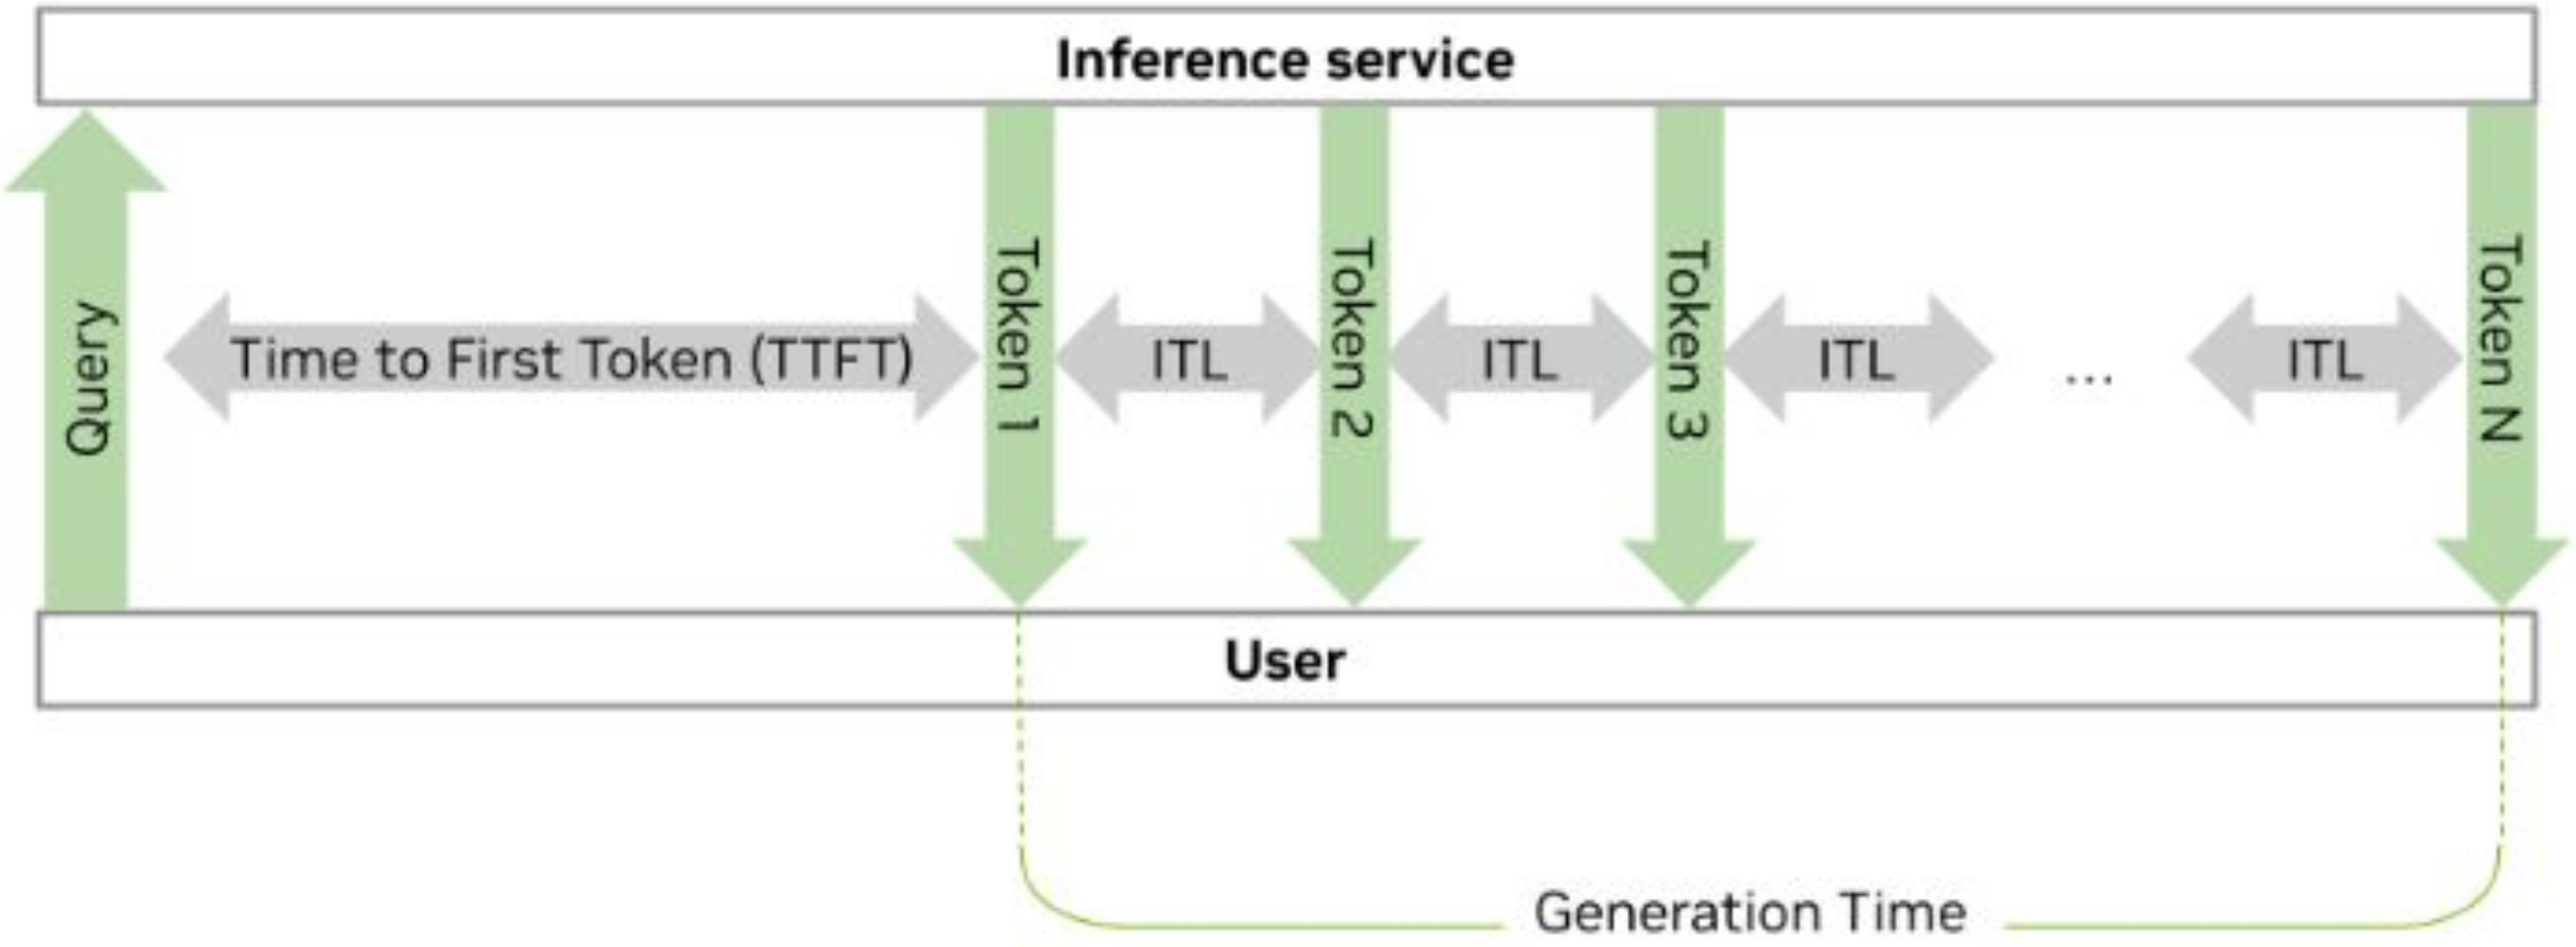

GPU Infrastructure

How to Calculate the Computational Cost of a Model

# of Parameter * Tensor Type Size = Total Model Size

ex) 3B Model * int8 = Total 3GB

KIMI

2.8T Model * fp32 = Total 11.2TB
(104B Activated → 416GB)

Can AI models run on a CPU?
→ Yes, but they're slow. Very, very, VERY slow.

GPU Infrastructure

How to Calculate the Computational Cost of a Model

https://www.youtube.com/watch?v=1BAZf3PsjWA

GPU Infrastructure

How to Calculate the Computational Cost of a Model

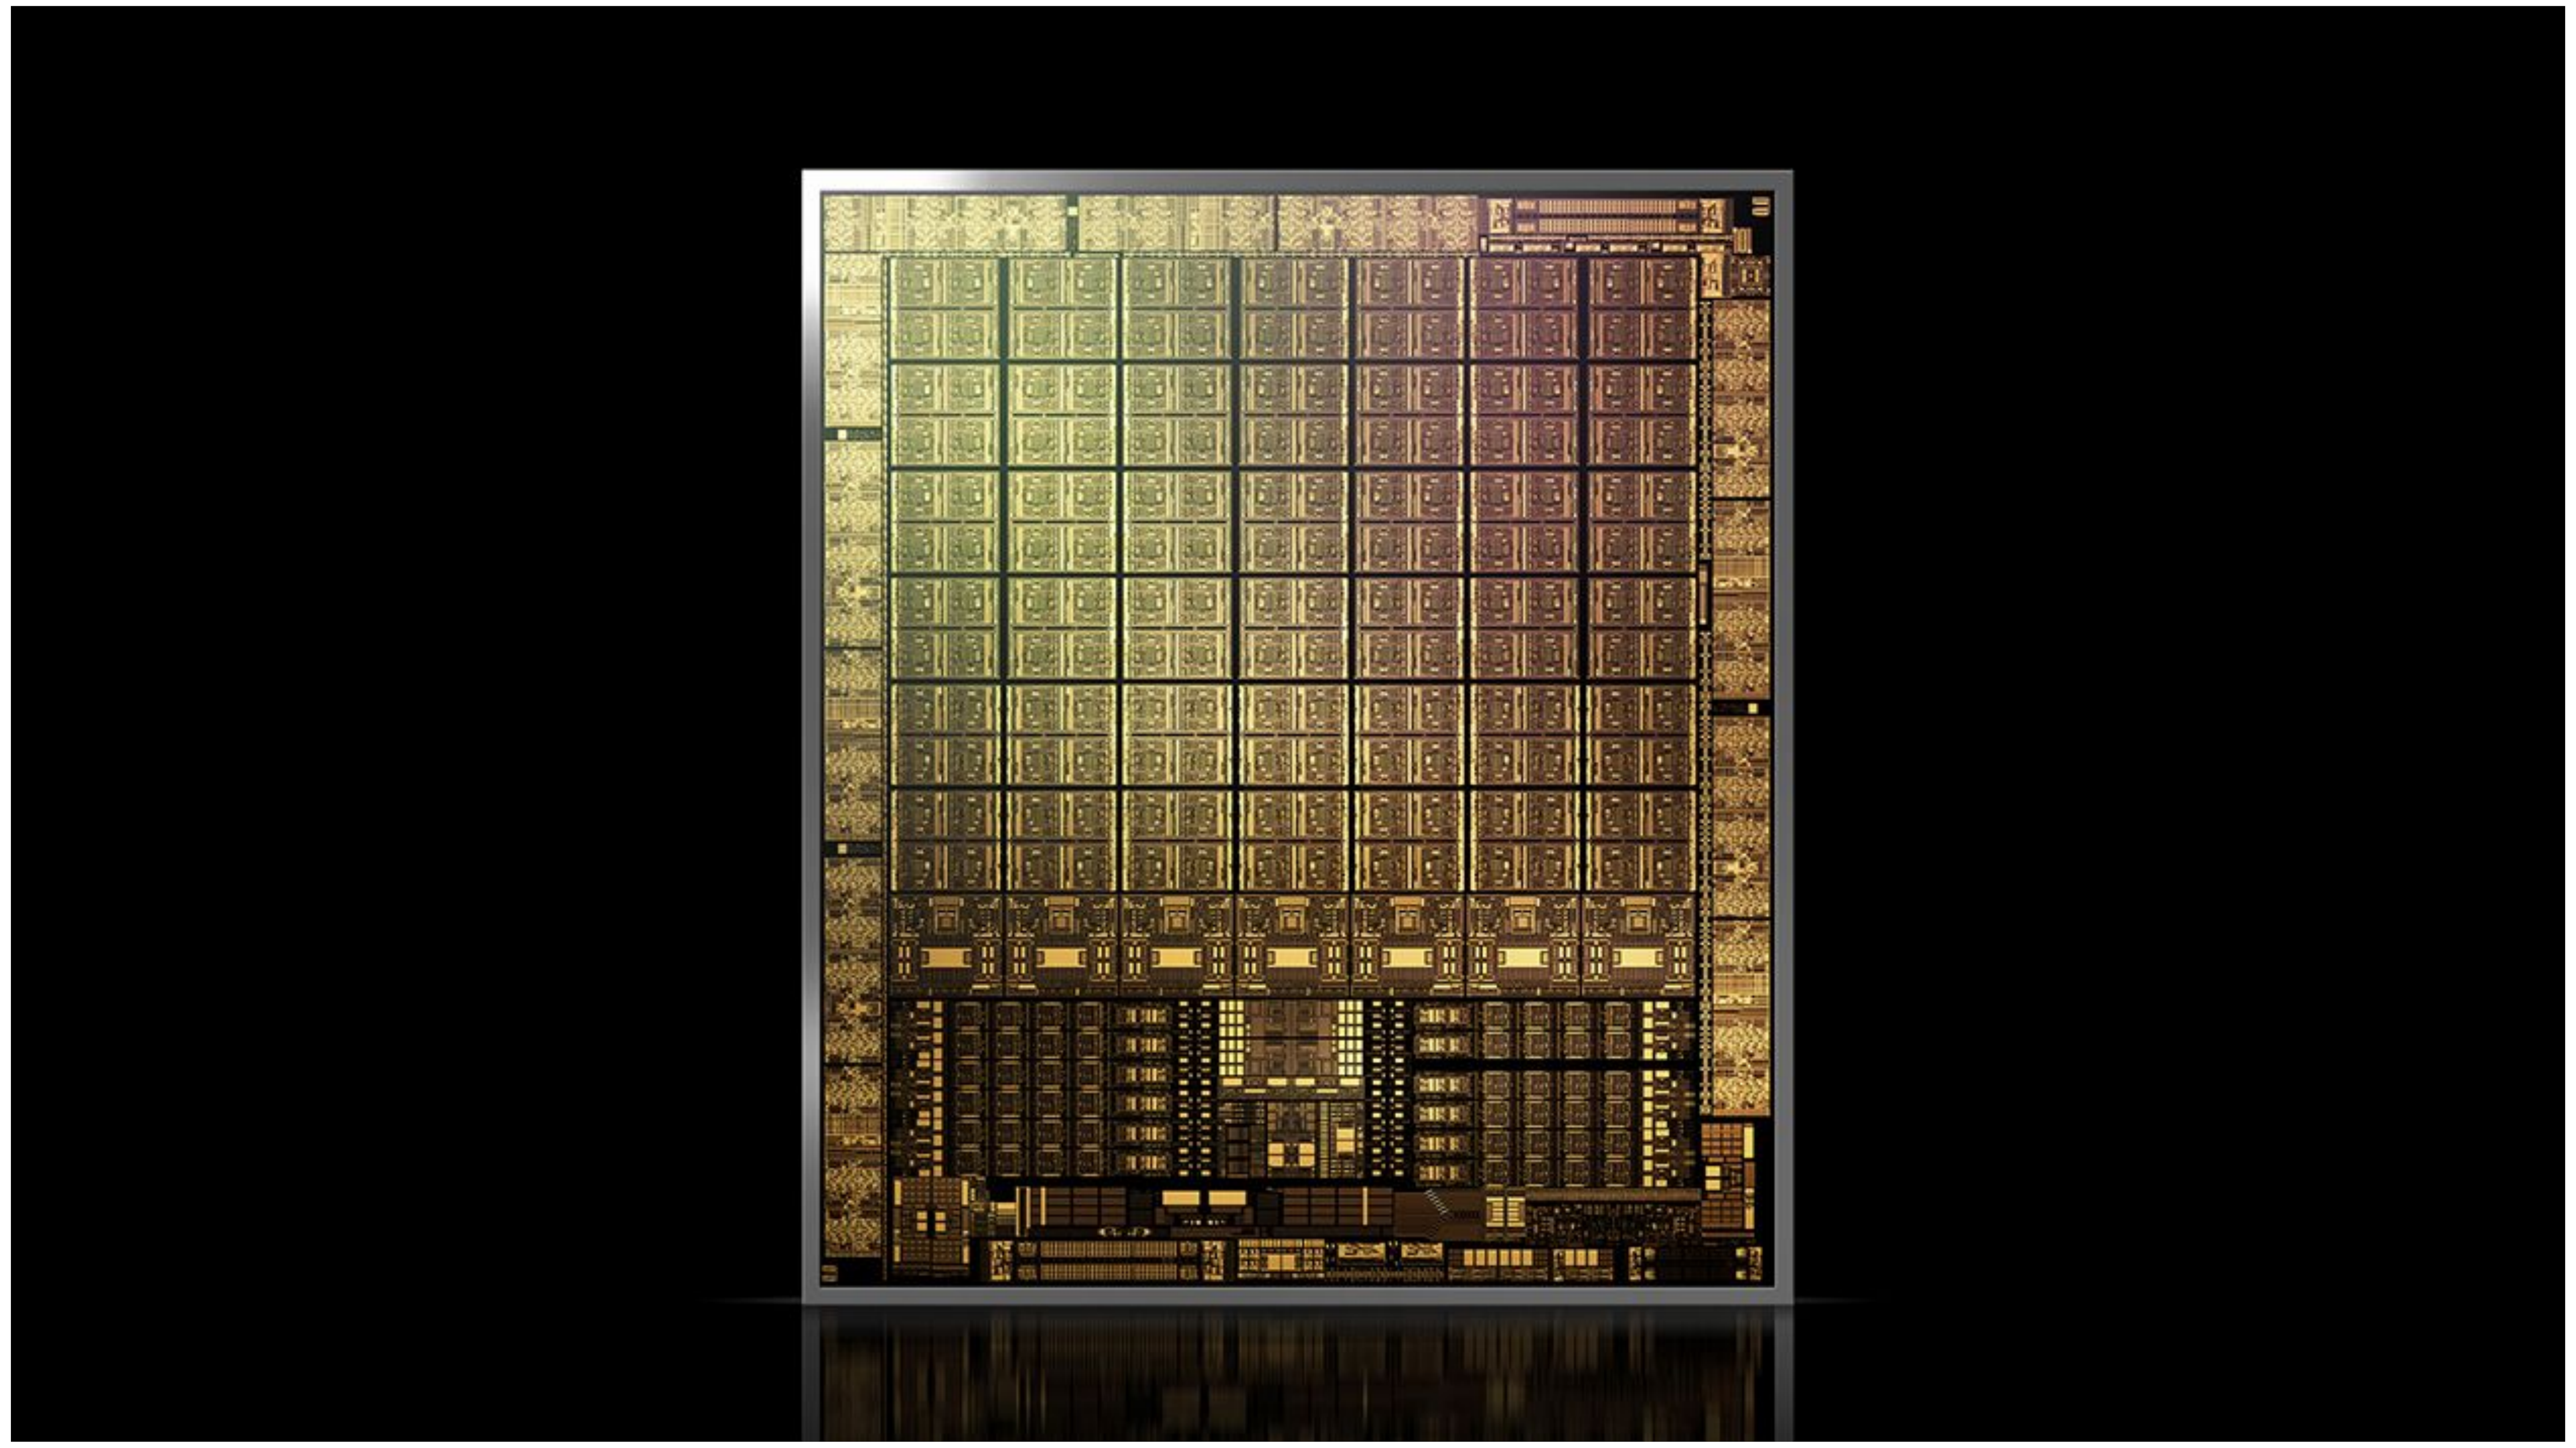

A Powerful Processing Unit — Is It Sufficient?

Memory → GPU
Data Transfer
⇒ New Bottleneck!

GPU Infrastructure

How to Calculate the Computational Cost of a Model

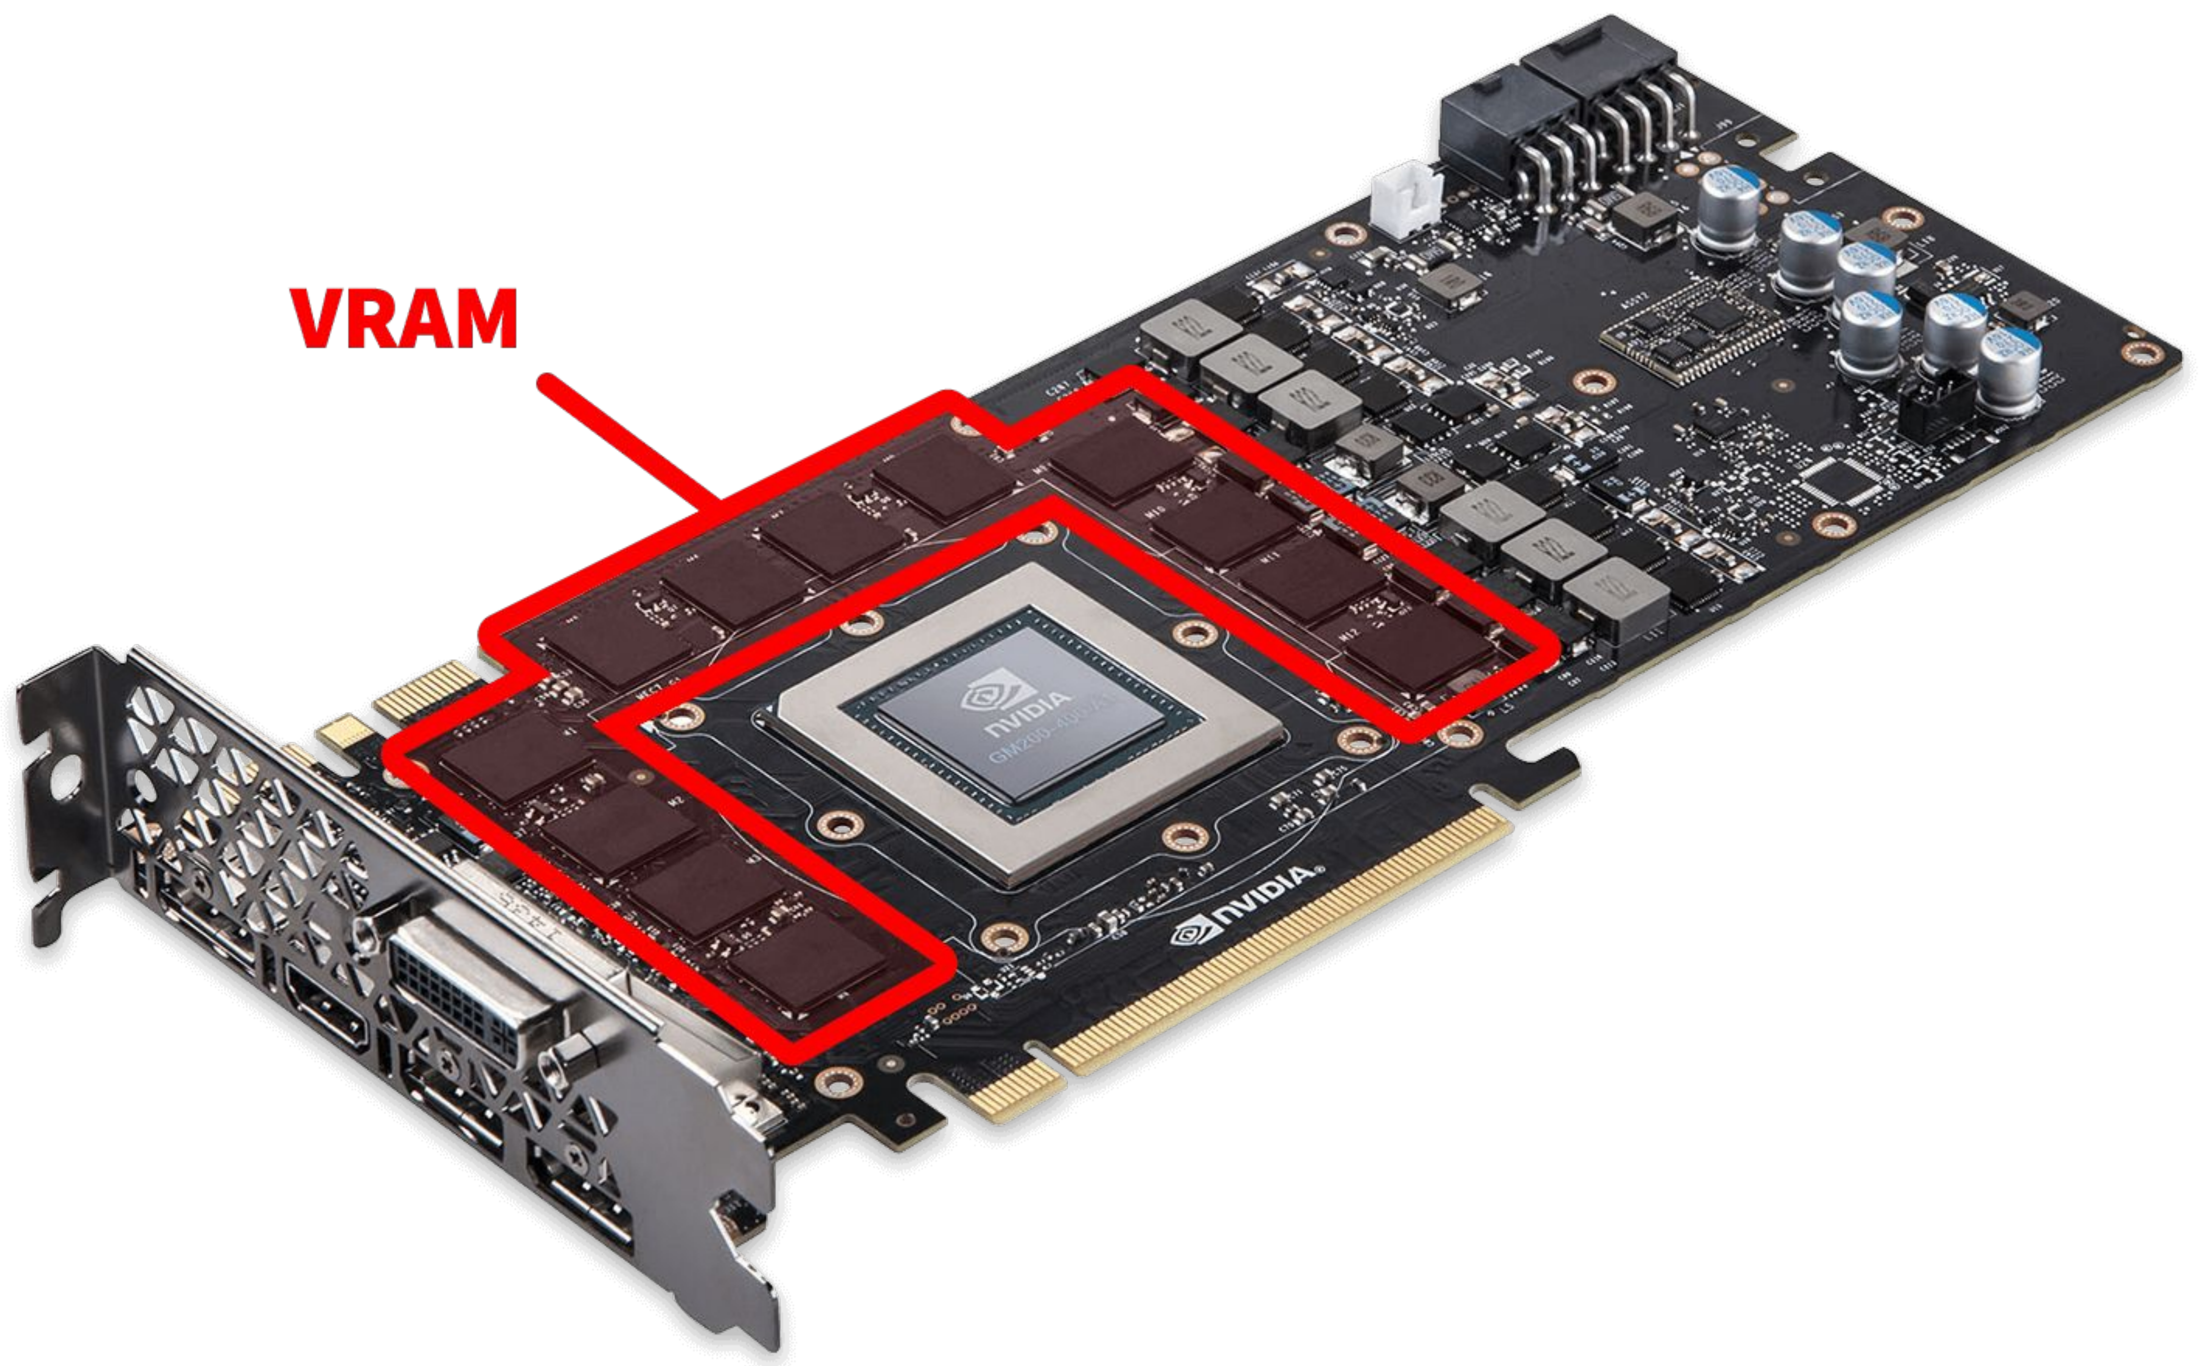

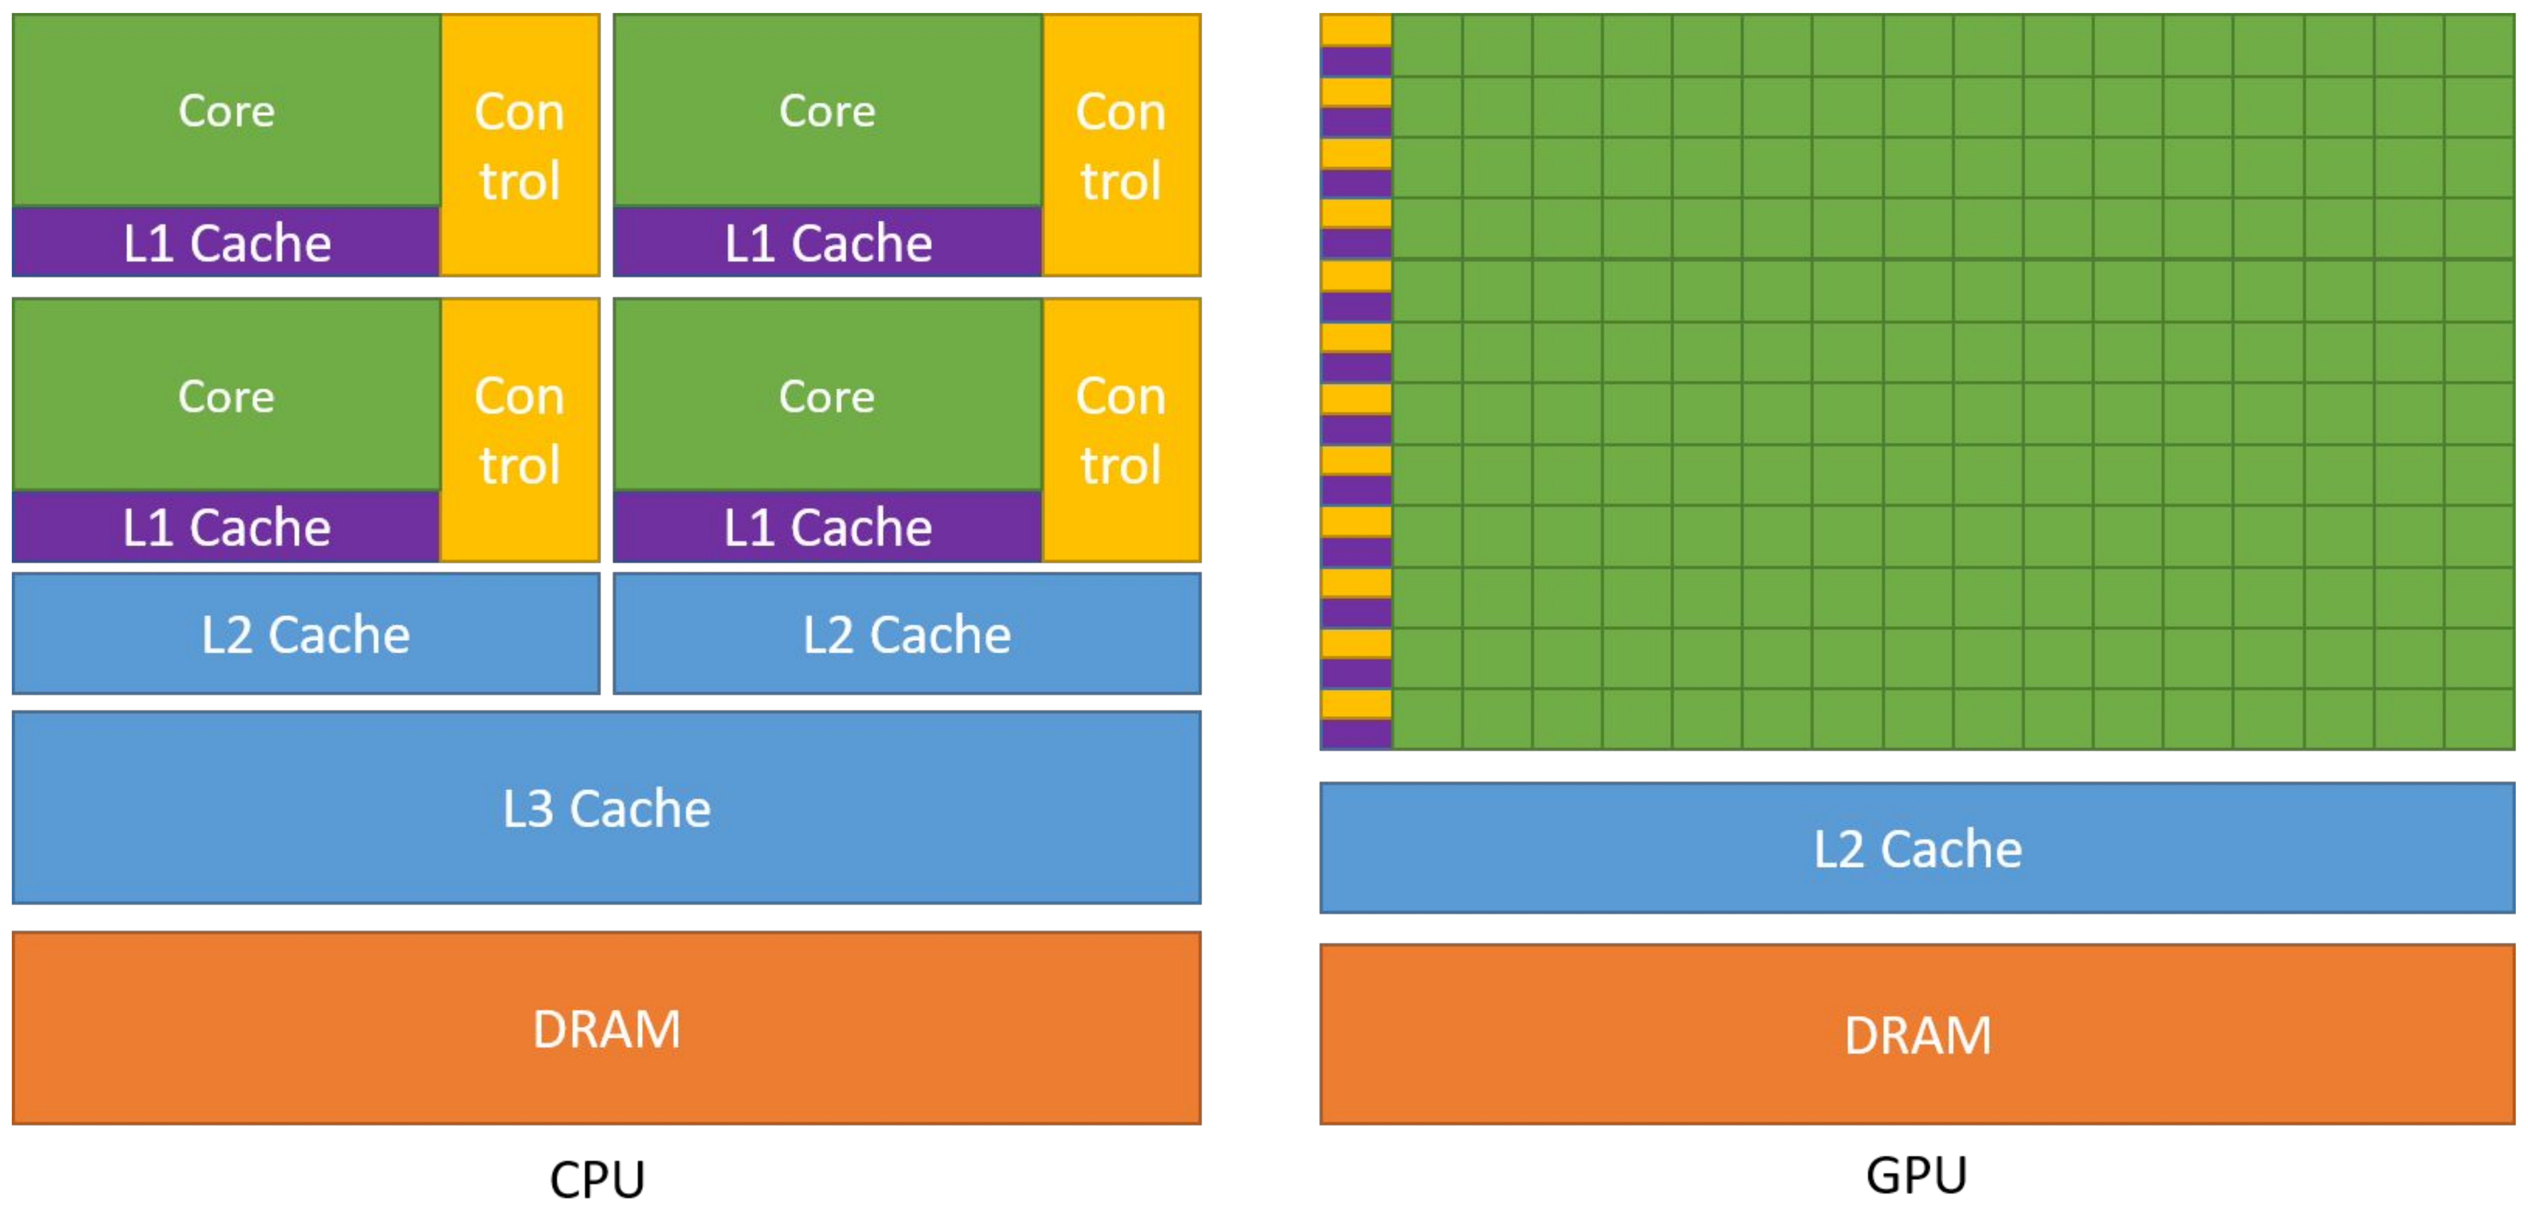

Solution — Store model parameters in GPU memory

GPU Infrastructure

How to Calculate the Computational Cost of a Model

VRAM Matters (On-board GPU Memory)

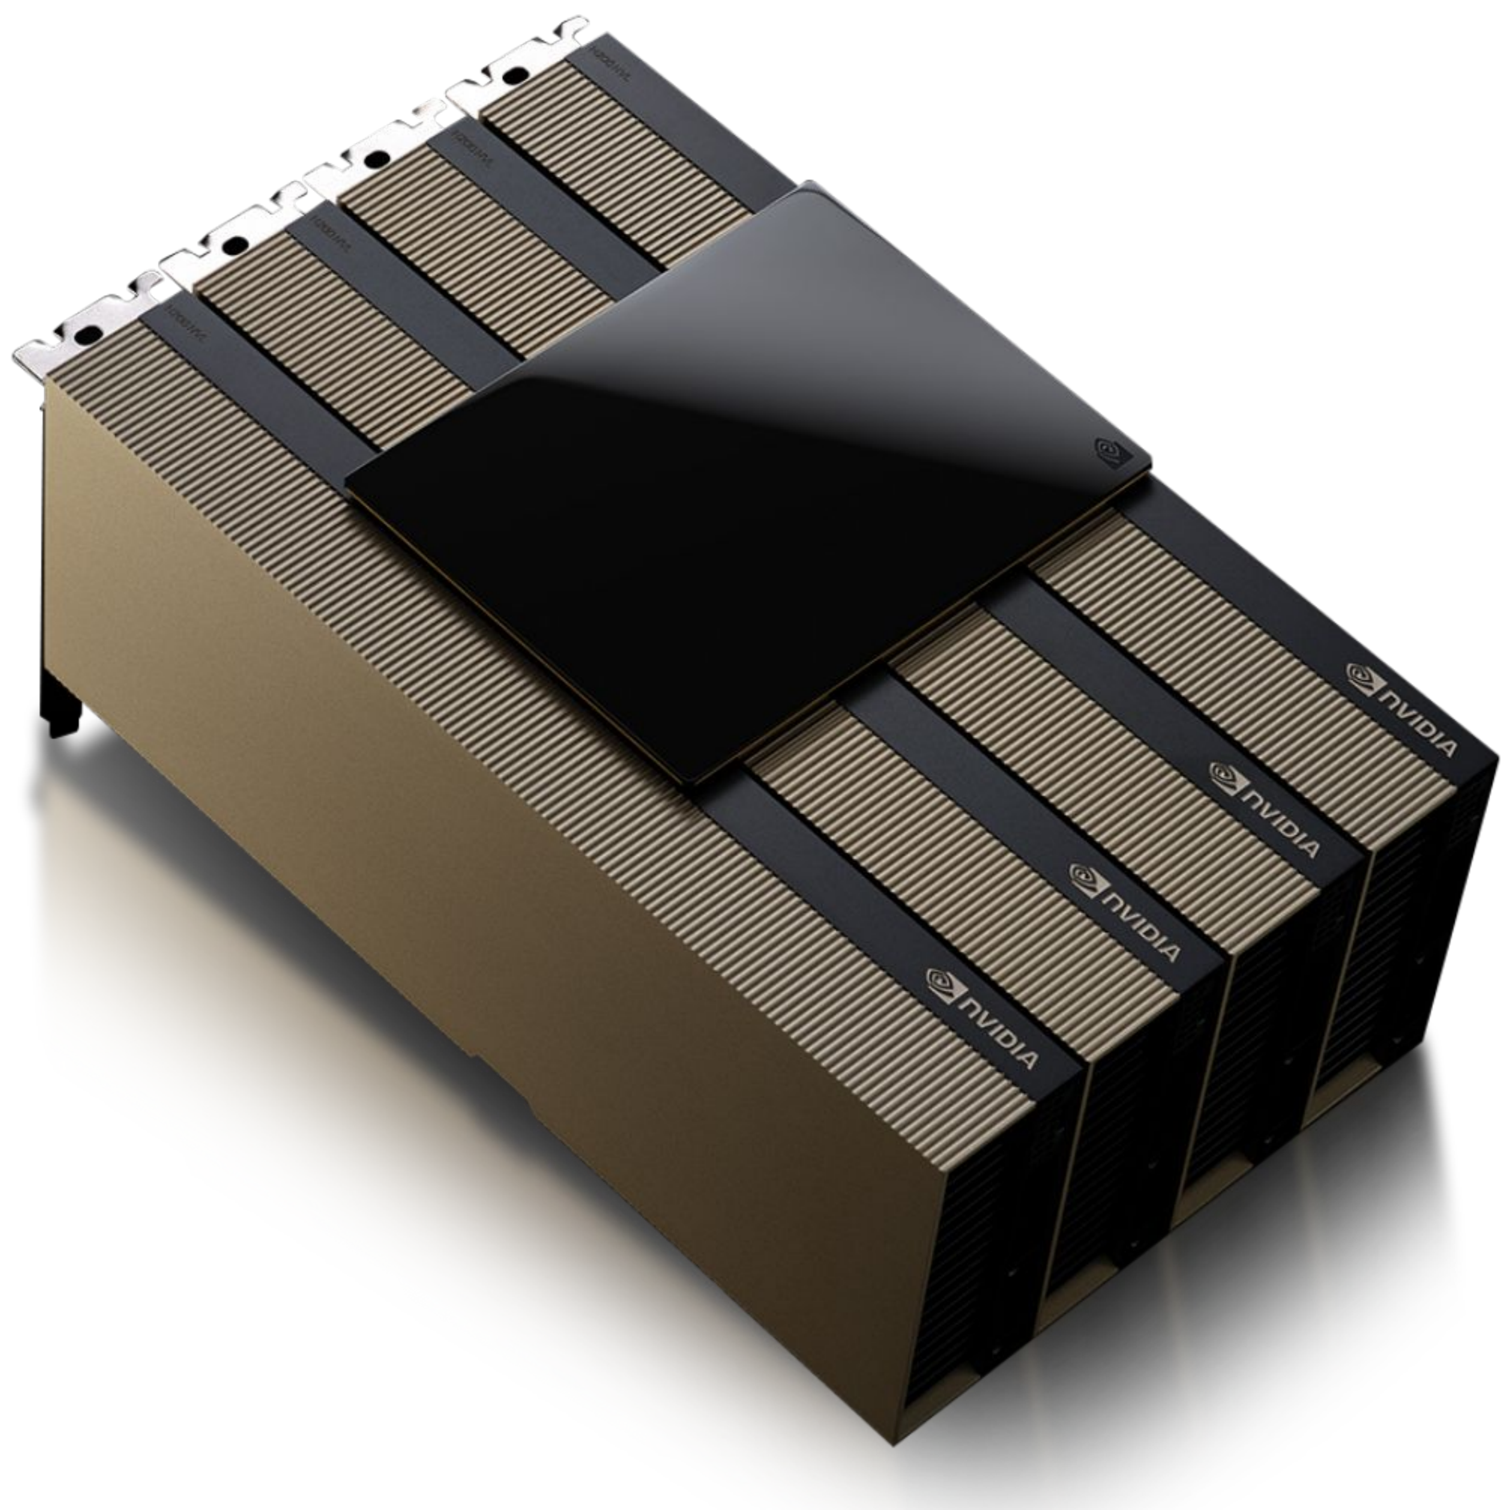

10 NVIDIA H100 HBM2e 80GB PCIe

H100 / PCIe x16 / 최대 AI TOPS: 3026(sparsity) / HBM(2세대) ECC / VRAM 대역폭: 2000 GB/s / 지원정보: 멀티VGA 지원 / 사용전력: 350W

24.04. 등록 | 의견 2 | 관심 | 대량구매

45,314,600원

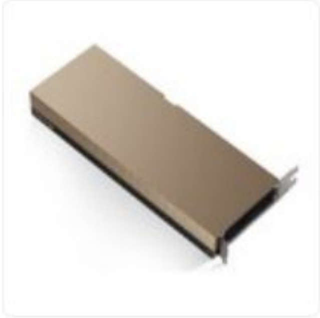

7 NVIDIA H100 HBM3 94GB NVL

H100 / PCIe5.0×16 / 스트림 프로세서: 16896 / 최대 AI TOPS: 3958(x2, sparsity) / HBM3 / VRAM 대역폭: 3900 GB/s / 지원정보: 멀티VGA 지원 / 사용전력: 400W

24.05. 등록 | 의견 20 | 관심 | 대량구매

44,738,950원

14 NVIDIA A100 HBM2e 80GB PCIe

A100 / PCIe4.0 / 스트림 프로세서: 6912 / 최대 AI TOPS: 1248(sparsity) / HBM2 / VRAM 대역폭: 2000 GB/s / 지원정보: 멀티VGA 지원 / 사용전력: 300W / FP16 Tensor Core: 312 TFLOPS / FP32: 19.5 TFLOPs

23.07. 등록 | 관심 | 대량구매

40,459,760원

12 NVIDIA RTX PRO 6000 Blackwell 서버 에디션 D7 96GB (병행수입)

RTX PRO 6000 Blackwell / PCIe5.0×16 / 가로(길이): 266.7mm / 스트림 프로세서: 24064 / GDDR7 ECC / 출력단자: DP2.1 / 지원정보: 8K 지원 / 사용전력: 600W / 두께: 39mm

25.07. 등록 | 관심 | 대량구매

27,499,990원

5 NVIDIA RTX PRO 6000 Blackwell 서버 에디션 D7 96GB 서린 (벌크)

RTX PRO 6000 Blackwell / PCIe5.0×16 / 가로(길이): 266.7mm / 스트림 프로세서: 24064 / GDDR7 ECC / 출력단자: DP2.1 / 지원정보: 8K 지원 / 사용전력: 600W / 무팬 / 두께: 39mm

25.07. 등록 | 브랜드로그 | 관심 | 대량구매

27,200,000원

Building RAG Agents with LLMs

THANK YOU SO MUCH!

Jae Y. CHOI

NVIDIA DLI Certified Instructor & University Ambassador